

DETERMINISTIC PUSHDOWN NORMALIZATION

From Completion-Aware DPDAs
to a General Theory

ALP EREN BÜTÜN

THEORY • CONSTRUCTION • PROOFS • GENERALIZATION

Independent Researcher

2023–2026

AUTOMATA THEORY • NORMALIZATION • COMPLETION • STRUCTURE

Research Monograph

Deterministic Pushdown Normalization

From Completion-Aware DPDAs to a General Theory

Author

Alp Eren Bütün
Independent Researcher

Research period

2023–2026

Edition note

This volume consolidates a sequence of research manuscripts developed between 2023 and 2026. The monograph preserves the chronological logic of discovery while consolidating repeated arguments, separating representation-specific results from representation-independent ones, and recording computational certificates beside the symbolic results they support.

Proof and validation note

Unless a result is explicitly identified as computer-assisted, the mathematical claims in this volume rest on symbolic arguments. Executable audits are used as independent consistency and reproducibility checks; reconstructed validation artifacts are identified as such in Appendix D.

Version

Content-locked research edition, August 2026.

No publisher, ISBN, or distribution licence is asserted on this page. Those publication details, if any, belong to a later publication edition.

Preface

This book grew from a single deterministic pushdown automaton.

In 2023 I was studying the binary counting language

$$L_{2,1} = \{w \in \{0, 1\}^* : N_1(w) = 2N_0(w)\}.$$

The first object of interest was not a general formalism but a concrete machine: a center state, two side states, and a stack manipulated by cancellation and replacement rules. The design was initially understood operationally. What made it the starting point of a larger research program was the realization that its state and stack were carrying more information than acceptance alone. After a prefix had been read, the configuration encoded a precise residual quantity and, after completion normalization, a shortest way of repairing that residual.

That observation supplied the first transition from a machine diagram to a mathematical theory. The stack ceased to be merely storage. It became a representation of completion demand. The side states ceased to be auxiliary control. They became finite corrections to the meaning of the stack. The cancellation transitions became local steps in a directed shortest-completion geometry.

The chapters that follow preserve this order of discovery. The original completion-aware DPDA is first analyzed on its own terms. Its reachable stores and state potentials lead to an exact residual invariant. Residuals then give a quotient coordinate for prefixes and a directed geometry of shortest continuations. That geometry supports progressively more economical realizations and, after the representation class is fixed carefully, exact resource frontiers. A related resource question then leaves the original coordinate system and becomes a representation-independent stack-alphabet hierarchy problem.

The research subsequently branches. Higher-dimensional counting systems turn shortest completions into cubical and local-action complexes. A different branch follows one cancellation pattern into Fuss–Catalan defect geometry and then into chronological and p -adic arithmetic. These branches are retained because they record genuine consequences of the completion viewpoint, but they are not allowed to obscure the central line of the book.

The central line returns to pushdown construction. For coprime positive p, q , a centered real-time literal-demand DPDA generalizes the original

2:1 architecture. Its finite-control rail and literal stack expose a WRITE–SEARCH–CANCEL normalization calculus. At that point a distinction that had been implicit from the beginning becomes explicit: the semantic process can be unique while its physical pushdown realization is not. This leads to the theory of exact semantic lifts, sections, representation debt, LIFO locality, normalization schedules, and causal obstruction.

The final abstraction is deliberately broader than weighted counting languages. In the General Theory, a deterministic semantic system is given first, together with a physical policy domain and an exact projection from physical configurations to semantic states. Questions about canonical sections, local realizability, safe and fatal debt, schedule conflicts, synthesis, and verification can then be asked without assuming that the semantic state space is the integer residual line or that the physical machine is one of the centered $p:q$ constructions. The weighted-balance family serves as a demanding test laboratory for the abstract theory, not as its definition.

Two editorial principles govern this monograph.

First, discovery history is preserved where it clarifies the mathematics. A later theorem is not retroactively presented as if it had motivated an earlier construction. In particular, the original 2023 machine remains the causal starting object, even when later chapters reveal cleaner coordinates or smaller realizations. The compact machines are conclusions of analysis, not substitutes for the original motivation.

Second, scope is treated as part of the theorem. Several different notions of “optimality” occur in this program: raw residual-faithful resource optimality, optimality inside a deferred-SEARCH literal-demand model, and lower bounds for broader DPDA representation classes. They are not interchangeable. Likewise, normalization to a common terminal representative is not promoted to a general Church–Rosser statement, and a fixed-policy causal obstruction is not silently upgraded to an unrestricted impossibility theorem. The book repeatedly marks these boundaries because many of the strongest results are strong precisely because their model assumptions are explicit.

The same principle applies to computation. Exhaustive programs were useful in three roles: discovering finite patterns, auditing symbolic formulas over large bounded regions, and certifying explicitly finite reductions. The first two roles do not replace proof. Where a theorem contains a genuinely finite computer-assisted component, the certificate and its scope are stated as such. Appendix D records the validation ledger and distinguishes recovered artifacts, reconstructed regression scripts, independently rerun checks, and purely symbolic results.

The book therefore has two simultaneous purposes. It is a consolidation of a sequence of papers, with repeated proofs merged and notation normalized. It is also a single mathematical narrative: a path from one completion-aware

DPDA to a general language for separating semantic evolution from physical normalization.

The final chapters return to the original $C/L/R$ machine after the abstract theory has been developed. This return is intentional. It allows the reader to see, in the same small automaton, the residual split between state and stack, literal completion demand, SEARCH debt, kernel materialization, schedule choice, and semantic lift structure. The aim of the abstraction is not to make the machine disappear. It is to explain why the machine was doing more than was initially visible.

Alp Eren Bütün
August 2026

Abstract

This monograph develops a theory of deterministic pushdown normalization from a concrete completion-aware DPDA for the binary language

$$L_{2,1} = \{w : N_1(w) = 2N_0(w)\}$$

to an abstract framework for exact semantic lifts, bounded LIFO locality, representation debt, causal scheduling, synthesis, and verification.

The first part analyzes the original 2:1 machine through the signed residual

$$\rho_{p,q}(w) = qN_1(w) - pN_0(w).$$

Its control state and logical stack satisfy an exact state-potential invariant, and end-of-input normalization converts the residual into a canonical shortest completion. The resulting completion problem has a directed-geodesic interpretation and, for general coprime p, q , a unique minimum Parikh demand characterized by the primitive kernel vector $K = (q, p)$. These coordinates lead to exact results on reachable stores, visibility, residual quotients, compact realizations, and resource frontiers. The monograph carefully separates raw residual-faithful optimality, model-relative deferred-SEARCH optimality, and representation-independent DPDA lower bounds.

The completion viewpoint is then extended in two directions. Higher-rank counting systems yield rectangular, cubical, and local-action completion complexes, including examples with nontrivial topology. A separate cancellation branch produces Fuss–Catalan defect geometry and a chronological carry refinement whose coefficient structure admits an exact Lagrange–Bürmann derivation and a p -adic interpretation.

Returning to machine construction, the book gives a centered real-time literal-demand DPDA for every coprime positive ratio $p:q$. The construction reveals WRITE, SEARCH, and CANCEL as a normalization calculus. At the semantic level, the outstanding demand evolves uniquely; at the physical level, LIFO order and finite-control sheets create representation debt. Quantity debt, order debt, and control/sheet debt are separated explicitly. The semantic quotient is unique while physical sectioned lifts need not be. In the 2:1 case, six

admissible reduced normalization schedules are pairwise non-isomorphic, a minimal causal dependency $A \rightsquigarrow B$ is exposed, and common terminal normalization is distinguished from unrestricted confluence.

The final theory abstracts away from weighted counting. A deterministic semantic system $(M, \Sigma, \{U_a\})$ is related to a physical policy domain by an exact projection Θ satisfying

$$\Theta \delta_a = U_a \Theta.$$

A canonical section induces an idempotent retraction, and bounded local observation leads to capacity and obstruction theorems. Behavioral safety is characterized by a greatest forward-invariant hull; fatal schedules are controlled by reachable unsafe fibers and robust conflict supports. Exact 3:2 and infinite odd- p :2 schedule classifications provide stress tests for the abstract theory. The framework also gives section-synthesis criteria, forced-debt lower bounds, regular saturation methods, and a deterministic pushdown equivalence fallback for exact fixed-policy verification.

Throughout, symbolic proofs remain primary. Large executable audits and finite certificates are recorded separately as reproducibility evidence. The result is a unified account of how completion semantics, pushdown representation, and normalization chronology interact—from a single 2023 DPDA to a general theory of semantic and physical normalization.

Contents

Preface	iii
Abstract	vi
Notation and Conventions	xxviii
Prologue — The Machine Before the Theory	1
I The Origin	3
1 The 2023 Completion-Aware DPDA	4
1.1 The motivating language	4
1.2 Two formal variants of the same core design	5
1.3 Recognition and completion are different operations	9
1.4 Removing the wrapper from the semantic analysis	10
1.5 The reachable stack is much more restricted than it looks . . .	11
1.6 The state carries part of the arithmetic	12
1.7 End-of-input completion transfers the state correction into the stack	14
1.8 A worked run: the prefix 111	15
1.9 The first structural surprise: two stack cells are enough	16
1.10 The cancellation transition that will reappear much later	17
1.11 What has been learned without leaving the original machine . .	18
2 What the Machine Is Really Storing	19
2.1 From the 2:1 imbalance to the general $p:q$ residual	19
2.2 The completion equation	20
2.3 The primitive balanced cycle	22
2.4 Every completion lies on one affine ray	23
2.5 Minimum completion demand	24
2.6 A canonical shortest word	26
2.7 The original DPDA implements this decoder at 2:1	28

2.8	Physical information versus semantic information	29
2.9	The residual transition system	30
2.10	Completion length as distance	31
2.11	How many shortest completions are there?	32
2.12	Why the minimum pair is a better semantic object than the canonical word	33
2.13	A first exact notion of redundancy	33
2.14	The original machine viewed through the new semantics	34
2.15	The first major conceptual compression	35
II	Completion Becomes Geometry	37
3	Minimum Demand and the Completion Ladder	38
3.1	Geodesic input symbols	38
3.2	The left branch as an exact shortest-path process	41
3.3	The ladder sources are actually reachable	42
3.4	Every position of the required one is shortest	43
3.5	Local commuting diamonds	46
3.6	Parikh synchronization	47
3.7	Completion and cancellation are opposite boundaries of one object	48
3.8	A commutation potential	49
3.9	Why this is more than a counting argument	50
3.10	The canonical completion is a boundary choice	51
3.11	What the ladder does not yet tell us	51
4	Quotients, Geodesics, and What Can Be Seen Locally	53
4.1	The residual classifies all future behavior	53
4.2	The quotient graph	55
4.3	Completion length is exactly directed distance	56
4.4	The fundamental balanced cycle	57
4.5	Balanced-cycle reduction	58
4.6	Which words are geodesic?	60
4.7	Exact quotient spheres	60
4.8	What the original two-cell observation really reveals	62
4.9	General canonical action visibility	64
4.10	A second bounded quantity: the completion frontier	66
4.11	One extremal word, two different bounds	69
4.12	The next representation idea is now forced	70

5	General $p:q$ Completion Rectangles	72
5.1	Remaining demand as a grid coordinate	72
5.2	The shortest interval is an actual rectangle	74
5.3	Exact number of shortest completions	76
5.4	Exact number of commuting squares	76
5.5	A worked 3:2 rectangle	77
5.6	Canonical order as a boundary normal form	78
5.7	Exactly when does a geodesic fork occur?	80
5.8	First obstruction at length $p + q$	81
5.9	Exact geodesic-word growth	82
5.10	Distinct paths versus distinct quotient vertices	84
5.11	The 2:1 ladder is now completely explained	85
5.12	From geometry back to representation	86
III	Better Coordinates and Exact Resource Frontiers	87
6	The First General Physical Lift and Its Defects	88
6.1	Finite phase plus unary magnitude	88
6.2	The stable four-symbol encoding	89
6.3	Stable rules away from the boundary	90
6.4	Forced boundary normalization	91
6.5	Exact normalized residual conjugacy	92
6.6	The abstract rectangles lift to physical configurations	93
6.7	A concrete 3:2 boundary crossing	94
6.8	Raw residual conjugacy	95
6.9	The defect set is finite and exact	96
6.10	Why the defects cannot be patched locally in the same coordinates	97
6.11	Which shortest rectangles are already raw?	98
6.12	The defect is coordinate-relative	99
7	From Boundary Defects to a Raw Three-Symbol DPDA	101
7.1	Before using two work symbols: the unary baseline	102
7.2	Boundary-compatible coordinates	103
7.3	Positive-tail arithmetic	104
7.4	The repaired sign boundary	105
7.5	Negative-tail arithmetic and the second repaired boundary . . .	106
7.6	Global raw conjugacy	106
7.7	The 3:2 boundary is now literal	109
7.8	Raw rectangles	110
7.9	What has actually been optimized?	110
7.10	The motivating 2:1 ratio survives arbitrary three-symbol encodings	111
7.11	Further exact special ratios	113

7.12	The right endmarker saves one control state in the scoped models	115
7.13	The coordinate-shift improvement	115
7.14	Three different meanings of “better”	116
7.15	The next question removes the representation shape	117
8	The Complete Raw Residual-Faithful Frontier	118
8.1	Model R	119
8.2	Why bounded physical height is impossible	119
8.3	The last occurrence of a height creates a stable suffix	120
8.4	Positive residue classes synchronize	122
8.5	Why one work symbol is worse	123
8.6	The complete raw frontier	124
8.7	What this theorem still does not say	125
8.8	Acceptance without a right endmarker	126
8.9	Can the optimum survive bounded pushes?	127
8.10	A stronger semantic-visibility budget	128
8.11	A superlinear exact family	130
8.12	Direct lossless factorization: counting semantic cells	131
8.13	Exact normalization-rule cost under push-at-most-two	131
8.14	Four resource questions, four different answers	133
8.15	What is now genuinely settled	134
IV	When the Resource Question Leaves Our Representation	136
9	From Residual Frontiers to a Representation-Independent Open Problem	137
9.1	The methodological jump	137
9.2	The external hierarchy question	138
9.3	Why realtime separation does not automatically survive ε -moves	139
9.4	Lineage from the original DPDA	140
9.5	The later 2:1 improved DPDA as the immediate precursor . . .	140
9.6	Why the witness development did not end with a direct encoding trick	141
9.7	The general DPDA model used in the lower bound	142
9.8	Modes	142
9.9	The lower-bound target	143
10	Completion Slopes and ε-Robust Pumping	145
10.1	The binary realtime witness	145
10.2	Four rule families	146
10.3	Upper-bound theorem	148

10.4	The exact completion-slope slice	148
10.5	The smallest new case	150
10.6	The lower-bound opponent is arbitrary	151
10.7	Raw continuation principle	152
10.8	No raw configuration can repeat on a live slope ray	152
10.9	Raw stack height must escape	153
10.10	Suffix-minimum positions	153
10.11	The ε -robust step pump	154
10.12	What remains to prove	156
11	One Mode, One Slope	157
11.1	Unary completion uniqueness	157
11.2	Slope separation	158
11.3	The product lower bound	161
11.4	The hierarchy theorem	163
11.5	What exactly has been resolved	164
11.6	Why this is not a realtime lower bound in disguise	164
11.7	Why the product, rather than states or symbols separately, appears	165
11.8	What the theorem does not prove	165
11.9	The methodological meaning for the monograph	166
11.10	Where the story goes next	167
V	Completion Geometry Beyond the Binary Line	169
12	When Minimum Demand Becomes a Fiber	170
12.1	Matrix counting languages	170
12.2	Completion totality	171
12.3	Residuals still classify right quotients	173
12.4	Minimum completion length	174
12.5	The optimal Parikh fiber	174
12.6	Equal-length trades	175
12.7	Affine-semigroup viewpoint	176
12.8	Continuous optimum geometry is not the same as integer multi- plicity	176
12.9	The shortest-prefix down-set	177
12.10	Exact geodesic-prefix criterion	178
12.11	First actions and compatible action sets	180
12.12	The conceptual transition	180

13 The Filled Shortest-Prefix Quotient	182
13.1 Why vertices alone are not enough	182
13.2 Cubical realization of a down-set	183
13.3 Equal-length identifications	183
13.4 The quotient theorem	184
13.5 Recovery of the binary rectangle	186
13.6 When is the quotient literally a box?	186
13.7 Unique optimum versus unique shortest word	188
13.8 Local compatibility and global identification are different . . .	189
13.9 The rank of equal-length trade	189
14 Local and Global Topology	191
14.1 The local action complex	191
14.2 Squarefree-divisor interpretation	192
14.3 Rank zero: one simplex	193
14.4 Rank one: the optimum fiber is an arithmetic interval	194
14.5 Rank two already permits a local hole	195
14.6 Rank-one global geometry	197
14.7 Overlap complexes	197
14.8 Periodic saturation	198
14.9 The global circle	200
14.10 Local triviality, global nontriviality	201
14.11 What the topology means operationally	201
15 Scalar Rank Hierarchy and Exact Three-Letter Geometry	203
15.1 Scalar counting systems	203
15.2 The $d = 2, 3, 4+$ hierarchy	204
15.3 Three-letter normal form	205
15.4 Parameterizing the optimum fiber	206
15.5 Exact three-letter quotient vertex count	207
15.6 Exact cubical f -vector	208
15.7 Worked example	209
15.8 What changed relative to the binary rectangle?	210
15.9 Why this hierarchy belongs in the same monograph	210
16 Eventual Geodesic Arithmetic	212
16.1 The continuous directional slope	212
16.2 Dual certificate and reduced costs	213
16.3 Zero-defect shifts	214
16.4 Stabilized defect profile	215
16.5 Why τ_b already contains the slope denominator	217
16.6 Fundamental defect period	217
16.7 Eventual local action geometry	219

16.8	What is and is not being claimed	220
16.9	The higher-dimensional branch in one picture	221
16.10	Returning to the original machine from the opposite direction	221
VI	One Cancellation Transition Becomes a Geometry	223
17	Zooming Back into the Original DPDA	224
17.1	The machine facts we reuse, not re-prove	224
17.2	The nonnegative path sector	225
17.3	Exact landing-height correspondence	226
17.4	From one machine event to a d -Dyck statistic	227
17.5	Chronological correction signatures	228
17.6	Path defect	229
17.7	The new inverse problem	229
18	Correction Signatures and Minimum Defect	230
18.1	The threshold function	230
18.2	Local Gap Cost	231
18.3	Terminal Cost	232
18.4	Intrinsic signature defect	233
18.5	Canonical blocks	233
18.6	Defect–Signature Characterization	234
18.7	Canonical does not mean unique	236
18.8	Signature persistence	237
18.9	Maximum number of corrections	237
18.10	Zero defect	237
18.11	Defect as penalized overshoot	238
19	The Catalan Boundary and the Binary Cancellation Skeleton	239
19.1	Two unavoidable unmarked nodes	239
19.2	The d entrance choices	240
19.3	Why correction nodes become binary	240
19.4	Binary Cancellation Skeleton	241
19.5	Tree coordinates	242
19.6	Correction height from the binary address	243
19.7	Inorder chronology	243
19.8	The signature is lossless	244
19.9	Area statistic	245
19.10	What the original transition has become	245

20 Exact Correction and Defect Enumeration	247
20.1 Scalar correction generating function	247
20.2 The auxiliary kernel and A371395	248
20.3 Exact correction triangle	249
20.4 Row sums and boundary columns	250
20.5 Intrinsic signature enumeration	251
20.6 Translation and first-zero decomposition	251
20.7 Defect–Signature Master Equation	252
20.8 Exact defect coefficients	252
20.9 Cumulative realizability	253
20.10 Three arrays, three questions	254
21 The Asymptotic Defect Hierarchy	255
21.1 Fixed-defect generating functions	255
21.2 Algebraicity and singularity location	256
21.3 Critical scaling	257
21.4 Pairwise leading singularities	258
21.5 Asymptotic Defect Hierarchy	258
21.6 The defect pairs	260
21.7 Geometric interpretation	260
21.8 Cumulative asymptotics	261
21.9 What the cancellation branch has established	261
21.10 The next arithmetic offshoot	263
VII Chronological Carry and the Arithmetic Offshoot	264
22 From Ordered Corrections to CCR-3	266
22.1 Scientific boundary	266
22.2 Connected supports through order three	267
22.3 The CCR-3 recurrence	267
22.4 Hessenberg realization	269
22.5 Operator form	270
22.6 What this chapter has and has not established	271
23 Fixed Columns and the Universal Stirling Shadow	272
23.1 Column coefficients	272
23.2 Transport equations	272
23.3 First three fixed columns	273
23.4 Reduction modulo p	274
23.5 Normalized shadow polynomials	275
23.6 Why the third column is the first singular one	276

24 The Third-Column p-Adic Split	277
24.1 The simple branch	277
24.2 The double shadow branch	279
24.3 Generic residue classes	280
24.4 The scaled third column	280
24.5 Independent consistency checks	281
24.6 The unresolved algebraic question	281
25 Ramified Recovery	282
25.1 Exact discriminant factorization	282
25.2 Removing the integral root	283
25.3 Half-integral roots	283
25.4 Ramified square class by $p \bmod 8$	284
25.5 The complete third-order arithmetic picture	284
25.6 Novelty boundary	285
25.7 Do not extrapolate CCR-3 blindly	286
25.8 Re-entry to the main DPDA story	286
VIII Returning to the Machine: General Real-Time $p:q$	288
26 Why Static Two-Visibility Fails	289
26.1 We do not re-prove minimum demand	289
26.2 A two-visible canonical word	290
26.3 The 3:2 obstruction	291
26.4 Static two-visible obstruction	291
26.5 SEARCH	292
27 Debt, Virtual Demand, and Minimal WRITE	293
27.1 Virtual demands	293
27.2 Virtual-demand projection	294
27.3 The WRITE equation	295
27.4 Minimal kernel carry	295
27.5 Why the historical long writes existed	296
27.6 From arbitrary debt to a finite rail	297
28 The Centered Real-Time $p:q$ DPDA	298
28.1 The centered rail	299
28.2 Residual invariant	301
28.3 Why the debt rail can stop at h_{ctr}	302
28.4 Boundary guard	304
28.5 Canonical segments	306
28.6 Two-cell qualitative visibility	307

28.7	Exact WRITE normalization	307
28.8	Zero residual	308
28.9	Main construction theorem	309
28.10	Shortest-demand quotient	310
29	What “Optimal” Means in the Deferred-SEARCH Architecture	311
29.1	Canonical deferred-SEARCH model	311
29.2	Residual distinguishability	312
29.3	One witness containing enough of both symbols	312
29.4	Exhaust the same stack, leave different debts	314
29.5	Exact deferred width	315
29.6	Scope — place this beside the theorem	316
29.7	Why Chapter 8 gives a different optimum	316
29.8	Why this chapter belongs before the 2:1 recovery	317
30	Closing the First Loop: Recovering the Original 2:1 Machine	318
30.1	Center mismatch directions	319
30.2	Left cancellation rail	319
30.3	Right locality-repair rail	319
30.4	Historical WRITE arithmetic	320
30.5	Specialization and recovery theorem	320
30.6	Lazy versus eager discharge	321
30.7	Canonical materialization map	322
30.8	Eager-lazy semantic equivalence	322
30.9	Historical machine as an eager refinement	323
30.10	Computational audit	323
30.11	The real discovery at the end of the $p:q$ generalization	324
IX	Write–Search–Cancel as a Structural Theory	326
31	Semantic Arithmetic Before Representation	327
31.1	Dominating kernel level	327
31.2	One input symbol: CANCEL plus optional carry	329
31.3	Semantic WRITE count	330
31.4	Demand-length clock	331
31.5	Semantic chronology is unique	331
32	What One-Top LIFO Locality Forbids	332
32.1	Exact literal-demand lift	332
32.2	Fixed block orientation	333
32.3	Top-rewrite suffix criterion	333
32.4	Complete obstruction classification	334

32.5	Two fundamentally different failures	335
33	Quantity, Order, and Control Debt	337
33.1	Quantity defect	337
33.2	Universal representation conservation	338
33.3	Order defect	338
33.4	Three-channel canonicity	339
33.5	Kernel economy	340
33.6	Forced unit SEARCH	340
33.7	Carry obstruction with zero quantity debt	342
33.8	Minimum order debt	342
34	The WSC Count Shadow	344
34.1	Phase coherence	344
34.2	WSC defect lattice	345
34.3	Defect coordinates	346
34.4	Defect quotient	346
34.5	WSC lattice is not the full LIFO fiber	347
34.6	Kernel economy in lattice coordinates	347
34.7	Abstract count-shadow fiber region	348
35	Physical Materialization, Backlog, and WSC Factorization	349
35.1	Materialization level	349
35.2	Physical length formula	350
35.3	Materialization amount	350
35.4	Backlog law	351
35.5	Canonical carry dichotomy	352
35.6	Unique physical WSC role factorization	352
35.7	Reconstructing the charges	354
35.8	Global materialization budget	354
35.9	Deferral capacity	355
36	Unique Semantic Quotient, Nonunique Physical Presentation	356
36.1	Semantic fibers	356
36.2	Fiber congruence	357
36.3	Canonical retraction	357
36.4	What is actually unique	358
36.5	Cheap nonuniqueness already exists	359
36.6	Why “confluence” is not yet justified	360

37 The Original 2:1 Machine as WSC	361
37.1 Historical states as SEARCH sheets	361
37.2 Historical stack-length changes	362
37.3 The global identity $b + 2c = n$	363
37.4 What Structure Theory has accomplished	363
37.5 The Structure Theory endpoint	364
 X Semantic Uniqueness and Physical Lift Freedom	 365
 38 The Semantic Core Belongs to the Language	 366
38.1 Right residuals	366
38.2 Universal semantic quotient	367
38.3 Minimum-demand semantics is a coordinate change	368
38.4 Three levels of description	368
 39 Exact Sectioned Lifts and Category-Relative Minimality	 370
39.1 Reachable versus ambient configurations	370
39.2 Exact sectioned lift	370
39.3 Section-preserving morphisms	371
39.4 Minimality is category-relative	372
 40 When LIFO Locality Forces Fiber Branching	 373
40.1 Uniform local rule	373
40.2 Canonical-section closure dichotomy	373
40.3 Essential physical fiber branching	374
40.4 Quotient-protected branching	374
 41 Two Reduced 2:1 Lifts	 376
41.1 Historical eager lift	376
41.2 Centered lazy lift	376
41.3 Section incomparability	376
41.4 Coordinate reducedness	377
41.5 Morphism-relative nonuniqueness	377
 42 The Six 2:1 Normalization Schedules	 379
42.1 Sites	379
42.2 Exact schedule classification	379
42.3 Schedule geometry	380
42.4 Six reduced lifts	380
42.5 Terminal normalization coherence	381

43 Why A Causes B	383
43.1 Local carry visibility	383
43.2 Portability	383
43.3 A expands the B -fiber	384
43.4 Direct canonicalization becomes impossible	384
43.5 Forced kernel excess	385
44 Directed Causal Scheduling	386
44.1 First-hit exposure	386
44.2 Exact 2:1 primary dependency	386
44.3 Concrete causal certificate	387
44.4 Coordinate invariance	388
44.5 But semantics does not determine causality	389
45 Which Obstructions Are Really Forced?	390
45.1 Nested representation classes	390
45.2 Obstruction monotonicity	390
45.3 Branching profile	391
45.4 Current status of $A \rightarrow B$	391
45.5 Representation–obstruction separation	392
45.6 Open problems retained	392
45.7 Reproducibility note	392
45.8 Transition to General Theory	393
XI General Theory of Exact Pushdown Normalization	394
46 Semantic Fibers and Canonical Retraction	395
46.1 Exact semantic lift	395
46.2 Canonical fibers	396
46.3 Canonical absorption	397
46.4 Zero debt and raw conjugacy	397
46.5 Three levels now separated permanently	398
47 Bounded-Local Realizability and the Three-Stage Obstruction	399
47.1 Uniform one-top portability	399
47.2 Finite observation capacity	400
47.3 Bounded- d -top pushdown capacity	400
47.4 The three stages	401
47.5 One-top specialization	401
47.6 Locality–debt–recoordination trilemma	402

48 Representation Debt and Behavioral Safety	403
48.1 Terminal semantics	403
48.2 Immediate disagreement	404
48.3 Behavioral hull	405
48.4 Forward invariance	405
48.5 Debt birth dichotomy	405
48.6 Two opposite examples	406
49 Semantic Containment Cores	407
49.1 Semantic trap containment	407
49.2 Absorbing-fiber harmlessness	408
49.3 Maximal full-fiber containment core	408
49.4 Reachable-fiber refinement	408
49.5 Root correctness via containment	409
50 Safe Reset and the Fatal Frontier	410
50.1 Wordwise safe-reset domain	410
50.2 Structural unsafe targets	410
50.3 Fatal Frontier	411
50.4 Activation tightness	411
50.5 Site-free diagnostics	411
50.6 Fatal causality	412
51 Conflict Complexes and Weighted-Balance Stress Tests	413
51.1 Minimal conflicts	413
51.2 Robust fatal supports	413
51.3 Reset-ray fork	414
51.4 Imported weighted-balance foundations	414
51.5 Exact 3:2 case	414
51.6 Infinite odd- p :2 family	415
52 Rooted Section Synthesis	419
52.1 Rooted path-outcome fibers	419
52.2 Path-outcome coherence	420
52.3 Finite synthesis algorithm	420
52.4 Short obstruction certificate	421
52.5 Counting sections	421
52.6 Infinite semantic cores require symbolic synthesis	421
53 Symbolic Re-Coordination and Forced-Debt Frontiers	422
53.1 Modular quotient/remainder synthesis	422
53.2 Finite symbolic skeleton	423
53.3 Partial canonical sections	423

53.4	Canonical-hit capacity	424
53.5	Quantitative forced debt	424
53.6	Multi-type Dyck frontier	424
54	Effective Verification Hierarchy	426
54.1	Level I — Structural safe-reset saturation	426
54.2	Level II — Exact immediate-disagreement saturation	426
54.3	Reachable-containment shortcut	427
54.4	Finite semantic containment computation	427
54.5	Level III — DPDA-equivalence fallback	428
54.6	Verification hierarchy	428
54.7	General Theory synthesis	429
XII	Closing the Loop	431
55	The 2023 DPDA Revisited	432
55.1	The two origin figures remain distinct	432
55.2	What the stack was storing	436
55.3	What the completion rules were doing	436
55.4	Why depth two appeared	437
55.5	The cancellation transition	437
55.6	Why the improved raw DPDA looked so different	438
55.7	Why the general $p:q$ machine returned to the original architecture	438
55.8	Why the historical table was not literally the centered specialization	439
55.9	C/L/R after Structure Theory	439
55.10	C/L/R after Uniqueness Theory	440
55.11	The original machine after General Theory	440
56	What Was Forced, What Was Chosen, What Became General	443
56.1	Language-semantic necessities	443
56.2	Necessities of a chosen locality category	444
56.3	Coordinate artifacts	444
56.4	Normalization-schedule choices	445
56.5	Special 2:1 features	445
56.6	What generalized to all coprime $p:q$	446
56.7	What generalized beyond rational binary counting	446
56.8	The research story in one chain	447
57	Open Problems and the Continuing Research Program	449
57.1	Unrestricted recognition complexity of $L_{p,q}$	449
57.2	Broader resource hierarchies	449
57.3	General causal schedule theory for $p:q$	450

57.4	Fixed-policy status of $A \rightarrow B$	450
57.5	Unrestricted lift minimality	450
57.6	Confluence and schedule rewriting	451
57.7	Higher-dimensional completion topology	451
57.8	Higher-order correction geometry	451
57.9	Fourth-order chronological arithmetic	452
57.10	General section synthesis	452
57.11	Verification complexity	452
57.12	Reproducibility tasks before publication	453
57.13	Final perspective	453
A	Complete Transition Audit of the Original 2023 DPDA	455
A.1	Purpose and scope	455
A.2	Wrapper projection	456
A.3	Center-state recognition transitions	457
A.4	Left-state recognition transitions	459
A.5	Right-state recognition transitions	461
A.6	Complete recognition-phase audit table	463
A.7	Reachable-store closure	464
A.8	Audit of the end-of-input completion rules	465
A.9	From residual preservation to the canonical shortest completion	467
A.10	Two representative runs	468
A.11	Determinism of the completion extension	469
A.12	Exhaustive computational cross-check	469
A.13	Conclusion	469
B	General Construction Transition Schemas	471
B.1	Purpose and model separation	471
B.2	Four-symbol normalized paired-sign lift	472
B.3	One-work-symbol raw conjugacy	474
B.4	Boundary-compatible three-symbol raw DPDA	475
B.5	Centered real-time literal-demand construction	477
B.6	Minimal kernel-carry compiler identity	479
B.7	The 2:1 specialization and the eager–lazy distinction	479
B.8	Determinism and implementation conventions	480
B.9	Summary table	480
C	Scoped Lower Bounds and Finite Certificates	482
C.1	Why this appendix is necessary	482
C.2	Lower-bound scope table	482
C.3	Raw residual-conjugacy frontier	483
C.4	The 3:2 finite reduction	484
C.5	Normalization of the minimum configuration	485

C.6	Legal branching and pruning rules	485
C.7	Exact survivor vectors	486
C.8	Certificate reproducibility	487
C.9	Architecture-relative deferred-SEARCH width	487
C.10	Representation-independent hierarchy product bound	488
C.11	Final scope rule	489
D	Reproducibility and Validation Ledger	490
D.1	Purpose	490
D.2	Provenance classes	490
D.3	Recovered and independently rerun suites	491
D.4	Reconstructed P10 validation suite	493
D.5	Reconstructed P11 validation suite	494
D.6	P12 and executable support	494
D.7	Preserved monograph code package	495
D.8	Interpretation rule	495
E	Lagrange–Bürmann Derivation of the Exact Correction Triangle	496
E.1	The marked first-return grammar	496
E.2	The auxiliary kernel variable	496
E.3	Coefficient formula for the auxiliary kernel	497
E.4	Recovering the correction generating function	498
E.5	Lagrange–Bürmann extraction for a fixed summand	499
E.6	Exact correction triangle	500
E.7	Equivalent elimination through $U = A - 1$	500
E.8	Row sum	501
E.9	Why the kernel and correction arrays differ	501
F	General-Theory Stress-Test Catalogues	503
F.1	Purpose and scope	503
F.2	The exact centered 3:2 certificate	503
F.2.1	Finite reset repertoires	504
F.2.2	Direct verification of reset closure	505
F.2.3	Exact schedule classification and fatal witnesses	506
F.3	The infinite odd- p :2 family	507
F.3.1	Arithmetic catalogue	507
F.3.2	Source-geometry catalogue	508
F.3.3	Reset-target catalogue and closure	509
F.3.4	Uniform D -induced source-fiber expansion	510
F.3.5	Uniform reachability of the length-three frontier	510
F.3.6	Uniform kernel-suffix fatality	511
F.3.7	Exact odd- p :2 classification	511

F.4 Interpretive boundary	512
Bibliography	513

List of Figures

1.1	Formalized redrawing of the original 2:1 DPDA variant with the empty input excluded. Solid arrows are ordinary recognition transitions; dashed callout boxes contain only right-endmarker completion rules. The listed left-state completion rule with top 1 is retained for fidelity but has an unreachable source condition.	6
1.2	Formalized redrawing of the original 2:1 DPDA variant with the empty input included. The center/left/right recognition logic and end-of-input completion rules are identical to Figure 1.1; only the initialization/final-state wrapper changes so that ε is accepted.	8
3.1	The exact configuration ladder for $k = 2$. The source $(L, 00)$ has residual 7 and minimum demand $(4, 1)$. Its five shortest continuations are 10000, 01000, 00100, 00010, 00001. The upper boundary is canonical completion; the lower boundary postpones the unique required 1 maximally.	46
5.1	The 3:2, $r = 2$ shortest-completion rectangle. Horizontal edges read 0 and subtract 3 from the residual; vertical edges read 1 and add 2. The nine residual vertices form four commuting squares, and every monotone path from residual 2 to residual 0 is one of the six shortest completions with Parikh pair $(2, 2)$.	79
7.1	The improved two-state, three-symbol raw residual-conjugate DPDA for the motivating 2:1 residual graph. Labels have the form input, top/replacement; only transition cells occurring on the encoded residual orbit are shown. The right endmarker is accepted exactly at (R_0, Z_0) .	113
14.1	The local action complex for $\omega = (-11, -2, 1, 10)$ at residual 7. Its maximal faces are $\{2, 3\}$, $\{1, 3\}$, and the filled triangle $\{1, 2, 4\}$. The cycle $1 - 2 - 3 - 1$ remains unfilled, giving $H_1 \cong \mathbb{Z}$.	196

15.1	The complete rank-one quotient for $\omega = (-5, -3, 1)$ and $r = 9$. The down-set below $(2, 0, 1)$ contributes two filled squares; the path below $(0, 3, 0)$ shares the origin, and its far endpoint is identified with $(2, 0, 1)$ by the equal-length kernel translation. The quotient has $(f_0, f_1, f_2, f_3) = (8, 10, 2, 0)$ and is homotopy equivalent to S^1	210
28.1	Schematic centered debt rail. Labels 1/0 and 0/1 denote input/top mismatches. Interior mismatches pop the visible cell and move debt by one; an outward mismatch at $\pm h_{\text{ctr}}$ triggers WRITE back to the center.	300
42.1	Hasse diagram of the six admissible 2:1 schedules. The schedule bits are ordered by inclusion of eager sites A, B, C . The two omitted cube vertices 110 and 111 are shown separately as forbidden because A and B cannot be simultaneously eager. . .	381
55.1	The empty-input-excluded original DPDA revisited after the general theory. This is the same formalized machine as Figure 1.1, deliberately repeated so that the final interpretation can be read against the original object.	433
55.2	The empty-input-included original DPDA revisited. This is the same recognition/completion core as Figure 1.2; the wrapper accepts ε	435
55.3	The original $C/L/R$ architecture annotated by the complete theory. The figure is interpretive rather than a replacement transition table: it overlays the residual/state-potential invariant, WSC SEARCH-sheet coordinates, literal-demand magnitude, kernel materialization, end-of-input completion, schedule freedom, and the semantic-lift viewpoint on the original machine architecture.	441

Notation and Conventions

This page records the notation used across several layers of the monograph. Symbols that are intentionally local to a chapter are defined again where they are used.

Words and weighted balance

$N_a(w)$ Number of occurrences of the symbol a in the word w .

$L_{p,q}$ For coprime positive integers p, q ,

$$L_{p,q} = \{w \in \{0, 1\}^* : qN_1(w) = pN_0(w)\}.$$

$\rho_{p,q}(w)$ Signed residual

$$\rho_{p,q}(w) = qN_1(w) - pN_0(w).$$

Thus reading 0 changes the residual by $-p$, while reading 1 changes it by $+q$.

m The total weight $m = p + q$ when the $p:q$ family is under discussion.

K Primitive nonnegative kernel vector

$$K = (q, p),$$

which preserves the weighted completion equation.

$\mu(r)$ The canonical minimum nonnegative completion-demand vector associated with residual r ; when written (x, y) , x counts appended zeros and y appended ones.

$\eta(x, y)$ The canonical literal stack word associated with demand (x, y) in the relevant orientation, typically a fixed two-block representative such as $1^y 0^x$.

Original and centered machines

C, L, R Center, left, and right states of the original 2:1 completion-aware DPDA. In the original-machine chapters their state potentials are $\Phi(C) = 0$, $\Phi(L) = 3$, and $\Phi(R) = -3$.

$D(S)$ Original-machine logical-stack weight

$$D(S) = 2N_0(S) - N_1(S).$$

Together with the state potential it represents the prefix residual.

h_{ctr} Centered-rail half-width

$$h_{\text{ctr}} = \left\lfloor \frac{p+q-1}{2} \right\rfloor.$$

The subscript distinguishes this scalar from the WSC exchange vector below.

D_d State at centered-rail coordinate d ; $D_0 = C$. In the centered construction the state potential is proportional to dm .

WRITE–SEARCH–CANCEL coordinates

P Physical count shadow of a literal stack or physical representation.

Θ Canonical semantic demand/count target represented by a physical configuration. In the abstract General Theory the same letter denotes the semantic projection from physical configurations to semantic states.

Q Quantity defect, typically

$$Q = \Theta - P.$$

Ω Order defect: the part of physical noncanonicity that is invisible to Parikh counts but visible to LIFO order.

$\mathbf{h}$ WSC exchange direction

$$\mathbf{h} = (1, -1).$$

It is typographically distinct from h_{ctr} .

M_t Kernel-materialization count at time t in a physical WSC factorization.

General Theory

M	Semantic state space of a deterministic semantic system. This use of M is local to the General Theory and should not be confused with the WSC scalar M_t .
Σ	Input alphabet.
$U_a : M \rightarrow M$	Semantic update induced by input symbol a .
C	Physical policy/configuration domain in the abstract General Theory. This is a context-local reuse of the letter C ; it is not the center state of the earlier machine chapters.
$\delta_a : C \rightarrow C$	Physical update under input symbol a , on the policy domain where it is defined.
$\Theta : C \rightarrow M$	Exact semantic projection satisfying $\Theta \delta_a = U_a \Theta.$
$\sigma : M \rightarrow C$	Chosen canonical section, with $\Theta \sigma = \text{id}_M$.
N	Canonical retraction $N = \sigma \Theta.$ It satisfies $N^2 = N$ on its domain.
$\mathcal{D}_H$	Immediate-disagreement or bad set associated with a fixed physical policy or schedule H , as defined in the relevant chapter.
$\mathcal{B}_H$	Behavioral safe hull for policy H , equivalently the greatest forward-invariant subset that avoids eventual observable disagreement under the stated hypotheses.
$\text{Adm}(\mathcal{E})$	Admissible schedule family for a schedule system or environment $\mathcal{E}$.

Scope conventions

“Optimal”, “minimal”, and “unique” are always read relative to the model specified in the theorem. In particular, the monograph distinguishes raw residual-faithful resource frontiers, deferred-SEARCH literal-demand width, and unrestricted DPDA recognition resources. A common terminal normal form is not, by itself, a claim of Church–Rosser confluence.

The symbol $\dashv$ denotes an explicit right endmarker. Completion callouts selected by $\dashv$ are not ordinary competing ε -moves.

Prologue — The Machine Before the Theory

This monograph begins with a machine, not with an abstract theory.

The starting problem is the binary counting language

$$L_{2,1} = \{w \in \{0,1\}^* : N_1(w) = 2N_0(w)\}.$$

At the level of arithmetic, the condition is elementary. One can measure the imbalance between the numbers of zeros and ones by a signed linear expression. At the level of language recognition, this already suggests counter-like memory.

The research program developed in this monograph began somewhere more concrete. In 2023, a deterministic pushdown design was constructed for this language using a center state, two side states, and a stack whose contents were manipulated by cancellation and replacement rules. The machine was not originally designed as an instance of a general theory of semantic fibers, local realizability, normalization debt, or causal schedules. Those notions did not yet exist in the project.

What made the design mathematically productive was a more elementary observation: after a prefix had been processed, the machine appeared to remember not only whether the prefix was currently balanced, but also **what was still needed to make it balanced**.

That observation changed the direction of the work.

The first question was therefore not whether $L_{2,1}$ was deterministic context-free. The more useful questions were:

- Which stack contents can actually occur?
- Why do the side states matter arithmetically?
- What exactly is the cancellation operation doing?
- When the input ends away from balance, does the machine encode a continuation that repairs the deficit?
- Is that continuation shortest?
- Why does looking only a small distance into the stack reveal information about the entire store?

- Can the same structure be built for a general ratio $p:q$?
- If several machines realize the same counting semantics, which aspects of their design are forced and which are representational choices?

The answers do not arrive all at once. They are the subject of the book.

The development will repeatedly return to the original machine. First it will be analyzed as a concrete DPDA. Its state and stack behavior will then be reinterpreted arithmetically. The arithmetic will become a geometry of shortest completions. That geometry will lead to more efficient pushdown representations and exact resource frontiers. A resource question will then escape the original representation class and lead to a stack-alphabet hierarchy theorem for general deterministic pushdown automata. The completion geometry will branch into higher-dimensional topology, while one local cancellation transition will generate a separate Fuss–Catalan defect geometry and, later, an arithmetic p -adic offshoot.

After those branches, the book will return to the machine-construction problem. A uniform real-time $p:q$ DPDA will explain the arithmetic hidden inside the original center/left/right architecture. WRITE, SEARCH, and CANCEL will then cease to be merely names for transition patterns and become components of a normalization calculus. Finally, the distinction between a unique semantic process and its nonunique physical pushdown realizations will lead to a general theory of locality, representation debt, causal schedules, synthesis, and verification.

The final chapters return once more to the 2023 design.

The aim is therefore not to replace the original machine by abstraction. It is to understand how much theory was already latent in it.

PART I

The Origin

Research provenance. Parts I–II begin from the author’s 2023 completion-aware machine and synthesize the formal analyses developed in the first two papers of the research sequence [1, 2]. Recognition by counter-like deterministic devices and related one-counter questions are classical background rather than novelty claims here [64, 51, 29].

Chapter 1

The 2023 Completion-Aware DPDA

1.1 The motivating language

For a binary word w , let $N_0(w)$ and $N_1(w)$ denote the numbers of zeros and ones in w . The motivating language is

$$L_{2,1} = \{w \in \{0, 1\}^* : N_1(w) = 2N_0(w)\}.$$

Equivalently,

$$2N_0(w) - N_1(w) = 0.$$

It is useful to define already, without yet turning it into a general theory, the imbalance

$$\rho_{2,1}(w) = N_1(w) - 2N_0(w).$$

Thus $\rho_{2,1}(w) = 0$ exactly when $w \in L_{2,1}$.

The original 2023 DPDA was built directly at the machine level. Its characteristic structure consists of three semantic core states:

$$C, \quad L, \quad R,$$

which we call the **center**, **left**, and **right** states.

The center is the balanced operational hub. The side states are not independent counters. They record a finite correction to the meaning of the physical stack. This interpretation will be proved rather than assumed.

Before doing so, we must distinguish two formal versions of the original design.

1.2 Two formal variants of the same core design

There are two versions of the 2023 machine that we will preserve throughout the monograph.

The first excludes the empty input and recognizes

$$L_{2,1} \setminus \{\varepsilon\}.$$

The second includes the empty input and recognizes the standard language

$$L_{2,1}.$$

Their center/left/right recognition logic is the same. Their difference lies in the initialization and acceptance wrapper.

This distinction is worth preserving for two reasons.

First, it records the actual development of the design rather than silently collapsing two formal choices into one picture. Second, it separates a semantic question—how the *C/L/R* core represents imbalance—from an acceptance-convention question—whether ε is accepted by the wrapper.




Figure 1.1: Formalized redrawing of the original 2:1 DPDA variant with the empty input excluded. Solid arrows are ordinary recognition transitions; dashed callout boxes contain only right-endmarker completion rules. The listed left-state completion rule with top 1 is retained for fidelity but has an unreachable source condition.

The logical core has states C, L, R . The wrapper contains the source state q_s and accepting state q_f .

The principal recognition transitions are the following.

From C :

$$1, 0/\varepsilon : C \rightarrow L,$$

$$0, 1/\varepsilon : C \rightarrow R,$$

and the center cancellation loops

$$1, 1/\varepsilon, \quad 0, 0/\varepsilon.$$

The left state has the characteristic cancellation loop

$$\boxed{0, 0/\varepsilon}$$

and returns toward the center through the replacement rules

$$0, Z_0/10Z_0,$$

$$1, 0/000,$$

$$1, Z_0/00Z_0.$$

The right state returns toward the center through

$$0, 1/111111,$$

$$0, Z_0/111111Z_0,$$

$$0, 0/111,$$

$$1, 0/\varepsilon,$$

$$1, Z_0/11Z_0,$$

$$1, 1/111.$$

The wrapper moves between C and q_f at empty logical stack and initializes nonempty inputs appropriately.

This version does not accept ε .




Figure 1.2: Formalized redrawing of the original 2:1 DPDA variant with the empty input included. The center/left/right recognition logic and end-of-input completion rules are identical to Figure 1.1; only the initialization/final-state wrapper changes so that ε is accepted.

The $C/L/R$ transition system is unchanged.

The difference is the initialization/final-state wrapper. In particular, an initial

$$\varepsilon, \varepsilon/Z_0$$

wrapper move allows the physical machine to represent the empty prefix at the accepting bottom-marker configuration.

Consequently this version recognizes

$$L_{2,1},$$

including ε .

The two figures will remain separate in the final monograph, but their common recognition core will be analyzed only once.

1.3 Recognition and completion are different operations

The original design also contains rules whose purpose is not to consume another binary input symbol, but to expose the continuation represented by a nonbalanced final configuration.

These are **completion rules**.

In the formal monograph model, they are enabled only when a distinguished right endmarker

$$\neg$$

is the next unread symbol.

For the left state, the completion list is

$$\neg, Z_0/100Z_0,$$

$$\neg, 1/0,$$

$$\neg, 0/1000.$$

For the right state, it is

$$\neg, Z_0/111Z_0,$$

$$\neg, 1/1111,$$

$\neg, 0/1$.

The rule with left-state top 1 is retained as part of the original completion list, but we will prove shortly that its source condition is unreachable.

The endmarker convention is essential. If these operations were interpreted as ordinary ε -moves, some state/top pairs would also have input-consuming moves, producing a conflict with standard DPDA determinism. With the right endmarker, recognition on 0 or 1 and completion on $\neg$ are disjoint cases.

Lemma 1.1 — Endmarker determinism.

If every post-input completion rule is interpreted as a $\neg$ -labeled transition to a completion sink with no outgoing transitions, then the completion-extended machine is deterministic in the standard DPDA sense.

Proof.

Fix a state/top pair.

During binary recognition, the next input symbol is 0 or 1. A completion rule is unavailable.

After binary input has been exhausted, the next input symbol is $\neg$. The binary recognition rules are unavailable, and at most one completion rule applies to the current state/top pair.

Thus every instantaneous configuration has at most one successor. The completion sink introduces no further conflict. $\square$

This formalization will be used consistently in the rest of the monograph.

1.4 Removing the wrapper from the semantic analysis

The physical states q_s and q_f serve initialization and acceptance. They should not be confused with the semantic $C/L/R$ core.

Write a physical stack as

$$SZ_0, \quad S \in \{0, 1\}^*,$$

where S is the **logical stack**.

Define the core projection

$$\pi(C, SZ_0) = (C, S),$$

$$\pi(L, SZ_0) = (L, S),$$

$$\pi(R, SZ_0) = (R, S),$$

together with

$$\pi(q_s, \varepsilon) = (C, \varepsilon),$$

$$\pi(q_f, Z_0) = (C, \varepsilon).$$

Thus the empty prefix and every balanced prefix accepted at q_f receive the same zero-residual semantic core meaning.

From now on, unless the wrapper itself matters, “the configuration after a prefix” means its core-projected configuration.

1.5 The reachable stack is much more restricted than it looks

The transition table contains replacements of several different lengths. One might therefore expect the machine to generate a complicated set of stack words.

It does not.

Let

$$\mathcal{S}_C = 0^* \cup 1^* \cup \{10\},$$

$$\mathcal{S}_L = 0^*,$$

$$\mathcal{S}_R = 1^* \cup \{0\}.$$

Theorem 1.2 — Reachable-store normal form.

For every binary prefix processed by either formal variant, if

$$(q, S)$$

is its core-projected configuration, then

$$S \in \mathcal{S}_q.$$

In particular:

- every reachable left-state stack is unary 0^k ;
- every reachable right-state stack is either unary 1^k or the exceptional word 0;
- every reachable center stack is unary or the exceptional mixed word 10.

Proof.

We proceed by induction on the number of consumed input symbols.
The projected empty configuration is

$$(C, \varepsilon),$$

and

$$\varepsilon \in 0^* \cap 1^*.$$

Assume the current projected configuration belongs to the stated family.

The center loops remove matching top symbols. Hence they preserve unary forms.

A center mismatch on top 0 moves to L and removes that top. The resulting left-state stack is therefore still in 0^* .

A center mismatch on top 1 moves to R and removes that top. The resulting right-state stack lies in 1^* , with the special bottom-related construction accounting for the exceptional reachable 0 case.

The left-state loop

$$0, 0/\varepsilon$$

preserves 0^* . Its return transitions create only the allowed center forms 0^* , 1^* , or 10 .

The right-state return transitions similarly produce only the allowed center families.

No recognition transition enters L with top 1. Therefore a left-state source condition with logical top 1 is unreachable.

All reachable configurations remain in the stated families. $\square$

The theorem is stronger than a convenient regular-language description. It tells us that the apparently global question

“Does a zero or a one occur somewhere in the stack?”

may have a bounded local answer.

Before proving that, we need to understand why the side states exist.

1.6 The state carries part of the arithmetic

Define the stack weight

$$D(S) = 2N_0(S) - N_1(S).$$

Define the state potential

$$\Phi(C) = 0,$$

$$\Phi(L) = 3,$$

$$\Phi(R) = -3.$$

Theorem 1.3 — Original-machine state-potential invariant.

After every consumed prefix w , let

$$(q, S)$$

be the core-projected configuration. Then

$$\boxed{\rho_{2,1}(w) = D(S) + \Phi(q).}$$

Proof.

For the empty prefix,

$$\rho_{2,1}(\varepsilon) = 0,$$

and the projected configuration is

$$(C, \varepsilon),$$

so

$$D(\varepsilon) + \Phi(C) = 0.$$

It remains to compare the change of the two sides under each recognition transition.

Reading 0 decreases $\rho_{2,1}$ by 2. Reading 1 increases it by 1.

Each stack replacement changes $D(S)$, and every state change changes Φ . The transition table is constructed so that

$$\Delta(D(S) + \Phi(q)) = \Delta\rho_{2,1}.$$

For the transition that will later generate the cancellation geometry,

$$(L, 0S) \xrightarrow{0} (L, S),$$

the stack weight decreases by 2, while the state potential remains 3. Thus the right-hand side changes by -2 , exactly matching the change induced by consuming the relevant symbol under the selected residual sign convention.

The remaining transition types satisfy the same identity by direct finite verification. A complete transition-by-transition audit table will be retained in the technical appendix of the complete research version. $\square$

Interpretation

The logical stack alone does **not** always equal the current arithmetic deficit.
 The side state stores a finite correction:

$$+3 \quad \text{in } L,$$

$$-3 \quad \text{in } R.$$

This explains a configuration such as

$$(L, \varepsilon).$$

Its empty logical stack does not mean zero residual. The missing value is carried by the state.

This is the first place where the machine reveals a distinction that will become much more important later:

semantic information = stack information + finite-control information.
--

At this point, however, we keep the observation concrete.

1.7 End-of-input completion transfers the state correction into the stack

Suppose the binary input has been exhausted.

If the projected configuration is already in C , no side-state correction is required.

If it is in L or R , the unique endmarker completion rule converts the side-state potential into a normalized stack representation.

The result has an exceptionally simple form.

Theorem 1.4 — Original-machine shortest-completion extraction.

After end-of-input normalization, the logical stack is

$$\widehat{S}(w) = \begin{cases} 0^m, & \rho_{2,1}(w) = 2m, \quad m \geq 0, \\ 10^{m+1}, & \rho_{2,1}(w) = 2m + 1, \quad m \geq 0, \\ 1^m, & \rho_{2,1}(w) = -m, \quad m \geq 1. \end{cases}$$

When interpreted as a continuation word, $\widehat{S}(w)$ is a shortest suffix that completes w to $L_{2,1}$.

Proof.

By Theorems 1.2 and 1.3, every side-state configuration belongs to a one-parameter stack family and differs from its stack weight by exactly the fixed state potential.

The endmarker completion rule removes that finite state correction and returns the representation to the zero-potential center convention.

Consequently, after normalization,

$$D(\widehat{S}(w)) = \rho_{2,1}(w).$$

The reachable normalized center family is

$$0^* \cup 1^* \cup \{10^+\},$$

so the displayed residual-to-stack map is injective.

Now let a completion contain x zeros and y ones. The completed word must satisfy

$$N_1(w) + y = 2(N_0(w) + x).$$

Since

$$\rho_{2,1}(w) = N_1(w) - 2N_0(w),$$

this is equivalent to

$$\boxed{2x - y = \rho_{2,1}(w)}.$$

The three displayed normalized forms give the unique minimum nonnegative Parikh pair for this equation. Hence the continuation selected by the normalized stack has minimum possible length. $\square$

The important discovery is independent of notation:

the machine's end configuration can be decoded as a shortest repair word.

This is the observation from which the completion theory begins.

1.8 A worked run: the prefix 111

Consider the empty-input-excluded presentation and ignore Z_0 in the logical-stack notation.

Starting in the center:

$$(C, 111 \dashv, \varepsilon) \vdash (C, 11 \dashv, 10) \vdash (C, 1 \dashv, 0) \vdash (L, \dashv, \varepsilon).$$

The consumed prefix has

$$\rho_{2,1}(111) = 3.$$

The physical logical stack is empty, but the machine is in L , whose potential carries the missing correction.

The endmarker completion rule

$$\neg, Z_0/100Z_0$$

materializes the continuation

$$100.$$

Indeed,

$$111100$$

contains four ones and two zeros, so

$$111100 \in L_{2,1}.$$

No shorter continuation can repair the count imbalance.

Thus a side state that initially looks like an implementation detail has a precise role: it stores a finite amount of the completion demand outside the stack until the end-of-input normalizer materializes it.

1.9 The first structural surprise: two stack cells are enough

For a word v , let

$$\text{alph}(v) = \{a \in \{0, 1\} : a \text{ occurs in } v\}.$$

Let

$$\text{pref}_d(v)$$

denote the prefix of v of length $\min(d, |v|)$.

Theorem 1.5 — Tight depth-two visibility.

For every completion-normalized stack S ,

$$\boxed{\text{alph}(S) = \text{alph}(\text{pref}_2(S)).}$$

Depth two is necessary.

Proof.

By Theorem 1.4, every normalized stack belongs to one of the families

$$0^*, \quad 1^*, \quad 10^+.$$

Every symbol that occurs anywhere in one of those words therefore occurs within the first two positions.

Depth one does not suffice. The normalized words

$$1$$

and

$$10$$

have the same top symbol but different global symbol supports.

Hence the exact visibility depth is two. $\square$

This theorem formalizes an intuition present in the original design: a question that appears to require searching the entire stack can, on the reachable normalized family, be answered from only the first two cells.

The reason is not that arbitrary pushdown stacks are locally transparent. It is that the machine's reachable stores occupy a highly restricted language.

That distinction will matter repeatedly later.

1.10 The cancellation transition that will reappear much later

One transition deserves to be marked now even though its deeper significance will not be developed until Part V:

$$(L, 0S) \xrightarrow{0} (L, S).$$

At the machine level it is simply a cancellation operation.

At the arithmetic level it participates in the state-potential invariant.

Much later, when prefixes are represented as nonnegative 2-Dyck paths, the same transition will be characterized by an exact family of landing heights. That observation will generate correction signatures, an intrinsic defect, a Catalan boundary, and a binary cancellation skeleton.

We do not need any of that theory yet.

For now we record only the transition and its exact place in the original machine.

1.11 What has been learned without leaving the original machine

We can now answer several questions that were not visible from the transition diagram alone.

First, the reachable stack language is extremely small.

Second, the side states carry an exact finite arithmetic potential.

Third, the post-input completion rules can be formalized deterministically with a right endmarker.

Fourth, the final state/stack configuration can be normalized into a shortest continuation.

Fifth, global symbol support of the normalized store is exactly visible at depth two.

These facts suggest that the original DPDA is not merely storing an arbitrary history of the input. Its configuration is representing a much smaller mathematical object.

The next chapter identifies that object.

The progression will be:

physical state and stack $\longrightarrow$ residual
 $\longrightarrow$ minimum completion demand $\longrightarrow$ directed completion geometry.

Only after that geometry is understood will we return to the engineering question:

Can the representation itself be improved?

That question will eventually produce the improved and scoped-optimal DPDA constructions, the exact resource frontiers, and the later general $p:q$ architectures.

Chapter 2

What the Machine Is Really Storing

Chapter 1 stayed deliberately close to the 2023 transition diagram. We saw that the reachable stores are highly restricted, that the side states carry a finite arithmetic potential, and that an end-of-input normalization converts a nonbalanced configuration into a shortest continuation.

Those observations suggest a simpler object behind the physical machine.

The purpose of this chapter is to identify it.

The central distinction will be:

$$\begin{array}{l} \text{physical configuration} \longrightarrow \text{arithmetic residual} \\ \longrightarrow \text{minimum completion demand.} \end{array}$$

The first object depends on a particular DPDA representation. The second and third do not.

That separation is the first decisive abstraction in the monograph.

2.1 From the 2:1 imbalance to the general $p:q$ residual

The original language satisfies

$$N_1(w) = 2N_0(w).$$

The natural general family is

$$L_{p,q} = \{w \in \{0,1\}^* : qN_1(w) = pN_0(w)\},$$

where

$$p, q \geq 1, \quad \gcd(p, q) = 1.$$

The coprimality assumption removes a redundant common scale. Nothing about the language changes if p and q are divided by their greatest common divisor.

Define the residual

$$\rho_{p,q}(w) = qN_1(w) - pN_0(w).$$

Thus

$$w \in L_{p,q} \iff \rho_{p,q}(w) = 0.$$

Reading a symbol changes the residual by a fixed amount:

$$\rho_{p,q}(w0) = \rho_{p,q}(w) - p,$$

$$\rho_{p,q}(w1) = \rho_{p,q}(w) + q.$$

For the motivating ratio $(p, q) = (2, 1)$,

$$\rho_{2,1}(w) = N_1(w) - 2N_0(w),$$

exactly the convention used in Chapter 1.

The residual is already a substantial simplification. Instead of retaining the whole Parikh pair

$$(N_0(w), N_1(w)),$$

we retain one integer.

But the stronger observation is that this one integer determines exactly what is needed on the right to reach the language.

2.2 The completion equation

Let w be a prefix with residual

$$r = \rho_{p,q}(w).$$

Suppose a continuation z contains

$$x = N_0(z)$$

zeros and

$$y = N_1(z)$$

ones.

Then

$$wz \in L_{p,q}$$

exactly when

$$q(N_1(w) + y) = p(N_0(w) + x).$$

Rearranging gives the basic completion equation.

Lemma 2.1 — Completion equation.

For every prefix w and continuation z with Parikh pair (x, y) ,

$$z \in w^{-1}L_{p,q} \iff px - qy = \rho_{p,q}(w).$$

Proof.

Expand the defining equality

$$qN_1(wz) = pN_0(wz).$$

Since

$$N_0(wz) = N_0(w) + x, \quad N_1(wz) = N_1(w) + y,$$

we obtain

$$qN_1(w) + qy = pN_0(w) + px.$$

Therefore

$$px - qy = qN_1(w) - pN_0(w) = \rho_{p,q}(w).$$

The steps are reversible. $\square$

This equation is the semantic heart of the original completion mechanism.

For $p = 2, q = 1$, it becomes

$$2x - y = r.$$

Thus the Chapter 1 example

$$w = 111, \quad r = 3,$$

requires

$$2x - y = 3.$$

The minimum nonnegative solution is

$$(x, y) = (2, 1).$$

Hence every shortest completion has two zeros and one one.

There are three such words:

$$100, \quad 010, \quad 001.$$

The original DPDA selects 100.

This distinction will be important throughout the book:

minimum Parikh demand is unique; a shortest word need not be unique.

2.3 The primitive balanced cycle

The completion equation

$$px - qy = r$$

has an elementary but fundamental symmetry.

If (x, y) is a solution, then

$$(x + q, y + p)$$

is another solution, because

$$p(x + q) - q(y + p) = px - qy.$$

The vector

$$K = (q, p)$$

is therefore the primitive nonnegative kernel direction.

Its meaning is transparent:

- adding q zeros contributes pq ;
- adding p ones contributes $-pq$;
- the two contributions cancel.

At the word level, any word containing exactly q zeros and p ones has net residual change zero.

Its length is

$$p + q.$$

Later this will become the fundamental directed cycle length of the residual graph.

At this stage we need only its Diophantine role.

2.4 Every completion lies on one affine ray

Because $\gcd(p, q) = 1$, the integer solutions of

$$px - qy = r$$

form a one-dimensional affine lattice.

Proposition 2.2 — Affine completion ray.

For every

$$r \in \mathbb{Z},$$

there is a unique nonnegative solution

$$\mu(r) = (x_r, y_r)$$

minimizing

$$x + y$$

subject to

$$px - qy = r.$$

Every nonnegative solution is

$$(x_r + qk, y_r + pk), \quad k \in \mathbb{N}_0.$$

Consequently the possible completion lengths are exactly

$$\Delta(r) + (p + q)k, \quad k \in \mathbb{N}_0,$$

where

$$\Delta(r) = x_r + y_r.$$

Proof.

Since p and q are coprime, Bézout's identity guarantees an integer solution

$$(x_0, y_0)$$

to

$$px - qy = r.$$

Every integer solution has the form

$$(x_0 + qk, y_0 + pk), \quad k \in \mathbb{Z}.$$

For sufficiently large k , both coordinates are nonnegative.

Let k_* be the least integer for which this happens, and define

$$(x_r, y_r) = (x_0 + qk_*, y_0 + pk_*).$$

Every other nonnegative solution is obtained by replacing k_* with

$$k_* + j, \quad j \in \mathbb{N}_0.$$

Therefore every other nonnegative solution adds

$$j(q, p).$$

Its length increases by

$$j(p + q).$$

Hence (x_r, y_r) is the unique minimum-length nonnegative Parikh pair. $\square$

This proposition gives the first exact picture of the completion space:

$$\boxed{\mu(r), \quad \mu(r) + K, \quad \mu(r) + 2K, \quad \mu(r) + 3K, \dots}$$

The minimum demand is the first lattice point of an affine ray inside the nonnegative quadrant.

The later geometry of the monograph begins here.

2.5 Minimum completion demand

We now give the central object a name.

Definition 2.3 — Minimum completion demand.

For residual

$$r \in \mathbb{Z},$$

let

$$\boxed{\mu_{p,q}(r) = (x_r, y_r)}$$

be the unique nonnegative solution of

$$px - qy = r$$

minimizing

$$x + y.$$

Define the minimum completion length

$$\Delta_{p,q}(r) = x_r + y_r.$$

When p, q are fixed, we write simply

$$\mu(r), \quad \Delta(r).$$

The word **demand** is intentional.

The pair

$$(x_r, y_r)$$

does not describe the prefix already read. It describes what the prefix still requires:

- x_r additional zeros;
- y_r additional ones.

The physical DPDA need not store those two integers literally.

The semantic claim is only that a configuration representing residual r implicitly represents the unique pair $\mu(r)$.

For $L_{2,1}$,

$$2x - y = r.$$

The minimum demand has the closed form

$$\mu(2m) = (m, 0), \quad m \geq 0,$$

$$\mu(2m + 1) = (m + 1, 1), \quad m \geq 0,$$

and

$$\mu(-m) = (0, m), \quad m \geq 1.$$

These are exactly the three families already materialized by the original 2023 DPDA:

$$0^m, \quad 10^{m+1}, \quad 1^m.$$

So the stack forms discovered operationally in Chapter 1 are not arbitrary normal forms.

They are physical encodings of $\mu(r)$.

2.6 A canonical shortest word

The pair

$$\mu(r) = (x_r, y_r)$$

is unique, but its letters can usually be ordered in several ways.

For semantic decoding it is convenient to choose one fixed representative.

The minimum pair has a simple modular description.

Theorem 2.4 — Canonical two-block shortest completion.

Let $r > 0$.

There is a unique

$$y_r \in \{0, \dots, p-1\}$$

such that

$$qy_r \equiv -r \pmod{p}.$$

Set

$$x_r = \frac{r + qy_r}{p}.$$

Then

$$\boxed{C(r) = 1^{y_r} 0^{x_r}}$$

is a shortest completion word.

For $r < 0$, let

$$x_r \in \{0, \dots, q-1\}$$

be the unique solution of

$$px_r \equiv r \pmod{q}.$$

Set

$$y_r = \frac{px_r - r}{q}.$$

Then

$$\boxed{C(r) = 0^{x_r} 1^{y_r}}$$

is a shortest completion.

Finally,

$$C(0) = \varepsilon.$$

Proof.

Suppose first that $r > 0$.

The congruence

$$qy \equiv -r \pmod{p}$$

has a unique solution modulo p because q is invertible modulo p .

Selecting the representative

$$0 \leq y_r < p$$

makes

$$x_r = \frac{r + qy_r}{p}$$

a nonnegative integer.

Any other integer solution differs by the kernel vector

$$(q, p).$$

Subtracting one copy of K would make the bounded coordinate y_r negative. Therefore this is the first nonnegative point on the affine completion ray and hence the unique minimum Parikh pair.

The argument for $r < 0$ is symmetric, now bounding the zero coordinate modulo q .

The displayed block word is simply a fixed ordering of the unique minimum pair. $\square$

The canonical word is a **decoder convention**, not a claim that other shortest orderings are invalid.

For example, for

$$p = 3, \quad q = 2, \quad r = 4,$$

we solve

$$2y \equiv -4 \pmod{3}.$$

Thus

$$y = 1, \quad x = \frac{4 + 2}{3} = 2.$$

Hence

$$\mu(4) = (2, 1)$$

and

$$C(4) = 100.$$

The other shortest words with the same Parikh pair are equally valid as completions.

The canonical map simply chooses one of them reproducibly.

2.7 The original DPDA implements this decoder at 2:1

We can now restate the Chapter 1 completion theorem in semantic language.
Let

$$c = (q, S)$$

be any reachable core configuration of the original 2:1 DPDA.
Its residual is decoded by

$$\boxed{\Theta_{\text{orig}}(c) = D(S) + \Phi(q).}$$

Thus

$$\Theta_{\text{orig}}(c) = \rho_{2,1}(w)$$

for the prefix w that reached c .

The endmarker normalizer removes the side-state correction and produces

$$\widehat{S}(c).$$

Chapter 1 proved that

$$\widehat{S}(c) = \begin{cases} 0^m, & r = 2m, \\ 10^{m+1}, & r = 2m + 1, \\ 1^m, & r = -m, \end{cases}$$

where

$$r = \Theta_{\text{orig}}(c).$$

But these are exactly the canonical two-block words

$$C(r)$$

for

$$p = 2, \quad q = 1.$$

Therefore the original machine realizes the diagram

$$\boxed{c \xrightarrow{\Theta_{\text{orig}}} r \xrightarrow{\mu} (x_r, y_r) \xrightarrow{C} C(r)}.$$

The machine is not being claimed to be an output transducer.

Rather, its configuration is a mathematical certificate from which the shortest canonical continuation can be decoded.

This distinction between **stored information** and **explicit output** will remain important later.

2.8 Physical information versus semantic information

We can now answer the question posed at the end of Chapter 1.

What is the machine really storing?

Not the entire prefix.

Not even the complete Parikh pair of the prefix.

At the semantic level, it stores only the residual

$$r.$$

From r , the minimum completion demand

$$\mu(r)$$

is uniquely determined.

From $\mu(r)$, the canonical shortest continuation

$$C(r)$$

is determined by convention.

This gives three layers.

Layer 1 — Physical representation

A concrete configuration such as

$$(L, 0^k),$$

$$(R, 1^k),$$

or

$$(C, 10)$$

belongs to the chosen 2023 machine.

Layer 2 — Semantic residual

The decoder

$$\Theta_{\text{orig}}$$

maps that configuration to one integer

$$r.$$

Layer 3 — Completion semantics

The residual determines

$$\mu(r), \quad \Delta(r), \quad C(r).$$

The crucial fact is that Layer 1 is not unique in principle.

Later papers will construct different physical DPDAs representing the same Layer 2 and Layer 3 objects.

That is the beginning of the distinction between semantics and representation.

We will not formalize “exact lifts,” “sections,” or “semantic fibers” yet. Those concepts belong much later in the book.

For now it is enough to observe that a complicated state/stack configuration may be merely a coordinate system for a much simpler arithmetic state.

2.9 The residual transition system

Because the residual changes by fixed amounts,

$$r \xrightarrow{0} r - p,$$

$$r \xrightarrow{1} r + q.$$

Thus every binary prefix traces a path on the integer line.

For the motivating language,

$$r \xrightarrow{0} r - 2,$$

$$r \xrightarrow{1} r + 1.$$

This graph is independent of the original DPDA's stack representation.

The 2023 machine is one physical realization of motion over this arithmetic system.

That observation immediately suggests a new question:

If the semantic state is only one integer, why does the physical machine need the particular stack words and side states that it uses?

There are two very different possible answers.

One possibility is that the physical complexity is largely accidental and can be reduced by choosing better coordinates.

The other is that LIFO locality forces some of it.

Both possibilities will eventually occur.

Before studying resources, however, we need to understand the shortest paths in the residual graph.

2.10 Completion length as distance

The quantity

$$\Delta(r) = x_r + y_r$$

has an immediate graph interpretation.

It is the length of a shortest directed path from residual r to residual 0.

Indeed:

- appending a zero follows the edge

$$r \rightarrow r - p;$$

- appending a one follows

$$r \rightarrow r + q;$$

- a continuation reaches the language exactly when the resulting residual is zero.

Therefore

$$\boxed{\Delta(r) = \text{dist}^{\rightarrow}(r, 0).}$$

We will prove and exploit this metric statement systematically in later chapters.

Here it already changes the interpretation of the minimum demand.
The pair

$$\mu(r) = (x_r, y_r)$$

does not merely count missing symbols.
It counts the two types of edges used by a shortest path to zero.
The canonical completion

$$C(r)$$

selects one distinguished shortest path among all orderings of those edges.
Thus the completion problem has become a geodesic problem.

2.11 How many shortest completions are there?

Because every shortest completion has the same Parikh pair

$$(x_r, y_r),$$

the number of shortest completion words is

$$\boxed{\binom{x_r + y_r}{x_r}}.$$

This simple count is the first indication that shortest completion has an internal geometry.

If both coordinates are positive, there are several ways to interleave the required zeros and ones.

For the prefix 111,

$$\mu(3) = (2, 1),$$

so there are

$$\binom{3}{1} = 3$$

shortest completions.

The original machine chooses the boundary ordering

100.

Later we will see that different shortest orderings are connected by elementary commuting moves, and the left branch of the original machine realizes a complete ladder of such shortest schedules at configuration level.

The multiplicity of shortest words is therefore not noise.

It is the first object that turns completion arithmetic into geometry.

2.12 Why the minimum pair is a better semantic object than the canonical word

It would be possible to define the semantics directly by the canonical word

$$C(r).$$

For the later theory, however, the pair

$$\mu(r) = (x_r, y_r)$$

is more fundamental.

There are several reasons.

First, the completion equation is linear in (x, y) .

Second, the kernel direction

$$K = (q, p)$$

acts naturally on Parikh pairs.

Third, shortest-path descent becomes coordinate deletion:

- consuming a geodesic zero removes one unit of the zero demand;
- consuming a geodesic one removes one unit of the one demand.

Fourth, later physical representations will use different word orderings while preserving the same minimum pair.

Finally, in higher-dimensional counting languages the minimum demand will cease to be a single pair and become a fiber of optimal Parikh vectors. That generalization is visible only if we separate the Parikh object from its chosen word ordering now.

Thus our semantic hierarchy is

$$\boxed{r \longrightarrow \mu(r) \longrightarrow C(r),}$$

with the residual and minimum demand carrying the mathematics, and the canonical word supplying a convenient physical or decoding coordinate.

2.13 A first exact notion of redundancy

Suppose two different prefixes u and v have the same residual:

$$\rho_{p,q}(u) = \rho_{p,q}(v) = r.$$

Then both prefixes have exactly the same completion equation

$$px - qy = r.$$

Consequently they have:

- the same minimum demand;
- the same minimum completion length;
- the same set of feasible completion Parikh pairs;
- the same canonical shortest continuation.

Already, therefore, their past histories have become irrelevant to the completion problem.

This foreshadows the quotient theorem proved later:

$$u^{-1}L_{p,q} = v^{-1}L_{p,q} \iff \rho_{p,q}(u) = \rho_{p,q}(v).$$

We postpone the full quotient interpretation to Chapter 4, where it will be connected to directed geometry.

For the present chapter, the important message is simpler:

the residual is a sufficient statistic for all right-completion arithmetic.

2.14 The original machine viewed through the new semantics

We can now revisit several Chapter 1 observations without introducing any new transition rule.

The side states

The state potential means that a side state stores part of the residual outside the logical stack.

For example,

$$(L, \varepsilon)$$

represents

$$r = 3.$$

This is why the stack may be physically empty even though the completion demand is nonzero.

Endmarker materialization

The completion rule does not create new semantic information.

It changes coordinates.

Before the endmarker, the residual may be split between:

stack weight + state potential.

After normalization, the state potential is zero and the entire minimum demand is visible in the canonical stack word.

Tight depth-two visibility

The normalized words

$$0^*, \quad 1^*, \quad 10^+$$

are exactly the 2:1 canonical completion words.

Their global support is visible in the first two positions because the canonical decoder has only these three structural forms.


Thus the depth-two theorem is not an isolated stack-language coincidence.

It is a property of the chosen canonical ordering of the minimum demand.

This distinction will later become crucial when we compare alternative physical representations.

2.15 The first major conceptual compression

The entire original configuration can now be summarized by the following commuting picture:



The physical DPDA is rich.

The semantic state is one-dimensional.

The completion demand is exact.

This is the first compression of the research program.

It also creates the next problem.

If

$$\Delta(r)$$

is the shortest distance to balance, what does a single input symbol do to that distance?

When does reading a zero move one step closer to acceptance?

When does reading a one do so?

Can both actions be shortest simultaneously?

And do the original DPDA's cancellation and completion moves physically realize those shortest paths?

Those questions lead directly to the **minimum-demand geodesic descent** and the **completion ladder**.

That is the subject of Chapter 3.

PART II

Completion Becomes Geometry

Research provenance. The configuration ladder, minimum-demand geometry, and rectangular shortest-completion theory in this Part are developed from the exact structural analysis in [2], with the quotient/completion viewpoint cross-checked against [1]. The use of shortest-path structure in pushdown graphs has prior literature; the claim here is the exact arithmetic specialization and configuration-level geometry of this family [23, 25].

Chapter 3

Minimum Demand and the Completion Ladder

Chapter 2 replaced the original machine's physical configuration by a simpler semantic object:

$$r \mapsto \mu(r) = (x_r, y_r).$$

For the motivating language,

$$L_{2,1} = \{w : N_1(w) = 2N_0(w)\},$$

the residual evolves by

$$r \xrightarrow{0} r - 2, \quad r \xrightarrow{1} r + 1.$$

The minimum completion demand tells us how many zeros and ones are still required. We now ask what happens to that demand when another input symbol is read.

The answer reveals that the original DPDA is traversing shortest paths.

3.1 Geodesic input symbols

Fix

$$p = 2, \quad q = 1.$$

For residual r , write

$$\mu(r) = (x_r, y_r),$$

and

$$\Delta(r) = x_r + y_r.$$

Recall:

$$\mu(2m) = (m, 0),$$

$$\mu(2m + 1) = (m + 1, 1),$$

and

$$\mu(-m) = (0, m).$$

Define the residual updates

$$T_0(r) = r - 2,$$

$$T_1(r) = r + 1.$$

A symbol $a \in \{0, 1\}$ will be called **geodesic at residual r** if reading it decreases the remaining shortest-completion distance by exactly one:

$$\boxed{\Delta(T_a(r)) = \Delta(r) - 1.}$$

This is the semantic version of the question:

Does this next input symbol perform one unit of the completion work that was already required?

Lemma 3.1 — Minimum-demand geodesic descent.

Let

$$\mu(r) = (x_r, y_r).$$

If

$$x_r > 0,$$

then

$$\boxed{\mu(r - 2) = (x_r - 1, y_r)}$$

and

$$\Delta(r - 2) = \Delta(r) - 1.$$

If

$$y_r > 0,$$

then

$$\boxed{\mu(r+1) = (x_r, y_r - 1)}$$

and

$$\Delta(r+1) = \Delta(r) - 1.$$

Therefore the set of geodesic first symbols is

$$\boxed{\text{GeoFirst}(r) = \{0 : x_r > 0\} \cup \{1 : y_r > 0\}.$$

Proof.

Suppose

$$x_r > 0.$$

Since

$$2x_r - y_r = r,$$

we have

$$2(x_r - 1) - y_r = r - 2.$$

Thus

$$(x_r - 1, y_r)$$

is a completion demand for residual $r - 2$.

If $r - 2$ had a completion pair of length strictly less than

$$x_r + y_r - 1,$$

adding one zero to that pair would produce a completion of r shorter than

$$\mu(r),$$

contradicting minimality.

Hence

$$\mu(r-2) = (x_r - 1, y_r).$$

The argument for $y_r > 0$ is symmetric:

$$2x_r - (y_r - 1) = r + 1.$$

Again, any shorter completion after reading 1 would yield a shorter completion before reading it.

Thus precisely those symbols whose coordinate in $\mu(r)$ is positive decrease the minimum distance by one. $\square$

This lemma has a useful interpretation.

If the minimum demand is

$$(x, y),$$

then every geodesic input symbol simply deletes one required unit:

$$0 : (x, y) \mapsto (x - 1, y),$$

$$1 : (x, y) \mapsto (x, y - 1).$$

At the semantic level, shortest completion is therefore almost trivial.

The interesting question is whether the physical DPDA realizes the same simple coordinate deletion.

3.2 The left branch as an exact shortest-path process

Consider the reachable left configurations

$$c_k = (L, 0^k), \quad k \geq 0.$$

From the state-potential invariant,

$$\rho(c_k) = 2k + 3.$$

Therefore

$$\boxed{\mu(c_k) = (k + 2, 1)}$$

and

$$\boxed{\Delta(c_k) = k + 3.}$$

So every such configuration requires:

- $k + 2$ zeros;
- one one.

The left cancellation transition

$$(L, 0^k) \xrightarrow{0} (L, 0^{k-1})$$

does exactly what Lemma 3.1 predicts:

$$(k + 2, 1) \mapsto (k + 1, 1).$$

It decreases the shortest-completion distance by exactly one.

This gives the original cancellation loop a new meaning.

It is not merely deleting matching physical stack symbols.

It is executing one edge of a shortest path in completion space.

The boundary sequence

$$(L, \varepsilon) \xrightarrow{0} (C, 10) \xrightarrow{0} (R, 0) \xrightarrow{1} (C, \varepsilon)$$

continues the same exact descent.

Thus the entire route from a left configuration to balance can be read geodesically.

3.3 The ladder sources are actually reachable

The configurations

$$(L, 0^k)$$

used in the next theorem are not hypothetical legal stack patterns.

They are reachable.

The prefix

$$111$$

reaches

$$(L, \varepsilon).$$

Two additional ones increase the left unary tail by one unit through the machine's center-side transition pattern.

Inductively,

$$\boxed{1^{2k+3} \text{ reaches } (L, 0^k).}$$

This reachability matters.

The geometry we are about to describe is therefore a geometry of actual machine computations, not merely of syntactically possible configurations.

3.4 Every position of the required one is shortest

From

$$(L, 0^k),$$

the minimum demand is

$$(k + 2, 1).$$

Hence every shortest continuation contains exactly one 1 and $k + 2$ zeros. There are

$$k + 3$$

possible positions for that unique one:

$$1 0^{k+2},$$

$$0 1 0^{k+1},$$

$$0^2 1 0^k,$$

$$\dots,$$

$$0^{k+2} 1.$$

The remarkable fact is that the original DPDA realizes **every one of these shortest schedules**, and it realizes them in a rectangular ladder of configurations.

Theorem 3.2 — Exact geodesic ladder.

For every

$$k \geq 0,$$

define the lower row

$$\Psi_k(i, 0) = \begin{cases} (L, 0^{k-i}), & 0 \leq i \leq k, \\ (C, 10), & i = k + 1, \\ (R, 0), & i = k + 2, \end{cases}$$

and the upper row

$$\Psi_k(i, 1) = (C, 0^{k+2-i}), \quad 0 \leq i \leq k+2.$$

Then:

1. every horizontal ladder edge is a solid input-0 transition;
2. every vertical ladder edge is a solid input-1 transition;
3. every monotone path from

$$(L, 0^k)$$

to

$$(C, \varepsilon)$$

consumes the minimum demand

$$(k+2, 1);$$

4. therefore every such path is shortest;
5. the shortest continuation words are exactly

$$\boxed{0^i 1 0^{k+2-i}, \quad 0 \leq i \leq k+2.}$$

There are exactly

$$\boxed{k+3}$$

shortest continuations.

Proof.

Along the lower row, while a zero remains on the left stack, the horizontal edge is the cancellation loop

$$(L, 0^m) \xrightarrow{0} (L, 0^{m-1}).$$

At the lower boundary,

$$(L, \varepsilon) \xrightarrow{0} (C, 10),$$

followed by

$$(C, 10) \xrightarrow{0} (R, 0).$$

Along the upper row, each horizontal edge is a center cancellation

$$(C, 0^m) \xrightarrow{0} (C, 0^{m-1}).$$

The vertical edges consume the unique required 1.

For an interior left configuration,

$$(L, 0^m) \xrightarrow{1} (C, 0^{m+2}).$$

At the lower boundary,

$$(L, \varepsilon) \xrightarrow{1} (C, 00).$$

At

$$(C, 10),$$

the input 1 removes the visible 1 and reaches

$$(C, 0).$$

At

$$(R, 0),$$

the input 1 reaches

$$(C, \varepsilon).$$

Thus the entire two-row grid is physically realized by the DPDA.

Every source-to-target monotone path contains exactly:

$$k + 2$$

horizontal zero edges and one vertical one edge.

Hence every path consumes the unique minimum Parikh demand

$$(k + 2, 1).$$

By Chapter 2, no shorter continuation exists.

The position of the vertical edge determines the position of the unique 1 in the continuation, giving exactly

$$0^i 1 0^{k+2-i}, \quad 0 \leq i \leq k + 2.$$

□

The source is

$$(L, 00),$$

with residual

$$7$$

and minimum demand




Figure 3.1: The exact configuration ladder for $k = 2$. The source $(L, 00)$ has residual 7 and minimum demand $(4, 1)$. Its five shortest continuations are 10000, 01000, 00100, 00010, 00001. The upper boundary is canonical completion; the lower boundary postpones the unique required 1 maximally.

$(4, 1)$.

Its five shortest completions are

10000,

01000,

00100,

00010,

00001.

The upper route materializes the required 1 as early as possible.

The lower route postpones it as long as possible.

Both are shortest.

Everything between them is shortest as well.

3.5 Local commuting diamonds

A ladder square gives two ways to consume one zero and one one:

01

or

10.

The machine reaches the same configuration either way.

Lemma 3.3 — Local commuting diamonds.

At every square of the geodesic ladder,

$$\delta^*(c, 01) = \delta^*(c, 10).$$

Proof.

For an interior left configuration

$$(L, 0^m), \quad m \geq 1,$$

one route is

$$(L, 0^m) \xrightarrow{0} (L, 0^{m-1}) \xrightarrow{1} (C, 0^{m+1}).$$

The other is

$$(L, 0^m) \xrightarrow{1} (C, 0^{m+2}) \xrightarrow{0} (C, 0^{m+1}).$$

Both end at the same configuration.

The two boundary squares are checked directly from the corresponding transition rules:

- at (L, ε) , both 01 and 10 reach the same center configuration;
- at $(C, 10)$, both orders reach (C, ε) .

These are exactly the ladder squares. $\square$

The point is not that 0 and 1 commute as input symbols globally.

They do not.

The point is much more specific:

Within this shortest-completion interval, adjacent geodesic actions commute physically.

3.6 Parikh synchronization

Because adjacent 01 and 10 factors can be exchanged locally, any two shortest prefixes that have consumed the same numbers of zeros and ones reach the same ladder position.

Corollary 3.4 — Geodesic Parikh synchronization.

Let u and v be prefixes of shortest continuations from

$$(L, 0^k).$$

If

$$N_0(u) = N_0(v)$$

and

$$N_1(u) = N_1(v),$$

then

$$\boxed{\delta^*((L, 0^k), u) = \delta^*((L, 0^k), v).}$$

In particular, every shortest continuation reaches

$$(C, \varepsilon).$$

Proof.

Inside the ladder, repeatedly exchange adjacent factors

$$01 \leftrightarrow 10$$

using Lemma 3.3.

The consumed Parikh pair uniquely determines the grid coordinate. $\square$

This is the first place where the physical machine's computation depends only on the consumed Parikh pair **inside a geodesic interval**, even though the machine is still a stack automaton processing an ordered word.

That phenomenon will later generalize.

3.7 Completion and cancellation are opposite boundaries of one object

The original machine suggested two seemingly different strategies.

One strategy is to expose or materialize the canonical completion as early as possible.

The other is to continue cancelling zeros on the left before finally consuming the required one.

The ladder shows that these are not different semantic processes.

They are boundary schedules of the same shortest-path rectangle.

Corollary 3.5 — Completion–cancellation boundary duality.

For the source

$$(L, 0^k),$$

the canonical completion

$$10^{k+2}$$

and the maximal-cancellation continuation

$$0^{k+2}1$$

are the two extremal boundary paths of the same geodesic ladder.

The word **duality** here is operational.

One boundary performs the unique required 1 first.

The other performs it last.

Both consume the same minimum demand.

This is a major reinterpretation of the original design:

completion and cancellation are scheduling choices
inside one geodesic object.

3.8 A commutation potential

For residual r , define

$$\Xi(r) = x_r y_r$$

when

$$\mu(r) = (x_r, y_r).$$

For the left-ladder source,

$$\mu(2k+3) = (k+2, 1),$$

so

$$\Xi(2k+3) = k+2.$$

This number has several simultaneous meanings.

It is:

1. the number of elementary squares in the two-row ladder;
2. the maximum number of geodesic zeros that can be consumed before the unique one becomes unavoidable;
3. the number of adjacent $01 \leftrightarrow 10$ swaps needed to move the unique one from the first position to the last.

Corollary 3.6 — Cancellation–diamond identity.

For every reachable left configuration,

$$\Xi = \text{diamond count} = \text{cancellation horizon} \\ = \text{extremal adjacent-swap distance.}$$

This identity is small, but it is conceptually useful.

A quantity defined purely from minimum-demand arithmetic,

$$x_r y_r,$$

simultaneously measures:

- a piece of configuration-space geometry;
- a cancellation horizon in the DPDA;
- a word-reordering distance.

The different descriptions are beginning to converge.

3.9 Why this is more than a counting argument

At the arithmetic level, it is easy to observe that a demand

$$(k + 2, 1)$$

has

$$k + 3$$

possible shortest word orderings.

The DPDA theorem says more.

It says that the machine realizes these orderings in a coherent configuration-level structure:

$$\text{shortest words} \longleftrightarrow \text{monotone paths} \longleftrightarrow \text{reachable configurations.}$$

Adjacent word commutations become literal commuting squares of machine transitions.

This is why the completion ladder matters.

It is the first exact bridge between:

- Diophantine completion arithmetic;
- shortest-path geometry;
- pushdown-store evolution.

3.10 The canonical completion is a boundary choice

The original endmarker normalizer selects

$$10^{k+2}$$

from the $k + 3$ shortest completions of a left source.

Nothing about the minimum Parikh demand itself forces the 1 to occur first.

The ordering

$$10^{k+2}$$

is a canonical representation choice.

The geodesic ladder makes this distinction visible:

semantic minimum demand $\neq$ unique physical schedule.

This sentence will acquire much deeper meanings later.

In P9, the original and centered machines will differ by eager versus lazy normalization chronology.

In P10, semantic carry chronology will be separated from physical materialization chronology.

In P11, several reduced physical lifts with the same semantics will coexist.

At this early stage, however, the ladder already contains the seed of that distinction.

3.11 What the ladder does not yet tell us

The ladder gives a complete operational picture for the motivating left branch.

It does not yet answer three broader questions.

First, when we look at an arbitrary residual, how much of its shortest-path structure can be inferred from a bounded part of the physical stack?

Second, do equal residuals really represent identical future language behavior, not merely identical completion lengths?

Third, what replaces this two-row ladder for a general ratio

$$p:q,$$

where a shortest demand may contain several copies of both symbols?

These questions lead in two directions:

quotient and visibility

and

general $p:q$ rectangles.

The next chapter takes the first direction.

Chapter 4

Quotients, Geodesics, and What Can Be Seen Locally

Chapter 2 showed that the residual

$$\rho_{p,q}(w) = qN_1(w) - pN_0(w)$$

determines the entire arithmetic of right completion. Chapter 3 then showed, in the original 2:1 machine, that shortest completions are realized by an actual configuration-level ladder of geodesic schedules.

We now strengthen both observations.

The residual does not merely determine the length or Parikh vector of a shortest repair. It determines the **entire future right language** of the prefix.

At the same time, the canonical shortest-completion word has a bounded visible frontier: one need not inspect its unbounded tail to know which types of geodesic action are available.

These two facts connect three levels:

prefix quotient $\longleftrightarrow$ residual vertex $\longleftrightarrow$ shortest-completion geometry.

4.1 The residual classifies all future behavior

For a language L and prefix u , the left quotient

$$u^{-1}L = \{z : uz \in L\}$$

is the set of all continuations that lead from u into L .

For $L_{p,q}$, Chapter 2 already gave the completion equation

$$z \in u^{-1}L_{p,q} \iff pN_0(z) - qN_1(z) = \rho_{p,q}(u).$$

The right-hand side depends on u only through its residual.
That observation is exact.

Lemma 4.1 — Every integer residual is reachable.

For every

$$r \in \mathbb{Z},$$

there exists a word

$$w \in \{0, 1\}^*$$

such that

$$\rho_{p,q}(w) = r.$$

Proof.

Since

$$\gcd(p, q) = 1,$$

Bézout's identity gives integers x_0, y_0 satisfying

$$qy_0 - px_0 = r.$$

Every pair

$$(x_0 + qk, y_0 + pk)$$

has the same value of $qy - px$.

For sufficiently large k , both coordinates are nonnegative. Any binary word with those numbers of zeros and ones therefore has residual r . $\square$

Thus the semantic residual graph will genuinely contain every integer, not merely a subset reached by special prefixes.

Theorem 4.2 — Prefix quotient classification.

For all

$$u, v \in \{0, 1\}^*,$$

$$\boxed{u^{-1}L_{p,q} = v^{-1}L_{p,q} \iff \rho_{p,q}(u) = \rho_{p,q}(v).}$$

Proof.

If the residuals agree, the completion equation is identical for u and v . Hence every continuation completes one prefix exactly when it completes the other.

Conversely, suppose

$$\rho_{p,q}(u) = r, \quad \rho_{p,q}(v) = s, \quad r \neq s.$$

Choose any nonnegative solution

$$(x, y)$$

of

$$px - qy = r,$$

for example the minimum completion demand $\mu(r)$. Let z be any word with that Parikh pair.

Then

$$z \in u^{-1}L_{p,q}.$$

But the same pair cannot also satisfy

$$px - qy = s.$$

Therefore

$$z \notin v^{-1}L_{p,q}.$$

The quotients differ. $\square$

This theorem gives the residual a stronger interpretation than the one used so far.

It is not merely a useful numeric invariant.

It is an exact coordinate for right-continuation behavior:

same residual $\iff$ same future language.
--

Left quotients themselves are classical objects. The role of the theorem in this monograph is not to claim a new general quotient theory. Its importance is that the physical DPDA configurations developed from the original machine can later be shown to encode this quotient coordinate exactly.

4.2 The quotient graph

Define the directed labeled graph

$$\mathcal{G}_{p,q}$$

with vertex set

$$\mathbb{Z}$$

and edges

$$r \xrightarrow{0} r - p,$$

$$r \xrightarrow{1} r + q.$$

By Theorem 4.2 and Lemma 4.1, the vertices of this graph may equally well be regarded as the prefix left quotients of $L_{p,q}$.

Thus the arithmetic picture

$$\dots, -2, -1, 0, 1, 2, \dots$$

is simultaneously a quotient graph.

A word labels a path by applying its symbols from left to right.

The acceptance quotient is the vertex

$$0.$$

4.3 Completion length is exactly directed distance

Chapter 2 interpreted the minimum completion length

$$\Delta(r)$$

as a shortest path to zero.

The quotient graph allows a slightly stronger translation-invariant form.

Theorem 4.3 — Directed completion metric.

Let

$$d_{\mathcal{G}}(r, s)$$

denote directed shortest-path distance in $\mathcal{G}_{p,q}$. Then

$$\boxed{d_{\mathcal{G}}(r, s) = \Delta(r - s).}$$

Proof.

A path from r to s containing x zeros and y ones has displacement

$$s - r = qy - px.$$

Equivalently,

$$px - qy = r - s.$$

Thus the Parikh vectors of directed r -to- s paths are exactly the nonnegative solutions of the completion equation for residual $r - s$.

Minimizing path length

$$x + y$$

is therefore exactly the minimum-completion problem. $\square$

The geometry is homogeneous:

$$d_{\mathcal{G}}(r, s)$$

depends only on the difference

$$r - s.$$

The language-specific shortest-completion problem is therefore one distinguished family of distances in a translation-invariant directed graph.

4.4 The fundamental balanced cycle

A directed closed walk has zero net residual displacement.

If it contains x zeros and y ones, then

$$qy - px = 0.$$

Coprimality forces a rigid form.

Theorem 4.4 — Fundamental directed cycle.

Every nonempty directed closed walk satisfies

$$N_0 = qk, \quad N_1 = pk$$

for some integer

$$k \geq 1.$$

Consequently every directed cycle length is a positive multiple of

$$p + q,$$

and the shortest possible directed cycle length is

$$\boxed{g_{p,q} = p + q.}$$

Proof.

A closed walk has

$$qN_1 = pN_0.$$

Since p and q are coprime,

$$N_0 = qk, \quad N_1 = pk$$

for some nonnegative integer k .

A nonempty walk has $k \geq 1$.

Conversely, any ordering of q zeros and p ones has net residual displacement zero and therefore gives a closed walk of length $p + q$. $\square$

The primitive kernel vector

$$K = (q, p)$$

from Chapter 2 has now become a geometric object:

$K = \text{Parikh vector of a shortest directed balanced cycle.}$

This is the global analogue of cancellation.

4.5 Balanced-cycle reduction

Suppose a word z has Parikh pair

$$(x, y).$$

If

$$x \geq q \quad \text{and} \quad y \geq p,$$

then one full balanced kernel packet

$$(q, p)$$

can be removed without changing the displacement of the word.

Repeat this as often as possible.

Theorem 4.5 — Balanced-cycle reduction.

Let a word have Parikh pair

$$(x, y).$$

Define

$$k = \min \left(\left\lfloor \frac{x}{q} \right\rfloor, \left\lfloor \frac{y}{p} \right\rfloor \right).$$

Then

$$(x', y') = (x - kq, y - kp)$$

has the same residual displacement as (x, y) and is the unique minimum Parikh pair for that displacement.

Proof.

Subtracting

$$(q, p)$$

preserves

$$px - qy.$$

By maximality of k ,

$$x' < q \quad \text{or} \quad y' < p.$$

Every integer solution with the same displacement differs by an integer multiple of

$$(q, p).$$

No further positive multiple can be removed while remaining nonnegative.

Therefore (x', y') is the first nonnegative point on the affine solution ray and hence the unique minimum pair. $\square$

This theorem gives a global quotient-graph interpretation of the minimum demand:

minimum demand = Parikh vector after deleting
every removable balanced cycle.

The original machine's cancellation behavior was local and representation-specific.

Balanced-cycle reduction is semantic and representation-independent.

4.6 Which words are geodesic?

A word is geodesic between its endpoints precisely when no balanced kernel packet can be deleted.

Corollary 4.6 — Exact geodesic domain.

A word z is shortest among all words with the same residual displacement if and only if

$$\boxed{N_0(z) < q \quad \text{or} \quad N_1(z) < p.}$$

Equivalently,

$$\text{Geo}_{p,q} = \{z : N_0(z) < q\} \cup \{z : N_1(z) < p\}.$$

The geodesic-word language is therefore regular.

The criterion is strikingly simple.

A word ceases to be shortest exactly when it contains enough of **both** symbol types to hide one complete balanced cycle in its Parikh vector.

Consequently every binary word of length less than

$$p + q$$

is geodesic.

The first possible failure occurs exactly at the fundamental cycle length.

4.7 Exact quotient spheres

Let

$$S_n = \{r \in \mathbb{Z} : d_{\mathcal{G}}(r, 0) = n\}.$$

A length- n Parikh pair has the form

$$(x, n - x).$$

Its residual displacement is

$$(p + q)x - qn,$$

which is strictly increasing in x . Hence distinct values of x give distinct residual vertices.

Combining this with the geodesic criterion gives an exact growth law.

Proposition 4.7 — Exact quotient spheres.

For every

$$n \geq 0,$$

$$|S_n| = \min(n+1, p+q).$$

Therefore

$$\sum_{n \geq 0} |S_n| z^n = \frac{1 - z^{p+q}}{(1 - z)^2}.$$

Proof.

If

$$n < p + q,$$

every length- n Parikh pair is geodesic, so all

$$n + 1$$

possible zero counts give distinct vertices at distance n .

If

$$n \geq p + q,$$

a length- n pair is geodesic precisely when

$$x < q$$

or

$$n - x < p.$$

These two regions are disjoint and contain respectively q and p pairs.

Hence exactly

$$p + q$$

residuals occur at distance n . $\square$

The graph therefore has linear growth only until the first balanced cycle appears. After that, every directed sphere has constant size

$$p + q.$$

This small exact formula will later help distinguish semantic growth from the sometimes much larger physical state/stack resources used by a particular lift.

4.8 What the original two-cell observation really reveals

Chapter 1 proved that the original 2:1 normalized completion word exposes its global symbol support within two stack cells.

Chapter 3 gave a more operational interpretation.

For a reachable configuration c , define

$$\text{GeoFirst}(c)$$

to be the set of symbols that begin some shortest continuation from the residual represented by c .

If

$$\mu(r) = (x_r, y_r),$$

then

$$\boxed{\text{GeoFirst}(r) = \{0 : x_r > 0\} \cup \{1 : y_r > 0\}.$$

For the original machine, the reachable qualitative roles are:

$$(C, \varepsilon) \mapsto \emptyset,$$

$$(C, 0^m) \mapsto \{0\},$$

$$(C, 1^m) \mapsto \{1\},$$

$$(C, 10) \mapsto \{0, 1\},$$

every reachable left configuration maps to

$$\{0, 1\},$$

and every reachable right configuration maps to

$$\{1\}.$$

Thus the old “bounded stack search” intuition can be stated exactly.

Theorem 4.8 — Exact depth-two geodesic locality of the original machine.

For the motivating 2:1 DPDA:

1. the complete geodesic first-action set is determined by the control state and the first two logical stack symbols;
2. depth one is insufficient;
3. after end-of-input normalization,

$$\boxed{\text{alph}(\widehat{S}(c)) = \text{GeoFirst}(c);}$$

4. no fixed finite observation depth determines the full quantitative demand vector on all reachable configurations.

Proof.

The reachable-store normal form from Chapter 1 leaves only the center forms

$$\varepsilon, \quad 0^m, \quad 1^m, \quad 10,$$

together with the fixed qualitative left and right roles.

At the center, two visible symbols distinguish exactly the four possible geodesic-action sets:

$$\emptyset, \quad \{0\}, \quad \{1\}, \quad \{0, 1\}.$$

Depth one fails because

$$(C, 1)$$

and

$$(C, 10)$$

have the same state and top symbol but different geodesic first-action sets.

End-of-input normalization produces the canonical shortest continuation represented by the residual, and its symbol support is exactly the set of positive coordinates of the minimum demand.

Finally, arbitrarily large reachable unary center stacks

$$(C, 0^m)$$

share every fixed top- d observation once $m \geq d$, yet their demands

$$(m, 0)$$

have different magnitudes.

Thus qualitative geometry is exactly two-local, while quantitative demand is not fixed-depth local. $\square$

This is the precise form of an important distinction:

bounded local observation can reveal the shape of the shortest-path choice

without revealing

the unbounded distance still remaining.

4.9 General canonical action visibility

For general $p:q$, recall the canonical shortest word

$$C(r) = \begin{cases} 1^{y_r} 0^{x_r}, & r > 0, \\ 0^{x_r} 1^{y_r}, & r < 0, \\ \varepsilon, & r = 0. \end{cases}$$

The minimum pair satisfies the bounded-coordinate property:

$$r > 0 \implies 0 \leq y_r < p,$$

$$r < 0 \implies 0 \leq x_r < q.$$

Therefore the first change of symbol, if it occurs, cannot be arbitrarily deep.

Definition 4.9 — Canonical visibility depth.

Let $d_{p,q}$ be the least integer d such that for every residual r ,

$$\text{alph}(C(r)) = \text{alph}(\text{pref}_d(C(r))).$$

Theorem 4.10 — Exact bounded visibility.

$$d_{p,q} = \max(p, q).$$

Proof.

Suppose first that

$$r > 0.$$

Then

$$C(r) = 1^{y_r} 0^{x_r}$$

with

$$y_r < p.$$

If a zero occurs, the first zero appears no later than position

$$p.$$

Thus depth p suffices on the positive side.

For

$$r < 0,$$

the symmetric argument shows that depth q suffices.

Hence

$$d_{p,q} \leq \max(p, q).$$

For tightness, suppose

$$p > q.$$

Take residual

$$r = q.$$

The minimum solution is

$$(x_r, y_r) = (q, p - 1),$$

so

$$C(q) = 1^{p-1}0^q.$$

Depth $p - 1$ sees only ones even though zeros occur later. Therefore

$$d_{p,q} \geq p.$$

If

$$q > p,$$

the symmetric residual

$$r = -p$$

gives

$$C(-p) = 0^{q-1}1^p,$$

forcing

$$d_{p,q} \geq q.$$

For

$$p = q = 1,$$

depth one is plainly necessary and sufficient.
Therefore

$$d_{p,q} = \max(p, q).$$

□

Since

$$\text{alph}(C(r)) = \text{GeoFirst}(r),$$

the same theorem can be read operationally:

The complete set of locally shortest first actions is visible within
 $\max(p, q)$ canonical symbols.

For the original machine,

$$(p, q) = (2, 1),$$

so

$$d_{2,1} = 2,$$

exactly reproducing the original tight depth-two theorem.

4.10 A second bounded quantity: the completion frontier

Visibility asks one question:

How far into a canonical shortest completion must we inspect before every symbol type that occurs anywhere in it has appeared?

There is another bounded quantity that will become important when we redesign the physical DPDA.

Write positive and negative residuals in Euclidean form:

$$r = pm + s, \quad 0 \leq s < p,$$

and

$$r = -(qm + t), \quad 0 \leq t < q.$$

Define the finite residue-dependent bases

$$A_s = C(s),$$

$$B_t = C(-t),$$

with

$$A_0 = B_0 = \varepsilon.$$

Then

$$\boxed{C(pm + s) = A_s 0^m,}$$

and

$$\boxed{C(-(qm + t)) = B_t 1^m.}$$

Thus every canonical completion decomposes into:

$$\boxed{\text{bounded residue-dependent frontier} + \text{unbounded unary tail}.}$$

Definition 4.11 — Completion frontier.

Define

$$K_{p,q} = \max \left(\max_{0 \leq s < p} |A_s|, \max_{0 \leq t < q} |B_t| \right).$$

Theorem 4.12 — Exact almost-cycle frontier.

For every coprime pair (p, q) ,

$$\boxed{K_{p,q} = \begin{cases} 0, & (p, q) = (1, 1), \\ p + q - 1, & \text{otherwise.} \end{cases}}$$

Proof.

For

$$(p, q) = (1, 1),$$

the only residue bases are

$$A_0 = B_0 = \varepsilon,$$

so

$$K_{1,1} = 0.$$

Now assume the coprime pair is nontrivial. Then

$$p \neq q.$$

The bounded-coordinate formulas for $C(r)$ imply

$$K_{p,q} \leq p + q - 1.$$

If

$$p > q,$$

the base residual

$$q$$

has

$$A_q = C(q) = 1^{p-1}0^q,$$

whose length is

$$p + q - 1.$$

If

$$q > p,$$

the symmetric base is

$$B_p = C(-p) = 0^{q-1}1^p,$$

again of length

$$p + q - 1.$$

Hence the upper bound is attained. $\square$

Since the shortest balanced cycle length is

$$g_{p,q} = p + q,$$

we obtain, for every nontrivial coprime pair,

$$K_{p,q} = g_{p,q} - 1.$$

The largest finite frontier is literally one symbol short of a complete balanced cycle.

4.11 One extremal word, two different bounds

The same almost-cycle word makes both visibility and frontier tight.

If

$$p > q,$$

the witness is

$$C(q) = 1^{p-1}0^q.$$

If

$$q > p,$$

the witness is

$$C(-p) = 0^{q-1}1^p.$$

But the two parameters measure different things.

Visibility depth

$$d_{p,q} = \max(p, q)$$

asks:

How many leading symbols must be inspected before the **qualitative support** of the canonical shortest completion is known?

Completion frontier

$$K_{p,q} = p + q - 1$$

for nontrivial pairs asks:

How large can the complete **residue-dependent non-unary base** be before only an unbounded unary magnitude remains?

These should not be conflated.

In particular, their numerical values are generally different even though one extremal word witnesses both lower bounds.

This distinction is crucial for later resource arguments.

A finite observation depth is not the same thing as a finite amount of residue information that can be moved into control.

4.12 The next representation idea is now forced

We have reached a useful decomposition:

$$C(r) = \text{finite residue-dependent base} + \text{unbounded unary tail}.$$

The finite base has exact worst-case size

$$K_{p,q}.$$

The unary tail stores the unbounded magnitude.

This immediately suggests a new physical representation:

store the bounded frontier in finite control

and

leave only the unbounded unary magnitude on the pushdown.

Such a factorization can be lossless because the pair

$$\text{control residue role} + \text{unary tail length}$$

reconstructs the exact residual.

Once the residual is known, Theorem 4.2 gives the exact prefix quotient, Theorem 4.3 gives the exact distance to acceptance, and Chapter 2 gives a canonical shortest continuation.

This will eventually yield a compact physical realization.

But before we redesign the machine, one more piece of geometry is needed.

The 2:1 ladder of Chapter 3 was only two rows high because its minimum demand contained a single 1.

For a general residual,

$$\mu(r) = (x_r, y_r)$$

may have several copies of both symbols.

What is the full shortest-path interval then?

The answer is a rectangle.

That is the subject of Chapter 5.

Chapter 5

General $p:q$ Completion Rectangles

The 2:1 ladder of Chapter 3 had two rows because every left-ladder source had minimum demand

$$(k + 2, 1).$$

Only one 1 was required, so the position of that single action was the only scheduling choice.

For a general ratio $p:q$, the minimum demand

$$\mu(r) = (x_r, y_r)$$

may contain several copies of both symbols.

The ladder then becomes a rectangle.

This chapter identifies that rectangle exactly and shows that three objects coincide:

minimum-demand coordinates $\longleftrightarrow$ directed shortest-path interval
 $\longleftrightarrow$ commuting schedule geometry.

5.1 Remaining demand as a grid coordinate

Fix coprime positive integers

$$p, q$$

and residual

$$r.$$

Let

$$\mu(r) = (x_r, y_r).$$

Suppose a shortest completion has so far consumed

$$i$$

zeros and

$$j$$

ones.

Its current residual is

$$r - pi + qj.$$

If the prefix is still geodesic, the remaining demand should be

$$(x_r - i, y_r - j).$$

Thus the natural coordinate set is

$$[0, x_r] \times [0, y_r],$$

where the point

$$(i, j)$$

records how much of the unique minimum Parikh demand has already been consumed.

A horizontal step consumes a zero:

$$(i, j) \rightarrow (i + 1, j).$$

A vertical step consumes a one:

$$(i, j) \rightarrow (i, j + 1).$$

The source is

$$(0, 0),$$

and the balanced target is

$$(x_r, y_r).$$

The only possible problem is that two different grid coordinates might correspond to the same residual before the target is reached.

The minimum-demand property prevents exactly that.

5.2 The shortest interval is an actual rectangle

Theorem 5.1 — Rectangular shortest-completion geometry.

Let

$$\mu(r) = (x_r, y_r).$$

The map

$$\boxed{\varphi_r(i, j) = r - pi + qj}$$

is injective on

$$[0, x_r] \times [0, y_r].$$

Consequently the directed shortest-path interval from r to 0 in the residual graph is exactly the rectangular grid

$$\boxed{[0, x_r] \times [0, y_r].}$$

Every monotone source-to-target path is a shortest completion, and every shortest completion occurs in this way.

Proof.

Suppose two rectangle coordinates satisfy

$$r - pi + qj = r - pi' + qj'.$$

Then

$$p(i - i') = q(j - j').$$

Since

$$\gcd(p, q) = 1,$$

there is an integer t such that

$$i - i' = qt, \quad j - j' = pt.$$

If

$$t \neq 0,$$

then the rectangle contains two points separated by at least one full kernel vector

$$(q, p).$$

In particular, the side lengths must satisfy

$$x_r \geq q, \quad y_r \geq p.$$

But then

$$(x_r - q, y_r - p)$$

would be a shorter nonnegative solution of the same completion equation

$$px - qy = r,$$

contradicting the minimality of

$$\mu(r).$$

Hence

$$t = 0,$$

so

$$i = i', \quad j = j'.$$

Thus φ_r is injective.

Every monotone path from

$$(0, 0)$$

to

$$(x_r, y_r)$$

consumes exactly x_r zeros and y_r ones, hence exactly the minimum demand.

Conversely, every shortest completion has that same unique minimum Parikh vector, so its sequence of consumed zero/one counts traces one monotone path through the rectangle. $\square$

This theorem generalizes the Chapter 3 ladder without changing its meaning. When

$$y_r = 1,$$

the rectangle has exactly two rows.

When

$$x_r = 0$$

or

$$y_r = 0,$$

it degenerates to a one-dimensional corridor.

When both coordinates are positive, the residual has a genuine scheduling fork.

5.3 Exact number of shortest completions

A monotone path in an

$$x_r \times y_r$$

rectangle consists of:

- x_r horizontal moves;
- y_r vertical moves.

Therefore the number of source-to-target paths is immediate.

Corollary 5.2 — Exact shortest-completion count.

For every residual r ,

$$|\text{GeoComp}(r)| = \binom{x_r + y_r}{x_r} = \binom{x_r + y_r}{y_r}.$$

Equivalently,

$$\text{GeoComp}(r) = \{z : N_0(z) = x_r, N_1(z) = y_r\}.$$

The statement is stronger than “there exists a shortest completion.”

It says that the unique minimum demand completely classifies **all** shortest words.

Their only freedom is order.

5.4 Exact number of commuting squares

Every elementary unit square of the rectangle corresponds to one local choice between

01

and

10.

There are

$$x_r y_r$$

such squares.

Define

$$\Xi(r) = x_r y_r.$$

Corollary 5.3 — Exact commutation potential.

The number of elementary commuting squares in the shortest-completion rectangle of residual r is

$$\Xi(r) = x_r y_r.$$

Thus the same arithmetic quantity introduced in the 2:1 ladder has a general meaning.

It measures the number of elementary places at which the two required symbol types can interchange their local chronological order while remaining inside the shortest interval.

For the Chapter 3 sources,

$$\mu(2k+3) = (k+2, 1),$$

so

$$\Xi(2k+3) = k+2,$$

exactly recovering the number of ladder diamonds.

5.5 A worked 3:2 rectangle

Take

$$p = 3, \quad q = 2, \quad r = 2.$$

The completion equation is

$$3x - 2y = 2.$$

The minimum nonnegative solution is

$$\mu(2) = (2, 2).$$

Thus

$$\Delta(2) = 4.$$

The shortest interval is

$$[0, 2] \times [0, 2].$$

It contains

$$(2 + 1)(2 + 1) = 9$$

residual vertices.

There are

$$\binom{4}{2} = 6$$

shortest completion words:

$$0011, \quad 0101, \quad 0110, \quad 1001, \quad 1010, \quad 1100.$$

The canonical two-block completion is

$$\boxed{C(2) = 1100.}$$

The rectangle contains

$$2 \cdot 2 = 4$$

elementary commuting squares.

A useful feature of this example is that neither symbol type occurs only once. The two-row ladder has become a genuinely two-dimensional scheduling space.

5.6 Canonical order as a boundary normal form

For

$$r > 0,$$

our canonical representative is

$$C(r) = 1^{y_r} 0^{x_r}.$$

Among all shortest words with Parikh pair

$$(x_r, y_r),$$




Figure 5.1: The 3:2, $r = 2$ shortest-completion rectangle. Horizontal edges read 0 and subtract 3 from the residual; vertical edges read 1 and add 2. The nine residual vertices form four commuting squares, and every monotone path from residual 2 to residual 0 is one of the six shortest completions with Parikh pair $(2, 2)$.

this is one boundary ordering: every required 1 is performed before every required 0.

For

$$r < 0,$$

the canonical representative

$$C(r) = 0^{x_r} 1^{y_r}$$

uses the opposite boundary orientation.

The rest of the shortest words lie between these extremal orderings.

For positive residuals, repeatedly apply the adjacent rewrite

$$01 \longrightarrow 10.$$

Each rewrite preserves the Parikh pair and therefore remains inside the same shortest-completion rectangle.

If we count inversions of the form “a 0 occurring before a 1,” every rewrite decreases that count by one.

The unique normal form is

$$1^{y_r} 0^{x_r} = C(r).$$

The maximum possible inversion count is

$$x_r y_r.$$

Thus:

$$\begin{aligned} \Xi(r) &= \text{rectangle-square count} \\ &= \text{maximum adjacent-swap normalization distance.} \end{aligned}$$

For negative residuals the symmetric rewrite

$$10 \longrightarrow 01$$

normalizes toward

$$0^{x_r} 1^{y_r}.$$

The canonical shortest word is therefore not arbitrary.

It is a boundary normal form for a finite commuting geometry.

5.7 Exactly when does a geodesic fork occur?

A residual has a genuine first-action fork exactly when both minimum-demand coordinates are positive.

For positive residuals, the bounded coordinate satisfies

$$0 \leq y_r < p.$$

The completion congruence gives

$$q y_r \equiv -r \pmod{p}.$$

Since q is invertible modulo p ,

$$y_r = 0 \iff p \mid r.$$

The negative case is symmetric.

Corollary 5.4 — Exact geodesic-fork classification.

For

$$r > 0,$$

$$\text{GeoFirst}(r) = \begin{cases} \{0\}, & p \mid r, \\ \{0, 1\}, & p \nmid r. \end{cases}$$

For

$$r < 0,$$

$$\text{GeoFirst}(r) = \begin{cases} \{1\}, & q \mid r, \\ \{0, 1\}, & q \nmid r. \end{cases}$$

At

$$r = 0,$$

the set is empty.

This gives an exact arithmetic classification of local shortest-path shape.

There are only three nonterminal qualitative possibilities:

zero corridor,	one corridor,	commuting fork.
----------------	---------------	-----------------

The original 2:1 machine has forks only on the positive side because

$$q = 1.$$

Every negative residual is divisible by q , so its shortest path is a one-way 1-corridor.

The one-sided appearance of the original transition diagram is therefore arithmetic, not merely graphical.

5.8 First obstruction at length $p + q$

Chapter 4 proved that a word is geodesic precisely when

$$N_0 < q \quad \text{or} \quad N_1 < p.$$

This immediately gives the first scale at which a word can fail to be shortest. If

$$n < p + q,$$

a length- n word cannot simultaneously contain at least q zeros and at least p ones.

Therefore:

every binary word of length $n < p + q$ is geodesic.
--

At length

$$p + q,$$

the first failures appear.

They are exactly the words with Parikh vector

$$(q, p),$$

that is, the orderings of one primitive balanced cycle.

There are

$$\boxed{\binom{p+q}{q}}$$

such first non-geodesic words.

This is another appearance of the same threshold:

$$\boxed{\begin{aligned} p + q &= \text{kernel packet length} = \text{fundamental cycle length} \\ &= \text{first possible non-geodesic length.} \end{aligned}}$$

5.9 Exact geodesic-word growth

The quotient-sphere count of Chapter 4 counted **distinct residual endpoints** at a given directed distance.

We now count something different:

How many length- n binary words themselves are geodesic between their endpoints?

Let

$$G_{p,q}(n)$$

denote that number.

Theorem 5.5 — Exact geodesic-word growth.

For

$$0 \leq n < p + q,$$

$$\boxed{G_{p,q}(n) = 2^n.}$$

For

$$n \geq p + q,$$

$$G_{p,q}(n) = \sum_{i=0}^{q-1} \binom{n}{i} + \sum_{j=0}^{p-1} \binom{n}{j}.$$

The first non-geodesic words occur at length

$$p + q$$

and are exactly the

$$\binom{p+q}{q}$$

orderings of the primitive balanced Parikh vector

$$(q, p).$$

If

$$(p, q) \neq (1, 1)$$

and

$$M = \max(p, q),$$

then

$$G_{p,q}(n) \sim \frac{n^{M-1}}{(M-1)!}.$$

For

$$(p, q) = (1, 1),$$

$$G_{1,1}(0) = 1, \quad G_{1,1}(n) = 2 \quad (n \geq 1).$$

Proof.

A length- n word with i zeros is geodesic exactly when

$$i < q$$

or

$$n - i < p.$$

For

$$n < p + q,$$

these inequalities cannot both fail, so every word is geodesic.
For

$$n \geq p + q,$$

the two geodesic regions are disjoint:

- words with fewer than q zeros;
- words with fewer than p ones.

Counting them gives the two fixed binomial tails.
For a nontrivial coprime pair,

$$p \neq q.$$

The larger fixed binomial tail has highest index

$$M - 1.$$

Hence the dominant term is

$$\binom{n}{M-1} \sim \frac{n^{M-1}}{(M-1)!}.$$

The balanced case $(1, 1)$ reduces to the two unary words at every positive length. $\square$

This theorem shows a sharp change of growth.

Before the first balanced cycle, geodesic words grow exponentially:

$$2^n.$$

After the threshold, only two bounded-width binomial tails survive, and growth becomes polynomial.

Again, the primitive balanced cycle is the obstruction.

5.10 Distinct paths versus distinct quotient vertices

Two counting functions have now appeared:

$$|S_n| = \min(n + 1, p + q)$$

for quotient vertices at directed distance n , and

$$G_{p,q}(n)$$

for geodesic words of length n .

They measure different kinds of multiplicity.

The first forgets word order and remembers only the endpoint residual.

The second counts all geodesic orderings.

Inside a fixed shortest-completion rectangle, the same distinction appears as:

grid vertices versus monotone paths.

The semantic quotient collapses many word histories into one residual coordinate.

The rectangle restores exactly the scheduling information relevant to a shortest interval.

This interplay between quotienting and lifting will become a recurring theme later in the monograph.

5.11 The 2:1 ladder is now completely explained

Return to the Chapter 3 source

$$(L, 0^k).$$

Its residual is

$$2k + 3,$$

and its minimum demand is

$$(k + 2, 1).$$

The general rectangle is therefore

$$[0, k + 2] \times [0, 1].$$

That is precisely a two-row ladder.

Its path count is

$$\binom{k + 3}{1} = k + 3.$$

Its square count is

$$(k + 2) \cdot 1 = k + 2.$$

The positive residual is odd, hence not divisible by $p = 2$, so it has a fork:

$$\text{GeoFirst}(2k + 3) = \{0, 1\}.$$

Every feature of the original operational ladder is therefore a specialization of the general minimum-demand rectangle.

This is an important methodological point.

We did not begin by defining a rectangle and then force the original machine into it.

The rectangle was discovered by first analyzing a concrete cancellation/completion ladder in the machine and then identifying the arithmetic that generalized it.

That discovery order will be preserved in the monograph.

5.12 From geometry back to representation

We now have two complementary pieces of structure.

Chapter 4 gave the canonical completion decomposition

$$C(r) = \text{finite residue frontier} + \text{unbounded unary tail}.$$

This suggests moving the finite frontier into control.

Chapter 5 gives the full shortest-path rectangle

$$[0, x_r] \times [0, y_r].$$

This gives an exact semantic graph that a physical DPDA representation ought to realize.

These facts create a new engineering objective.

Can we construct a pushdown representation in which:

1. the exact residual is decodable;
2. the exact prefix quotient is therefore known;
3. the unbounded magnitude remains stack-like;
4. the finite residue information is moved into control;
5. ordinary input transitions reproduce the residual graph as faithfully as possible?

The first attempt will work semantically but reveal a boundary defect.

That defect will turn out to be useful.

It will teach us that the problem is not the residual semantics.

The problem is the coordinate system.

The next chapter begins that redesign.

PART III

Better Coordinates and Exact Resource Frontiers

Research provenance. This Part consolidates the coordinate constructions and exact residual-faithful resource frontiers developed across [1, 2, 4]. State/stack-symbol tradeoffs and limited pushdown alphabets have a substantial descriptional-complexity literature [37, 36, 54, 60]; the frontier statements below are therefore kept explicitly scoped to their representation classes.

Chapter 6

The First General Physical Lift and Its Defects

The previous chapters separated two objects that were entangled in the original 2023 machine:

semantic residual and physical pushdown representation.

The residual graph is now completely explicit:

$$r \xrightarrow{0} r - p, \quad r \xrightarrow{1} r + q.$$

Its shortest intervals are rectangles, its quotient coordinate is one integer, and its canonical completion has a bounded finite frontier followed by an unbounded unary tail.

This makes a new design problem unavoidable.

Can the residual graph itself be lifted into a compact DPDA configuration graph?

The first answer is yes—but only after a deterministic normalization step at a finite boundary.

That apparent imperfection will be useful. It will identify exactly what the first coordinates got wrong.

6.1 Finite phase plus unary magnitude

Let

$$m = \max(p, q).$$

The idea is to move the bounded residue information into finite control while leaving only an unbounded magnitude on the stack.

Use stable control states

$$R_0, \dots, R_{m-1}$$

and stack alphabet

$$\Gamma_4 = \{0, 1, X, Z_0\}.$$

The positive and negative ends are represented by different unary symbols:

- 0 for the positive end;
- 1 for the negative end.

The additional symbol X is a negative-side sentinel.

For

$$0 \leq s < p,$$

define

$$a_s = \begin{cases} 0, & s = 0, \\ s - p, & 1 \leq s < p. \end{cases}$$

For

$$0 \leq t < q,$$

define

$$b_t = \begin{cases} 0, & t = 0, \\ q - t, & 1 \leq t < q. \end{cases}$$

These are not arbitrary constants. They choose one finite base residual for each positive and negative residue class.

6.2 The stable four-symbol encoding

Define

$$\eta(0) = (R_0, Z_0).$$

If

$$r > 0,$$

let

$$s = r \bmod p$$

and write uniquely

$$r = pk + a_s, \quad k \geq 1.$$

Set

$$\boxed{\eta(r) = (R_s, 0^k Z_0).}$$

If

$$r < 0,$$

let

$$t = (-r) \bmod q$$

and write uniquely

$$r = -qk + b_t, \quad k \geq 1.$$

Set

$$\boxed{\eta(r) = (R_t, 1^k X).}$$

The representation is injective.

The control state records the finite phase. The unary tail records the unbounded quotient magnitude. The sign is visible physically from the top/sentinel convention.

This is the first systematic realization of the Chapter 4 factorization

finite boundary role + unbounded unary magnitude.

6.3 Stable rules away from the boundary

On the positive side, input 0 should subtract p . Deep inside the unary tail, this is simply a pop:

$$\delta(R_s, 0, 0) = (R_s, \varepsilon).$$

Input 1 adds q . Let

$$s' = (s + q) \bmod p$$

and define

$$c_s = \frac{a_s + q - a_{s'}}{p}.$$

Then

$$pk + a_s + q = p(k + c_s) + a_{s'}.$$

Thus the correct stable rule is

$$\delta(R_s, 1, 0) = (R_{s'}, 0^{c_s+1}).$$

The negative side is symmetric.

Let

$$t' = (t + p) \bmod q$$

and

$$d_t = \frac{p - b_t + b_{t'}}{q}.$$

Then

$$\delta(R_t, 1, 1) = (R_t, \varepsilon)$$

and

$$\delta(R_t, 0, 1) = (R_{t'}, 1^{d_t+1}).$$

Away from the finite boundary, the design is already raw: one input symbol produces the exact next stable residual encoding.

The only difficulty appears when a unary tail is completely consumed.

6.4 Forced boundary normalization

Suppose a positive configuration has only one unary 0 left.

Consuming that 0 exposes

$$Z_0.$$

For residue class $s = 0$, this is already the correct zero configuration.

For

$$1 \leq s < p,$$

the configuration

$$(R_s, Z_0)$$

is only transient. A unique forced normalization moves it to the stable encoding of

$$a_s = s - p < 0.$$

Likewise, on the negative side, consuming the last 1 exposes

$$X,$$

and a unique forced normalization moves to the stable encoding of the corresponding nonnegative base residual b_t .

This motivates a semantic macro-step.

Definition 6.1 — Normalized input macro-step.

Write

$$c \xRightarrow{a} c'$$

when the machine:

1. consumes the binary symbol a ;
2. performs the unique forced finite normalization chain, if one is enabled;
3. stops at the next stable encoded configuration.

The normalization consumes no additional binary input and introduces no choice.

6.5 Exact normalized residual conjugacy

Theorem 6.2 — Normalized residual–configuration conjugacy.

For every

$$r \in \mathbb{Z},$$

$$\boxed{\eta(r) \xRightarrow{0} \eta(r - p), \quad \eta(r) \xRightarrow{1} \eta(r + q).}$$

Consequently, for every binary word u ,

$$\boxed{\widehat{\delta}^*(\eta(r), u) = \eta(r - pN_0(u) + qN_1(u)).}$$

Thus η is a label-preserving graph isomorphism between the residual graph and the stable configuration graph **modulo deterministic normalization**.

Proof.

Let

$$r = pk + a_s > 0.$$

On input 0, if $k > 1$, one unary 0 is popped and the result is immediately

$$\eta(r - p).$$

If $k = 1$, the pop reaches the finite boundary. The uniquely forced normalization sends the exposed boundary pair to the stable encoding of $r - p$.

On input 1,

$$pk + a_s + q = p(k + c_s) + a_{s'},$$

so the positive stable rule lands directly in

$$\eta(r + q).$$

The negative case is symmetric. Input 1 pops one unary 1, with the $k = 1$ case followed by its unique sentinel normalization; input 0 uses the identity defining d_t .

At residual zero the two binary moves are defined directly toward the encodings of $-p$ and q .

The one-step identities imply the word formula by induction. $\square$

This theorem is already a major improvement in understanding.

The original $C/L/R$ machine encoded the residual through a more specialized stack/state interaction.

The new construction makes the quotient coordinate explicit:

finite residual phase in control
+ unbounded quotient magnitude on a unary tail.

6.6 The abstract rectangles lift to physical configurations

Let

$$\mu(r) = (x_r, y_r).$$

Define

$$\widehat{\Psi}_r(i, j) = \eta(r - pi + qj).$$

Corollary 6.3 — Exact normalized configuration rectangle.

The map

$$\widehat{\Psi}_r$$

is a label-preserving isomorphism from

$$[0, x_r] \times [0, y_r]$$

onto the normalized shortest-path interval from

$$\eta(r)$$

to

$$\eta(0).$$

Every elementary rectangle square is a genuine normalized commuting diamond:

$$\widehat{\delta}^*(\widehat{\Psi}_r(i, j), 01) = \widehat{\Psi}_r(i + 1, j + 1) = \widehat{\delta}^*(\widehat{\Psi}_r(i, j), 10).$$

The shortest-completion rectangle is therefore not merely an arithmetic diagram.

It is realized by actual DPDA configurations.

The qualification is important:

$$\boxed{\text{one rectangle edge may be one input move + forced normalization.}}$$

The original 2:1 ladder was stronger: its relevant edges were ordinary binary moves.

The question is now whether the general construction can be made equally literal.

6.7 A concrete 3:2 boundary crossing

Take

$$p = 3, \quad q = 2, \quad r = 4.$$

Then

$$\mu(4) = (2, 1).$$

The normalized rectangle contains

$$\eta(4) = (R_1, 00Z_0),$$

$$\eta(1) = (R_1, 0Z_0),$$

$$\eta(-2) = (R_0, 1X),$$

$$\eta(6) = (R_0, 00Z_0),$$

$$\eta(3) = (R_0, 0Z_0),$$

$$\eta(0) = (R_0, Z_0).$$

The lower horizontal edge

$$1 \xrightarrow{0} -2$$

is semantically correct, but in these coordinates the binary pop first exposes a transient boundary pair before the forced normalization reaches

$$\eta(-2).$$

The rectangle remains exact.

The edge is simply not yet **raw**.

6.8 Raw residual conjugacy

We now strengthen the representation requirement.

Definition 6.4 — Raw residual–configuration conjugacy.

An injective encoding

$$\zeta : \mathbb{Z} \rightarrow \text{DPDA configurations}$$

is a raw residual conjugacy if

$$\boxed{\delta(\zeta(r), 0) = \zeta(r - p), \quad \delta(\zeta(r), 1) = \zeta(r + q)}$$

for every residual r , using exactly one ordinary binary-input transition and no normalization step.

A raw defect of an encoding is therefore an edge on which this stronger identity fails.

The four-symbol construction is normalized-exact.

We now ask exactly where it fails to be raw.

6.9 The defect set is finite and exact

Proposition 6.5 — Exact four-symbol raw-defect set.

For the four-symbol encoding η ,

$$\mathcal{D}_{p,q} = \{(s, 0) : 1 \leq s < p\} \cup \{(-j, 1) : 1 \leq j \leq q\}.$$

Therefore

$$|\mathcal{D}_{p,q}| = p + q - 1.$$

Every residual edge outside this set is already one raw binary transition between stable configurations.

Proof.

On the positive side, input 1 always uses a stable rule.

Input 0 is also raw whenever at least two unary zeros are present. The only problem occurs when the source residual is one of

$$1, 2, \dots, p - 1.$$

The last zero is popped, leaving the transient state/bottom pair

$$(R_s, Z_0)$$

instead of the required negative encoding.

These are exactly the $p - 1$ positive defects.

On the negative side, input 0 is always raw.

Input 1 is raw while at least two unary ones remain. At the final unary one, the pop exposes X , and the unique normalization is still required.

The affected residuals are

$$-q, -(q - 1), \dots, -1,$$

giving exactly q negative defects.

The two families are disjoint. $\square$

The count

$$p + q - 1$$

is numerically the same as the completion-frontier scale of Chapter 4 for nontrivial coprime pairs.

The meanings are different.

One counts the largest bounded canonical completion base.

The other counts the raw boundary defects of this particular four-symbol encoding.

The equality is evidence that the same finite arithmetic boundary is being seen from two physical viewpoints—not that the notions are identical.

6.10 Why the defects cannot be patched locally in the same coordinates

It is tempting to keep the encoding η and merely change the boundary transition rules.

That does not work.

Proposition 6.6 — Boundary obstruction for the fixed encoding.

Keeping the four-symbol physical representations $\eta(r)$, no standard top-rewrite PDA transition function can make all edges in

$$\mathcal{D}_{p,q}$$

raw while preserving the already correct behavior on the adjacent deeper unary tails.

Hence every defect of Proposition 6.5 is unavoidable **within these coordinates**.

Proof.

Fix

$$1 \leq s < p.$$

Compare

$$\eta(p + s) = (R_s, 00Z_0)$$

with

$$\eta(s) = (R_s, 0Z_0).$$

The two sources have the same state and visible top symbol.
Correct realization of

$$p + s \xrightarrow{0} s$$

forces the transition on the local pair

$$(R_s, 0)$$

to pop one zero and remain in state R_s .

But the same deterministic rule applied to

$$(R_s, 0Z_0)$$

must produce

$$(R_s, Z_0),$$

whereas the correct target is the negative encoding

$$\eta(s - p).$$

The transition cannot distinguish the two hidden suffixes.

Thus the positive boundary edge cannot be raw without breaking the already correct deeper edge.

The negative case is analogous. Compare

$$(R_t, 11X)$$

and

$$(R_t, 1X).$$

Correct deep behavior forces the visible-1 rule to pop. Applied at the boundary, the same pop leaves

$$(R_t, X),$$

not the required nonnegative target.

The obstruction depends only on the fixed state/top encoding and one-top locality. $\square$

This is the real lesson of the defect theorem.

The problem is not that we failed to invent a clever enough boundary rule.

The local rule is already determined by the deeper tail.

6.11 Which shortest rectangles are already raw?

Corollary 6.7 — Raw rectangle criterion in the four-symbol orientation.

The entire shortest-completion rectangle from residual r is raw under η exactly when

$$\boxed{r = 0 \quad \text{or} \quad (r > 0 \text{ and } p \mid r).}$$

Interpretation. If

$$r > 0$$

and

$$p \mid r,$$

then

$$\mu(r) = (r/p, 0),$$

so the rectangle degenerates to a pure zero corridor and never crosses a defective boundary.

Every positive residual with a genuine two-symbol fork eventually exposes one of the positive boundary defects.

Every negative nonzero rectangle contains the final edge

$$-q \xrightarrow{1} 0,$$

which is defective in this orientation.

The asymmetry is representational.

Zero shares the positive sentinel Z_0 , while the negative side uses X .

That observation points directly toward the cure.

6.12 The defect is coordinate-relative

The fixed-coordinate obstruction proves:

these coordinates cannot be made globally raw.

It does **not** prove:

four stack symbols are necessary.

Those are very different statements.

The first is a theorem.

The second would be an unjustified extrapolation.

The right design question is therefore not:

What extra symbol should we add?

It is:

Can we align the finite residual bases differently so that the sign-changing boundary is itself an ordinary top rewrite?

The answer is yes.

Even better, the new coordinates remove the extra sentinel X .

The next chapter performs that coordinate change.

Chapter 7

From Boundary Defects to a Raw Three-Symbol DPDA

Chapter 6 ended with a precise failure.

The four-symbol representation

$$\eta$$

was semantically exact, but at

$$p + q - 1$$

boundary edges an ordinary binary transition could not land directly in the next stable residual encoding.

The obstruction was not global.

It was local:

the same state/top pair had to serve both
a deep unary tail and a sign boundary.

The first coordinates placed the finite residual bases on the wrong side of that boundary.

We now shift those bases.

The result is stronger in three ways:

1. every residual edge becomes one raw input transition;
2. the extra sentinel X disappears;
3. the construction still uses only

$$\max(p, q)$$

control states.

This is the first major example in the monograph where understanding the mathematics of a DPDA leads directly to a better DPDA.

7.1 Before using two work symbols: the unary baseline

It is useful to know what can be done with only one work-stack symbol.

Let

$$\Gamma_2 = \{A, Z_0\},$$

where Z_0 is a permanent bottom marker.

Every encoded stack then has the form

$$A^h Z_0.$$

The state must therefore distinguish every finite role that cannot be carried by stack-symbol type.

Theorem 7.1 — Exact one-work-symbol raw-conjugacy complexity.

For every coprime

$$p, q \geq 1,$$

any ε -free permanent-bottom raw residual conjugacy using only one work symbol requires

$$\boxed{|Q| \geq p + q.}$$

There is a right-endmarker construction attaining

$$\boxed{|Q| = p + q, \quad |\Gamma| = 2.}$$

Hence $p + q$ is exact in this model.

Proof idea

Input 0 partitions the residual line into p bi-infinite orbits modulo p .

Far out on the positive end, injectivity forces the stack height to escape every bounded region. Since a single top rewrite can decrease work height by at most one, the eventual input-0 dynamics must contain enough negative-drift state phases to distinguish all p positive residue classes.

Thus at least p state roles are required.

Symmetrically, input 1 forces at least q state roles on the negative end.

With only one work symbol, those two tail families cannot safely share the same control roles: otherwise repeated negative drift would send infinitely many distinct residuals into finitely many bottom configurations.

Therefore the roles are disjoint:

$$|Q| \geq p + q.$$

A matching construction uses state families

$$P_0, \dots, P_{p-1}$$

and

$$N_0, \dots, N_{q-1},$$

with the single work symbol storing only unary magnitude.

The full transition table will be retained in the construction appendix.

The theorem establishes an exact baseline:

$$\boxed{\text{one work symbol} \implies p + q \text{ states.}}$$

The next question is whether a second work symbol can let the positive and negative phase roles share control states.

It can.

7.2 Boundary-compatible coordinates

Assume first

$$p \geq q.$$

Use only

$$\boxed{\Gamma_3 = \{0, 1, Z_0\}.}$$

Use control states

$$\boxed{Q_3 = \{R_0, \dots, R_{p-1}\}.}$$

There are exactly

$$p = \max(p, q)$$

states.

For nonnegative residuals, use ordinary Euclidean division:

$$r = pk + s, \quad k \geq 0, \quad 0 \leq s < p.$$

Define

$$\zeta(r) = (R_s, 0^k Z_0).$$

For negative residuals, write uniquely

$$r = u - qh, \quad h \geq 1, \quad 0 \leq u < q.$$

Define

$$\zeta(r) = (R_u, 1^h Z_0).$$

This is the decisive shift.

The negative side no longer terminates in a separate sentinel X .

Its last unary 1 sits directly above the ordinary bottom marker.

The states

$$R_0, \dots, R_{q-1}$$

are deliberately shared between:

- nonnegative boundary roles;
- negative residue phases.

The sign remains visible because:

- a negative configuration has top 1;
- a positive nonboundary configuration has top 0;
- a nonnegative boundary configuration has top Z_0 .

Thus state sharing does not destroy injectivity.

7.3 Positive-tail arithmetic

For

$$0 \leq s < p,$$

define

$$s^+ = (s + q) \bmod p$$

and

$$c_s = \frac{s + q - s^+}{p}.$$

Since

$$0 < q \leq p,$$

we have

$$c_s \in \{0, 1\}.$$

On a positive unary tail:

$$\boxed{\delta(R_s, 0, 0) = (R_s, \varepsilon),}$$

and

$$\boxed{\delta(R_s, 1, 0) = (R_{s^+}, 0^{c_s+1}).}$$

The arithmetic identity is

$$pk + s + q = p(k + c_s) + s^+.$$

So the visible rewrite is exactly the residual update.

7.4 The repaired sign boundary

For the nonnegative boundary, define

$$u_s = (s - p) \bmod q$$

and

$$h_s = \frac{u_s + p - s}{q}.$$

Then

$$s - p = u_s - qh_s.$$

Therefore input 0 at the bottom marker can cross directly into the negative encoding:

$$\boxed{\delta(R_s, 0, Z_0) = (R_{u_s}, 1^{h_s} Z_0).}$$

No transient pair is created.

No normalization follows.

The boundary itself is now raw.

Input 1 at a nonnegative boundary uses

$$\boxed{\delta(R_s, 1, Z_0) = (R_{s^+}, 0^{c_s} Z_0).}$$

Again the target is already the exact stable encoding.

7.5 Negative-tail arithmetic and the second repaired boundary

For

$$0 \leq u < q,$$

define

$$u^- = (u - p) \bmod q$$

and

$$d_u = \frac{u^- - u + p}{q}.$$

On the negative tail:

$$\boxed{\delta(R_u, 1, 1) = (R_u, \varepsilon),}$$

and

$$\boxed{\delta(R_u, 0, 1) = (R_{u-}, 1^{d_u+1}).}$$

Now observe what happens when the negative stack has only one 1:

$$(R_u, 1Z_0) \xrightarrow{1} (R_u, Z_0).$$

But by definition

$$(R_u, Z_0) = \zeta(u).$$

The final negative 1-pop therefore crosses the sign boundary automatically.

The very rule that was forced by the deep negative tail is now also the correct boundary rule.

This is exactly what the four-symbol coordinates failed to achieve.

7.6 Global raw conjugacy

Theorem 7.2 — Three-symbol raw residual–configuration conjugacy.

For every coprime

$$p, q \geq 1,$$

there is an ε -free right-endmarker DPDA with

$$|Q| = \max(p, q), \quad |\Gamma| = 3,$$

whose stack alphabet is

$$\{0, 1, Z_0\}$$

and whose injective encoding

$$\zeta : \mathbb{Z} \rightarrow Q \times \Gamma^*$$

satisfies

$$\delta(\zeta(r), 0) = \zeta(r - p), \quad \delta(\zeta(r), 1) = \zeta(r + q)$$

for every residual r .

Hence, for every word w ,

$$\delta^*(\zeta(r), w) = \zeta(r - pN_0(w) + qN_1(w)).$$

The right endmarker is accepted exactly at

$$\zeta(0).$$

Proof.

Assume

$$p \geq q.$$

The nonnegative and negative encodings are injective within their own Euclidean decompositions.

They are mutually disjoint physically because the negative family has top 1, whereas the nonnegative family has top 0 or Z_0 .

If

$$r = pk + s$$

with

$$k \geq 1,$$

input 0 pops one zero:

$$pk + s - p = p(k - 1) + s.$$

Input 1 uses

$$pk + s + q = p(k + c_s) + s^+.$$

If

$$k = 0,$$

the same input-1 identity works at Z_0 , while the boundary identity

$$s - p = u_s - qh_s$$

makes input 0 land directly in the negative encoding.

Now let

$$r = u - qh < 0.$$

Input 1 pops one unary 1.

If

$$h > 1,$$

the result is still negative.

If

$$h = 1,$$

the pop leaves

$$(R_u, Z_0) = \zeta(u),$$

which is exactly the nonnegative target.

Finally,

$$u - p = u^- - qd_u$$

gives

$$r - p = u^- - q(h + d_u),$$

so the input-0 rewrite produces exactly the required deeper negative encoding.

No case uses an ε -move.

For

$$q > p,$$

exchange the roles of 0 and 1, exchange p and q , and reverse the residual sign. $\square$

This theorem removes both weaknesses of Chapter 6 at once:

$(\max(p, q), 4, \text{normalized macro-edges})$

becomes

$$(\max(p, q), 3, \text{raw input edges}).$$

7.7 The 3:2 boundary is now literal

Return to

$$p = 3, \quad q = 2, \quad r = 4.$$

The new encodings include

$$\zeta(4) = (R_1, 0Z_0),$$

$$\zeta(1) = (R_1, Z_0),$$

$$\zeta(-2) = (R_0, 1Z_0),$$

$$\zeta(0) = (R_0, Z_0).$$

The old defective edge

$$1 \xrightarrow{0} -2$$

is now exactly

$$\delta(R_1, 0, Z_0) = (R_0, 1Z_0).$$

The old negative normalization edge

$$-2 \xrightarrow{1} 0$$

is now the ordinary pop

$$\delta(R_0, 1, 1) = (R_0, \varepsilon),$$

which exposes Z_0 .

The same semantic rectangle now exists literally in the raw configuration graph.

7.8 Raw rectangles

Corollary 7.3 — Three-symbol raw configuration rectangles.

Let

$$\mu(r) = (x_r, y_r).$$

Define

$$\Psi_r^{\text{raw}}(i, j) = \zeta(r - pi + qj).$$

Then

$$\Psi_r^{\text{raw}}$$

is a label-preserving isomorphism from

$$[0, x_r] \times [0, y_r]$$

onto the ordinary DPDA shortest-path interval from

$$\zeta(r)$$

to

$$\zeta(0).$$

Every elementary square is a raw commuting diamond:

$$\boxed{\delta^*(\Psi_r^{\text{raw}}(i, j), 01) = \Psi_r^{\text{raw}}(i + 1, j + 1) = \delta^*(\Psi_r^{\text{raw}}(i, j), 10).}$$

The Chapter 5 rectangle has now been fully lifted into one-symbol DPDA dynamics.

No quotient by normalization is needed.

7.9 What has actually been optimized?

The construction uses

$$\max(p, q)$$

states and two work symbols 0, 1 above the bottom marker.

At this point it is tempting to call it simply “the optimal DPDA.”

That wording is too broad.

We must specify the representation category.

The Chapter 7 machine is optimal under several increasingly strong statements, but not yet under every conceivable notion of DPDA recognition complexity.

Theorem 7.4 — Exact homogeneous-tail frontier.

In the permanent-bottom, ε -free raw-conjugacy model with asymptotically homogeneous positive and negative tails:

$$S_{\text{hom}}(k) = \begin{cases} p + q, & k = 1, \\ \max(p, q), & k \geq 2, \end{cases}$$

where k is the number of work symbols.

Thus a second work symbol gives the entire possible state reduction in this model.

Additional homogeneous-tail work symbols cannot reduce the state count below

$$\max(p, q).$$

The proof uses preserved deep suffixes and the cyclic action of:

- addition by q on positive classes modulo p ;
- subtraction of p on negative classes modulo q .

Coprimality forces one common positive tail symbol and one common negative tail symbol. With one work symbol the two ends must remain state-disjoint; with two symbols they may share control phases.

This is already an exact state/stack tradeoff.

But it still assumes a homogeneous-tail representation shape.

7.10 The motivating 2:1 ratio survives arbitrary three-symbol encodings

The original research program began at

$$(p, q) = (2, 1).$$

Here the stronger lower bound can be proved without assuming homogeneous tails or sign separation.

Theorem 7.5 — Exact three-symbol state complexity for 2:1.

Consider the full permanent-bottom stack alphabet

$$\Gamma_3 = \{A, B, Z_0\}$$

with:

- no ε -moves;
- injective raw residual conjugacy;
- no homogeneous-tail assumption;
- no prescribed sign-separated encoding.

Then

$$\boxed{|Q| \geq 2.}$$

The construction of Theorem 7.2 uses exactly two states, so

$$\boxed{|Q|_{\min} = 2}$$

for three-symbol raw conjugacy of the motivating 2:1 residual graph in this model.

Proof strategy

Assume one state.

Let

$$G : S_r \mapsto S_{r+1}$$

be input 1 and

$$F : S_r \mapsto S_{r-2}$$

be input 0.

On the encoded orbit,

$$\boxed{FG^2 = G^2F = \text{id.}}$$

Choose an encoded residual of minimum work-stack height.

The inverse identity bounds the lengths of all possible top replacements, reducing the problem to finitely many local rewrite patterns.

If the minimum height is positive, a common inert suffix can be stripped. The remaining possibilities for the first few G -iterates are finite, and each is eliminated by one of:

- preserved-suffix impossibility;
- forced rule reuse;
- violation of minimum height;
- collision of two residuals.

If the minimum height is zero, the bottom write is likewise bounded and only finitely many cases remain. Every case again contradicts injectivity or the inverse relation.

The complete case table will be retained in the lower-bound appendix.

Thus one state is impossible.

This theorem is particularly important for the monograph because it ties the improved machine back to the original ratio:

The analysis of the 2023 2:1 DPDA leads to an exact two-state, three-symbol raw residual-conjugate realization.




Figure 7.1: The improved two-state, three-symbol raw residual-conjugate DPDA for the motivating 2:1 residual graph. Labels have the form input, top/replacement; only transition cells occurring on the encoded residual orbit are shown. The right endmarker is accepted exactly at (R_0, Z_0) .

Figure 7.1 is exact in the permanent-bottom, ε -free raw residual-conjugacy model. It is not a claim of unrestricted global optimality among all DPDAs recognizing $L_{2,1}$.

The scope words remain essential.

7.11 Further exact special ratios

The same unrestricted permanent-bottom three-symbol question can be solved at two further small ratios.

Theorem 7.6 — Exact 3:1 three-symbol state complexity.

For

$$(p, q) = (3, 1),$$

every ε -free permanent-bottom three-symbol raw conjugacy requires

$$\boxed{|Q| \geq 3.}$$

The construction of Theorem 7.2 attains three states.
Hence

$$\boxed{|Q|_{\min} = 3.}$$

The symmetric result holds for

$$(1, 3).$$

The proof again reduces an assumed smaller-state conjugacy to finitely many bounded top-rewrite possibilities, but the case structure is substantially larger than at 2:1. The detailed finite reduction belongs in the appendix of the complete research version.

Theorem 7.7 — Computer-assisted exact 3:2 result.

For

$$(p, q) = (3, 2),$$

a finite mathematically bounded search eliminates every two-state permanent-bottom three-symbol raw conjugacy.

Therefore

$$\boxed{|Q| \geq 3,}$$

and the three-state construction is exact:

$$\boxed{|Q|_{\min} = 3.}$$

The symmetric result holds for

$$(2, 3).$$

This theorem is explicitly **computer-assisted**.

The computer is used only after the proof reduces the unrestricted-looking problem to a finite bounded search space. The survivor vectors and branch counts belong in the reproducibility appendix; the theorem should never be presented as a purely symbolic proof.

At this stage of the story, the corresponding arbitrary-encoding theorem for every coprime pair remains unresolved.

Chapter 8 will close the full raw residual-faithful frontier by a different, representation-shape-free argument.

7.12 The right endmarker saves one control state in the scoped models

The constructions above use a distinguished right endmarker and accept exactly at the zero residual configuration.

If ordinary final-state acceptance on the binary alphabet is required instead, the zero configuration must use a final state that does not also occur on a nonzero residual.

In the scoped raw-conjugacy models this costs exactly one state.

Corollary 7.8 — Exact no-endmarker cost in the early raw models.

For the one-work-symbol model:

$$|Q| = p + q + 1, \quad |\Gamma| = 2.$$

For the sign-separated homogeneous-tail three-symbol model:

$$|Q| = \max(p, q) + 1, \quad |\Gamma| = 3.$$

Both bounds are tight.

This acceptance cost is separate from the stack-alphabet tradeoff.

It changes only how residual zero is represented.

7.13 The coordinate-shift improvement

The conclusion of Chapters 6 and 7 can now be stated cleanly.

Corollary 7.9 — Coordinate-shift improvement.

The $p + q - 1$ raw defects of the first four-symbol encoding are not a lower bound on the number of stack symbols.

They are properties of those coordinates.

A boundary-compatible shift gives

$$(\max(p, q), 4, \text{normalized macro-steps}) \longrightarrow (\max(p, q), 3, \text{raw binary steps}).$$

The missing ingredient was not an additional endpoint bit.

The finite bases had been aligned incorrectly with the sign-changing unary pops.

After moving the bases to adjacent residual classes:

- the same state can serve the post-pop boundary target;
- the ordinary bottom marker is sufficient;
- the extra sentinel X disappears;

- every residual edge becomes literal.

This is the first major design lesson of the monograph:

a locality defect may be a coordinate defect
rather than a semantic necessity.

7.14 Three different meanings of “better”

The evolution from the original machine to the raw three-symbol construction improves the representation in several non-equivalent ways.

Semantic transparency

The residual is decoded by a simple Euclidean state/tail coordinate.

Transition literalness

Every binary residual edge is exactly one DPDA input transition.

Stack alphabet

The general raw construction needs only

$$\{0, 1, Z_0\}.$$

Control width

It uses

$$\max(p, q)$$

states.

But the exactness claims established so far belong to stated raw-conjugacy categories.

We have not proved that no completely different DPDA recognizing the same language can use another representation and beat those counts.

That distinction will become one of the central methodological themes of the book.

7.15 The next question removes the representation shape

The current lower bounds still come from several different levels of assumption:

- one work symbol;
- homogeneous positive/negative tails;
- special-ratio arbitrary three-symbol encodings.

The obvious next question is:

Are the resource points

$$p + q \quad \text{and} \quad \max(p, q)$$

accidents of these structured coordinates, or are they forced by raw residual conjugacy itself?

To answer that, we must stop assuming homogeneous tails.

We will keep only the semantic requirement:

Every integer residual has one injective stable physical encoding, and every residual edge is one raw input transition.

From that weak assumption alone, the next chapter derives the complete exact frontier.

That result will turn the improved construction discovered here into a mature resource theorem.

Chapter 8

The Complete Raw Residual-Faithful Frontier

Chapter 7 produced a sharp construction:

$$|Q| = \max(p, q), \quad |\Gamma| = 3,$$

with one raw binary transition for every residual edge.

It also produced exact lower bounds in several structured settings:

- one work symbol;
- homogeneous positive/negative tails;
- unrestricted three-symbol encodings for several small ratios.

The remaining question is whether the same resource points are forced **without assuming a particular stack shape**.

This chapter removes that assumption.

The result is the mature form of the early resource story:

$$S_{\text{raw}}(g) = \begin{cases} \infty, & g = 0, \\ p + q, & g = 1, \\ \max(p, q), & g \geq 2. \end{cases}$$

The proof is deliberately representation-shape-free.

It does not assume:

- homogeneous tails;
- sign separation;
- a canonical coordinate formula;
- fixed stack depth;
- bounded replacement length.

Only raw residual conjugacy remains fixed.

8.1 Model R

Let

$$\Gamma = \Delta \cup \{Z_0\},$$

where Z_0 is a permanent bottom marker and

$$g = |\Delta|$$

is the number of work symbols.

Definition 8.1 — Model R.

A raw residual conjugacy is an injective map

$$\eta : \mathbb{Z} \rightarrow Q \times \Delta^* Z_0$$

such that, for every residual r ,

$$\boxed{\delta(\eta(r), 0) = \eta(r - p), \quad \delta(\eta(r), 1) = \eta(r + q),}$$

with no ε -normalization on encoded residual configurations.

Let

$$S_{\text{raw}}(g)$$

be the minimum number of control states among such realizations, with value ∞ if no finite-state realization exists.

The semantic restriction is exact but narrow:

$$\boxed{\text{one residual} \leftrightarrow \text{one stable physical configuration.}}$$

The physical shape of that configuration is otherwise unrestricted.

This is the right setting for deciding whether the Chapter 7 optimum was merely a feature of its chosen coordinates.

8.2 Why bounded physical height is impossible

Fix a positive residue class

$$s \bmod p.$$

Consider the encoded chain

$$\eta(s), \quad \eta(s + p), \quad \eta(s + 2p), \quad \dots$$

or, if a finite initial segment crosses zero, begin sufficiently far into the positive tail.

Let

$$h(c)$$

denote the number of work symbols above Z_0 .

Lemma 8.2 — Tail escape.

For every positive residue class $s \bmod p$,

$$\boxed{h(\eta(s + pk)) \rightarrow \infty \quad (k \rightarrow \infty).}$$

The analogous statement holds on every negative residue class modulo q .

Proof.

Suppose infinitely many residuals on the chain were encoded at height at most H .

There are only finitely many physical configurations of bounded height because:

- Q is finite;
- Δ is finite;
- only finitely many stack words have length at most H .

Injectivity of η would then fail.

Therefore the height must escape every finite bound. $\square$

This elementary argument replaces every earlier homogeneous-tail assumption.

We do not know what the stack looks like.

We know only that it must grow.

8.3 The last occurrence of a height creates a stable suffix

Write

$$c_k = \eta(s + pk)$$

and

$$h_k = h(c_k).$$

Because

$$c_{k+1} \xrightarrow{0} c_k,$$

one ordinary PDA transition can remove at most the visible top work symbol.

Hence

$$h_{k+1} \leq h_k + 1.$$

This one-sided height bound, together with tail escape, creates a canonical cross-section.

Lemma 8.3 — Last-height crossing.

For every sufficiently large integer H , the sequence (h_k) has a last occurrence of height H .

If this last occurrence is at index k , then:

1.

$$h_{k+1} = H + 1;$$

2. the transition

$$c_{k+1} \xrightarrow{0} c_k$$

is a pure one-symbol pop;

3. every later configuration on that residue-class tail has height $> H$;
 4. all later configurations share the same bottom- H work-stack suffix.

Proof.

Tail escape implies that sufficiently large heights are eventually exceeded forever.

Because

$$h_{k+1} \leq h_k + 1,$$

the forward sequence cannot jump upward over a height H .

Thus every sufficiently large H is attained.

Choose its last occurrence.

The next height cannot be at most H , otherwise either H occurs again immediately or a later climb to values above H must pass through H once more.

Therefore

$$h_{k+1} = H + 1.$$

A one-step PDA transition that lowers height from $H + 1$ to H can only delete the visible work symbol. The lower H work symbols are untouched.

After the last height- H occurrence, every later configuration stays above H . Since subsequent transitions rewrite only the visible top region, that bottom- H suffix can never again be exposed or altered.

Hence it stabilizes forever. $\square$

This is the key physical lemma.

No regular tail shape was assumed.

A stable suffix emerged from injectivity and one-top locality alone.

8.4 Positive residue classes synchronize

For each positive residue class

$$s \in \mathbb{Z}_p,$$

choose the last-height- H cross-section.

Let

$$\beta_s \in \Delta^H$$

be the resulting work-stack word above Z_0 .

Lemma 8.4 — Positive synchronization.

For sufficiently large H ,

$$\boxed{\beta_s = \beta_{s+q}}$$

for every residue class $s \bmod p$.

Consequently all p positive residue classes have the same last-height- H stack word, and their control states must be pairwise distinct.

Hence

$$\boxed{|Q| \geq p}$$

on the positive side.

Proof.

Take sufficiently deep positive configurations so that both source and target lie beyond their last-height- H crossings.

Input 1 maps residual class

$$s$$

to

$$s + q \bmod p.$$

A single top rewrite cannot modify the already stabilized bottom- H suffix. Therefore

$$\beta_s = \beta_{s+q}.$$

Because

$$\gcd(p, q) = 1,$$

addition by q is one cycle on

$$\mathbb{Z}_p.$$

Thus all β_s are equal.

At the selected height- H cross-sections, the complete stack word is now identical across all p classes.

Since the residuals are distinct and η is injective, their control states must be distinct. $\square$

The negative side is symmetric.

Lemma 8.5 — Negative synchronization.

At sufficiently large synchronized cross-sections on the negative tail,

$$|Q| \geq q.$$

Therefore every nonempty work alphabet already forces

$$|Q| \geq \max(p, q).$$

This is the first shape-free lower bound matching the Chapter 7 construction.

8.5 Why one work symbol is worse

Suppose now

$$g = 1.$$

Write

$$\Delta = \{A\}.$$

At a fixed height H , every work stack is literally

$$A^H Z_0.$$

The positive and negative synchronized cross-sections can therefore be chosen at the same height with exactly the same physical stack.

Lemma 8.6 — Cross-sign unary separation.

If

$$g = 1,$$

the p synchronized positive states and the q synchronized negative states must all be distinct.

Thus

$$|Q| \geq p + q.$$

Proof.

Choose the synchronized positive and negative cross-sections sufficiently far from zero.

All selected configurations have identical stack

$$A^H Z_0.$$

They represent

$$p + q$$

distinct residuals.

Injectivity can distinguish them only through control state.

Hence all $p + q$ states are distinct. $\square$

The reason the second work symbol is powerful is now exact.

With one work symbol, sign cannot be carried by the tail symbol, so positive and negative residue roles cannot share states.

With two work symbols, the stack itself can separate the ends, allowing state reuse.

8.6 The complete raw frontier

Theorem 8.7 — Exact raw residual-faithful frontier.

For every coprime

$$p, q \geq 1,$$

$$S_{\text{raw}}(g) = \begin{cases} \infty, & g = 0, \\ p + q, & g = 1, \\ \max(p, q), & g \geq 2. \end{cases}$$

Proof.

If

$$g = 0,$$

the only stack is the permanent bottom marker. With finitely many control states, only finitely many physical configurations exist, so no injective encoding of all

$$r \in \mathbb{Z}$$

is possible.

For

$$g \geq 1,$$

positive and negative synchronization give

$$|Q| \geq \max(p, q).$$

For

$$g = 1,$$

cross-sign unary separation strengthens the bound to

$$p + q.$$

The matching upper constructions are exactly the Chapter 7 machines:

- one work symbol with $p + q$ states;
- two work symbols with $\max(p, q)$ states.

Extra work symbols may simply remain unused. $\square$

This theorem changes the status of several earlier results.

The Chapter 7 homogeneous-tail lower bound was not merely an artifact of homogeneous tails.

The exact same resource points survive after all such physical-shape assumptions are removed.

Thus the improved three-symbol machine discovered from the original DPDA lies on the complete raw residual-faithful frontier.

8.7 What this theorem still does not say

The theorem is strong, but its category remains explicit.

It proves exact resource complexity for **raw residual conjugacy**.

It does not prove unrestricted recognition complexity for every DPDA recognizing

$$L_{p,q}.$$

A completely different recognizer may:

- fail to give every residual one stable configuration;
- encode several residuals indirectly through longer physical histories;
- use normalization;
- use a different semantic quotient;
- leave Model R entirely.

Therefore the correct statement is:

the raw residual-faithful frontier is exact.

Not:

the unrestricted DPDA state/stack frontier is solved.

This distinction will become even more important in Chapter 9, where we leave residual-faithful representation entirely and attack a language-level hierarchy problem.

8.8 Acceptance without a right endmarker

In the raw right-endmarker constructions, residual zero is accepted only when the endmarker arrives.

Without a right endmarker, ordinary final-state acceptance requires a state that is final exactly at residual zero.

The synchronized lower-bound states all occur on nonzero residuals, so none may be final.

Corollary 8.8 — Exact no-endmarker raw cost.

Under standard final-state acceptance without a right endmarker,

$$S_{\text{raw}}^{\text{noend}}(g) = \begin{cases} \infty, & g = 0, \\ p + q + 1, & g = 1, \\ \max(p, q) + 1, & g \geq 2. \end{cases}$$

The extra state is an acceptance-wrapper cost.

It does not change the nonzero residual geometry.

This is another place where the two original Chapter 1 variants become conceptually useful: acceptance conventions and semantic core complexity should not be conflated.

8.9 Can the optimum survive bounded pushes?

The raw three-symbol construction may use fixed replacement words longer than two symbols.

Suppose we impose the standard bounded-push convention

$$\max |\alpha| \leq 2.$$

A naive compilation could add many states.

The remarkable fact is that no new state is necessary.

Assume

$$p \geq q.$$

The positive-side writes are already of length at most two because

$$c_s \in \{0, 1\}.$$

Only some negative unary writes are long.

For each negative target phase v , define a family of suffix roles

$$H_{v,j},$$

meaning:

append exactly j additional unary 1's, then enter stable state R_v .

The total number of such transient roles is exactly

$$\boxed{p - q}.$$

Now inspect the stable three-symbol construction.

Negative configurations use top 1 only in states

$$R_0, \dots, R_{q-1}.$$

Therefore the state/top cells

$$(R_i, 1), \quad q \leq i < p,$$

are unused by stable negative semantics.

There are exactly

$$p - q$$

such cells.

The transient suffix roles fit into them perfectly.

Theorem 8.9 — State-preserving bounded-push compiler.

For every coprime p, q , there is a deterministic right-endmarker DPDA whose stable configurations realize the residual graph and satisfy

$$\boxed{|Q| = \max(p, q), \quad |\Gamma| = 3, \quad \max |\alpha| \leq 2.}$$

The only ε -moves are uniquely forced terminating suffix-chain moves.

Proof idea. Split every long unary write into:

1. one input-consuming push of length at most two;
2. a strictly decreasing suffix-completion chain.

The total transient role count is

$$|p - q|.$$

Exactly that many stable state/top cells are unused.

Overlay the transient roles on those cells.

Since the base machine has no consuming transition on the overlaid cells, the forced ε -moves introduce no input/ ε conflict.

Every chain terminates and reconstructs exactly the macro-result of the original raw transition.

The case

$$q > p$$

is symmetric. $\square$

This theorem is a useful warning against counting “temporary states” too early.

A transient normalization role need not require a new control state if an unused state/top cell already exists.

8.10 A stronger semantic-visibility budget

Model R asks only for exact residual faithfulness.

We now intentionally demand more.

Fix

$$B \geq \max(p, q)$$

and write every residual as

$$r = s + Bz, \quad 0 \leq s < B.$$

Define the positive and negative phase periods

$$h_+(B) = \frac{p}{\gcd(B, p)},$$

$$h_-(B) = \frac{q}{\gcd(B, q)},$$

and

$$\boxed{H(B) = h_+(B) + h_-(B).}$$

In the **phase-local architecture**, each remainder s must expose its currently active signed quotient phase through:

- the control state;
- one visible work symbol.

Let

$$C$$

be the number of available work symbols.

For one fixed remainder, one state can provide at most C distinct state/top cells.

Therefore $H(B)$ phase roles require at least

$$\left\lceil \frac{H(B)}{C} \right\rceil$$

states.

Lemma 8.10 — Phase capacity.

For each remainder class s ,

$$\boxed{|Q_s| \geq \left\lceil \frac{H(B)}{C} \right\rceil.}$$

The B remainder classes are semantically distinct in this architecture. Summing gives the exact global frontier.

Theorem 8.11 — Exact normalized phase-local frontier.

For

$$B \geq \max(p, q),$$

$$\boxed{Q_{\min}^{\text{phase}}(B, C) = B \left\lceil \frac{\frac{p}{\gcd(B, p)} + \frac{q}{\gcd(B, q)}}{C} \right\rceil.}$$

For

$$C \geq 2,$$

the matching construction can be ε -free and raw.

For

$$C = 1,$$

the same state bound is attained with one forced bottom-cell normalization at the negative-to-zero crossing.

This theorem should not be confused with Theorem 8.7.

Theorem 8.7 asks for exact residual faithfulness only.

Theorem 8.11 additionally requires a chosen signed phase to be locally visible.

A stronger representation contract can require more states.

This is not a contradiction.

It is the beginning of a general principle:

Resource complexity depends on what physical information the representation is required to expose locally.

8.11 A superlinear exact family

Take

$$p = k^2, \quad q = k^2 - 1, \quad B = k(k + 1).$$

Then

$$H(B) = 2k - 1.$$

Hence

$$Q_{\min}^{\text{phase}}(B, C) = k(k + 1) \left\lceil \frac{2k - 1}{C} \right\rceil.$$

For fixed C ,

$$Q_{\min}^{\text{phase}} = \Theta(k^3).$$

Since

$$\max(p, q) = \Theta(k^2),$$

this is

$$\Theta\left(\max(p, q)^{3/2}\right).$$

The raw residual-faithful machine still needs only

$$\Theta(\max(p, q))$$

states.

The superlinear growth comes entirely from the extra phase-local visibility requirement.

This example makes category-relative resource claims impossible to ignore.

8.12 Direct lossless factorization: counting semantic cells

A still more structured model keeps the residual split into explicit semantic roles.

The stable roles are:

1. $p + q - 1$ boundary/bottom roles;
2. p deep positive roles with visible top 0;
3. q deep negative roles with visible top 1.

These role classes are disjoint.

Theorem 8.12 — Exact semantic-cell count.

Every exact implementation of this direct lossless unary-tail factorization requires

$$2p + 2q - 1$$

reachable stable state/top cells.

The canonical factorization attains this bound exactly.

The theorem counts semantic cells, not merely control states.

This is the correct unit when the representation contract explicitly requires different local roles to remain distinguishable.

8.13 Exact normalization-rule cost under push-at-most-two

Assume

$$p > q.$$

For each deep negative remainder t , the bounded-push normalization required by input 0 has a net stack-height increase

$$D_t = \left\lfloor \frac{p+t}{q} \right\rfloor - 1.$$

A terminating deterministic ε -computation that gains D persistent stack levels needs at least D distinct persistent-push rules.

Why?

If the same state/top rule created two different persistent levels, the intervening positive-drift ε -macro could repeat forever, contradicting termination.

Different source remainders also cannot share the same persistent rule, because after the common rule the lower context remains hidden and deterministic continuation would be identical, while the semantic target remainder is different.

The arithmetic sum is exact:

$$\sum_{t=0}^{q-1} D_t = p - q.$$

Theorem 8.13 — Exact ε -transition count.

In the direct lossless unary-tail factorization with:

- push-at-most-two;
- right endmarker;

the minimum number of forced normalization rules is

$$\boxed{|\delta_\varepsilon|_{\min} = |p - q|}.$$

Theorem 8.14 — Exact total transition count.

Under the same model,

$$\boxed{|\delta|_{\min} = 4(p + q) - 1 + |p - q|}.$$

Reason. The stable semantic graph forces:

$$4(p + q) - 2$$

binary input rules.

The endmarker contributes one accepting rule.

The bounded-push normalization requires exactly

$$|p - q|$$

forced ε -rules.

All three lower bounds are attained simultaneously.

This is the most explicit descriptional-complexity layer so far.

We are no longer counting only states or stack symbols.

We are counting:

- semantic state/top cells;
- binary rules;
- forced normalization rules;
- total transition description size.

8.14 Four resource questions, four different answers

We can now place the resource results in a hierarchy.

Level 1 — Raw residual conjugacy

Only the residual semantics is fixed.

$$S_{\text{raw}}(g) = \infty, \ p + q, \ \max(p, q).$$

Level 2 — Bounded-push normalized realization

The optimal three-symbol/state point survives:

$$|Q| = \max(p, q), \quad |\Gamma| = 3, \quad \max |\alpha| \leq 2.$$

The price is uniquely forced normalization.

Level 3 — Phase-local architecture

A selected signed phase must be visible from one state/top cell.

$$Q_{\min}^{\text{phase}}(B, C) = B \left\lceil \frac{H(B)}{C} \right\rceil.$$

Level 4 — Direct lossless unary-tail factorization

The semantic roles themselves are fixed and counted.

$$C_{\text{sem}, \min} = 2p + 2q - 1,$$

$$|\delta_\varepsilon|_{\min} = |p - q|,$$

$$|\delta|_{\min} = 4(p + q) - 1 + |p - q|.$$

These theorems do not compete.


They answer increasingly structured physical questions.

8.15 What is now genuinely settled

The early DPDA analysis has produced a complete statement:

Within raw residual conjugacy, the improved three-symbol machine is not merely good—it is on the exact frontier.

The discovery sequence is now complete:



three-symbol raw DPDA

↓

complete raw residual-faithful frontier.
--

At this point the natural resource question becomes more radical.

The entire analysis still assumes that a competing machine faithfully represents our residual semantics.

But a DPDA recognizing the same language has no obligation to use our residual coordinate at all.

It may use a completely different stack representation.

It may use deterministic ε -moves.

It may hide the arithmetic in a way our residual-faithful lower bounds do not see.

So we now ask:

Can a stack-alphabet lower bound survive against **arbitrary equivalent deterministic pushdown representations**?

That question takes the monograph outside the original counting family and into a previously stated hierarchy problem.

That is the subject of the next Part.

PART IV

When the Resource Question Leaves Our Representation

Research provenance. The representation-independent hierarchy in this Part is the subject of [3]; its scope is deliberately kept separate from the raw residual-faithful frontiers of the preceding Part. Related descriptonal-complexity work on pushdown alphabets and determinization provides the surrounding context [54, 53, 36].

Chapter 9

From Residual Frontiers to a Representation-Independent Open Problem

The resource story of the previous chapters has reached a natural stopping point.

Within raw residual conjugacy, the frontier is exact:

$$S_{\text{raw}}(g) = \begin{cases} \infty, & g = 0, \\ p + q, & g = 1, \\ \max(p, q), & g \geq 2. \end{cases}$$

The result is strong because the lower bound no longer assumes homogeneous tails, sign-separated coordinates, or a particular stack formula.

But it still assumes something fundamental:

the competing machine is required to realize the residual graph itself.

A DPDA recognizing the same language need not do that.

It may encode the future language in a completely different way. It may normalize internally. It may use deterministic ε -moves between input symbols. It may never expose anything resembling our residual coordinate.

The next step in the research program was therefore not another attempt to optimize the same representation.

It was to remove the representation from the lower bound.

9.1 The methodological jump

The first two resource stages asked:

Given the semantic object we have discovered, how efficiently can it be represented?

The new question is different:

Can the **continuation behavior of the language itself** force a state–stack lower bound, even if an equivalent DPDA chooses an unrelated internal representation?

This distinction is the reason Chapter 8 and the present Part must not be merged.

Chapter 8 is about **representation-faithful optimality**.

This Part is about **arbitrary equivalent recognizers**.

The lower-bound invariant can therefore no longer refer to:

- our residual stack shape;
- our canonical completion word;
- our sign symbol;
- our chosen state phases.

It must be externally visible through the accepted continuation language. Completion lengths provide exactly such an observable.

9.2 The external hierarchy question

Let

$$\mathcal{D}_{m,k}$$

denote the class of languages recognized by deterministic pushdown automata with at most:

$$m$$

control states and at most

$$k$$

pushdown symbols, in the general deterministic model where deterministic ε -moves are allowed.

The inclusion

$$\mathcal{D}_{m,k} \subseteq \mathcal{D}_{m,k+1}$$

is trivial: an extra stack symbol may simply remain unused.

The question is whether the inclusion is strict.

Earlier work had already established:

- a binary stack-alphabet hierarchy for stateless DPDAs;
- an extension to m -state **realtime** DPDAs for every $m \geq 1$.

The corresponding hierarchy for general m -state DPDAs with deterministic ε -moves was left open.

Open Problem 9.1 — Fixed-state stack-alphabet hierarchy

For fixed state budget m , does

$$\mathcal{D}_{m,k} \subsetneq \mathcal{D}_{m,k+1}$$

hold for general deterministic pushdown automata, even when deterministic ε -moves are allowed?

Can the strictness be witnessed over a fixed binary input alphabet?

The new result will answer this for

$$m \geq 2, \quad k \geq 1,$$

under the general single-initial-symbol convention used in this Part.

The stateless case $m = 1$ is kept separate because the earlier stateless formulation uses a slightly different initial-store convention.

That model distinction must remain explicit.

9.3 Why realtime separation does not automatically survive ε -moves

A realtime DPDA processes one input symbol per transition.

Its visible state/top pair is therefore synchronized with the input position.

A general DPDA can insert a deterministic ε -computation between two input symbols.

Such an internal computation can:

- rewrite the stack;
- move information between stack and control;
- expose a deeper symbol;
- normalize a temporary encoding;
- alter the local mode before the next symbol is consumed.

Therefore a lower bound based only on what a realtime machine must expose at input boundaries may collapse when the adversary is allowed internal stack work.

This is precisely why the lower bound must reason about **raw computation positions**, including positions inside forced ε -blocks.

The proof developed in this Part will do that.

9.4 Lineage from the original DPDA

The hierarchy theorem is formally self-contained.

It does not depend on the original 2:1 transition table.

Nevertheless, its mathematical route is directly connected to the earlier program.

The progression was:

original completion-aware DPDA $\rightarrow$ shortest continuation semantics
 $\rightarrow$ exact state/stack tradeoffs $\rightarrow$ question of arbitrary alternative representations.

The key methodological lesson from the earlier machines was:

Stack contents can be probed indirectly by the continuations they require.

The hierarchy witness turns that observation into an adversarial lower-bound tool.

Instead of assuming that a competitor stores a residual, we construct many prefix families whose required unary completion lengths grow at pairwise different rates.

Those rates are visible in the language itself.

The competitor may encode them however it wants.

9.5 The later 2:1 improved DPDA as the immediate precursor

For the motivating ratio (2, 1), Chapter 7 obtained a later raw residual-conjugate construction with:

$$|Q| = 2, \quad \Gamma = \{0, 1, Z_0\},$$

and no ε -normalization on binary residual edges.

This representation is more compact and more literal than the original center/left/right realization for the residual-conjugacy objective.

But that result still does not imply an unrestricted recognition lower bound.

An adversarial DPDA recognizing the same language is not required to use:

- the same two states;
- the same three stack symbols;
- the same residual coordinate;
- any stable residual encoding at all.

This is the conceptual bridge into the hierarchy problem.

9.6 Why the witness development did not end with a direct encoding trick

The route to the final binary hierarchy witness passed through two intermediate prototypes.

Four-letter prototype

The first general prototype used semantic modes

$$(i, t) \in \mathbb{Z}_m \times \mathbb{Z}_k$$

with distinct weights

$$\omega(i, t) = tm + i + 1 \in \{1, \dots, mk\}.$$

Its purpose was to isolate the product phenomenon:

many externally different completion rates
 $\implies$ many recurrent state/top roles.

Three-letter prototype

A mixed-radix control symbol merged two semantic navigation coordinates.

This reduced the input alphabet but still retained a separate delimiter/cleanup phase.

Binary witness

A direct binary code of the three-letter system would have required decoder memory.

Under an exact m -state upper budget, that decoder memory is itself a resource and can ruin the comparison.

The final construction therefore changes the invariant instead of coding the alphabet.

Rather than forcing

$$mk \text{ slopes} + 1 \text{ cleanup role,}$$

it forces directly

$$mk + 1 \text{ distinct completion slopes.}$$

The cleanup role disappears as a separate semantic mode.

The witness becomes binary without needing a phase decoder.

This is an important design lesson in its own right:

When an encoding adds resource overhead, change the observable rather than merely recoding symbols.

9.7 The general DPDA model used in the lower bound

A DPDA is

$$M = (Q, \Sigma, \Gamma, \delta, q_0, Z_0, F),$$

with partial transition function

$$\delta : Q \times \Gamma \times (\Sigma \cup \{\varepsilon\}) \rightarrow Q \times \Gamma^*.$$

We write the stack top on the right.

Thus if

$$\delta(q, X, a) = (p, \beta),$$

then

$$(q, \gamma X, av) \vdash (p, \gamma \beta, v).$$

No transition is available from an empty stack.

Determinism requires that for each state/top pair:

- there is at most one transition for each input symbol;
- if an ε -transition is enabled, then no input-consuming transition is enabled from that same state/top pair.

Acceptance is by:

all input consumed + final state + empty stack.

A realtime DPDA is the special case with no ε -moves.

The upper witness will be realtime.

The lower-bound opponent is allowed to be fully general.

9.8 Modes

A nonempty control-store configuration has the form

$$(q, \eta X),$$

where X is the visible top symbol.

Its **mode** is the pair

$$[q, X].$$

A DPDA has at most

$$|Q||\Gamma|$$

nonempty modes.

The product lower bound will ultimately be a mode-counting theorem.

But merely counting modes is not enough.

We must first prove that different externally visible completion slopes cannot share one recurrent mode.

That is the hard part.

9.9 The lower-bound target

For every

$$m \geq 2, \quad k \geq 1,$$

we will construct a binary language

$$B_{m,k}$$

such that:

1. $B_{m,k}$ is recognized by a realtime DPDA with exactly

$$m$$

states and

$$k + 1$$

stack symbols;

2. every general DPDA recognizing $B_{m,k}$, even with arbitrary deterministic ε -moves, satisfies

$$|Q||\Gamma| \geq mk + 1.$$

The hierarchy corollary will then be immediate.

If a machine used at most m states and at most k stack symbols, then

$$|Q||\Gamma| \leq mk,$$

contradicting the product theorem.

The real work is therefore to build a binary witness exposing exactly

$$mk + 1$$

pairwise incompatible continuation slopes.

That construction is the subject of Chapter 10.

Chapter 10

Completion Slopes and ε -Robust Pumping

Fix

$$m \geq 2, \quad k \geq 1,$$

and write

$$\boxed{\ell = mk + 1.}$$

The witness will have ℓ distinguished live prefix families.

Each family will admit exactly one all-zero completion length after an accumulation block of 1's.

Those lengths will be affine functions with slopes

$$1, 2, \dots, \ell.$$

The upper machine uses exactly:

$$m \text{ states,} \quad k + 1 \text{ stack symbols.}$$

The lower bound then asks whether an arbitrary equivalent DPDA can hide those ℓ externally different rates inside fewer than ℓ state/top modes.

It cannot.

10.1 The binary realtime witness

Let

$$Q = \{q_0, \dots, q_{m-1}\}$$

and

$$\Gamma = \{A, M_0, \dots, M_{k-1}\}.$$

Hence

$$|Q| = m, \quad |\Gamma| = k + 1.$$

Define the permutation

$$\pi(0) = 1, \quad \pi(1) = 0,$$

and

$$\pi(r) = r \quad (2 \leq r < m).$$

Choose ℓ pairwise distinct slope modes

$$x_s = (r_s, X_s), \quad 0 \leq s < \ell,$$

by

$$x_0 = (q_1, A),$$

and, for

$$0 \leq r < m, \quad 0 \leq t < k,$$

$$x_{1+rk+t} = (q_{\pi(r)}, M_t).$$

Thus

$$\{x_1, \dots, x_{\ell-1}\} = Q \times \{M_0, \dots, M_{k-1}\}.$$

The initial pair

$$(q_0, A)$$

is deliberately not a slope mode.

Define the fixed table word

$$\boxed{T = X_{\ell-1}X_{\ell-2} \cdots X_1X_0},$$

with top on the right, so X_0 is initially visible.

10.2 Four rule families

The witness DPDA $U_{m,k}$ is realtime.

Header

$$\delta(q_0, A, 1) = (q_1, T).$$

The first 1 installs the entire finite mode table.

Generic zero-pop on counter symbol

For every state q_i ,

$$\delta(q_i, A, 0) = (q_i, \varepsilon).$$

Selector transitions

For

$$0 \leq s < \ell - 1,$$

$$\delta(r_s, X_s, 0) = (r_{s+1}, \varepsilon).$$

At the last table mode,

$$\delta(r_{\ell-1}, X_{\ell-1}, 0) = (q_0, \varepsilon).$$

Thus a zero can walk down the installed table.

Slope transitions

For every

$$0 \leq s < \ell,$$

$$\delta(r_s, X_s, 1) = (r_s, A^{s+1} X_s).$$

One accumulation 1 inserts exactly $s + 1$ copies of A immediately below the visible mode symbol.

The same mode remains visible.

This is the entire slope mechanism.

10.3 Upper-bound theorem

Proposition 10.1 — Realtime upper construction.

$U_{m,k}$ is a deterministic realtime DPDA with exactly

$$m \text{ states}$$

and

$$k + 1 \text{ stack symbols.}$$

Proof.

There are no ε -moves.

The slope modes are pairwise distinct, so their 1-rules and selector 0-rules do not conflict.

The only apparent overlap is at

$$x_0 = (q_1, A),$$

where the selector rule agrees exactly with the generic A -pop rule.

The header uses

$$(q_0, A, 1),$$

which is not one of the slope modes.

Thus the transition function is deterministic. $\square$

Define

$$B_{m,k} = L(U_{m,k}).$$

The lower bound will not require a closed combinatorial description of the entire language.

One exact slice is enough.

10.4 The exact completion-slope slice

For

$$0 \leq s < \ell,$$

define

$$N_s(n) = \ell - s + (s + 1)n.$$

Let

$$p_s = 10^s.$$

Lemma 10.2 — Exact pure-zero slice.

For every

$$0 \leq s < \ell$$

and every

$$n, r \geq 0,$$

$$\boxed{10^s 1^n 0^r \in B_{m,k} \iff r = N_s(n).}$$

Proof.

After the initial 1, the machine is in mode x_0 with table T .

Each of the next s zeros removes the current table symbol and exposes the next designated slope mode.

Thus after

$$10^s$$

the visible mode is

$$x_s$$

and the remaining table contains exactly

$$\ell - s$$

symbols.

Each subsequent 1 inserts

$$s + 1$$

copies of A below the visible mode symbol.

After n such ones, the total stack height is therefore

$$\ell - s + (s + 1)n.$$

Now read zeros.

The first zero removes the visible mode symbol.

Inserted A 's are then popped one by one, and whenever the next table symbol is exposed the machine is already in its designated state.

Thus every remaining zero removes exactly one stack symbol.

Because acceptance requires empty stack after the input has been exhausted, the all-zero suffix is accepted exactly when its length equals the current stack height. $\square$

Hence the ℓ prefix families

$$p_s 1^n$$

have unique pure-zero completion slopes

$$\boxed{w_s = s + 1.}$$

So the language exposes:

$$\boxed{1, 2, \dots, mk + 1}$$

as pairwise distinct affine continuation rates.

10.5 The smallest new case

Take

$$m = 2, \quad k = 1.$$

Then

$$\ell = 3.$$

The slope modes are

$$x_0 = (q_1, A),$$

$$x_1 = (q_1, M_0),$$

$$x_2 = (q_0, M_0).$$

The exact slice becomes:

$$11^n 0^{3+n},$$

$$101^n 0^{2+2n},$$

$$1001^n 0^{1+3n}.$$

The three slopes are:

1, 2, 3.

The theorem of the next chapter will imply that no DPDA with:

2 states

and

1 stack symbol

can recognize the same language—even if that competitor uses deterministic ε -moves.

10.6 The lower-bound opponent is arbitrary

Let

$$M = (Q, \{0, 1\}, \Gamma, \delta, q_0, Z_0, F)$$

be any general DPDA with

$$L(M) = B_{m,k}.$$

Fix one slope index s and consider the unique computation on

$$p_s 1^\omega.$$

Every finite prefix

$$p_s 1^n$$

is live because it has the accepted continuation

$$0^{N_s(n)}.$$

Therefore the computation cannot become permanently trapped in an ε -divergence before the next input symbol is consumed.

We call **raw positions** all individual computation boundaries, including positions inside forced ε -blocks.

This is essential.

The proof must not look only at input-boundary configurations.

10.7 Raw continuation principle

At any raw configuration

$$(q, \gamma),$$

determinism fixes the future computation once the unread input is fixed.

If an ε -transition is enabled, it is forced and independent of whether the next unread binary symbol is 0 or 1.

Therefore an equal raw control-store configuration can be reused with a changed future suffix.

This trivial-looking fact is the bridge from equality of physical configurations to equality of continuation behavior.

10.8 No raw configuration can repeat on a live slope ray

Lemma 10.3 — No raw configuration repetition.

Along the computation on

$$p_s 1^\omega,$$

the same full control-store configuration cannot occur twice at distinct raw positions.

Proof.

Suppose first that the two occurrences have consumed different numbers

$$n < n'$$

of accumulation ones.

The later prefix has accepted continuation

$$0^{N_s(n')}.$$

If the earlier and later raw control-store configurations are identical, determinism gives the same future computation from both.

Thus M would also accept

$$p_s 1^n 0^{N_s(n')}.$$

But the exact slice says that after

$$p_s 1^n$$

the unique all-zero completion length is

$$N_s(n),$$

and

$$N_s(n') \neq N_s(n).$$

Contradiction.

If the two repeated raw configurations occur after the same amount of input has been consumed, then the segment between them contains only ε -moves.

Determinism would repeat the same ε -cycle forever, preventing the next live input symbol from being read.

Again contradiction. $\square$

This lemma absorbs arbitrary deterministic internal normalization into the argument.

Repeated physical configurations are impossible because the language exposes a unique continuation length.

10.9 Raw stack height must escape

Lemma 10.4 — Raw-height escape.

Along every slope ray,

raw stack height tends to infinity.

More explicitly, for every H , after some raw position the stack height is always greater than H .

Proof.

There are only finitely many control-store configurations of stack height at most H .

By Lemma 10.3, none can occur twice.

Therefore only finitely many raw positions have height at most H . $\square$

The adversary may use any representation it likes.

But the infinitely many distinguishable live prefixes force its physical store to escape every bounded region.

10.10 Suffix-minimum positions

A raw position t is a **step** if its stack height is a minimum over the entire future suffix of the computation:

$$h(t) \leq h(t') \quad \text{for every } t' \geq t.$$

Lemma 10.5 — Infinitely many steps.

Every slope ray contains infinitely many step positions.

Proof.

The height sequence tends to infinity.

For every starting position, the future tail has a minimum.

Choose a position where that minimum occurs.

As the starting point moves arbitrarily far to the right, these minima cannot all lie in a fixed finite set. $\square$

Step positions are the correct pumping locations because the computation never descends below their initial stack height.

Therefore the lower stack context remains hidden.

10.11 The ε -robust step pump

Lemma 10.6 — Step-pump lemma.

On every slope ray there exist:

- a mode

$$[q, X];$$

- an integer

$$c \geq 1;$$

- a nonempty word

$$\rho \in \Gamma^+;$$

such that for every lower context

$$\eta \in \Gamma^*,$$

$$(q, \eta X) \xRightarrow{1^c} (q, \eta \rho X),$$

where the macro-computation:

1. consumes exactly c input symbols 1;
2. may contain forced ε -moves;
3. never goes below the initial height

$$|\eta X|.$$

The mode $[q, X]$ is called an **up-mode**.

Proof.

There are infinitely many step positions but only

$$|Q||\Gamma|$$

possible modes.

Choose two step positions

$$t_1 < t_2$$

with the same mode

$$[q, X].$$

Write the earlier stack as

$$\eta X.$$

Because t_1 is a suffix minimum, the computation between t_1 and t_2 never goes below height

$$|\eta X|.$$

Therefore the lower prefix η is never exposed.

The later stack has form

$$\eta\rho X.$$

If

$$\rho = \varepsilon,$$

the two positions would have the same full configuration, contradicting Lemma 10.3.

Hence

$$\rho \in \Gamma^+.$$

Let c be the number of input 1's consumed between the two steps.

Because the lower context is never inspected, the same deterministic transition sequence works for every replacement lower context.

Thus the pump is context-independent.

Finally,

$$c \neq 0.$$

If it consumed no input, it would be a height-increasing ε -macro from the same mode back to itself. Determinism would repeat it forever and prevent the next live symbol from being consumed.

Therefore

$$c \geq 1.$$

□

The pump iterates:

$$(q, \eta X) \xrightarrow{1^{jc}} (q, \eta \rho^j X) \quad (j \geq 0).$$

This is the technical point at which the lower bound becomes genuinely ε -robust.

The proof did not forbid internal normalization.

It found a recurrent growth mechanism inside it.

10.12 What remains to prove

Every slope ray now has an up-mode.

To obtain the product lower bound, we need one final structural fact:

one up-mode cannot serve two different completion slopes.

If that holds, then the $\ell = mk + 1$ slopes force ℓ different modes.

The proof requires one more determinism lemma about unary completion.

That is where the next chapter begins.

Chapter 11

One Mode, One Slope

The upper witness exposes

$$\ell = mk + 1$$

different affine completion slopes.

Chapter 10 showed that every arbitrary general DPDA recognizing the witness must, along each live slope ray, develop a context-independent height-increasing up-mode.

We now prove the separation theorem:

$\text{one up-mode} \implies \text{at most one completion slope.}$

Once this is established, the product bound and the hierarchy follow immediately.

11.1 Unary completion uniqueness

The witness slice uses all-zero continuations.

Determinism makes those continuations rigid.

Lemma 11.1 — Unary completion uniqueness.

Fix any raw control-store configuration

$$(p, \gamma)$$

of a deterministic pushdown automaton in the present model.

There is at most one integer

$$r \geq 0$$

such that input

0^r

leads from (p, γ) to acceptance.

Proof.

Suppose

$$r < r'$$

and both

0^r

and

$0^{r'}$

were accepted.

The two deterministic computations are identical until the shorter input has consumed all r zeros, including every forced ε -move enabled independently of the presence of further input.

The accepting run on 0^r eventually reaches:

- empty unread input;
- a final state;
- empty stack.

The run on $0^{r'}$ reaches the same empty-stack configuration while

$$r' - r > 0$$

zeros remain unread.

No transition is available from the empty stack.

Thus the longer word cannot be accepted.

Contradiction. $\square$

This is a very small lemma, but it is exactly what converts a recurrent physical mode into a constraint on completion length.

11.2 Slope separation

Lemma 11.2 — Slope separation.

Let

$$s \neq t.$$

The slope rays

$$p_s 1^\omega \quad \text{and} \quad p_t 1^\omega$$

cannot have the same up-mode in any DPDA recognizing

$$B_{m,k}.$$

Proof.

Assume both rays share an up-mode

$$[q, X].$$

Take the context-independent pump

$$(q, \eta X) \xRightarrow{1^c} (q, \eta \rho X),$$

with

$$c \geq 1, \quad \rho \in \Gamma^+.$$

Suppose the mode occurs on the s -ray after n_s accumulation ones with lower context γ_s , and on the t -ray after n_t accumulation ones with lower context γ_t .

After j pump iterations the two rays reach:

$$(q, \gamma_s \rho^j X)$$

and

$$(q, \gamma_t \rho^j X)$$

after respectively

$$n_s + jc$$

and

$$n_t + jc$$

accumulation ones.

The exact witness slice provides accepted zero completions of lengths

$$N_s(n_s + jc)$$

and

$$N_t(n_t + jc).$$

Now run those two zero computations in parallel.

As long as the common upper segment

$$\rho^j X$$

has not been completely removed, the two computations have:

- the same state;
- the same visible active stack suffix;
- the same deterministic local behavior.

The lower contexts γ_s and γ_t remain hidden.

Thus the common upper segment is removed after the same number

$$e_j$$

of consumed zeros.

Immediately after the final active symbol of that segment disappears, the two raw configurations have the form

$$(p_j, \gamma_s) \quad \text{and} \quad (p_j, \gamma_t)$$

for the same state p_j .

Because Q is finite, there is an infinite set of indices J on which

$$p_j = p$$

is constant.

For $j \in J$, the remaining accepted zero lengths are

$$R_s(j) = N_s(n_s + jc) - e_j$$

from the fixed raw configuration

$$(p, \gamma_s),$$

and

$$R_t(j) = N_t(n_t + jc) - e_j$$

from

$$(p, \gamma_t).$$

Unary completion uniqueness forces each of these two quantities to be constant on J .

Therefore

$$N_s(n_s + jc) - N_t(n_t + jc)$$

must be constant on J .

But

$$N_s(n) = \ell - s + (s + 1)n,$$

so

$$\begin{aligned} N_s(n_s + jc) - N_t(n_t + jc) \\ = (s - t)cj + O(1). \end{aligned}$$

Because

$$s \neq t$$

and

$$c \geq 1,$$

this expression is unbounded along the infinite set J .

Contradiction. $\square$

This lemma is the structural center of the hierarchy proof.

It says that a recurrent state/top mode has an externally measurable asymptotic identity:

its continuation slope.

Two different slopes cannot be hidden in the same up-mode.

11.3 The product lower bound

There are

$$\ell = mk + 1$$

slope rays.

Every slope ray has an up-mode.

Distinct slopes require distinct up-modes.

A DPDA with state set Q and stack alphabet Γ has at most

$$|Q||\Gamma|$$

nonempty modes.

Everything is now ready.

Theorem 11.3 — Product lower bound.

Let

$$m \geq 2, \quad k \geq 1.$$

If a general deterministic pushdown automaton M , with arbitrary deterministic ε -moves allowed, recognizes

$$B_{m,k},$$

then

$$\boxed{|Q_M||\Gamma_M| \geq mk + 1.}$$

Proof.

There are

$$mk + 1$$

slope rays.

By the step-pump lemma, each has an up-mode.

By slope separation, no two slope rays share an up-mode.

Therefore M has at least

$$mk + 1$$

distinct nonempty state/top-symbol modes.

But there are at most

$$|Q_M||\Gamma_M|$$

such modes.

Hence

$$|Q_M||\Gamma_M| \geq mk + 1.$$

□

This is stronger than the adjacent hierarchy statement.

It excludes an entire hyperbolic resource region:

$$\boxed{|Q||\Gamma| \leq mk.}$$

11.4 The hierarchy theorem

Corollary 11.4 — Binary adjacent stack-alphabet hierarchy.

For every

$$m \geq 2, \quad k \geq 1,$$

$$\boxed{\mathcal{D}_{m,k} \subsetneq \mathcal{D}_{m,k+1}.}$$

The strictness is witnessed over the binary alphabet

$$\{0, 1\}.$$

Proof.

The realtime witness $U_{m,k}$ uses exactly:

$$m$$

states and

$$k + 1$$

stack symbols.

Therefore

$$B_{m,k} \in \mathcal{D}_{m,k+1}.$$

If

$$B_{m,k} \in \mathcal{D}_{m,k},$$

some recognizing DPDA would satisfy

$$|Q| \leq m, \quad |\Gamma| \leq k.$$

Hence

$$|Q||\Gamma| \leq mk,$$

contradicting Theorem 11.3.

The non-strict inclusion follows by allowing an unused extra stack symbol.

□

Thus deterministic ε -moves do not collapse the fixed-state stack-alphabet hierarchy in the stated general model for

$$m \geq 2.$$

11.5 What exactly has been resolved

The bibliographic claim must remain narrow.

The theorem answers the fixed-state stack-alphabet hierarchy question left open after the earlier realtime result, under:

- the general deterministic model of this Part;
- the single-initial-symbol convention;
- fixed state budget

$$m \geq 2;$$

- binary witness alphabet.

The stateless case is not silently folded into the same theorem because its earlier formulation uses a different initial-store convention.

Thus the monograph should say:

the general-model fixed-state hierarchy is resolved here for $m \geq 2$

and separately acknowledge the known stateless result.

11.6 Why this is not a realtime lower bound in disguise

The upper witness is realtime.

The lower-bound opponent is not.

The arbitrary competitor may use deterministic ε -moves throughout.

The proof survives because it deliberately reasons at raw positions.

The chain is:

exact continuation slice

⇓

no raw configuration repetition

⇓

raw-height escape

⇓

suffix-minimum step pump

$$\Downarrow$$

slope separation

$$\Downarrow$$

$|Q||\Gamma| \geq mk + 1.$

The first box is a property of the explicit binary witness.

Every later box concerns an arbitrary equivalent DPDA.

That is why the theorem is representation-independent in a way the earlier residual-frontier results were not.

11.7 Why the product, rather than states or symbols separately, appears

The proof never forces a slope into a unique state alone.

It forces it into a unique **mode**:

$$[q, X].$$

Therefore the natural capacity is:

$|Q||\Gamma|.$

This is exactly the resource unit the proof can count.

A machine may trade states against stack symbols.

The product theorem allows that trade.

What it forbids is packing more than one distinct completion slope into one recurrent up-mode.

Thus the proof yields a quantitative state–stack tradeoff rather than only a yes/no alphabet separation.

11.8 What the theorem does not prove

Several stronger statements remain outside the theorem.

Not a full Pareto frontier

The bound

$$|Q||\Gamma| \geq mk + 1$$

does not prove that every point above the hyperbola is attainable.

It excludes one region.

It does not determine the complete state/stack Pareto frontier of the witness.

Not a bounded-push hierarchy theorem

The upper witness installs a fixed table word of length

$$\ell = mk + 1$$

in one header transition.

The general PDA model permits replacement by an arbitrary fixed stack word.

If one imposed a push-length bound, compiling the header would change the resource accounting.

That is a separate problem.

Not an unrestricted lower bound for $L_{p,q}$

The witness was engineered so that its continuation slopes are maximally separable.

The theorem does not automatically transfer to the original rational counting languages.

Whether related continuation observables can give equally strong unrestricted recognition lower bounds for specific $L_{p,q}$ remains a separate question.

Not a claim that mode arguments are themselves new

State/top modes and state–stack–symbol tradeoffs have a pre-existing literature.

The contribution here is the particular binary witness, the ε -robust slope argument, the product theorem, and the resulting resolution in the stated model.

11.9 The methodological meaning for the monograph

This Part began with a problem created by our own DPDA analysis.

The original machine and its successors showed that the same semantic quantity can be redistributed between:

- control state;
- stack symbol;
- stack height;
- normalization.

That made representation-scoped lower bounds suspicious.

The hierarchy proof answers that suspicion by abandoning the representation.

It asks only what continuations the language forces.

This produces a new lower-bound pattern:

externally visible continuation rates $\rightarrow$ recurrent physical modes
 $\rightarrow$ representation-independent resource count.

This method is conceptually distinct from the raw residual frontier.

Both belong in the monograph because they answer different questions.

11.10 Where the story goes next

We have now completed the first major resource arc.

Starting from the 2023 2:1 machine, we obtained:

completion semantics

$\downarrow$

geodesic geometry

$\downarrow$

better coordinates

$\downarrow$

exact residual-faithful frontier

$\downarrow$

representation-independent hierarchy theorem.

The next Part changes direction.

Instead of asking how many resources are needed to realize a one-dimensional residual geometry, we ask what happens when the counting constraint itself becomes higher-dimensional.

The binary rectangle will become an optimal Parikh fiber.

The absence of equal-length trade directions will become a kernel-rank condition.

And the first nontrivial equal-length trade will create topology.

We now leave resource complexity and return to completion geometry.

PART V

Completion Geometry Beyond the Binary Line

Research provenance. The higher-dimensional minimum-completion fibers, local action complexes, and cubical quotient geometry are developed in [5]. Affine-semigroup factorization, Markov/Graver moves, squarefree-divisor complexes, cubical state/rewrite complexes, and vector-partition functions are established neighboring theories [34, 38, 28, 17, 12, 62, 39, 63]; the new object here is their restriction to minimum-completion prefixes and the equal-length residual quotient.

Chapter 12

When Minimum Demand Becomes a Fiber

The first half of the monograph was built around one scalar residual:

$$\rho_{p,q}(w) = qN_1(w) - pN_0(w).$$

That scalar residual had a unique minimum completion Parikh vector

$$\mu(r) = (x_r, y_r).$$

As a consequence, every shortest-completion interval was a rectangle. The natural higher-dimensional question is:

What survives if the counting constraint is no longer one scalar equation?

The answer preserves the completion viewpoint but changes the geometry fundamentally.

The minimum demand need no longer be one vector.

It can become a finite **fiber of equally short Parikh vectors**.

That is where higher-dimensional completion geometry begins.

12.1 Matrix counting languages

Let

$$\Sigma = \{a_1, \dots, a_d\}$$

and let

$$\Psi : \Sigma^* \rightarrow \mathbb{N}^d$$

be the Parikh map.
Fix an integer matrix

$$A \in \mathbb{Z}^{k \times d}.$$

Define the vector residual

$$\rho(w) = A\Psi(w)$$

and the matrix counting language

$$L_A = \{w : \rho(w) = 0\}.$$

Residuals lie in the lattice

$$A\mathbb{Z}^d.$$

The binary rational counting family is a rank-one special case.
For

$$L_{p,q},$$

take

$$A = \begin{bmatrix} -p & q \end{bmatrix}.$$

Then

$$A\Psi(w) = qN_1(w) - pN_0(w).$$

Thus the higher-dimensional theory genuinely contains the earlier one.

12.2 Completion totality

A matrix counting language need not allow every residual to be completed by appending letters.

We therefore isolate a convenient hypothesis.

Definition 12.1 — Completion totality.

We call A **completion-total** if there exists

$$h \in \mathbb{N}_{>0}^d$$

such that

$$Ah = 0.$$

The vector h is a strictly positive balanced Parikh vector.
It plays the role of the primitive balanced packet

$$(q, p)$$

from the binary theory, except that it need not generate the entire integer kernel.

Lemma 12.2 — Every lattice residual is completable.

Assume completion totality.

For every

$$r \in A\mathbb{Z}^d,$$

there exists

$$x \in \mathbb{N}^d$$

such that

$$\boxed{Ax = -r.}$$

Proof.

Write

$$r = Az$$

with

$$z \in \mathbb{Z}^d.$$

Because every coordinate of h is positive, for sufficiently large t ,

$$th - z \in \mathbb{N}^d.$$

Since

$$Ah = 0,$$

$$A(th - z) = -Az = -r.$$

Thus $th - z$ is a nonnegative completion vector. $\square$

This hypothesis is deliberately strong enough to make completion semantics total on the entire residual lattice.

It lets us generalize the quotient theory cleanly.

12.3 Residuals still classify right quotients

Proposition 12.3 — Matrix residual quotient classification.

Under completion totality, for all prefixes

$$u, v \in \Sigma^*,$$

$$u^{-1}L_A = v^{-1}L_A \iff \rho(u) = \rho(v).$$

Proof.

A continuation z completes u exactly when

$$A\Psi(z) = -\rho(u).$$

Therefore equal residuals give equal quotients.

Conversely, suppose

$$\rho(u) \neq \rho(v).$$

By Lemma 12.2, choose

$$x \in \mathbb{N}^d$$

with

$$Ax = -\rho(u).$$

Choose a word z with

$$\Psi(z) = x.$$

Then

$$uz \in L_A.$$

But if

$$vz \in L_A,$$

we would have

$$Ax = -\rho(v)$$

as well, forcing

$$\rho(u) = \rho(v).$$

Contradiction. $\square$

The semantic lesson from Chapter 4 survives unchanged:

the residual is an exact coordinate for future right behavior.

What changes is the geometry of shortest return to residual zero.

12.4 Minimum completion length

For residual

$$r \in A\mathbb{Z}^d,$$

define

$$\ell_A(r) = \min\{\mathbf{1}^\top x : x \in \mathbb{N}^d, Ax = -r\}.$$

This is the shortest number of appended letters required to cancel r .

In the binary case,

$$\ell_A(r) = \Delta_{p,q}(r).$$

So far the generalization looks formal.

The decisive change appears when we ask which Parikh vectors attain this minimum.

12.5 The optimal Parikh fiber

Definition 12.4 — Optimal Parikh fiber.

Define

$$O_A(r) = \{x \in \mathbb{N}^d : Ax = -r, \mathbf{1}^\top x = \ell_A(r)\}.$$

In the binary $p:q$ theory,

$$O_A(r)$$

was always a singleton:

$$O_A(r) = \{\mu(r)\}.$$

In higher dimensions, several distinct nonnegative Parikh vectors can complete the same residual at the same minimum length.

Therefore the semantic object

$$\mu(r)$$

must be replaced by the entire set

$$O_A(r).$$

This is the first major qualitative change.

12.6 Equal-length trades

Introduce the augmented matrix

$$C = \begin{bmatrix} A \\ \mathbf{1}^\top \end{bmatrix}.$$

Define the integer equal-length kernel

$$K_A = \ker_{\mathbb{Z}} C = \left\{ g \in \mathbb{Z}^d : Ag = 0, \mathbf{1}^\top g = 0 \right\}.$$

A vector

$$g \in K_A$$

changes one completion Parikh vector into another while preserving:

1. residual contribution;
2. total completion length.

Thus K_A is the lattice of **equal-length trades**.

This is the higher-dimensional replacement for the binary kernel direction.

But there is a crucial difference.

For the binary minimum demand, subtracting one nonzero kernel packet always shortened the completion, so two distinct minimum Parikh vectors could not coexist.

In higher dimensions, a nonzero vector can have positive and negative coordinates while still satisfying

$$\mathbf{1}^\top g = 0.$$

Such a trade changes the distribution of letters without changing total length.

That is how an optimal fiber acquires multiple points.

12.7 Affine-semigroup viewpoint

The columns

$$C_i = \begin{bmatrix} A_i \\ 1 \end{bmatrix}$$

generate an affine semigroup.

For residual r , define the augmented optimum element

$$b_r = (-r, \ell_A(r)).$$

Then

$$O_A(r)$$

is exactly the factorization fiber of b_r into the generators C_i .

Every such factorization has the same last coordinate:

$$\ell_A(r).$$

This places the theory next to established affine-semigroup and factorization geometry.

The monograph uses that literature as background.

The new object of interest remains specifically the geometry of **minimum-completion schedules**.

12.8 Continuous optimum geometry is not the same as integer multiplicity

Define the continuous optimum polytope

$$P_r^A = \{x \in \mathbb{R}_{\geq 0}^d : Ax = -r, \mathbf{1}^\top x = \ell_A(r)\}.$$

Let

$$S(r) = \{i : \exists x \in P_r^A \text{ with } x_i > 0\}$$

be its active support.

Proposition 12.5 — Dimension of the continuous optimum polytope.

If

$$P_r^A \neq \emptyset,$$

then

$$\dim P_r^A = |S(r)| - \text{rank}_{\mathbb{R}} \begin{bmatrix} A^{S(r)} \\ \mathbf{1}^\top \end{bmatrix}.$$

Proof.

For every active coordinate choose a feasible point positive in that coordinate.

The average of these points is positive in every active coordinate and lies in the relative interior of the nonnegative orthant restricted to $S(r)$.

Thus no additional coordinate inequality is active in the relative interior.

The dimension is therefore the number of active variables minus the rank of the affine equality system on those variables. $\square$

The proposition is useful mainly because it prevents a common confusion.

A positive-dimensional real optimum polytope does not automatically imply many integer optimum points.

Conversely, discrete equal-length trade structure can be significant even when a simple continuous dimension count does not capture its combinatorics.

The integer fiber

$$O_A(r)$$

is the primary shortest-completion object.

12.9 The shortest-prefix down-set

In the binary rectangle, a shortest-completion prefix was described by how many required zeros and ones had already been consumed.

For a higher-dimensional optimum vector

$$x \in O_A(r),$$

every coordinatewise subvector

$$0 \leq u \leq x$$

is the Parikh vector of some prefix of an ordering of that optimum completion. Define

$$B_x = \{u \in \mathbb{N}^d : 0 \leq u \leq x\}.$$

Now take the union over all optimum vectors.

Definition 12.6 — Shortest-prefix down-set.

$$\mathcal{C}_A(r) = \downarrow O_A(r) = \bigcup_{x \in O_A(r)} B_x.$$

This is the higher-dimensional replacement for one rectangle.
 If the optimum is unique,

$$O_A(r) = \{x\},$$

then

$$\mathcal{C}_A(r) = B_x,$$

a literal coordinate box.

If the optimum fiber contains several vectors, the shortest-prefix region becomes a union of overlapping boxes.

The next theorem says this down-set is not merely a convenient overapproximation.

It is exact.

12.10 Exact geodesic-prefix criterion

Theorem 12.7 — Exact geodesic-prefix criterion.

For every

$$u \in \mathbb{N}^d,$$

$$u \in \mathcal{C}_A(r) \iff \ell_A(r + Au) = \ell_A(r) - |u|_1.$$

Equivalently,

$$\mathcal{C}_A(r)$$

is exactly the set of Parikh vectors that occur as prefixes of shortest completions of residual r .

Proof.

Suppose

$$u \in \mathcal{C}_A(r).$$

Then

$$u \leq x$$

for some

$$x \in O_A(r).$$

The vector

$$x - u$$

completes residual

$$r + Au$$

and has length

$$\ell_A(r) - |u|_1.$$

If a shorter completion existed, appending it after u would produce a completion of r shorter than $\ell_A(r)$.

Therefore

$$\ell_A(r + Au) = \ell_A(r) - |u|_1.$$

Conversely, assume the displayed equality.

Let y be an optimal completion of

$$r + Au.$$

Then

$$A(u + y) = -r,$$

and

$$|u + y|_1 = |u|_1 + \ell_A(r + Au) = \ell_A(r).$$

Thus

$$u + y \in O_A(r),$$

so

$$u \leq u + y.$$

Hence

$$u \in \mathcal{C}_A(r).$$

□

This theorem is the higher-dimensional form of minimum-demand descent.

The difference is that there may be several optimal endpoints x , so the geodesic-prefix set is a union of coordinate boxes rather than one box.

12.11 First actions and compatible action sets

Let e_i denote the i -th unit vector.

Corollary 12.8 — First actions.

Letter type i can begin a shortest completion exactly when

$$\ell_A(r + Ae_i) = \ell_A(r) - 1.$$

More generally, for

$$J \subseteq [d],$$

there exists an optimal completion using every letter type in J at least once if and only if

$$\ell_A(r + A\mathbf{1}_J) = \ell_A(r) - |J|.$$

Thus shortest-action compatibility can be tested intrinsically from the distance function.

This is the higher-dimensional analogue of the binary rule:

$$x_r > 0 \iff 0 \text{ is geodesic,}$$

$$y_r > 0 \iff 1 \text{ is geodesic.}$$

12.12 The conceptual transition

The completion object has now changed from

$$\boxed{\text{one minimum vector}}$$

to

$$\boxed{\text{a finite optimum fiber.}}$$

Accordingly, the shortest-prefix geometry has changed from

$$\boxed{\text{one rectangle}}$$

to

$$\boxed{\text{a union of boxes.}}$$

But the residual quotient still identifies prefixes that reach the same semantic state.

If two different shortest-prefix Parikh vectors reach the same residual at the same geodesic depth, they must differ by an equal-length trade in

$$K_A.$$

That observation leads to the next object:

$$\boxed{|\mathcal{C}_A(r)|/K_A.}$$

The next chapter will show that this quotient is exactly the filled shortest-path interval.

Chapter 13

The Filled Shortest-Prefix Quotient

Chapter 12 produced two objects:

$$\mathcal{C}_A(r) = \downarrow O_A(r),$$

the exact Parikh down-set of shortest prefixes, and

$$K_A = \ker_{\mathbb{Z}} \begin{bmatrix} A \\ \mathbf{1}^\top \end{bmatrix},$$

the lattice of equal-length residual-preserving trades.

We now combine them.

The result is the higher-dimensional completion object that replaces the binary shortest rectangle.

13.1 Why vertices alone are not enough

A shortest prefix is not only a point.

If several distinct next letters can be taken in any order while remaining geodesic, the geometry contains a commuting square, cube, or higher-dimensional cell.

The binary rectangle already had this property.

A pair of actions 0 and 1 produced an elementary commuting square.

In d letters, a compatible set

$$J \subseteq [d]$$

should produce a $|J|$ -dimensional cube.

Thus we should fill the shortest-prefix down-set cubically.

13.2 Cubical realization of a down-set

Let

$$D \subseteq \mathbb{N}^d$$

be a finite coordinatewise down-set.

Define

$$|D|$$

to be the cubical complex whose vertices are the elements of D .

For

$$J \subseteq [d],$$

an elementary J -oriented cube based at u is present exactly when

$$\boxed{u + \mathbf{1}_J \in D.}$$

Because D is a down-set, all intermediate vertices of that cube are automatically present.

Applied to

$$D = \mathcal{C}_A(r),$$

every cube records a set of distinct letter actions that can be executed in any order while remaining on shortest paths.

This is the filled schedule geometry before residual identifications.

13.3 Equal-length identifications

Two shortest-prefix Parikh vectors can be different while representing the same reached residual.

Let

$$u, v \in \mathcal{C}_A(r).$$

They reach the same residual exactly when

$$A(u - v) = 0.$$

But because both are shortest prefixes reaching the same residual, they also have the same geodesic depth.

Thus

$$\mathbf{1}^\top(u - v) = 0.$$

Therefore

$$u - v \in K_A.$$

This gives the correct equivalence relation.

Two equally oriented cells are identified when their base vertices differ by an element of

$$K_A.$$

Definition 13.1 — Filled shortest-geodesic quotient.

Define

$$\mathfrak{I}_A(r) = |\mathcal{C}_A(r)|/K_A.$$

This quotient keeps:

- every shortest prefix;
- every compatible commuting cell;

but identifies prefixes that reach the same residual at the same shortest depth.

13.4 The quotient theorem

Theorem 13.2 — Shortest-prefix quotient theorem.

The 0-skeleton of

$$\mathfrak{I}_A(r)$$

is label-preserving isomorphic to the complete directed shortest-path interval from residual r to residual 0.

Under this identification, elementary cubes are exactly compatible commuting sets of geodesic one-letter actions.

Proof.

Define

$$\pi_r : \mathcal{C}_A(r) \rightarrow A\mathbb{Z}^d$$

by

$$\pi_r(u) = r + Au.$$

By Theorem 12.7, every $u \in \mathcal{C}_A(r)$ is a geodesic prefix, so every image vertex lies on a shortest path from r to 0.

Conversely, every shortest path prefix has some Parikh vector u , and Theorem 12.7 places that vector in $\mathcal{C}_A(r)$.

Now suppose

$$\pi_r(u) = \pi_r(v).$$

Then

$$A(u - v) = 0.$$

Because u and v reach the same residual and both are geodesic,

$$\ell_A(r) - |u|_1 = \ell_A(\pi_r(u)) = \ell_A(\pi_r(v)) = \ell_A(r) - |v|_1.$$

Hence

$$|u|_1 = |v|_1,$$

so

$$\mathbf{1}^\top(u - v) = 0.$$

Therefore

$$u - v \in K_A.$$

Conversely, if

$$u - v \in K_A,$$

then

$$A(u - v) = 0,$$

so the two vertices have the same residual label.

Thus the vertex fibers of π_r are exactly the K_A -classes.

Finally, a J -cube based at u exists exactly when

$$u + \mathbf{1}_J \in \mathcal{C}_A(r).$$

By the geodesic-prefix criterion, every subcollection of the actions in J remains shortest, and coordinate addition commutes.

Passing to the equal-length kernel quotient therefore gives exactly the filled commuting shortest-path structure. $\square$

This theorem is the central bridge from integer completion arithmetic to topology.

13.5 Recovery of the binary rectangle

For the binary scalar system

$$A = \begin{bmatrix} -p & q \end{bmatrix},$$

the augmented matrix is

$$\begin{bmatrix} -p & q \\ 1 & 1 \end{bmatrix}.$$

Its determinant is

$$-(p + q) \neq 0.$$

Therefore

$$\boxed{K_A = \{0\}}.$$

There are no nontrivial equal-length trades.

The optimum fiber is a singleton:

$$O_A(r) = \{\mu(r)\}.$$

Hence

$$\mathcal{C}_A(r) = [0, x_r] \times [0, y_r]$$

and no quotient identifications occur.

Thus

$$\boxed{\mathfrak{I}_A(r) = [0, x_r] \times [0, y_r]}.$$

The binary rectangle is therefore the **rank-zero member** of the general theory.

This is an important reinterpretation.

The rectangle did not disappear in higher dimensions.

It became the simplest possible fiber-and-quotient case.

13.6 When is the quotient literally a box?

Theorem 13.3 — Literal-box criterion.

$$\boxed{|O_A(r)| = 1 \iff \mathfrak{I}_A(r) \text{ is a literal coordinate box.}}$$

Here “literal coordinate box” means that the quotient contains no nontrivial kernel identifications among the vertices and cells of the prequotient box.

Proof.

Suppose

$$O_A(r) = \{x\}.$$

Then

$$\mathcal{C}_A(r) = B_x.$$

Assume distinct

$$u, v \leq x$$

were identified.

Then

$$u - v \in K_A.$$

Consider

$$x - u + v.$$

It is nonnegative.

Because

$$A(u - v) = 0,$$

it has the same residual contribution as x .

Because

$$\mathbf{1}^\top(u - v) = 0,$$

it has the same length as x .

It is different from x .

Therefore it would be a second element of

$$O_A(r),$$

contradicting uniqueness.

Hence no nontrivial kernel collision occurs inside the box.

Conversely, suppose

$$x \neq y$$

both belong to

$$O_A(r).$$

Then x and y are distinct maximal vertices of $\mathcal{C}_A(r)$, both reach residual zero, and

$$x - y \in K_A.$$

Thus the quotient identifies distinct shortest schedules.

It cannot be the original literal coordinate box with all coordinate vertices distinct. $\square$

This theorem identifies the first source of global complexity:

nonunique optimum $\iff$ nontrivial equal-length identification.
--

13.7 Unique optimum versus unique shortest word

The phrase “unique optimum” refers to the Parikh vector.

It does not mean there is only one shortest word.

Even when

$$O_A(r) = \{x\},$$

the letters represented by x can be ordered in many ways.

The literal box already records those schedules.

What uniqueness eliminates is a stronger ambiguity:

two different minimum Parikh endpoints.

Thus there are two levels of shortest-path multiplicity.

Ordering multiplicity

One Parikh vector, many permutations.

This produces a box.

Factorization multiplicity

Several distinct minimum Parikh vectors.

This creates kernel identifications among different boxes.

The second phenomenon is new beyond the binary line.

13.8 Local compatibility and global identification are different

The cubical complex

$$|\mathcal{C}_A(r)|$$

records which letter actions can coexist in one optimum and commute as schedule choices.

The quotient

$$\mathfrak{I}_A(r)$$

adds a different phenomenon:

globally different shortest prefixes can reach the same residual.

A later chapter will show that these two notions have different topology.

In rank one:

- every connected component of the local action-compatibility complex is contractible;
- the full global quotient can nevertheless be a circle.

Thus local coexistence and global schedule identification must not be conflated.

We are now ready to study that distinction.

13.9 The rank of equal-length trade

The integer lattice

$$K_A$$

measures the number of independent directions in which one optimal Parikh vector can trade letters for others without changing:

- residual;
- completion length.

Three regimes will become especially important.

Rank zero

$$K_A = \{0\}.$$

No equal-length trade exists.

Every optimum is unique.

The shortest interval is a literal box.

This is the binary $p:q$ regime.

Rank one

One primitive equal-length trade direction exists.

The optimum fiber becomes an arithmetic interval.

The global quotient may acquire its first cycle.

Rank at least two

Several independent trade directions exist.

Even the local action complex can develop holes.

The next chapters will make this hierarchy exact.

Chapter 14

Local and Global Topology

Chapter 13 produced the filled shortest-prefix quotient

$$\mathfrak{I}_A(r) = |\mathcal{C}_A(r)|/K_A.$$

Two different kinds of structure are present inside this object.

The first is **local**:

Which distinct letter types can coexist inside one shortest completion?

The second is **global**:

After all shortest prefixes are filled cubically and equal-length residual-preserving schedules are identified, what is the topology of the whole shortest interval?

These questions need not have the same answer.

The first genuinely new phenomenon occurs already when the equal-length kernel has rank one.

Locally, nothing topologically complicated happens.

Globally, a circle can appear.

That separation is one of the principal structural results of the higher-dimensional branch.

14.1 The local action complex

Recall the optimal Parikh fiber

$$O_A(r) = \{x \in \mathbb{N}^d : Ax = -r, |x|_1 = \ell_A(r)\}.$$

Define the simplicial complex

$$\mathcal{H}_A(r) = \{J \subseteq [d] : \exists x \in O_A(r) \text{ with } x_j > 0 \text{ for every } j \in J\}.$$

A face

$$J \in \mathcal{H}_A(r)$$

means that some shortest completion contains at least one copy of every action in J .

Equivalently, by Chapter 12,

$$J \in \mathcal{H}_A(r) \iff \ell_A(r + A\mathbf{1}_J) = \ell_A(r) - |J|.$$

Thus $\mathcal{H}_A(r)$ is the local compatibility complex of geodesic letter types. It deliberately forgets:

- how many times an action occurs;
- the order in which actions occur;
- which shortest prefixes are globally identified.

It records only coexistence.

14.2 Squarefree-divisor interpretation

Let

$$C_i = \begin{bmatrix} A_i \\ 1 \end{bmatrix}$$

and let

$$S_C = \langle C_1, \dots, C_d \rangle$$

be the affine semigroup generated by the augmented columns.

For residual r , recall

$$b_r = (-r, \ell_A(r)).$$

Proposition 14.1 — Squarefree divisor identification.

$$\mathcal{H}_A(r) = \left\{ J \subseteq [d] : b_r - \sum_{j \in J} C_j \in S_C \right\}.$$

Proof.

The semigroup condition means that there exists

$$y \in \mathbb{N}^d$$

with

$$Cy = b_r - \sum_{j \in J} C_j.$$

Then

$$x = y + \mathbf{1}_J$$

satisfies

$$Cx = b_r.$$

Hence

$$x \in O_A(r)$$

and every coordinate indexed by J is positive.

The converse is obtained by subtracting $\mathbf{1}_J$ from any optimum vector supporting J . $\square$

Squarefree divisor complexes are established objects in affine-semigroup topology [17].

Their role here is specific:

they are exactly the local shortest-action complexes
of the completion problem.

The topology of $\mathcal{H}_A(r)$ therefore measures local compatibility among geodesic actions.

14.3 Rank zero: one simplex

If

$$K_A = \{0\},$$

the optimum fiber is a singleton

$$O_A(r) = \{x\}.$$

Then

$$\mathcal{H}_A(r)$$

is simply the simplex on

$$\text{supp}(x).$$

Thus rank zero has no nontrivial local topology.

This includes every binary $p:q$ completion problem.

The local action set may have one or two symbols, but it never develops a hole.

14.4 Rank one: the optimum fiber is an arithmetic interval

Assume now

$$K_A = \mathbb{Z}g$$

for a primitive nonzero vector

$$g \in \mathbb{Z}^d.$$

Because the optimum fiber is finite, it has the form

$$\boxed{O_A(r) = \{x_t = x_0 + tg : 0 \leq t \leq N\}.$$

For each letter type i , define

$$T_i = \{t \in [0, N] \cap \mathbb{Z} : x_t(i) > 0\}.$$

Since

$$x_t(i) = x_0(i) + tg_i$$

is affine in t , the set T_i is an integer interval, possibly empty.

A set of actions

$$J$$

is locally compatible exactly when

$$\bigcap_{j \in J} T_j \neq \emptyset.$$

Thus the local action complex is the nerve of a finite family of intervals.

Theorem 14.2 — Rank-one local action topology.

If

$$\text{rank}_{\mathbb{Z}} K_A \leq 1,$$

then every connected component of

$$\mathcal{H}_A(r)$$

is contractible.

In particular,

$$\boxed{\widetilde{H}_j(\mathcal{H}_A(r); \mathbb{Z}) = 0 \quad (j \geq 1).}$$

Proof.

The rank-zero case was just described.

In rank one, $\mathcal{H}_A(r)$ is the nerve of the interval family

$$\{T_i\}.$$

Every connected component of the union of finitely many intervals is itself an interval.

By the nerve theorem [47], each connected component of the nerve is homotopy equivalent to the corresponding connected component of the union.

Hence every connected component is contractible. $\square$

This theorem says that one equal-length trade direction is not enough to create a local homology hole.

Yet it is enough to create a global cycle.

Before proving that, we show that rank two changes the local picture immediately.

14.5 Rank two already permits a local hole

Consider the scalar weight vector

$$\boxed{\omega = (-11, -2, 1, 10)}$$

and residual

$$r = 7.$$

The minimum completion length is 5, and the complete optimum fiber is

$$\boxed{(0, 4, 1, 0), \quad (1, 0, 4, 0), \quad (1, 3, 0, 1).}$$

The maximal local action faces are

$$\{2, 3\}, \quad \{1, 3\}, \quad \{1, 2, 4\}.$$

Therefore

$$f_0 = 4, \quad f_1 = 5, \quad f_2 = 1.$$

There is one filled triangle, but one independent graph cycle remains unfilled.
Hence

$$H_1(\mathcal{H}_\omega(7); \mathbb{Z}) \cong \mathbb{Z}.$$

The augmented equal-length kernel has rank two.

This example proves that the rank-one local theorem is sharp.

The first nontrivial local H_1 can appear as soon as two independent equal-length trades are available.




Figure 14.1: The local action complex for $\omega = (-11, -2, 1, 10)$ at residual 7. Its maximal faces are $\{2, 3\}$, $\{1, 3\}$, and the filled triangle $\{1, 2, 4\}$. The cycle $1 - 2 - 3 - 1$ remains unfilled, giving $H_1 \cong \mathbb{Z}$.

The figure should make the distinction visual:

- one triangular face is filled;
- another cycle is not.

This is not a global schedule-identification phenomenon.

The hole is already present in the local coexistence complex.

14.6 Rank-one global geometry

We now return to

$$K_A = \mathbb{Z}g.$$

Write the primitive generator as

$$g = g^+ - g^-,$$

where

$$g^+, g^- \in \mathbb{N}^d$$

have disjoint supports.

Let

$$O_A(r) = \{x_t = x_0 + tg : 0 \leq t \leq N\}.$$

Define the prequotient shortest-prefix complex

$$X = |\mathcal{C}_A(r)| = \left| \bigcup_{t=0}^N B_{x_t} \right|.$$

Because $\mathcal{C}_A(r)$ is a finite coordinate down-set,

$$\boxed{X \simeq *.}$$

So the topology of the final quotient cannot come from a hole already present in the prequotient region.

It must come from the kernel identification.

14.7 Overlap complexes

For

$$m \geq 1,$$

define

$$P_m = \{\sigma \subseteq X : \sigma + mg \subseteq X\},$$

where translation preserves cube orientation.

Lemma 14.3 — Exact m -step overlap.

For

$$1 \leq m \leq N,$$

translation by

$$-mg^-$$

identifies P_m with the anchored cubical down-set

$$\bigcup_{t=0}^{N-m} B_{x_t - mg^-}.$$

Hence P_m is nonempty and contractible.

For

$$m > N,$$

$$P_m = \emptyset.$$

Interpretation. The complete optimum fiber has parameter span

$$0, \dots, N.$$

Two shortest-prefix cells can overlap after an mg translation exactly while that translation stays inside the span.

Thus the global quotient glues only finitely many neighboring translates.

The overlap itself remains topologically simple.

14.8 Periodic saturation

Define the periodic saturation

$$\tilde{X} = \bigcup_{k \in \mathbb{Z}} (X + kg).$$

Let

$$X_k = X + kg.$$

Translation by g acts freely on $\tilde{X}$.

Moreover,

$$\tilde{X}/\mathbb{Z} \cong \mathfrak{I}_A(r).$$

The quotient by the integer translation action performs exactly the rank-one equal-length kernel identification.

To understand $\tilde{X}$, we study the cover

$$\{X_k\}_{k \in \mathbb{Z}}.$$

Lemma 14.4 — Good-cover intersections.

Let

$$k_0 < \cdots < k_s.$$

Then

$$X_{k_0} \cap \cdots \cap X_{k_s} \neq \emptyset \iff k_s - k_0 \leq N.$$

Every nonempty finite intersection is contractible.

Proof idea. For any fixed cell, the set of translate indices k for which the cell lies in X_k is an integer interval.

Thus membership in the two extreme translates implies membership in all intermediate translates.

After translating by $-k_0 g$, the intersection reduces to the $(k_s - k_0)$ -step overlap from Lemma 14.3.

That overlap is contractible exactly when the index gap is at most N . $\square$

The nerve of this cover has vertices $\mathbb{Z}$.

A finite set of vertices spans a simplex exactly when its diameter is at most N .

For

$$N \geq 1,$$

this is the nerve of the interval cover

$$\{[k, k + N]\}_{k \in \mathbb{Z}}$$

of the real line.

Hence the nerve is contractible.

Theorem 14.5 — Contractible rank-one saturation.

If

$$N \geq 1,$$

then

$$\boxed{\tilde{X} \simeq *}.$$

The periodic saturation is therefore a contractible universal-cover candidate.

14.9 The global circle

Theorem 14.6 — Rank-one global homotopy dichotomy.

If

$$\text{rank}_{\mathbb{Z}} K_A = 1,$$

then

$$\boxed{\mathfrak{I}_A(r) \simeq \begin{cases} *, & |O_A(r)| = 1, \\ S^1, & |O_A(r)| > 1. \end{cases}}$$

Proof.

If the optimum is unique, Chapter 13 gives a literal coordinate box, hence a contractible complex.

Now suppose

$$|O_A(r)| > 1.$$

Then

$$N \geq 1.$$

By Theorem 14.5,

$$\tilde{X}$$

is contractible.

The action of

$$\mathbb{Z}$$

by translation through g is free, properly discontinuous, and cellular.

Therefore

$$\tilde{X} \rightarrow \tilde{X}/\mathbb{Z} \cong \mathfrak{I}_A(r)$$

is a universal covering with deck group

$$\mathbb{Z}.$$

Hence

$$\pi_1(\mathfrak{I}_A(r)) \cong \mathbb{Z}$$

and all higher homotopy groups vanish.

Thus

$$\mathfrak{I}_A(r)$$

is a $K(\mathbb{Z}, 1)$, hence homotopy equivalent to

$$S^1.$$

□

This gives the first exact topological transition in the completion theory:

rank one + nonunique optimum $\implies$ global circle.
--

14.10 Local triviality, global nontriviality

Combining Theorems 14.2 and 14.6 gives the sharp distinction:

rank-one local action complex: every component contractible rank-one global shortest-prefix quotient: possibly S^1 .

The circle is not caused by a local incompatibility hole.

It is caused by globally identifying different shortest schedules through an equal-length trade.

This difference will matter later when the monograph returns to normalization schedules.

Local feasibility and global identification are fundamentally different layers of structure.

14.11 What the topology means operationally

A point of

$$\mathcal{C}_A(r)$$

records a shortest-prefix Parikh vector.

A path through the cubical complex records gradual execution of a shortest schedule.

Translation by g changes the Parikh bookkeeping without changing:

- residual;
- geodesic depth.

The quotient therefore says that two apparently different shortest schedule histories have reached the same semantic point.

When one primitive trade direction exists, those identifications wrap the contractible periodic saturation around one generator.

That is the geometric origin of the circle.

The result is not merely “some complex has homology.”

It is a statement about the topology created by schedule equivalence.

The next chapter turns kernel rank into an exact alphabet-size hierarchy for scalar counting languages.

Chapter 15

Scalar Rank Hierarchy and Exact Three-Letter Geometry

The previous chapter expressed the new topology in terms of

$$\text{rank}_{\mathbb{Z}} K_A.$$

For a general matrix A , this rank is an independent structural parameter.

For a scalar counting language, it has a remarkably simple form.

This produces a sharp hierarchy by alphabet size.

15.1 Scalar counting systems

Let

$$\omega = (\omega_1, \dots, \omega_d) \in \mathbb{Z}^d$$

have at least two distinct entries.

Treat

$$A = \omega$$

as a $1 \times d$ matrix.

The equal-length kernel is

$$K_\omega = \ker_{\mathbb{Z}} \begin{bmatrix} \omega \\ \mathbf{1}^\top \end{bmatrix}.$$

The two rows are linearly independent over $\mathbb{R}$.

Therefore:

$$\boxed{\text{rank}_{\mathbb{Z}} K_\omega = d - 2.}$$

This single formula turns the topology of Chapter 14 into an alphabet-size threshold.

15.2 The $d = 2, 3, 4+$ hierarchy

Two letters

If

$$d = 2,$$

then

$$\text{rank } K_\omega = 0.$$

Therefore:

- every optimum Parikh vector is unique;
- every shortest interval is a literal coordinate box;
- no equal-length schedule identification occurs.

This is exactly the binary theory of Chapters 2–5.

Three letters

If

$$d = 3,$$

then

$$\text{rank } K_\omega = 1.$$

Therefore:

- the optimum fiber is an arithmetic interval;
- every connected component of the local action complex is contractible;
- the global quotient is either contractible or homotopy equivalent to

$$S^1.$$

Thus three letters are the first scalar alphabet size capable of creating a global shortest-schedule cycle.

Four or more letters

If

$$d \geq 4,$$

then

$$\text{rank } K_\omega \geq 2.$$

Higher-rank equal-length trade geometry becomes possible.

The four-letter example of Chapter 14 already gives

$$H_1(\mathcal{H}; \mathbb{Z}) \cong \mathbb{Z}.$$

Thus four letters are the first scalar alphabet size at which a local shortest-action hole can occur.

The hierarchy is therefore:

d	shortest-completion geometry
2	rank 0: literal boxes
3	rank 1: first global S^1
≥ 4	higher rank: local holes possible

This gives a structural interpretation of alphabet size that is invisible in the original binary problem.

15.3 Three-letter normal form

Assume

$$\omega_1 < \omega_2 < \omega_3.$$

Define the adjacent gaps

$$a = \omega_3 - \omega_2, \quad b = \omega_2 - \omega_1.$$

Let

$$h = \gcd(a, b),$$

and set

$$\alpha = \frac{a}{h}, \quad \beta = \frac{b}{h}, \quad m = \alpha + \beta.$$

A primitive equal-length kernel generator is

$$g = (\alpha, -m, \beta).$$

Indeed,

$$\alpha - m + \beta = 0,$$

and

$$\alpha\omega_1 - m\omega_2 + \beta\omega_3 = 0.$$

Thus one equal-length trade replaces m copies of the middle letter by:

- α copies of the first;
- β copies of the third.

Or vice versa.

This is the exact three-letter analogue of a primitive commuting schedule trade.

15.4 Parameterizing the optimum fiber

Write the complete optimum fiber as

$$O_\omega(r) = \{x(t) = y + tg : 0 \leq t \leq N\},$$

where

$$y = (y_1, y_2, y_3).$$

Thus

$$x(t) = (y_1 + \alpha t, y_2 - mt, y_3 + \beta t).$$

Let

$$z = y_2 - mN, \quad 0 \leq z < m.$$

Define

$$P_t = (y_1 + \alpha t + 1)(y_3 + \beta t + 1).$$

Slicing the prequotient down-set by the middle coordinate gives an exact vertex count before quotienting:

$$|\mathcal{C}_\omega(r)| = m \sum_{t=0}^{N-1} P_t + (z+1)P_N.$$

The kernel quotient then collapses contiguous segments along the g -direction.

Because every rank-one orbit intersects the down-set in an integer interval, the quotient can be counted by:

$$\text{vertices} - \text{adjacent orbit pairs}.$$

15.5 Exact three-letter quotient vertex count

Proposition 15.1 — Exact rank-one quotient vertex count.

Let

$$\Omega(y, N)$$

be the number of vertex orbits in

$$\mathcal{C}_\omega(r)/\mathbb{Z}g.$$

If

$$N = 0,$$

then

$$\boxed{\Omega(y, 0) = (y_1 + 1)(y_2 + 1)(y_3 + 1).}$$

If

$$N \geq 1,$$

then

$$\boxed{\Omega(y, N) = mP_{N-1} + (z + 1)(P_N - P_{N-1}).}$$

Proof idea. Every kernel orbit meets the down-set in a finite integer interval.

An orbit containing s vertices contributes exactly $s - 1$ adjacent pairs

$$u, u + g.$$

Therefore:

$$\# \text{orbits} = \# \text{vertices} - \# \text{adjacent pairs}.$$

The same middle-coordinate slicing gives the adjacent-pair count, and subtraction yields the formula. $\square$

The result is an exact finite formula for the number of semantic shortest-prefix states after equal-length schedule identification.

15.6 Exact cubical f -vector

The same reduction works separately for every cube orientation.

For

$$J \subseteq \{1, 2, 3\},$$

define

$$T_J = \{t : x_i(t) \geq 1 \text{ for every } i \in J\}.$$

Because each coordinate is affine in t ,

$$T_J$$

is an integer interval, possibly empty.

If nonempty, write

$$T_J = [L_J, U_J] \cap \mathbb{Z},$$

set

$$N_J = U_J - L_J,$$

and define

$$y^{(J)} = x(L_J) - \mathbf{1}_J.$$

Proposition 15.2 — Exact rank-one cubical f -vector.

The number of quotient cubes oriented by J is

$$f_J(r) = \begin{cases} 0, & T_J = \emptyset, \\ \Omega(y^{(J)}, N_J), & T_J \neq \emptyset. \end{cases}$$

Hence

$$f_j(r) = \sum_{|J|=j} f_J(r), \quad 0 \leq j \leq 3.$$

Thus the entire finite cubical shortest-interval geometry is computable from one rank-one orbit formula.

15.7 Worked example

Take

$$\omega = (-5, -3, 1), \quad r = 9.$$

Then:

$$a = 4, \quad b = 2, \quad h = 2,$$

so

$$\alpha = 2, \quad \beta = 1, \quad m = 3.$$

The primitive equal-length trade is

$$g = (2, -3, 1).$$

The complete optimum fiber is

$$O_\omega(9) = \{(0, 3, 0), (2, 0, 1)\}.$$

Thus the minimum completion has two distinct Parikh realizations of the same length:

$$(0, 3, 0)$$

and

$$(2, 0, 1).$$

The quotient cubical f -vector is

$$(f_0, f_1, f_2, f_3) = (8, 10, 2, 0).$$

Therefore the filled quotient contains:

- 8 vertices;
- 10 labeled edges;
- 2 elementary commuting squares;
- no 3-cube.

Yet by Theorem 14.6,

$$\mathfrak{I}_\omega(9) \simeq S^1.$$

This is an ideal example for the monograph because the topology is nontrivial while the finite complex remains small enough to draw completely.




Figure 15.1: The complete rank-one quotient for $\omega = (-5, -3, 1)$ and $r = 9$. The down-set below $(2, 0, 1)$ contributes two filled squares; the path below $(0, 3, 0)$ shares the origin, and its far endpoint is identified with $(2, 0, 1)$ by the equal-length kernel translation. The quotient has $(f_0, f_1, f_2, f_3) = (8, 10, 2, 0)$ and is homotopy equivalent to S^1 .

15.8 What changed relative to the binary rectangle?

In two letters, the minimum completion is one Parikh vector.

Different shortest words differ only by ordering.

In three letters, there can be two different minimum Parikh vectors:

$$x_0 \quad \text{and} \quad x_0 + g.$$

The prequotient geometry contains boxes beneath both.

The quotient then identifies equal-depth prefixes differing by g .

That is what creates the global circle.

Thus the transition

$$d = 2 \rightarrow d = 3$$

does not merely add another coordinate.

It adds the first equal-length trade direction.

And the transition

$$d = 3 \rightarrow d = 4$$

does not merely add one more trade.

It permits two independent trades, enough for local action topology itself to develop a hole.

15.9 Why this hierarchy belongs in the same monograph

The scalar alphabet-size hierarchy may look distant from the original 2:1 DPDA.

Its origin is nevertheless direct.

The original machine led to:

shortest completion $\rightarrow$ minimum Parikh demand $\rightarrow$ rectangle.

The higher-dimensional question was:

Was the rectangle an accident of the binary alphabet?

The answer is no.

It was the rank-zero case of a general equal-length trade geometry.

The binary theory therefore sits at the first level of a topological hierarchy.

The next chapter studies a different direction: not how geometry changes with alphabet size, but how it behaves far out along a residual ray.

Chapter 16

Eventual Geodesic Arithmetic

The previous chapters studied the exact shortest-completion geometry of one fixed residual.

We now vary the residual systematically.

Fix two lattice residuals

$$r, b \in A\mathbb{Z}^d$$

and consider the ray

$$r_n = r + nb, \quad n \in \mathbb{N}.$$

The minimum completion length

$$\ell_A(r_n)$$

need not be exactly linear in n .

Integer constraints can create periodic corrections.

The main result of this chapter is that those corrections eventually stabilize into a finite defect profile.

The same finite mechanism controls both:

- numerical shortest-completion distance;
- local shortest-action geometry.

16.1 The continuous directional slope

Consider the linear program

$$\lambda_b = \min\{\mathbf{1}^\top x : Ax = -b, x \in \mathbb{R}_{\geq 0}^d\}.$$

Completion totality guarantees feasibility.

Because the data are rational,

$$\lambda_b \in \mathbb{Q}_{\geq 0}.$$

The number

$$\lambda_b$$

is the continuous asymptotic completion cost per unit step in direction b .
If integer effects disappeared completely, one would expect

$$\ell_A(r + nb) \approx n\lambda_b + \text{constant}.$$

The eventual defect theory measures the exact failure of this naive linear law.

16.2 Dual certificate and reduced costs

Choose a rational dual optimum

$$y^*$$

for

$$\max\{-b^\top y : A^\top y \leq \mathbf{1}\}.$$

Thus

$$-b^\top y^* = \lambda_b.$$

Choose an integer

$$L > 0$$

clearing all denominators needed below.

Define the reduced costs

$$\delta_i = L(1 - A_i^\top y^*).$$

Dual feasibility gives

$$\delta_i \in \mathbb{N}.$$

If

$$x_n$$

is any optimal integer completion of r_n , then

$$Ax_n = -r_n.$$

Therefore

$$\begin{aligned} L(\ell_A(r_n) + (y^*)^\top r - n\lambda_b) &= L\left(\mathbf{1}^\top x_n - (y^*)^\top Ax_n\right) \\ &= \sum_i \delta_i x_n(i). \end{aligned}$$

This motivates the exact integer defect.

Definition 16.1 — Scaled integrality defect.

$$\Delta_r(n) = L(\ell_A(r_n) + (y^*)^\top r - n\lambda_b).$$

Then

$$\Delta_r(n) \in \mathbb{N}.$$

The defect is not an asymptotic error term.

It is an exact nonnegative integer reduced-cost total.

That arithmetic discreteness is what makes eventual stabilization possible.

16.3 Zero-defect shifts

Define

$$\mathcal{S}_b = \left\{ t \geq 1 : \exists z \in \mathbb{N}^d, Az = -tb, |z|_1 = t\lambda_b \right\}.$$

An element

$$t \in \mathcal{S}_b$$

is a direction length for which the continuous optimum is attained integrally with zero defect.

Lemma 16.2 — Zero-defect semigroup.

$$\mathcal{S}_b$$

is nonempty and closed under addition.

Proof sketch. A rational primal optimum for direction b becomes integral after clearing denominators, giving at least one zero-defect multiple.

If

$$t_1, t_2 \in \mathcal{S}_b$$

with witnesses z_1, z_2 , then

$$z_1 + z_2$$

satisfies

$$A(z_1 + z_2) = -(t_1 + t_2)b$$

and

$$|z_1 + z_2|_1 = (t_1 + t_2)\lambda_b.$$

Thus

$$t_1 + t_2 \in \mathcal{S}_b.$$

□

Define the canonical arithmetic period

$$\tau_b = \gcd \mathcal{S}_b.$$

A zero-defect shift by any $t \in \mathcal{S}_b$ can be appended to an existing completion without increasing the defect.

Thus $\mathcal{S}_b$ transports good integer behavior forward along the ray.

16.4 Stabilized defect profile

For each residue class

$$s \in \mathbb{Z}/\tau_b\mathbb{Z},$$

define

$$M_s = \min\{\Delta_r(n) : n \geq 0, n \equiv s \pmod{\tau_b}\}.$$

Theorem 16.3 — Stabilized defect profile.

For every residue class s , there exists

$$N_s$$

such that

$$\boxed{\Delta_r(n) = M_s}$$

whenever

$$n \geq N_s$$

and

$$n \equiv s \pmod{\tau_b}.$$

Consequently, eventually,

$$\boxed{\ell_A(r_n) = -(y^*)^\top r + n\lambda_b + \frac{M_{n \bmod \tau_b}}{L}}.$$

Proof.

Choose finitely many elements

$$t_1, \dots, t_m \in \mathcal{S}_b$$

whose gcd is

$$\tau_b.$$

The additive numerical semigroup they generate contains every sufficiently large multiple of τ_b .

For each residue class s , choose an index n_s attaining the minimum defect M_s .

If

$$n \equiv s \pmod{\tau_b}$$

is sufficiently large, then

$$n - n_s$$

is a nonnegative combination of the t_i .

Append the corresponding zero-defect completions to an optimum at n_s .

This gives a feasible completion at n with the same defect M_s .

By definition of M_s , the defect cannot be smaller.

Hence it is exactly M_s . $\square$

The formula has a useful interpretation:

$$\boxed{\text{eventual distance} = \text{rational linear drift} + \text{finite periodic integer defect.}}$$

The eventual arithmetic is therefore controlled by a finite cyclic table.

16.5 Why τ_b already contains the slope denominator

Write

$$\lambda_b = \frac{a}{q}$$

in lowest terms.

For every

$$t \in \mathcal{S}_b,$$

the zero-defect completion has integer length

$$t\lambda_b.$$

Hence

$$q \mid t.$$

Therefore

$$\boxed{q \mid \tau_b.}$$

The canonical zero-defect period is thus automatically compatible with the denominator of the asymptotic slope.

But the **minimal eventual affine period** can be smaller than τ_b .

The remaining reduction comes from symmetry of the stabilized defect table.

16.6 Fundamental defect period

View

$$M = (M_0, \dots, M_{\tau_b-1})$$

as a cyclic word.

Define its fundamental circular period

$$\boxed{d_M = \min\{p \geq 1 : M_{s+p} = M_s \text{ for every } s\}.}$$

Then

$$d_M \mid \tau_b.$$

We can now determine the exact eventual distance period.

Theorem 16.4 — Exact minimal eventual distance period.

Let

$$\lambda_b = \frac{a}{q}$$

in lowest terms.

The smallest positive integer p for which there exists an integer D_p satisfying

$$\ell_A(r_{n+p}) = \ell_A(r_n) + D_p$$

for all sufficiently large n is

$$p_{\min}(r, b) = \text{lcm}(q, d_M).$$

Moreover,

$$p_{\min}(r, b) \mid \tau_b,$$

and

$$D_p = p\lambda_b.$$

Proof.

Suppose p is an eventual affine period.

Comparing the linear terms in Theorem 16.3 forces

$$D_p = p\lambda_b.$$

Since D_p is integral,

$$q \mid p.$$

The periodic correction term must also be invariant under shift by p , so

$$d_M \mid p.$$

Therefore

$$\text{lcm}(q, d_M) \mid p.$$

Conversely, let

$$p_0 = \text{lcm}(q, d_M).$$

Then

$$p_0\lambda_b$$

is integral and the stabilized defect table is invariant under shift by p_0 .

The formula of Theorem 16.3 gives

$$\ell_A(r_{n+p_0}) = \ell_A(r_n) + p_0 \lambda_b$$

for all sufficiently large n .

Thus p_0 is the minimal eventual affine period. $\square$

This theorem separates two arithmetic sources of periodicity:

1. the denominator of the continuous slope;
2. the intrinsic symmetry period of the integer defect profile.

The actual distance period is their least common multiple.

16.7 Eventual local action geometry

The same mechanism controls local shortest-action compatibility.

For

$$J \subseteq [d],$$

let

$$c_J = A\mathbf{1}_J.$$

Recall:

$$J \in \mathcal{H}_A(r_n) \iff \ell_A(r_n + c_J) = \ell_A(r_n) - |J|.$$

Apply Theorem 16.3 not only to the ray

$$r_n$$

but also to every finitely many shifted rays

$$r_n + c_J.$$

The same:

- dual optimum;
- zero-defect semigroup;
- canonical period τ_b ;

can be used.

Only the stabilized defect table changes with J .

Therefore the truth value of

$$J \in \mathcal{H}_A(r_n)$$

is eventually determined by

$$n \bmod \tau_b.$$

Let

$$\Theta_s$$

be the stabilized local action complex on residue class s .

Define

$$d_{\mathcal{H}} = \min\{p \geq 1 : \Theta_{s+p} = \Theta_s \text{ for every } s\}.$$

Corollary 16.5 — Exact minimal local-geometry period.

The smallest eventual period of the local action complex is

$$p_{\min}^{\mathcal{H}}(r, b) = d_{\mathcal{H}}.$$

Moreover,

$$d_{\mathcal{H}} \mid \tau_b.$$

The first-action set is eventually periodic for the same reason.

Thus one finite defect mechanism governs both:

$$\text{numerical shortest distance}$$

and

$$\text{local shortest-action geometry.}$$

16.8 What is and is not being claimed

Eventual periodicity in semilinear and parametric integer settings is not new in itself.

The point here is more specific.

The completion program produces one exact object:

$$\Delta_r(n).$$

From it we derive:

- eventual affine distance;

- exact minimal distance period;
- eventual local action complex;
- exact minimal local-geometry period.


Thus the numerical and geometric observables are not proved periodic by unrelated arguments.

They descend from the same minimum-completion defect profile.

This is the structural contribution relevant to the monograph.

16.9 The higher-dimensional branch in one picture

The path from the original binary machine to the present theorem is now complete:



The binary rectangle was therefore not a small isolated trick.

It was the rank-zero member of a much larger completion geometry.

16.10 Returning to the original machine from the opposite direction

The next Part does not continue outward.

It zooms back in.

In Chapter 1 we deliberately marked one local transition of the original 2:1 DPDA:

$$(L, 0S) \xrightarrow{0} (L, S).$$

So far that transition has been interpreted as:

- a stack cancellation;
- one unit of residual correction;
- one edge of a shortest-completion ladder.

But it contains another structure that the higher-dimensional completion theory did not use.

When the nonnegative 2:1 computations are drawn as lattice paths, the cancellation occurs at a precise arithmetic family of landing heights.

That observation leads to an entirely different branch:

pushdown cancellation $\rightarrow$ correction signatures $\rightarrow$ intrinsic defect $\rightarrow$
Catalan geometry.

We now return to the 2023 machine and follow that transition.

PART VI

One Cancellation Transition Becomes a Geometry

Research provenance. The Fuss–Catalan defect geometry and its connection back to the original cancellation transition are developed in [6]. Generalized Dyck paths, Fuss–Catalan enumeration, and related Catalan statistics are classical or independently studied background [18, 16, 41, 59]; singularity analysis is used in its standard analytic-combinatorics role [32].

Chapter 17

Zooming Back into the Original DPDA

The previous Part generalized shortest completion outward.

We now move in the opposite direction.

Return to the original 2:1 DPDA of Chapter 1 and isolate one local transition:

$$(L, 0S) \xrightarrow{0} (L, S).$$

At the machine level, this is a cancellation.

At the residual level, it decreases

$$\rho = N_1 - 2N_0$$

by 2.

At the geodesic level, Chapter 3 showed that it consumes one required zero along a shortest-completion path.

There is, however, another structure hidden in the same event.

If the input word is drawn as a nonnegative 2-Dyck path, the cancellation occurs at a precise arithmetic set of landing heights.

That observation starts a second geometry.

17.1 The machine facts we reuse, not re-prove

The cancellation analysis needs only two facts already established in Chapter 1:

Reachable-store normal form

$$\mathcal{S}_C = 0^* \cup 1^* \cup \{10\},$$

$$\mathcal{S}_L = 0^*,$$

$$\mathcal{S}_R = 1^* \cup \{0\}.$$

State-potential invariant

For every consumed prefix,

$$\rho(w) = D(S) + \Phi(q),$$

with

$$D(S) = 2N_0(S) - N_1(S),$$

$$\Phi(C) = 0, \quad \Phi(L) = 3, \quad \Phi(R) = -3.$$

We do not duplicate those proofs here.

Their consequence for positive odd residuals is immediate.

Corollary 17.1 — Positive odd residuals.

Every reachable core configuration with positive odd residual at least 3 is uniquely of the form

$$(L, 0^k)$$

for some

$$k \geq 0,$$

and its residual is

$$2k + 3.$$

The left state therefore has a very precise path-height meaning.

17.2 The nonnegative path sector

Encode input symbols as steps:

$$1 \longleftrightarrow U \quad \text{of height } +1,$$

$$0 \longleftrightarrow D \quad \text{of height } -2.$$

For a prefix w , the path height is

$$h(w) = N_1(w) - 2N_0(w) = \rho(w).$$

A word whose prefix heights are all nonnegative and whose final height is 0 is a 2-Dyck excursion.

The cancellation statistic will first be discovered inside this nonnegative sector.

This restriction matters.

The later combinatorial theory will be defined directly on generalized Dyck paths and will no longer depend on the remaining DPDA transition table.

17.3 Exact landing-height correspondence

Immediately before a distinguished cancellation,

$$(L, 0^k) \xrightarrow{0} (L, 0^{k-1}),$$

the logical top 0 must exist, so

$$k \geq 1.$$

The pre-step residual is

$$2k + 3.$$

Reading 0 decreases the height by 2, so the down-step lands at

$$2k + 1.$$

Therefore every cancellation lands at one of:

$$3, 5, 7, \dots$$

The converse is equally rigid.

Proposition 17.2 — Motivating $d = 2$ correspondence.

Within the nonnegative 2-Dyck sector of the original computation, the transition

$$\boxed{0, 0/\varepsilon}$$

occurs if and only if the corresponding down-step lands at height

$$\boxed{3, 5, 7, \dots}$$

Proof.

The forward direction was just computed.

Conversely, suppose a down-step lands at

$$2k + 1 \geq 3.$$

Its pre-step height is

$$2k + 3 \geq 5.$$

This is a positive odd residual.

Corollary 17.1 therefore forces the pre-step configuration to be

$$(L, 0^k)$$

with

$$k \geq 1.$$

The enabled input-0 transition is exactly the distinguished left cancellation.

□

This theorem is the only point at which the full original machine is needed to motivate the correction statistic.

Everything that follows will be stated intrinsically for paths.

17.4 From one machine event to a d -Dyck statistic

For

$$d \geq 2,$$

a d -Dyck excursion has:

- up-step U of height $+1$;
- down-step D of height $-d$;
- nonnegative prefix heights;
- final height 0.

A size- n excursion has:

$$n$$

down-steps and

$$dn$$

up-steps.

The original $d = 2$ cancellation landing heights were

$$3, 5, 7, \dots$$

or equivalently

$$h_{\text{end}} \equiv 1 \pmod{2}, \quad h_{\text{end}} \geq 3.$$

This suggests the intrinsic generalization.

Definition 17.3 — Correction.

A down-step is a **correction** if

$$h_{\text{end}} \equiv 1 \pmod{d}, \quad h_{\text{end}} \geq d + 1.$$

Equivalently,

$$h_{\text{end}} = d + 1 + dr$$

for a unique

$$r \in \mathbb{N}.$$

The parameter d now belongs to the combinatorial path system.

For $d > 2$, we are **not** claiming that the original DPDA itself implements this generalized rule.

The correct causal statement is:

$$\begin{aligned} &\text{original } d = 2 \text{ cancellation} \rightarrow \text{intrinsic landing-height condition} \\ &\quad \rightarrow \text{combinatorial } d\text{-generalization.} \end{aligned}$$

17.5 Chronological correction signatures

Suppose a path contains $m \geq 1$ corrections, in chronological order, landing at heights

$$c_1, \dots, c_m.$$

Write

$$c_i = d + 1 + dr_i.$$

Definition 17.4 — Normalized correction signature.

The ordered tuple

$$r(P) = (r_1, \dots, r_m)$$

is the normalized correction signature.

The order is essential.

Two paths can have the same number of corrections and even the same multiset of correction heights while having different ordered signatures.

This branch therefore remembers more than the scalar count statistic.

17.6 Path defect

If a size- n path has $m \geq 1$ corrections, define

$$\delta(P) = n - m - 2.$$

The form may look artificial before the realization theorem.

It is chosen so that:

- maximal-correction paths have defect 0;
- every proposed ordered signature has an exact minimum defect;
- the excess above that minimum can be increased freely by zero-to-zero padding.

The theorem in Chapter 18 will prove in particular that

$$\delta(P) \geq 0.$$

17.7 The new inverse problem

The original question was forward:

Given a path, where do cancellations/corrections occur?

The signature theory asks the inverse question:

Given an ordered sequence

$$r = (r_1, \dots, r_m),$$

can it occur as a correction history?

And, if so:

What is the smallest path defect at which it can occur?

This is much stronger than counting corrections.

The answer is one explicit functional

$$\Delta(r).$$

To derive it, we first need the exact local price of moving from one correction height to the next.

Chapter 18

Correction Signatures and Minimum Defect

Let

$$r = (r_1, \dots, r_m) \in \mathbb{N}^m, \quad m \geq 1.$$

The goal is to determine exactly when r is realizable as a normalized correction signature.

The answer is built from two local costs:

1. the price of moving between consecutive corrections;
2. the price of terminating after the final correction.

The first cost contains the key nonlinearity.

18.1 The threshold function

Define

$$\psi(a) = \begin{cases} 0, & a \leq 1, \\ a, & a \geq 2. \end{cases}$$

Why is there a jump from zero to a ?

Suppose consecutive normalized correction heights are

$$r \quad \text{and} \quad r',$$

and set

$$a = r - r'.$$

Let b be the number of non-correction down-steps strictly between the two corrections, and let u be the number of intervening up-steps.

Height balance gives

$$u = d(b + 1 - a).$$

This one equation explains the threshold.

18.2 Local Gap Cost

Lemma 18.1 — Local Gap Cost.

For two consecutive normalized correction heights r, r' , with

$$a = r - r',$$

the number b of intervening non-correction down-steps satisfies

$$b \geq \psi(a).$$

The bound is sharp.

Proof.

If

$$a \leq 1,$$

choose

$$b = 0.$$

Then

$$u = d(1 - a) \geq 0.$$

The block

$$U^{d(1-a)}D$$

moves directly to the next correction without an intervening down-step.

Now assume

$$a \geq 2.$$

The balance equation gives

$$b \geq a - 1.$$

If

$$b = a - 1,$$

then

$$u = 0.$$

But the first intervening down-step would land again in the correction residue class, creating an unintended correction between the two prescribed ones.

Therefore

$$b \geq a.$$

Equality is attainable by the detour block

$$\boxed{UD^aU^{d-1}D.}$$

The intermediate down-steps lie one residue class away from the correction class, so they are not corrections. The final D returns exactly to the prescribed correction height. $\square$

This theorem is the geometric origin of ψ .

A downward change of size at most 1 is free.

A drop of size at least 2 cannot be taken directly because the naive descent would create an unwanted correction.

The path must leave the correction residue class and return.

That detour costs exactly the full drop a .

18.3 Terminal Cost

Lemma 18.2 — Terminal Cost.

If the final correction has normalized height r , then at least

$$\boxed{r + 2}$$

non-correction down-steps are required after it.

The bound is sharp.

Proof.

Immediately after the final correction, the height is

$$d + 1 + dr.$$

If b down-steps and u up-steps remain, then

$$d + 1 + dr + u - db = 0.$$

Hence

$$u = d(b - r - 1) - 1.$$

Since

$$u \geq 0,$$

we must have

$$b \geq r + 2.$$

Equality is attained by

$$\boxed{U^{d-1}D^{r+2}}.$$

All its down-steps land at multiples of d , so none is a correction. $\square$

The final signature entry therefore contributes directly to the minimum defect.

18.4 Intrinsic signature defect

Define

$$\boxed{\Delta(r) = r_m + \sum_{i=1}^{m-1} \psi(r_i - r_{i+1}).}$$

This quantity contains no d .

That universality is already striking.

The physical correction heights scale as

$$d + 1 + dr_i,$$

but the normalized realizability cost is independent of the Dyck parameter.

18.5 Canonical blocks

To prove sufficiency constructively, define:

Initial block

$$I_r = U^{2d+1+dr} D.$$

Its unique down-step lands at correction height

$$d + 1 + dr.$$

Gap block

For $a \in \mathbb{Z}$,

$$G_a = \begin{cases} U^{d(1-a)} D, & a \leq 1, \\ U D^a U^{d-1} D, & a \geq 2. \end{cases}$$

Terminal block

$$T_r = U^{d-1} D^{r+2}.$$

Padding block

$$P = U^d D.$$

The padding block starts and ends at height zero and creates no correction.

Proposition 18.3 — Block Purity.

Each canonical block creates exactly the advertised correction behavior and no unintended correction.

Moreover, when started at its prescribed height, every intermediate height remains nonnegative.

This proposition is what makes the global realization theorem compositional.

18.6 Defect–Signature Characterization

Theorem 18.4 — Defect–Signature Characterization.

For every

$$d \geq 2$$

and every normalized signature

$$r = (r_1, \dots, r_m),$$

the following are equivalent:

1. r is realized by a d -Dyck excursion of defect δ ;
- 2.

$$\boxed{\Delta(r) \leq \delta.}$$

Equivalently,

$$\boxed{\delta_{\min}(r) = \Delta(r).}$$

If the signature has length m , then its minimum path size is

$$\boxed{n_{\min}(r) = m + 2 + \Delta(r).}$$

Proof — lower bound. A path realizing r contains:

$$n - m$$

non-correction down-steps.

Between correction i and correction $i + 1$, Lemma 18.1 forces at least

$$\psi(r_i - r_{i+1})$$

such steps.

After the last correction, Lemma 18.2 forces at least

$$r_m + 2.$$

Any non-correction down-steps before the first correction only increase the count.

Therefore

$$n - m \geq 2 + r_m + \sum_{i=1}^{m-1} \psi(r_i - r_{i+1}) = 2 + \Delta(r).$$

Hence

$$\delta(P) = n - m - 2 \geq \Delta(r).$$

Proof — construction. Let

$$a_i = r_i - r_{i+1}.$$

For

$$s \geq 0,$$

define

$$W_d(r; s) = I_{r_1} G_{a_1} \cdots G_{a_{m-1}} T_{r_m} P^s.$$

Block Purity shows that this is a nonnegative d -Dyck excursion whose normalized correction signature is exactly r .

Its number of down-steps is

$$m + 2 + \Delta(r) + s.$$

Thus

$$\delta(W_d(r; s)) = \Delta(r) + s.$$

Taking $s = 0$ gives a minimum realization.

Taking

$$s = \delta - \Delta(r)$$

realizes every larger defect. $\square$

This is the structural center of the cancellation branch.

It solves both directions:

$$\text{path} \rightarrow \text{signature}$$

and

$$\text{signature} \rightarrow \text{minimum realizing path}.$$

18.7 Canonical does not mean unique

Corollary 18.5 — Canonical minimum realization.

Every signature r has the explicit minimum witness

$$W_d^{\min}(r) = W_d(r; 0).$$

The word **canonical** means:

the distinguished block construction selected by the theory.

It does **not** mean:

every minimum realizing path is identical to this word.

This scope matters.

The theorem determines the minimum defect uniquely, not the complete set of minimum path orderings.

18.8 Signature persistence

Corollary 18.6 — Signature Persistence.

A signature r is realizable at exactly the defect levels

$$\boxed{\Delta(r), \Delta(r) + 1, \Delta(r) + 2, \dots}$$

because each copy of

$$P = U^d D$$

adds one defect unit without changing the correction signature.

Therefore the intrinsic defect is a **birth level**.

A signature appears for the first time at $\Delta(r)$, then persists forever.

18.9 Maximum number of corrections

Since

$$\Delta(r) \geq 0,$$

every path with $m \geq 1$ corrections satisfies

$$\boxed{m \leq n - 2.}$$

Thus defect zero is exactly the maximal-correction regime.

That boundary turns out to be Catalan.

18.10 Zero defect

Corollary 18.7 — Zero-defect signatures.

$$\boxed{\Delta(r) = 0}$$

if and only if

$$\boxed{r_m = 0}$$

and

$$\boxed{r_i - r_{i+1} \leq 1 \quad (1 \leq i < m).}$$

Reverse the sequence:

$$a_i = r_{m+1-i}.$$

Then

$$a_1 = 0$$

and

$$a_{i+1} \leq a_i + 1.$$

These are exactly the usual Dyck area-sequence inequalities. Therefore the number of zero-defect signatures of length m is

$$C_m.$$

The Catalan family is not the whole correction theory. It is its exact zero-cost boundary.

18.11 Defect as penalized overshoot

After reversal,

$$a_i = r_{m+1-i},$$

the defect becomes

$$\Delta^{\leftarrow}(a) = a_1 + \sum_{\substack{1 \leq i < m \\ a_{i+1} - a_i \geq 2}} (a_{i+1} - a_i).$$

Thus ordinary Dyck area sequences are precisely the nonnegative sequences with zero cost.

Positive defect allows overshoots above the Catalan nearest-neighbor condition, but every overshoot of size at least 2 is charged by its full magnitude.

This makes the geometry intuitive:

$$\text{Catalan behavior} = \text{zero-defect regime};$$

$$\text{non-Catalan jumps} = \text{quantified correction debt}.$$

The next chapter shows that the Catalan boundary is not only enumerative. Its maximal-correction paths themselves contain a full binary tree.

Chapter 19

The Catalan Boundary and the Binary Cancellation Skeleton

At defect zero,

$$m = n - 2.$$

Thus a size- n path contains the maximum possible number of corrections.
The signature count is Catalan:

$$C_{n-2}.$$

But the path count contains an additional factor d :

$$dC_{n-2}.$$

This chapter explains both factors structurally.

The intrinsic correction structure is a full binary tree.

The factor d records only how the path enters that binary regime.

19.1 Two unavoidable unmarked nodes

Use the standard first-return decomposition of a nonempty d -Dyck excursion:

$$P = UP_1UP_2 \cdots UP_dDP_{d+1}.$$

Each internal node corresponds to one down-step.

A node is marked exactly when its down-step is a correction.

A maximal-correction size- n path has:

$$n - 2$$

marked internal nodes and therefore exactly two unmarked internal nodes.

The root is necessarily unmarked because its baseline is zero.
Before a correction baseline

$$d + 1$$

can be reached, at least one further unmarked internal node is necessary.
In the maximal-correction regime, these are the only two unmarked nodes.
That rigid entrance forces the rest of the tree into a binary skeleton.

19.2 The d entrance choices

Choose the unique nonempty pre-down-step child of the root.

If it is slot

$$j \in [d],$$

its baseline is j .

Every other root child must be empty, otherwise another unmarked internal node would be created.

At this second unmarked node, the unique child that reaches correction baseline $d + 1$ without introducing an additional unmarked node is slot

$$d + 1 - j.$$

Thus there are exactly

$$\boxed{d}$$

possible entrance choices.

After those two forced unmarked nodes, the remaining structure is entirely inside the correction residue class.

19.3 Why correction nodes become binary

Suppose a correction node has baseline

$$h \equiv 1 \pmod{d}, \quad h \geq d + 1.$$

Its pre-down-step child baselines are

$$h + 1, \dots, h + d,$$

and its trailing child has baseline h .

Among these, the only baselines that remain in the correction class are:

$$h + d$$

and

$$h.$$

Therefore, once all unmarked nodes have been used up, a correction node has exactly two recursive correction slots:

- slot d ;
- slot $d + 1$.

All other child slots must be empty.

The correction structure is therefore binary.

19.4 Binary Cancellation Skeleton

Let

$$\mathcal{B}_m$$

be the set of plane full binary trees with m internal nodes.

Let

$$\mathcal{M}_{d,n}$$

be the class of size- n maximally corrected d -Dyck excursions.

Theorem 19.1 — Binary Cancellation Skeleton.

For every

$$n \geq 3,$$

$$\boxed{\mathcal{M}_{d,n} \cong [d] \times \mathcal{B}_{n-2}.$$

Consequently,

$$\boxed{Q_d(n, n-2) = dC_{n-2}.$$

Proof.

From a maximal-correction path:

1. record the root entrance slot

$$j \in [d];$$

2. delete the two forced unmarked entrance nodes;
3. retain the remaining marked subtree;
4. reinterpret its recursive slots d and $d + 1$ as the two ordered binary children.

This gives an element of

$$[d] \times \mathcal{B}_{n-2}.$$

Conversely, given:

- an entrance choice j ;
- a plane full binary tree;

reinsert the two forced unmarked nodes and place the binary tree recursively in the only two correction-preserving child slots.

The constructions are inverse. $\square$

The Catalan factor is therefore not a generating-function accident.

It is the count of an actual binary cancellation skeleton.

19.5 Tree coordinates

Assign child weights

$$\omega(j) = j \quad (1 \leq j \leq d),$$

and

$$\omega(d + 1) = 0.$$

Lemma 19.2 — Tree-coordinate lemma.

If an internal node has root-to-node child address

$$i_1 \cdots i_s,$$

then the baseline height of its associated down-step is

$$h(v) = \sum_{\nu=1}^s \omega(i_\nu).$$

For a correction node, the two recursive binary directions are:

- **high**: slot d , increment d ;
- **low**: slot $d + 1$, increment 0.

The two entrance weights always sum to

$$d + 1.$$

Therefore the binary root always begins at correction height

$$d + 1,$$

independently of the entrance choice.

19.6 Correction height from the binary address

Let

$$L(v)$$

be the number of high edges on the path from the binary root to a correction node v .

Then:

$$c(v) = d + 1 + dL(v),$$

and hence its normalized correction height is

$$r(v) = L(v).$$

Thus the normalized signature has a purely binary-tree interpretation.

The parameter d only rescales the physical path heights.

The normalized correction values themselves count high edges.

This is another manifestation of the universality first seen in Chapter 18.

19.7 Inorder chronology

In the d -ary path traversal:

- the slot- d correction subtree is visited before the node's closing down-step;
- the slot- $(d + 1)$ correction subtree is visited after it.

Therefore correction nodes appear in ordinary binary inorder:

$$\text{high subtree} \rightarrow \text{node} \rightarrow \text{low subtree.}$$

Let

$$v_1, \dots, v_m$$

be the internal binary nodes in inorder.
Then the normalized signature is

$$r_i = L(v_i).$$

So the chronological correction signature is exactly an inorder depth statistic on the binary skeleton.

19.8 The signature is lossless

A Catalan-counted statistic need not normally determine the full tree.

Here it does.

Theorem 19.3 — Lossless Signature Theorem.

In the zero-defect regime, the normalized correction signature uniquely determines the binary cancellation skeleton.

Proof.

Let

$$r = (r_1, \dots, r_m)$$

satisfy the zero-defect conditions.

Let j be the first index such that

$$r_j = 0.$$

In binary inorder, the first zero is the root.

Entries before j lie in the high subtree and each carries one additional high edge.

Therefore the normalized code of the high subtree is

$$(r_1 - 1, \dots, r_{j-1} - 1).$$

Entries after the root lie in the low subtree, whose code is

$$(r_{j+1}, \dots, r_m).$$

Both subsequences satisfy the same zero-defect conditions.

Recursively apply the same decomposition.

This reconstructs the tree uniquely. $\square$

Hence:

$$\begin{aligned} \text{binary cancellation skeleton} &\longleftrightarrow \text{zero-defect signature} \\ &\longleftrightarrow \text{Dyck area sequence.} \end{aligned}$$

The factor d is external to this intrinsic Catalan structure.

It remembers only the entrance choice.

19.9 Area statistic

If the zero-defect signature has length m , then

$$c_i = d + 1 + dr_i.$$

Therefore

$$\sum_{i=1}^m c_i = m(d + 1) + d \sum_{i=1}^m r_i.$$

After reversal,

$$\sum_i r_i$$

is the ordinary area statistic of the corresponding Dyck path.

If

$$\mathcal{A}_m(v) = \sum_D v^{\text{area}(D)}$$

is the Dyck area polynomial, then the maximal-correction paths satisfy

$$\sum_{P \in \mathcal{M}_{d,m+2}} v^{\sum_i c_i} = d v^{m(d+1)} \mathcal{A}_m(v^d).$$

Thus the Catalan boundary is structural at three levels:

1. number of signatures;
2. binary cancellation skeleton;
3. correction-height distribution.

19.10 What the original transition has become

We began with:

$$(L, 0S) \xrightarrow{0} (L, S).$$

At defect zero, the ordered history of occurrences of that transition becomes:

a full binary tree encoded losslessly by its normalized correction signature.

This is a much stronger statement than “Catalan numbers appear.”

The original transition has generated a complete combinatorial coordinate system.

The next chapter moves away from individual signatures and counts:

- paths by number of corrections;
- signatures by exact intrinsic defect.

Those are different arrays and must remain separate.

Chapter 20

Exact Correction and Defect Enumeration

The structural theory has produced three different kinds of objects.
They must not be conflated.

K_d auxiliary kernel coefficients,

$Q_d(n, k)$ paths with exactly k corrections,

and

$N_{m,\ell}$ normalized signatures born at exact defect ℓ .
--

The first appears naturally inside the scalar generating-function grammar.
The second is the actual path correction statistic.

The third belongs to the intrinsic signature geometry and is independent of d .

20.1 Scalar correction generating function

Let

$$G_h(z, t)$$

count d -Dyck excursions that begin and end at baseline h , never go below h , with:

- z marking down-steps;
- t marking corrections.

A down-step landing at baseline h has weight

$$\tau_h = \begin{cases} t, & h \in \{d+1, 2d+1, \dots\}, \\ 1, & \text{otherwise.} \end{cases}$$

The standard first-return decomposition gives

$$G_h = 1 + z\tau_h G_{h+1} \cdots G_{h+d} G_h.$$

Above the exceptional low boundary, the marking pattern is periodic modulo d .

This collapses the infinite family to three formal series:

- F : baseline zero;
- A : unmarked tail residue;
- B : marked correction residue.

They satisfy

$$F = 1 + zA^d F,$$

$$A = 1 + zA^d B,$$

$$B = 1 + ztA^{d-1}B^2.$$

Define

$$Q_d(n, k) = [z^n t^k] F(z, t).$$

Thus $Q_d(n, k)$ is the actual correction triangle.

20.2 The auxiliary kernel and A371395

Set

$$Y = zA^{d-1}B.$$

Then

$$A = \frac{1}{1-Y}, \quad B = \frac{1}{1-tY},$$

and

$$Y(1-Y)^{d-1}(1-tY) = z.$$

Define

$$K_d = \frac{Y}{z}.$$

Lagrange inversion gives

$$[z^n t^k] K_d = \frac{1}{n+1} \binom{n+k}{k} \binom{dn+d-k-2}{n-k}.$$

For

$$d = 2,$$

this becomes

$$[z^n t^k] K_2 = \frac{1}{n+1} \binom{n+k}{k} \binom{2n-k}{n-k}.$$

This is OEIS A371395 [58].

But K_2 is not the correction distribution.

The actual correction generating function is

$$F_d(z, t) = \frac{1 - Y}{1 - 2Y + tY^2}.$$

Proposition 20.1 — Kernel–correction separation.

For $d = 2$, A371395 is the coefficient triangle of the auxiliary kernel K_2 , not the path correction triangle Q_2 .

This distinction is important because the two arrays diverge immediately.

A371395 is a useful algebraic waypoint.

It is not the final statistic generated by the DPDA-derived correction event.

20.3 Exact correction triangle

Eliminating B from the three-class system and writing

$$U = A - 1$$

gives

$$U(1 + (1 - t)U) = z(1 + U)^{d+1}.$$

Lagrange–Bürmann inversion yields the exact path count.

Theorem 20.2 — Exact correction triangle.

For

$$n \geq 1, \quad k \geq 0,$$

$$Q_d(n, k) = \mathbf{1}_{\{k=0\}} + \sum_{j=k+1}^{n-1} \frac{d(n-j)}{j} \binom{j+k-1}{k} \binom{dn-k-1}{j-k-1}.$$

In particular,

$$Q_d(n, k) = 0 \quad (k > n-2).$$

The complete Lagrange–Bürmann coefficient derivation belongs in the technical appendix of the complete research version.

The theorem statement and its structural consequences remain in the main text.

20.4 Row sums and boundary columns

Setting

$$t = 1$$

forgets the correction marking.

The row sum is the ordinary Fuss–Catalan number:

$$\sum_{k \geq 0} Q_d(n, k) = \frac{1}{dn+1} \binom{(d+1)n}{n}.$$

At the zero-correction boundary,

$$Q_d(n, 0) = \frac{1}{n} \sum_{j=0}^{n-1} (n-j) \binom{dn}{j}.$$

For

$$d = 2,$$

this simplifies to

$$Q_2(n, 0) = \binom{2n-1}{n-1}.$$

At the maximal-correction boundary,

$$Q_d(n, n-2) = dC_{n-2}.$$

Chapter 19 already explained the last identity bijectively through the binary cancellation skeleton.

This is a useful example of two proof styles meeting:

- algebraic coefficient extraction;
- structural combinatorial decomposition.

20.5 Intrinsic signature enumeration

Now forget the paths and enumerate normalized signatures directly.

Define

$$N_{m,\ell} = \#\{r \in \mathbb{N}^m : \Delta(r) = \ell\}.$$

Let

$$S(x, q) = \sum_{m \geq 1} \sum_{\ell \geq 0} N_{m,\ell} x^m q^\ell.$$

Here:

- x marks signature length;
- q marks exact intrinsic minimum defect.

This enumeration is fundamentally different from $Q_d(n, k)$.

A signature of defect ℓ can be realized at every larger path defect.

Thus $N_{m,\ell}$ records the level at which a signature is **born**.

20.6 Translation and first-zero decomposition

Let $Z(x, q)$ count nonempty signatures whose minimum entry is zero.

Every nonempty signature has a unique decomposition

$$r = \bar{r} + s\mathbf{1},$$

where

$$s = \min_i r_i$$

and

$$\min_i \bar{r}_i = 0.$$

Adding 1 to every entry leaves all adjacent differences unchanged and raises the terminal term by 1.

Therefore

$$\Delta(r + \mathbf{1}) = \Delta(r) + 1.$$

This gives

$$Z = (1 - q)S.$$

Now cut a minimum-zero signature at its first zero.

A second auxiliary series E counts nonempty signatures ending in zero.

The first-zero decomposition yields a closed algebraic system, which simplifies to the master equation.

20.7 Defect–Signature Master Equation

Theorem 20.3 — Defect–Signature Master Equation.

The formal power series $S(x, q)$ with zero constant term is uniquely determined by

$$(1 - q)S(1 + qS) = x(1 + S)^2(1 + q^2S).$$

Equivalently,

$$S = x \frac{(1 + S)^2(1 + q^2S)}{(1 - q)(1 + qS)}.$$

The equation contains no d .

Thus the universality of the Defect–Signature theorem persists at the enumerative level.

The path parameter d changes physical height scaling and path multiplicity.

It does not change the intrinsic normalized signature birth counts.

20.8 Exact defect coefficients

Lagrange inversion gives a finite coefficient formula for every

$$N_{m,\ell}.$$

The first four defect layers simplify to:

$$N_{m,0} = C_m,$$

$$N_{m,1} = \frac{3m}{m+2}C_m = C_{m+1} - C_m,$$

$$N_{m,2} = \frac{m(m^2 + 10m + 1)}{(m+2)(m+3)}C_m,$$

and

$$N_{m,3} = \frac{5m(m^3 + 7m^2 + 2m + 2)}{(m+2)(m+3)(m+4)}C_m.$$

The zero layer is exactly Catalan.

The positive layers quantify how the signature class expands as non-Catalan overshoots are admitted.

20.9 Cumulative realizability

Define

$$\widehat{N}_{m,\delta} = \#\{r \in \mathbb{N}^m : r \text{ is realizable at path defect } \delta\}.$$

By signature persistence,

$$\widehat{N}_{m,\delta} = \sum_{\ell=0}^{\delta} N_{m,\ell}.$$

Thus:

- $N_{m,\ell}$ counts signatures born at exact intrinsic defect ℓ ;
- $\widehat{N}_{m,\delta}$ counts all signatures visible by defect δ .

At defect one,

$$\widehat{N}_{m,1} = C_{m+1}.$$

The first cumulative layer is still exactly Catalan, one size higher.

Defect two is the first level at which a genuinely non-Catalan local overshoot such as

$$(2, 0)$$

can appear.

20.10 Three arrays, three questions

The distinction can now be stated cleanly.

K_d

An auxiliary kernel series produced by the scalar grammar.

Question:

What coefficient array appears in the algebraic series-reversion core?

$Q_d(n, k)$

Path correction distribution.

Question:

How many size- n d -Dyck excursions contain exactly k correction events?

$N_{m,\ell}$

Intrinsic ordered signature distribution.

Question:

How many length- m correction histories are born at exact minimum defect ℓ ?

The DPDA-derived project passes through all three, but they are not interchangeable.

The next chapter studies the asymptotic hierarchy of the third array.

Chapter 21

The Asymptotic Defect Hierarchy

The master equation

$$(1 - q)S(1 + qS) = x(1 + S)^2(1 + q^2S)$$

provides exact coefficients.

It also reveals a surprisingly rigid asymptotic structure.

For fixed defect ℓ , the number of length- m signatures grows on the Catalan exponential scale:

$$4^m.$$

The defect level changes only the polynomial factor.

More precisely, defect levels occur in pairs:

$$(0, 1), \quad (2, 3), \quad (4, 5), \dots$$

and every two defect units increase the polynomial exponent by one.

21.1 Fixed-defect generating functions

Define

$$S_\ell(x) = [q^\ell]S(x, q) = \sum_{m \geq 1} N_{m, \ell} x^m.$$

Set

$$t = \sqrt{1 - 4x}.$$

At defect zero, the master equation gives

$$S_0(x) = C(x) - 1,$$

hence

$$S_0(x) = \frac{1-t}{1+t}.$$

Thus the dominant singularity is the Catalan point

$$x = \frac{1}{4}.$$

The same singular location controls every fixed defect layer.

21.2 Algebraicity and singularity location

Proposition 21.1 — Fixed-defect algebraicity.

For every fixed

$$\ell \geq 0,$$

$$S_\ell(x) \in \mathbb{Q}(t), \quad t = \sqrt{1-4x}.$$

More precisely,

$$S_\ell(x) \in \mathbb{Q}[t, t^{-1}, (1+t)^{-1}].$$

Therefore, on the branch $t(0) = 1$, the only possible dominant singularity is

$$x = \frac{1}{4}.$$

Proof idea. Let

$$\mathcal{F}(S, q, x) = (1-q)S(1+qS) - x(1+S)^2(1+q^2S).$$

At

$$q = 0,$$

the coefficient of the unknown S_ℓ in the extracted q^ℓ equation is

$$\frac{\partial \mathcal{F}}{\partial S} = 1 - 2x(1+S_0) = t.$$

Thus each new fixed-defect layer is obtained rationally from the earlier ones by dividing by powers of t .

Induction gives the claimed field. $\square$

This proposition justifies using singularity analysis uniformly across all fixed defect levels.

21.3 Critical scaling

To extract the leading singular behavior of all layers simultaneously, set

$$q = ty.$$

Seek an expansion

$$S(x, ty) = 1 + tf(y) + t^2g(y) + O(t^3).$$

The branch compatible with

$$S(x, 0) = \frac{1-t}{1+t}$$

must satisfy

$$f(0) = -2.$$

Substitution into the master equation gives

$$f^2 - 4yf + 8y^2 - 4 = 0.$$

The correct branch is

$$f(y) = 2y - 2\sqrt{1-y^2}.$$

At the next order,

$$g(y) = 2 \frac{4y^3 - 3y + \sqrt{1-y^2}}{\sqrt{1-y^2}}.$$

The even coefficients of f and odd coefficients of g determine the dominant singular terms of successive defect layers.

This is the analytic origin of the pair structure.

21.4 Pairwise leading singularities

For

$$j \geq 1,$$

one obtains

$$S_{2j}(x) = A_j t^{1-2j} + O(t^{2-2j}),$$

and

$$S_{2j+1}(x) = B_j t^{1-2j} + O(t^{2-2j}),$$

with

$$A_j = \frac{C_{j-1}}{4^{j-1}},$$

and

$$B_j = (2j + 3)A_j.$$

The key point is that defects $2j$ and $2j + 1$ have the **same singular exponent**.

Their leading constants differ, but their polynomial growth degree is identical.

21.5 Asymptotic Defect Hierarchy

Theorem 21.2 — Asymptotic Defect Hierarchy.

For every fixed

$$j \geq 0,$$

$$N_{m,2j} \sim \frac{m^j}{j!} C_m,$$

and

$$N_{m,2j+1} \sim \frac{2j+3}{j!} m^j C_m.$$

Equivalently,

$$N_{m,2j} \sim \frac{4^m}{j!\sqrt{\pi}} m^{j-\frac{3}{2}},$$

and

$$N_{m,2j+1} \sim \frac{2j+3}{j!\sqrt{\pi}} 4^m m^{j-\frac{3}{2}}.$$

Proof.

Apply the standard transfer theorem to the leading singular terms in powers of

$$t = \sqrt{1-4x}.$$

For $j \geq 1$, the singular exponent is

$$t^{1-2j} = (1-4x)^{-(j-\frac{1}{2})}.$$

The transfer coefficient, together with the explicit constants A_j, B_j , gives the displayed formulas.

The boundary case

$$j = 0$$

uses the exact identities

$$N_{m,0} = C_m$$

and

$$N_{m,1} = C_{m+1} - C_m \sim 3C_m.$$

□

Thus:

$$\frac{N_{m,\ell}}{C_m} = \Theta\left(m^{\lfloor \ell/2 \rfloor}\right)$$

for fixed ℓ .

21.6 The defect pairs

The first leading ratios to the Catalan scale are:

$$1, \quad 3, \quad m, \quad 5m, \quad \frac{1}{2}m^2, \quad \frac{7}{2}m^2, \quad \frac{1}{6}m^3, \quad \frac{3}{2}m^3, \dots$$

The hierarchy is therefore:

$$(0, 1)$$

on the Catalan polynomial scale,
then

$$(2, 3)$$

one power of m higher,
then

$$(4, 5)$$

two powers higher,
and so on.

Every two defect units add one asymptotically free polynomial degree.

21.7 Geometric interpretation

The theorem is analytic.

The following interpretation comes afterward.

The zero-defect Catalan condition permits adjacent reversed-signature rises of at most 1.

A true violation requires a rise of at least 2.

Such an overshoot costs at least two defect units.

Therefore the pair structure

$$(2j, 2j + 1)$$

is consistent with j asymptotically free non-Catalan defect locations among m signature positions.

The factor

$$\frac{m^j}{j!}$$

has exactly the scale expected from choosing j widely separated defect locations.

This is an interpretation of the proven singularity result, not a replacement for its proof.

The monograph should keep that distinction explicit.

21.8 Cumulative asymptotics

Since

$$\hat{N}_{m,\delta} = \sum_{\ell \leq \delta} N_{m,\ell},$$

the top defect layer dominates within each fixed cumulative cutoff.
Thus:

$$\hat{N}_{m,2j} \sim \frac{m^j}{j!} C_m,$$

and

$$\hat{N}_{m,2j+1} \sim \frac{2j+4}{j!} m^j C_m.$$

At

$$j = 0,$$

this agrees with the exact identity

$$\hat{N}_{m,1} = C_{m+1}.$$

21.9 What the cancellation branch has established


We can now summarize the complete branch generated by one transition of the original DPDA:

$$0, 0/\varepsilon$$

$\Downarrow$

odd landing-height event in the nonnegative 2-Dyck sector

$\Downarrow$



The DPDA was the discovery mechanism.

The final normalized signature theory is intrinsic and no longer depends on the machine's remaining states or stack alphabet.

21.10 The next arithmetic offshoot

The ordered correction signature remembers chronology.

The next paper asks what happens when we push that chronology one step further and attach an explicit third-order arithmetic interaction to:

- singleton correction positions;
- adjacent pairs;
- connected triples.

That construction produces the CCR-3 recurrence.

It then develops a modular Stirling shadow, a simple Hensel branch, singular extinction, and ramified p -adic recovery.

This arithmetic refinement is a genuine continuation of the correction branch, but it is not a prerequisite for the later general real-time $p:q$ DPDA theory.

Accordingly, the next Part will be marked as a specialized offshoot.

Readers interested only in the main normalization spine may skip it and rejoin at the later return to pushdown machines.

PART VII

Chronological Carry and the Arithmetic Offshoot

Research provenance. The chronological-carry refinement and p -adic root analysis in this Part are developed in [7]. The surrounding arithmetic literature includes carries and the p -adic behavior of Stirling numbers [27, 13, 42, 55].

Reader's guide.

Chapters 22–25 are a specialized arithmetic offshoot. They continue the ordered-correction branch of Chapters 17–21, but they are not prerequisites for the later real-time $p:q$ pushdown-normalization theory. Readers following only the main DPDA spine may skip directly to Chapter 26.

Chapter 22

From Ordered Corrections to CCR-3

The previous Part ended with an intrinsic correction-signature theory.

The original DPDA supplied a distinguished cancellation event.

The path interpretation ordered those events chronologically.

The Defect-Signature theorem then showed that the realization cost is local along that chronology:

$$\Delta(r) = r_m + \sum_{i=1}^{m-1} \psi(r_i - r_{i+1}).$$

Thus the correction history carries a natural path graph

$$1 - 2 - \cdots - n.$$

This is exactly the amount of pushdown-derived structure needed for the next step.

It is also the boundary of what the DPDA itself forces.

22.1 Scientific boundary

The arithmetic refinement introduced in this Part must be stated carefully.

What is inherited from the earlier DPDA/correction branch is:

1. a distinguished cancellation/correction event;
2. chronological ordering of those events;
3. nearest-neighbor locality along the ordered positions.

What is **introduced here** is a new arithmetic polarization:

1. a numerical weight for a singleton chronological slot;

2. a weight for an adjacent pair interaction;
3. a weight for the first connected triple collision.

Therefore the correct implication is

$$\begin{array}{c} \text{DPDA-derived chronological locality + specified arithmetic} \\ \text{polarization} \\ \implies \text{CCR-3.} \end{array}$$

The stronger statement

$$\text{original DPDA} \implies \text{unique CCR-3 coefficients}$$

is not claimed.

This distinction is not editorial caution; it marks the exact scientific scope of the arithmetic branch.

22.2 Connected supports through order three

A connected subset of a path graph is an interval.

Therefore, up to size three, there are only three local connected support types.

Lemma 22.1 — Third-order connected support.

Every connected chronological support of cardinality at most three is one of

$$\boxed{\{j\}, \quad \{j, j+1\}, \quad \{j, j+1, j+2\}.$$

These are, respectively:

- a singleton slot;
- an adjacent pair;
- a consecutive triple.

The lemma is elementary, but its role is structural: no other connected third-order motif exists on a one-dimensional chronology.

22.3 The CCR-3 recurrence

Fix a prime

$$p \geq 5$$

and define

$$\ell_n = (p+1)n - p.$$

Work in the truncated ring

$$\mathbb{Z}[u]/(u^4).$$

Set

$$P_{p,0}^{[3]}(u) = 1,$$

and

$$P_{p,-1}^{[3]}(u) = P_{p,-2}^{[3]}(u) = 0.$$

Definition 22.2 — CCR-3.

For

$$n \geq 1,$$

define

$$\begin{aligned} P_{p,n}^{[3]}(u) = & (1 + \ell_n u) P_{p,n-1}^{[3]}(u) \\ & + p(n-1) \ell_{n-1} u^2 P_{p,n-2}^{[3]}(u) \\ & - p(n-2)(n-1) u^3 P_{p,n-3}^{[3]}(u) \pmod{u^4}. \end{aligned}$$

The three nontrivial sources represent:

$$\text{singleton} + \text{adjacent pair carry} + \text{connected triple collision}.$$

The first rows are

$$P_{p,0}^{[3]} = 1,$$

$$P_{p,1}^{[3]} = 1 + u,$$

$$P_{p,2}^{[3]} = 1 + (p+3)u + 2(p+1)u^2,$$

and

$$P_{p,3}^{[3]} = 1 + (3p+6)u + (4p^2 + 15p + 11)u^2 + 6(p+1)^2 u^3.$$

The superscript [3] is important.

The recurrence is intentionally third-order.

It should not be extrapolated blindly into an all-column theory: at order four a genuinely new connected motif may enter.

22.4 Hessenberg realization

Let

$$H_{p,n}$$

be the $n \times n$ leading principal truncation of the infinite lower-Hessenberg matrix defined by

$$(H_p)_{i,i} = \ell_i,$$

$$(H_p)_{i,i+1} = 1,$$

$$(H_p)_{i+1,i} = -pi\ell_i,$$

$$(H_p)_{i+2,i} = -pi(i+1),$$

with all entries below the second subdiagonal equal to zero. Define

$$D_{p,n}(u) = \det(I_n + uH_{p,n}).$$

Proposition 22.3 — Hessenberg realization.

$$\begin{aligned} D_{p,n} &= (1 + \ell_n u) D_{p,n-1} \\ &\quad + p(n-1)\ell_{n-1}u^2 D_{p,n-2} \\ &\quad - p(n-2)(n-1)u^3 D_{p,n-3}. \end{aligned}$$

Hence

$$P_{p,n}^{[3]}(u) \equiv D_{p,n}(u) \pmod{u^4}.$$

Proof.

Expand the determinant by the last row.

The diagonal contribution is

$$(1 + \ell_n u) D_{p,n-1}.$$

The first subdiagonal closes with one superdiagonal entry and contributes

$$+p(n-1)\ell_{n-1}u^2 D_{p,n-2}.$$

The second subdiagonal closes with two consecutive superdiagonal entries and contributes

$$-p(n-2)(n-1)u^3D_{p,n-3}.$$

No deeper band is present. $\square$

Thus the three chronological interaction orders are encoded by a three-band Hessenberg structure.

22.5 Operator form

Use the exponential basis

$$e_i(x) = \frac{x^{i-1}}{(i-1)!}, \quad i \geq 1.$$

Let

$$D = \partial_x, \quad Xf = xf, \quad N = XD,$$

and define

$$L_p = I + (p+1)N.$$

Then

$$De_i = e_{i-1},$$

$$Xe_i = ie_{i+1},$$

$$X^2e_i = i(i+1)e_{i+2},$$

and

$$L_pe_i = \ell_ie_i.$$

Corollary 22.4 — Operator form.

The Hessenberg matrix H_p is the matrix, in the basis (e_i) , of

$$\boxed{\mathcal{H}_p = D + L_p - pXL_p - pX^2.}$$

Equivalently,

$$\boxed{\mathcal{H}_p = [1 + (p+1)x - p(p+1)x^2]\partial_x + 1 - px - px^2.}$$

The determinant recurrence and differential-operator form are therefore two representations of the same CCR-3 dynamics.

22.6 What this chapter has and has not established

We have not derived CCR-3 from the original DPDA transition table.

We have established the narrower chain

ordered correction chronology $\rightarrow$ path-local connected supports
 $\rightarrow$ specified third-order polarization $\rightarrow$ Hessenberg/operator
recurrence.

The next question is arithmetic.

What do the first three fixed columns of this recurrence look like?

The answer is unexpectedly rigid: they are polynomials, and modulo p the entire third jet collapses to a universal first-kind Stirling shadow.

Chapter 23

Fixed Columns and the Universal Stirling Shadow

The CCR-3 recurrence is nonlinear in its dependence on the chronological index but finite in interaction order.

To expose its arithmetic, we now freeze the power of u and let the row index vary.

This produces a transport system for the first three fixed columns.

23.1 Column coefficients

Define

$$C_{n,k}^{(p)} = [u^k] P_{p,n}^{[3]}(u), \quad 0 \leq k \leq 3.$$

For

$$m \geq k + 1,$$

normalize by the falling factorial:

$$\Pi_{p,k}(m) = \frac{C_{m-1,k}^{(p)}}{(m-1)_{\underline{k}}}.$$

The normalization removes the obvious chronological placement multiplicity and reveals polynomial structure.

23.2 Transport equations

Proposition 23.1 — Transport system.

For the indicated ranges,

$$m\Pi_{p,1}(m+1) = (m-1)\Pi_{p,1}(m) + \ell_m,$$

$$m\Pi_{p,2}(m+1) = (m-2)\Pi_{p,2}(m) + \ell_m\Pi_{p,1}(m) + p\ell_{m-1},$$

and

$$\begin{aligned} m\Pi_{p,3}(m+1) &= (m-3)\Pi_{p,3}(m) + \ell_m\Pi_{p,2}(m) \\ &\quad + p\ell_{m-1}\Pi_{p,1}(m-1) - p. \end{aligned}$$

Proof.

Extract the coefficients of

$$u, \quad u^2, \quad u^3$$

from CCR-3.

Then substitute

$$n = m$$

and divide by the corresponding falling factorial.

The chronological factors cancel exactly, leaving the displayed first-order transport equations. $\square$

The third source can be separated into a connected form.

Since

$$\Pi_{p,1}(m-1) = \frac{\ell_{m-1} - p}{2},$$

the genuinely new third-order source is

$$\mathfrak{T}_{p,3}(m) = p \left[\frac{1}{2}\ell_{m-1}(\ell_{m-1} - p) - 1 \right].$$

This isolates the connected triple correction from propagated lower-order contributions.

23.3 First three fixed columns

Theorem 23.2 — First three fixed columns.

For

$$p \geq 5,$$

$$\Pi_{p,1}(m) = \frac{(p+1)m - 2p}{2},$$

$$\Pi_{p,2}(m) = \frac{p+1}{24} [3(p+1)m^2 - (5p+1)m - 12p],$$

and

$$\Pi_{p,3}(m) = \frac{p+1}{48} F_p(m),$$

where

$$F_p(m) = (p+1)^2 m^3 + (p^2 - 1)m^2 - (34p^2 + 18p)m + 56p^2 - 8p.$$

The boundary values are

$$\Pi_{p,1}(2) = 1,$$

$$\Pi_{p,2}(3) = p + 1,$$

$$\Pi_{p,3}(4) = (p+1)^2.$$

Proof.

The first transport equation is solved directly by the displayed linear polynomial.

Substituting that result into the second equation gives a first-order recurrence with polynomial source; the displayed quadratic satisfies both the recurrence and its boundary value.

Substituting the first two solutions into the third transport equation gives a cubic source. The displayed cubic satisfies the recurrence and boundary value.

Uniqueness follows at each stage from the first-order recurrence. $\square$

The key methodological point is that these formulas arise from transport. They should not be presented as unexplained computer-algebra output.

23.4 Reduction modulo p

Modulo p ,

$$\ell_n = (p+1)n - p \equiv n \pmod{p}.$$

Every pair and triple interaction source carries an explicit factor of p .

Therefore those interaction terms vanish modulo p .
CCR-3 degenerates to the free recurrence

$$P_{p,n}^{[3]}(u) \equiv (1 + nu)P_{p,n-1}^{[3]}(u) \pmod{(p, u^4)}.$$

Iteration gives the universal product.

Theorem 23.3 — Stirling shadow.

For every

$$n \geq 0,$$

$$P_{p,n}^{[3]}(u) \equiv \prod_{j=1}^n (1 + ju) \pmod{(p, u^4)}.$$

Consequently, for $k \leq 3$,

$$C_{m-1,k}^{(p)} \equiv e_k(1, \dots, m-1) \equiv \left[\begin{matrix} m \\ m-k \end{matrix} \right] \pmod{p},$$

where the bracket denotes the unsigned Stirling number of the first kind.
Thus the arithmetic interaction deformation has a universal free shadow.

23.5 Normalized shadow polynomials

The normalized fixed columns reduce to:

$$S_1(m) = \frac{m}{2},$$

$$S_2(m) = \frac{m(3m-1)}{24},$$

$$S_3(m) = \frac{m^2(m-1)}{48}.$$

So, as polynomial congruences,

$$\Pi_{p,k}(m) \equiv S_k(m) \pmod{p} \quad (k = 1, 2, 3).$$

Important notation point

One must not derive these congruences by naively cancelling the falling factorial

$$(m-1)_{\underline{k}}$$

inside $\mathbb{F}_p$.

That denominator may vanish at exceptional residues.

The safe argument is:

1. derive the exact polynomial formulas over $\mathbb{Q}[m]$;
2. reduce those polynomial formulas modulo p .

Because

$$2, \quad 24, \quad 48$$

are units for $p \geq 5$, the reduction is legitimate.

This avoids illegal pointwise modular division.

23.6 Why the third column is the first singular one

The third normalized shadow is

$$S_3(m) = \frac{m^2(m-1)}{48}.$$

Therefore modulo p , the third-column zero set contains:

- a simple root at

$$m \equiv 1 \pmod{p};$$

- a double root at

$$m \equiv 0 \pmod{p}.$$

The exact cubic deformation $F_p(m)$ must decide what happens to these two shadow branches.

The next chapter shows that they behave completely differently.

The simple root survives integrally.

The double root becomes extinct in $\mathbb{Z}_p$.

Chapter 24

The Third-Column p -Adic Split

The third normalized column is

$$\Pi_{p,3}(m) = \frac{p+1}{48}F_p(m),$$

where

$$F_p(m) = (p+1)^2m^3 + (p^2-1)m^2 - (34p^2+18p)m + 56p^2 - 8p.$$

Modulo p ,

$$\boxed{F_p(m) \equiv m^2(m-1) \pmod{p}.$$

Thus the shadow has two different types of roots:

$$m \equiv 1 \pmod{p}$$

is simple, while

$$m \equiv 0 \pmod{p}$$

is double.

The exact deformation separates them sharply.

24.1 The simple branch

At

$$m = 1,$$

$$F_p(1) \equiv 0 \pmod{p},$$

and

$$\boxed{F'_p(1) \equiv 1 \pmod{p}.}$$

The derivative is a unit.

Hensel's lemma therefore applies in its simplest form.

Theorem 24.1 — Simple Hensel tower.

There exists a unique

$$\boxed{\alpha_p \in \mathbb{Z}_p}$$

such that

$$\alpha_p \equiv 1 \pmod{p}$$

and

$$F_p(\alpha_p) = 0.$$

Its expansion begins

$$\boxed{\begin{aligned} \alpha_p = & 1 + 24p - 888p^2 + 58008p^3 \\ & - 4715256p^4 + 428583192p^5 + O(p^6). \end{aligned}}$$

Moreover, throughout the residue class

$$m \equiv 1 \pmod{p},$$

$$\boxed{\nu_p(F_p(m)) = \nu_p(m - \alpha_p).}$$

Proof.

Hensel gives existence and uniqueness.

Factor in $\mathbb{Z}_p[m]$:

$$F_p(m) = (m - \alpha_p)G_p(m).$$

Since

$$G_p(\alpha_p) = F'_p(\alpha_p)$$

is a unit, $G_p(m)$ remains a unit throughout the residue class

$$m \equiv \alpha_p \equiv 1 \pmod{p}.$$

Therefore the valuation is exactly the distance to the Hensel root. $\square$

Thus the simple shadow branch survives exactly as expected.

24.2 The double shadow branch

Now set

$$m = pu.$$

Direct substitution gives

$$\frac{F_p(pu)}{p} \equiv -8 \pmod{p}.$$

Since $p \geq 5$,

$$-8$$

is nonzero modulo p .

Therefore:

Theorem 24.2 — Singular extinction.

For every

$$u \in \mathbb{Z}_p,$$

$$\nu_p(F_p(pu)) = 1.$$

Consequently, although

$$m \equiv 0 \pmod{p}$$

is a double root of the Stirling shadow, it has **no lift to a root modulo p^2** .

We call this phenomenon

singular extinction.

The term means exactly:

multiple root mod p but no integral branch survives mod p^2 .

It does not mean that the corresponding algebraic roots disappear from the algebraic closure.

Chapter 25 will show where they go.

24.3 Generic residue classes

If

$$m \not\equiv 0, 1 \pmod{p},$$

then

$$F_p(m) \not\equiv 0 \pmod{p}.$$

Hence

$$\boxed{\nu_p(F_p(m)) = 0.}$$

Thus the residue classes split into three exact arithmetic regimes:

1. generic;
2. simple Hensel;
3. singularly extinct.

24.4 The scaled third column

Define

$$\boxed{Q_{m,3}^{\text{CCR}} = p^{m-4} \Pi_{p,3}(m), \quad m \geq 4.}$$

Because

$$\frac{p+1}{48}$$

is a p -adic unit for $p \geq 5$,

$$\nu_p(\Pi_{p,3}(m)) = \nu_p(F_p(m)).$$

The previous three cases therefore give the complete valuation law.

Theorem 24.3 — Complete third-column valuation law.

For

$$m \geq 4,$$

$$\boxed{\nu_p(Q_{m,3}^{\text{CCR}}) = \begin{cases} m-4, & m \not\equiv 0, 1 \pmod{p}, \\ m-3, & m \equiv 0 \pmod{p}, \\ m-4 + \nu_p(m - \alpha_p), & m \equiv 1 \pmod{p}. \end{cases}}$$

This theorem is best read after the branch analysis.
 The three cases are not an arbitrary valuation table.
 They are exactly:

generic unit

one forced p -factor from singular extinction

Hensel distance to the unique integral root.
--

24.5 Independent consistency checks

The monograph audit independently checked the arithmetic at several levels:

- the fixed-column formulas were regenerated symbolically through several rows;
- the displayed formulas matched 15 tested symbolic column/boundary instances;
- substituting the displayed Hensel expansion into F_p cancels every coefficient through p^5 ;
- digitwise Hensel lifting for

$$p = 5, 7, 11, 13, 17, 19$$

agrees through modulus p^6 ;

- hundreds of sample valuation checks agree with both the simple-branch and extinction laws.

These computations support the proof chain.
 They do not replace it.

24.6 The unresolved algebraic question

The double shadow root has now disappeared from the integral p -adic locus:

$$0 \bmod p \longrightarrow \emptyset \bmod p^2.$$

But F_p is cubic.

One root has become the integral Hensel root α_p .

Two algebraic roots remain.

The discriminant will show that they have not disappeared.

They have moved into a ramified quadratic extension, with half-integral valuation.

Chapter 25

Ramified Recovery

The simple shadow root has survived inside

$$\mathbb{Z}_p.$$

The double shadow root has undergone singular extinction.
The final question of the arithmetic offshoot is:

Where are the two missing algebraic roots?

The answer is encoded by the cubic discriminant.

25.1 Exact discriminant factorization

Theorem 25.1 — Discriminant and ramified field.

The discriminant of F_p factors as

$$\text{Disc}(F_p) = 4p(p+1)^2 R_p,$$

where

$$R_p = 9801p^5 + 22672p^4 + 32502p^3 + 14360p^2 - 927p - 8.$$

Modulo p ,

$$R_p \equiv -8 \pmod{p}.$$

Therefore

$$\nu_p(\text{Disc}(F_p)) = 1.$$

The discriminant has odd valuation, so the residual quadratic factor after removing the simple Hensel root is ramified.

25.2 Removing the integral root

Factor

$$F_p(m) = (m - \alpha_p)Q_p(m)$$

over $\mathbb{Z}_p$.

The root α_p is simple and integral.

The remaining quadratic Q_p has discriminant with the same square class as the cubic discriminant, because the resultant contribution from the linear/quadratic splitting is a square.

Now

$$R_p/(-8) \equiv 1 \pmod{p}.$$

By Hensel's lemma,

$$R_p/(-8)$$

is a square in

$$\mathbb{Z}_p^\times.$$

The factors

$$4 \quad \text{and} \quad (p+1)^2$$

are also squares.

Hence the discriminant square class is

$$\boxed{[-2p]}.$$

The quadratic splitting field is therefore

$$\boxed{\mathbb{Q}_p(\sqrt{-2p})}.$$

25.3 Half-integral roots

The quadratic factor is Eisenstein after the integral branch is removed.

Its two roots

$$\beta_{p,+}, \quad \beta_{p,-}$$

therefore satisfy

$$\boxed{\nu_p(\beta_{p,\pm}) = \frac{1}{2}}.$$

Thus the double shadow root has not vanished algebraically.
Instead:

$$\boxed{\text{double root mod } p}$$

splits in the exact deformation into

$$\boxed{\text{two ramified roots of valuation } 1/2.}$$

This is the conceptual resolution of singular extinction.

Integral lifting fails because the two branches leave the integral valuation lattice.

25.4 Ramified square class by $p \bmod 8$

Corollary 25.2 — Ramified square class.

The ramified field is exactly

$$\boxed{\mathbb{Q}_p(\sqrt{-2p}).}$$

Its standard square-class description depends on

$$p \bmod 8.$$

For the residue classes in which -2 is a square unit, this is the standard p -ramified class.

For the remaining residue classes, it is the other ramified quadratic square class represented by a nonsquare unit times p .

This corollary is classical square-class arithmetic applied to the exact discriminant already proved.

25.5 The complete third-order arithmetic picture

The results of Chapters 22–25 can now be summarized without adding a new proof.

Corollary 25.3 — CCR-3 arithmetic geometry.

The third-order chronological carry refinement has the following structure.

1. Chronological locality allows only singleton, adjacent-pair, and consecutive-triple connected supports through order three.
2. The chosen arithmetic polarization produces the CCR-3 recurrence and its Hessenberg/operator realization.

3. The first three normalized columns are explicit polynomials.
4. Modulo p , all interaction terms disappear and

$$P_{p,n}^{[3]}(u) \equiv \prod_{j=1}^n (1 + ju) \pmod{(p, u^4)}.$$

5. The normalized columns reduce to near-diagonal first-kind Stirling polynomials.
6. In the third column:

- the simple shadow root

$$1 \pmod{p}$$

lifts uniquely to

$$\alpha_p \in \mathbb{Z}_p;$$

- the double shadow root

$$0 \pmod{p}$$

undergoes singular extinction;

- the two missing algebraic branches survive with valuation

$$1/2$$

over

$$\mathbb{Q}_p(\sqrt{-2p}).$$

This is the complete third-order arithmetic offshoot.

25.6 Novelty boundary

Nothing in this chapter requires claiming novelty for the ingredients individually.

The following are classical or established areas:

- carry processes;
- Holte-type carry matrices;
- Stirling numbers;
- p -adic valuations of Stirling numbers;
- Hensel lifting;
- Eisenstein irreducibility;
- discriminants;

- Newton polygons;
- ramified quadratic extensions;
- lower-Hessenberg production matrices.

The family-specific contribution is the chain

chronological local interactions $\rightarrow$ fixed-column polynomialization
 $\rightarrow$ Stirling shadow $\rightarrow$ simple lifting / singular extinction $\rightarrow$ exact
 ramified recovery.

A complete external priority audit of the full family-specific chain remains a publication-stage task.

25.7 Do not extrapolate CCR-3 blindly

The recurrence is intentionally third-order.

At order four, a new connected interaction can appear.

Therefore the next arithmetic research question is not:

“What happens if we simply continue the same three-band formula?”

It is:

“What is the genuine fourth-order connected chronological motif, and what arithmetic weight does the intended refinement assign to it?”

That question is open within this branch.

25.8 Re-entry to the main DPDA story

The arithmetic offshoot is now complete.

We return to the main machine problem.

The monograph began with the original 2:1 center/left/right DPDA.

Chapters 2–8 extracted its shortest-demand semantics and produced compact residual-faithful machines.

But those compact machines changed the physical representation.

A different question remained:

Can the actual operational idea hidden in the original $C/L/R$
 architecture be generalized to every $p:q$?

That question produces the real-time centered construction.

From Chapter 26 onward, the main line will remain uninterrupted:

general real-time $p:q \rightarrow$ Structure Theory $\rightarrow$ Uniqueness/Lift Theory
 $\rightarrow$ General Theory.

PART VIII

Returning to the Machine: General Real-Time $p:q$

Research provenance. The uniform centered real-time literal-demand construction for all coprime positive p, q is developed in [8].

Chapter 26

Why Static Two-Visibility Fails

We now return to pushdown machines.

This is not another residual-counter construction.

The goal is more specific:

Reconstruct, for every coprime $p:q$, the operational idea hidden in the original center/left/right machine.

The original 2:1 design had three striking features:

1. a literal completion demand was physically visible in the stack;
2. a mismatch could temporarily move information into finite control;
3. a later long replacement could rematerialize that information.

The general construction will preserve this style rather than replace it by the compact Euclidean residual encoding of Chapters 7–8.

That is why its resource count is different.

26.1 We do not re-prove minimum demand

The semantic foundation is already complete.

For

$$L_{p,q} = \{w : qN_1(w) = pN_0(w)\},$$

the residual is

$$\rho = qN_1 - pN_0.$$

For every residual r , the unique minimum completion demand is

$$\mu(r) = (x_r, y_r),$$

and the primitive zero-weight kernel is

$$K = (q, p).$$

Every nonnegative vector of the same weight differs from the minimum demand by a nonnegative multiple of K .

Those results are imported from the earlier chapters.

The new problem is physical ordering.

26.2 A two-visible canonical word

For a demand vector

$$(x, y),$$

define

$$\eta(x, y) = \begin{cases} (10)^y 0^{x-y}, & x \geq y, \\ (10)^x 1^{y-x}, & y > x. \end{cases}$$

This word has an attractive property.

If both symbols occur, the stack begins with

$$10.$$

More generally, every suffix that remains mixed begins with either

$$10$$

or

$$01.$$

Thus the presence of both demand types is always visible within two stack cells.

This is the general literal-demand analogue of the original 2:1 depth-two phenomenon.

But static visibility is not enough.

26.3 The 3:2 obstruction

Consider the canonical stack

1010.

Both demand symbols are visible immediately.

Suppose the next input is 0, but the required zero is the **second** visible cell.

A naive LIFO strategy could:

1. pop the top 1;
2. delete the 0;
3. immediately restore the displaced 1.

This produces

$$1010 \longrightarrow 110.$$

The surviving zero has moved to depth three.

Two-cell visibility has been destroyed.

The problem is not lack of static information.

It is chronology.

26.4 Static two-visible obstruction

Lemma 26.1 — Static two-visible obstruction.

Suppose a mixed binary stack contains at least two zeros and two ones and begins with

$$10.$$

No purely static strategy can support both deletion directions while:

1. deleting either visible demand type;
2. immediately restoring a displaced first-cell blocker;
3. guaranteeing that a surviving mixed suffix still exposes both symbol types in its first two cells.

Proof.

If deleting the top 1 must leave a mixed two-visible suffix, then the next two surviving symbols must differ. Therefore the third original symbol must be 1.

If instead the second-cell 0 is deleted and the first 1 is immediately restored, preserving two-visibility requires the next surviving symbol beneath it to be 0. Therefore the third original symbol must be 0.

The same third cell cannot satisfy both conditions. $\square$

This lemma isolates the exact reason the original architecture needs side-state memory.

A LIFO blocker sometimes must leave the stack temporarily.

26.5 SEARCH

The repair is:

SEARCH = move one displaced blocker into finite control.

The blocker is **not** immediately restored.

Its semantic effect is carried as finite-control debt while the remaining physical stack continues to be popped.

A later WRITE will materialize that debt.

This is the general version of the role played by L and R in the original machine.

The crucial design question becomes:

How much signed debt must finite control be able to carry before
WRITE becomes safe?

That will produce the centered state rail.

Before building the rail, we need to formalize what a debt state means semantically.

Chapter 27

Debt, Virtual Demand, and Minimal WRITE

Let

$$m = p + q.$$

A one-cell mismatch changes the relationship between the physical stack and the semantic demand by one kernel-length unit m .

This motivates a centered integer debt coordinate.

27.1 Virtual demands

For

$$d \in \mathbb{Z},$$

define the virtual demand

$$\boxed{V_d = \mu(dm).}$$

Its residual weight is

$$W(V_d) = dm.$$

A debt state D_d therefore carries finite-control potential

$$\boxed{\Phi(D_d) = dm.}$$

Let a physical logical stack S have Parikh vector

$$P(S) = (N_0(S), N_1(S)).$$

The stack and state together may contain more material than the minimum demand.

The semantic projection removes any complete kernel packets.

27.2 Virtual-demand projection

Define kernel reduction

$$\text{red}_{p,q}(x, y)$$

to mean repeated subtraction of

$$K = (q, p)$$

while both coordinates remain nonnegative.

The result is exactly the unique minimum demand of the same residual weight.

Definition 27.1 — Rich-to-demand projection.

For configuration

$$(D_d, S),$$

define

$$\boxed{\Theta(D_d, S) = \text{red}_{p,q}(P(S) + V_d).}$$

Theorem 27.2 — Virtual-demand projection.

If

$$\boxed{\rho = W(P(S)) + dm,}$$

then

$$\boxed{\Theta(D_d, S) = \mu(\rho).}$$

Proof.

The nonnegative vector

$$P(S) + V_d$$

has weight

$$W(P(S)) + W(V_d) = W(P(S)) + dm = \rho.$$

Kernel reduction of any nonnegative vector of weight ρ is the unique minimum demand $\mu(\rho)$. $\square$

This theorem is one of the main semantic bridges in the monograph.

It says:

rich state+stack configuration $\longrightarrow$ unique shortest-demand semantics.

The rich machine may retain normalization history.

Its semantic projection does not.

27.3 The WRITE equation

Suppose one transition:

- consumes input a ;
- removes visible top symbol z ;
- changes debt state from D_s to D_t ;
- writes replacement word u .

Let

$$e_0 = (1, 0), \quad e_1 = (0, 1).$$

Ignoring the unchanged stack suffix, residual preservation forces

$$P(u) = e_z - e_a + V_s - V_t + \kappa K$$

for some integer κ .

At the bottom marker, the term e_z is omitted.

Thus every invariant-preserving rewrite lies on one affine kernel ray.

27.4 Minimal kernel carry

Let

$$B = e_z - e_a + V_s - V_t.$$

The first kernel shift making both coordinates nonnegative is

$$\kappa_* = \max \left(\left\lceil -\frac{B_0}{q} \right\rceil, \left\lceil -\frac{B_1}{p} \right\rceil \right).$$

Theorem 27.3 — Minimal kernel-carry rewrite.

Among all nonnegative invariant-preserving replacement vectors, the unique shortest one is

$$R_* = B + \kappa_* K.$$

Every other solution has the form

$$R_* + \ell K, \quad \ell \geq 1,$$

and is longer by exactly

$$\ell(p + q)$$

symbols.

Proof.

All integer solutions differ by a multiple of the primitive kernel K .

The displayed κ_* is the least shift that places both coordinates in the nonnegative quadrant.

Every further kernel packet adds

$$q + p = m$$

symbols. $\square$

When a physical word is needed, we write the vector in canonical order using η .

27.5 Why the historical long writes existed

The original 2:1 table contained replacements such as:

$$10, \quad 11, \quad 000, \quad 111111.$$

Before the general theory, these can look ad hoc.

The WRITE equation changes their interpretation.

A long replacement is determined by:

1. the consumed input symbol;
2. the visible stack symbol removed;
3. the source debt potential;
4. the target debt potential;
5. the least kernel carry needed to make the replacement vector nonnegative.

Thus the long words are not arbitrary repair strings.

They are shortest solutions of a constrained residual-preservation equation.

This is one of the main retrospective explanations of the original machine.

27.6 From arbitrary debt to a finite rail

The virtual-demand definition allows every

$$d \in \mathbb{Z}.$$

A finite-state DPDA cannot.

The construction therefore needs a bounded debt interval.

Set

$$h_{\text{ctr}} = \left\lfloor \frac{m-1}{2} \right\rfloor.$$

The centered core will use exactly

$$D_{-h_{\text{ctr}}}, \dots, D_{-1}, D_0, D_1, \dots, D_{h_{\text{ctr}}}.$$

with

$$D_0 = C.$$

Why is this width sufficient?

Because the canonical word created at every WRITE has bounded minority count.

Why is it exact in the deferred architecture?

Because one common canonical witness can exhaust the physical stack while leaving every debt value in this interval purely in finite control.

Those are the two halves of Chapters 28 and 29.

Chapter 28

The Centered Real-Time $p:q$ DPDA

Set

$$m = p + q, \quad h_{\text{ctr}} = \left\lfloor \frac{m-1}{2} \right\rfloor.$$

The centered state set is

$$Q = \{D_{-h_{\text{ctr}}}, \dots, D_{-1}, D_0, D_1, \dots, D_{h_{\text{ctr}}}\},$$

with

$$D_0 = C.$$

The stack alphabet is

$$\{0, 1, Z_0\}.$$

The bottom marker Z_0 remains fixed throughout binary-input processing and is removed only by the final right-endmarker acceptance move.

The construction has four binary transition roles:

$$\text{MATCH, SEARCH, outer WRITE, bottom WRITE.}$$

Every binary move consumes exactly one input symbol.

There are no ε -moves during binary processing.

28.1 The centered rail

MATCH

For every

$$|d| \leq h_{\text{ctr}}$$

and $z \in \{0, 1\}$,

$$(D_d, zT) \xrightarrow{z} (D_d, T).$$

A matching input simply cancels one physically present demand cell.

Interior SEARCH

For

$$d < h_{\text{ctr}},$$

$$(D_d, 0T) \xrightarrow{1} (D_{d+1}, T).$$

For

$$d > -h_{\text{ctr}},$$

$$(D_d, 1T) \xrightarrow{0} (D_{d-1}, T).$$

A mismatch still pops the visible cell, but instead of restoring it the state debt moves one step.

Outer-boundary WRITE

Define

$$U_+ = \eta(\mu((h_{\text{ctr}} + 1)m)),$$

$$U_- = \eta(\mu(-(h_{\text{ctr}} + 1)m)).$$

Then

$$(D_{h_{\text{ctr}}}, 0T) \xrightarrow{1} (C, U_+T),$$

and

$$(D_{-h_{\text{ctr}}}, 1T) \xrightarrow{0} (C, U_T).$$

At the outer debt boundary, another outward mismatch triggers physical rematerialization and resets debt to zero.

Bottom WRITE

For every

$$|d| \leq h_{\text{ctr}},$$

$$(D_d, Z_0) \xrightarrow{1} (C, \eta(\mu(dm + q))Z_0),$$

and

$$(D_d, Z_0) \xrightarrow{0} (C, \eta(\mu(dm - p))Z_0).$$

When no pre-existing demand remains, the next input directly creates the exact new minimum-demand word.


Acceptance

A mandatory right endmarker is allowed only at

$$(C, Z_0).$$

It removes Z_0 and accepts by empty stack.

No other reachable configuration has a right-endmarker acceptance move.



Adjacent mismatch rules repeat along the rail.

Matching $0/0$ or $1/1$ pops and keeps the same debt state.

Figure 28.1: Schematic centered debt rail. Labels $1/0$ and $0/1$ denote input/top mismatches. Interior mismatches pop the visible cell and move debt by one; an outward mismatch at $\pm h_{\text{ctr}}$ triggers WRITE back to the center.

The figure should visually emphasize:

- adjacent mismatch SEARCH moves;
- center state C ;
- outer WRITE arcs back to C ;
- MATCH as state-preserving pops.

28.2 Residual invariant

Let the logical stack be S .

Theorem 28.1 — Residual invariant.

After every processed binary prefix, if the centered configuration is

$$(D_d, SZ_0),$$

then

$$\rho = W(P(S)) + dm.$$

Proof.

Initially,

$$(C, Z_0)$$

has residual zero.

A matching 0-pop decreases stack weight by p , exactly matching the input-0 residual update.

A matching 1-pop increases stack weight by q .

For mismatch 1/0, popping one physical zero changes stack weight by $-p$, while moving

$$d \rightarrow d + 1$$

adds state potential

$$m = p + q.$$

The net change is

$$-p + m = q,$$

exactly the input-1 residual update.

The 0/1 case is symmetric:

$$q - m = -p.$$

At an outer WRITE, the replacement word is defined to carry exactly the missing boundary potential.

At bottom WRITE, the target demand is defined directly from the new residual.

Therefore every binary transition preserves the invariant. $\square$

Combined with Chapter 27:

$$\boxed{\Theta(D_d, S) = \mu(\rho)}$$

for every reachable configuration.

28.3 Why the debt rail can stop at h_{ctr}

The key is the **minority size**

$$\text{minc}(x, y) = \min(x, y).$$

Lemma 28.2 — WRITE-frontier minority bound.

Let

$$h_{\text{ctr}} = \left\lfloor \frac{m-1}{2} \right\rfloor.$$

Then:

1. for every $|d| \leq h_{\text{ctr}}$, both bottom-WRITE demands

$$\mu(dm + q)$$

and

$$\mu(dm - p)$$

have minority size at most h_{ctr} ;

2. if

$$\mu((h_{\text{ctr}} + 1)m) = (x_+, y_+),$$

then

$$\boxed{y_+ \leq h_{\text{ctr}} < x_+;}$$

3. if

$$\mu(-(h_{\text{ctr}} + 1)m) = (x_-, y_-),$$

then

$$\boxed{x_- \leq h_{\text{ctr}} < y_-}.$$

Proof.

For the bottom-WRITE family, both residual forms can be written

$$r = jm + q$$

with

$$-h_{\text{ctr}} - 1 \leq j \leq h_{\text{ctr}}.$$

Let

$$\mu(r) = (x, y), \quad \ell = x + y.$$

Modulo m , since

$$p \equiv -q \pmod{m},$$

the completion equation gives

$$-q\ell \equiv q \pmod{m}.$$

Because

$$\gcd(q, m) = 1,$$

$$\ell \equiv -1 \pmod{m}.$$

Hence

$$\ell = km - 1.$$

If both x and y were at least $h_{\text{ctr}} + 1$, then $\ell > m - 1$, so $k \geq 2$.

But $\mu(r)$ is kernel-free, so either

$$x < q$$

or

$$y < p.$$

Each alternative, together with $x, y \geq h_{\text{ctr}} + 1$, forces the quotient index j outside the allowed interval.

Contradiction.

For the outer positive boundary, write $D = h_{\text{ctr}} + 1$. Integer solutions to weight Dm have form

$$x = D + qk, \quad y = pk - D.$$

The minimum nonnegative solution uses the first feasible k , giving

$$y \leq h_{\text{ctr}}$$

and

$$x \geq h_{\text{ctr}} + 1.$$

The negative case is symmetric. $\square$

This lemma explains the centered width geometrically.

Every freshly written canonical word contains at most h_{ctr} copies of its minority symbol.

The debt rail is exactly wide enough to search through that minority frontier.

28.4 Boundary guard

A subtle problem remains.

Suppose the machine reaches $D_{h_{\text{ctr}}}$.

Can the next outward mismatch occur while the surviving suffix is still mixed?

If so, writing a new canonical prefix in front of that suffix could destroy the normal form.

The minority bound prevents this.

Lemma 28.3 — Boundary guard.

Assume the current stack is a suffix of a canonical word

$$\eta(x, y)$$

with

$$\min(x, y) \leq h_{\text{ctr}}.$$

Then:

1. if the state is $D_{h_{\text{ctr}}}$ and the stack begins with 0, an outward mismatch 1/0 can occur only on a unary suffix

$$\boxed{0^{t+1};}$$

2. symmetrically, at $D_{-h_{\text{ctr}}}$, an outward 0/1 mismatch can occur only on a unary suffix

$$\boxed{1^{t+1}.}$$

Proof.

Consider the positive case.

Reaching debt h_{ctr} requires at least h_{ctr} positive mismatch events.

Each such mismatch has popped one zero from the original canonical stack.

Suppose the surviving suffix begins with 0 but still contains a later 1.

Then the current zero lies inside the alternating part of $\eta(x, y)$, not in its final unary tail.

Every previously popped zero in that alternating region has a distinct preceding one that must also have been popped.

Together with:

- the current surviving zero;
- the later surviving one;

the original word would contain at least

$$h_{\text{ctr}} + 1$$

zeros and at least

$$h_{\text{ctr}} + 1$$

ones.

This contradicts

$$\min(x, y) \leq h_{\text{ctr}}.$$

Therefore the suffix must be unary.

The negative case is symmetric. $\square$

This is the LIFO safety theorem of the construction.

28.5 Canonical segments

Theorem 28.4 — Canonical-segment invariant.

Immediately after every WRITE, the machine is in the center state with stack

$$\boxed{\eta(\mu(r))Z_0}$$

for its current residual r .

The corresponding demand has minority size at most h_{ctr} .

Between two WRITES, the logical stack is a suffix of that canonical word.

Proof.

Initially the logical stack is empty.

A bottom WRITE is canonical by definition, and Lemma 28.2 gives the minority bound.

Between WRITES, MATCH and interior SEARCH only pop one pre-existing work symbol, so the stack remains a suffix of the last canonical WRITE word.

Now consider a positive outer WRITE.

By Lemma 28.3, its source suffix is

$$0^{t+1}.$$

Write

$$\mu((h_{\text{ctr}} + 1)m) = (x, y).$$

Lemma 28.2 gives

$$y \leq h_{\text{ctr}} < x.$$

Thus the canonical word $\eta(x, y)$ ends in 0, and the WRITE produces

$$\eta(x, y)0^t = \eta(x + t, y).$$

The residual invariant shows that this vector has exactly the new residual weight.

Moreover, $y < p$, so no kernel packet can be subtracted; the vector is therefore already the minimum demand.

The negative boundary is symmetric. $\square$

The entire reachable-store theorem has now been reduced to episodes:

WRITE creates a canonical word;

MATCH/SEARCH consume a suffix of it;

boundary WRITE can occur only after the mixed part is gone.

28.6 Two-cell qualitative visibility

Corollary 28.5 — Two-cell visibility.

For every reachable logical stack S ,

$$\text{alph}(S) = \text{alph}(\text{pref}_2(S)).$$

Equivalently, if both 0 and 1 occur anywhere in S , then both occur among the first two cells.

Proof.

Every canonical word $\eta(x, y)$, and every suffix of it, avoids

$$001 \quad \text{and} \quad 110.$$

Therefore a mixed suffix cannot begin with two equal symbols. It begins with

$$10 \quad \text{or} \quad 01.$$

Apply Theorem 28.4. $\square$

This is qualitative, not quantitative.

The first two cells reveal which demand types remain.

They do not reveal the unbounded magnitudes.

28.7 Exact WRITE normalization

Theorem 28.6 — Exact WRITE normalization.

After every bottom or outer-boundary WRITE,

$$(D_d, S) = (C, \eta(\mu(\rho))),$$

where ρ is the residual after the consumed binary symbol.

WRITE therefore does more than preserve the invariant.

It returns the physical stack to the exact minimum-demand normal form.

This result is the normalization analogue of the original endmarker materialization theorem, but now it occurs online.

28.8 Zero residual

Theorem 28.7 — Zero-residual correctness.

At the end of every binary prefix,

$$\boxed{\rho = 0 \iff (D_d, S) = (C, \varepsilon).}$$

Proof.

The reverse implication follows from the invariant.

For the forward direction, consider a nonempty processed prefix.

Immediately after the most recent WRITE, the machine was at

$$(C, \eta(\mu(r_0)))$$

for some residual r_0 .

Write

$$\mu(r_0) = (x, y).$$

Since then, each MATCH or SEARCH step has popped exactly one pre-existing demand cell.

Suppose the intervening input contains:

$$u$$

zeros and

$$v$$

ones, so

$$u + v \leq x + y.$$

If the current residual is zero, this segment completes r_0 .

By minimality of $\mu(r_0)$, every completion requires length at least

$$x + y.$$

Hence

$$u + v = x + y.$$

The entire canonical stack has been consumed.

So

$$S = \varepsilon.$$

The residual invariant then gives

$$dm = 0,$$

hence

$$d = 0.$$

Therefore the state is C . $\square$

This proof is important because acceptance is obtained from shortest-demand minimality, not from a large case analysis of reachable stacks.

28.9 Main construction theorem

Theorem 28.8 — Real-time centered rich DPDA.

For every coprime

$$p, q \geq 1,$$

the centered schema defines a deterministic pushdown machine with:

$$|Q| = 2 \left\lfloor \frac{p+q-1}{2} \right\rfloor + 1$$

semantic core states and stack alphabet

$$\{0, 1, Z_0\}.$$

Every binary transition consumes exactly one symbol.

No ε -transition is used during binary processing.

With the mandatory right endmarker, the machine recognizes

$$L_{p,q}.$$

Furthermore:

1. every WRITE returns to the exact minimum-demand normal form;
2. every reachable stack has two-cell qualitative visibility;
3. the residual is exactly decoded by the state-plus-stack invariant.

The state count is

$$|Q| = \begin{cases} p+q, & p+q \text{ odd,} \\ p+q-1, & p+q \text{ even.} \end{cases}$$

This is the exact size of the centered construction.

It is not yet an unrestricted state lower bound.

28.10 Shortest-demand quotient

The centered machine is physically richer than a residual counter.

Two configurations can store the same semantic demand while distributing it differently between:

- physical stack;
- finite-control debt.

Corollary 28.9 — Shortest-demand quotient.

For reachable configurations, define

$$c_1 \sim_{\Theta} c_2 \iff \Theta(c_1) = \Theta(c_2).$$

Then the residual/minimum-demand transition system is a semantic quotient of the rich configuration graph.

The quotient forgets normalization chronology.

The physical fiber may remember it.

That observation will later become a central theorem family in the Uniqueness/Lift Theory.

For now, it is already visible concretely.

Chapter 29

What “Optimal” Means in the Deferred-SEARCH Architecture

The centered construction uses

$$2h_{\text{ctr}} + 1$$

debt states, where

$$h_{\text{ctr}} = \left\lfloor \frac{p + q - 1}{2} \right\rfloor.$$

Is that number necessary?

The answer is yes—but only for a precisely stated physical interface.

This theorem must not be confused with the raw residual-faithful frontier of Chapter 8.

The two results optimize different representation categories.

29.1 Canonical deferred-SEARCH model

Definition 29.1 — Canonical deferred-SEARCH realization.

A realization belongs to the canonical deferred-SEARCH model if:

1. the normalized minimum demand

$$\mu(r)$$

is physically represented by the canonical word

$$\eta(\mu(r));$$

2. during a SEARCH episode, each binary transition consuming a pre-existing demand cell removes exactly one such cell;
3. a mismatch is stored only in finite control—no auxiliary debt marker is inserted into the remaining pre-existing word;
4. in the witness episode used below, the transition consuming the final pre-existing demand cell performs no materializing push and leaves the same bottom marker Z_0 ;
5. debt materialization is deferred to a later transition.

The last condition is what makes the model **deferred**.
An eager implementation lies outside this model.

29.2 Residual distinguishability

Lemma 29.2 — Residual distinguishability.

If two prefixes have distinct residuals

$$r_1 \neq r_2,$$

then a common binary continuation followed by the right endmarker distinguishes them.

Proof.

Choose a word z whose Parikh vector is

$$\mu(r_1).$$

Appending z sends the first residual to zero.

The second residual becomes

$$r_2 - r_1 \neq 0.$$

The right endmarker is accepted after the first prefix and rejected after the second. $\square$

This is a language property.

It is not assumed by the architecture.

29.3 One witness containing enough of both symbols

Define

$$\nu_{h_{\text{ctr}}} = \begin{cases} (h_{\text{ctr}}, h_{\text{ctr}}), & m = p + q \text{ odd,} \\ (h_{\text{ctr}}, h_{\text{ctr}} + 1), & m \text{ even.} \end{cases}$$

Its canonical word is

$$(10)^{h_{\text{ctr}}}$$

in the odd case and

$$(10)^{h_{\text{ctr}}} 1$$

in the even case.

Lemma 29.3 — Canonical witness.

$$\nu_{h_{\text{ctr}}} = \mu(W(\nu_{h_{\text{ctr}}})).$$

Thus $\nu_{h_{\text{ctr}}}$ is kernel-free and contains at least h_{ctr} zeros and h_{ctr} ones.

Proof.

If m is odd and a kernel packet

$$(q, p)$$

could be subtracted from $(h_{\text{ctr}}, h_{\text{ctr}})$, then

$$p \leq h_{\text{ctr}}, \quad q \leq h_{\text{ctr}},$$

so

$$m = p + q \leq 2h_{\text{ctr}} = m - 1,$$

impossible.

If m is even, subtraction from $(h_{\text{ctr}}, h_{\text{ctr}} + 1)$ would require

$$q \leq h_{\text{ctr}}, \quad p \leq h_{\text{ctr}} + 1,$$

giving again

$$m \leq 2h_{\text{ctr}} + 1 = m - 1.$$

Contradiction. $\square$

This single canonical block contains enough of both symbol types to generate every debt in the centered interval.

29.4 Exhaust the same stack, leave different debts

Lemma 29.4 — Debt distinguishability witness.

Starting from the same canonical witness word $\eta(\nu_{h_{\text{ctr}}})$, a deferred SEARCH episode can consume the entire pre-existing work stack while leaving any debt

$$d \in \{-h_{\text{ctr}}, -h_{\text{ctr}} + 1, \dots, h_{\text{ctr}}\}.$$

The resulting residual is

$$dm.$$

Proof.

For

$$d \geq 0,$$

choose exactly d witness zero-cells and mismatch them by reading 1 instead of 0.

Match every other pre-existing cell.

Debt increases monotonically from 0 to d .

For

$$d < 0,$$

mismatch exactly $|d|$ witness one-cells by reading 0.

Because the witness contains at least h_{ctr} copies of both symbols, every debt in the interval is possible.

Let the witness contain x zeros and y ones.

If A zero-cells are mismatched by input 1 and B one-cells are mismatched by input 0, then

$$d = A - B.$$

The consumed input contains

$$x - d$$

zeros and

$$y + d$$

ones.

Starting from residual

$$px - qy,$$

the final residual is

$$px - qy + q(y + d) - p(x - d) = d(p + q) = dm.$$

□

By the deferred-model assumption, after the last witness cell is consumed, every case has the **same physical stack**:

$$\boxed{Z_0.}$$

Only finite control can distinguish the residuals.

29.5 Exact deferred width

Theorem 29.5 — Exact deferred-SEARCH control width.

Every canonical deferred-SEARCH realization requires at least

$$\boxed{2h_{\text{ctr}} + 1}$$

finite-control states.

The centered construction uses exactly $2h_{\text{ctr}} + 1$.

Therefore

$$\boxed{|Q|_{\min}^{\text{deferred}} = 2 \left\lfloor \frac{p + q - 1}{2} \right\rfloor + 1}$$

or equivalently

$$\boxed{|Q|_{\min}^{\text{deferred}} = \begin{cases} p + q, & p + q \text{ odd,} \\ p + q - 1, & p + q \text{ even.} \end{cases}}$$

Proof.

Lemma 29.4 produces

$$2h_{\text{ctr}} + 1$$

different residuals:

$$-hm, \dots, -m, 0, m, \dots, hm.$$

In every case the work stack has been exhausted and the remaining physical stack is the same bottom marker Z_0 .

If two distinct residuals shared the same finite-control state, the two instantaneous configurations would be physically identical.

But Lemma 29.2 gives a common continuation that must distinguish the residuals.

Determinism from identical configurations would force identical behavior.
Contradiction.

Thus the $2h_{\text{ctr}} + 1$ residual debts require $2h_{\text{ctr}} + 1$ distinct states.

The centered rail attains the bound. $\square$

29.6 Scope — place this beside the theorem

The theorem does **not** say:

every DPDA recognizing $L_{p,q}$ needs $2h_{\text{ctr}} + 1$ states.

A smaller machine may:

- abandon literal-demand storage;
- use another stack encoding;
- materialize debt eagerly;
- insert auxiliary physical markers;
- use another normalization schedule;
- leave the deferred-SEARCH interface entirely.

This scope is part of the theorem.

It is not a technical footnote.

29.7 Why Chapter 8 gives a different optimum

Chapter 8 proved the raw residual-faithful frontier:

$$S_{\text{raw}}(g) = \begin{cases} \infty, & g = 0, \\ p + q, & g = 1, \\ \max(p, q), & g \geq 2. \end{cases}$$

With two work symbols, the raw residual graph can therefore be realized using only

$$\max(p, q)$$

states.

The centered deferred machine can require

$$p + q \quad \text{or} \quad p + q - 1.$$

There is no contradiction.

Raw residual-faithful question

Give every residual one stable physical configuration and make every residual edge one raw input transition.

Deferred literal-demand question

Store the actual canonical minimum-demand word, search it by pure pops, and keep displaced mismatch information only in finite control until later materialization.

These are different representation contracts.

The correct conclusion is:

resource optimality is representation-category relative.

This observation is more important than either number individually.

It is one of the direct motivations for the later Structure and Uniqueness theories.

29.8 Why this chapter belongs before the 2:1 recovery

The width theorem tells us what the centered architecture itself costs.

Only after its category has been made explicit should we specialize back to the historical 2:1 machine.

Otherwise the reader could incorrectly interpret the three-state historical architecture as globally state-optimal.

The next chapter will instead say something more precise and more interesting:

The general centered theory reconstructs the unusual architecture of the original machine.

And the remaining transition differences will turn out to be schedule choices.

Chapter 30

Closing the First Loop: Recovering the Original 2:1 Machine

The strongest test of a proposed generalization is not that it accepts the right family.

It is that, when specialized back to the motivating case, it explains the machine from which the theory arose.

Set

$$\boxed{p = 2, \quad q = 1.}$$

Then

$$m = p + q = 3,$$

and

$$h_{\text{ctr}} = \left\lfloor \frac{3-1}{2} \right\rfloor = 1.$$

The centered rail has exactly three states:

$$D_{-1}, \quad D_0, \quad D_1.$$

Identify them as

$$\boxed{D_{-1} = R, \quad D_0 = C, \quad D_1 = L.}$$

The state potentials become

$$\boxed{-3, \quad 0, \quad +3.}$$

These are exactly the state potentials discovered directly in the original DPDA.

The coincidence is not imposed.

It falls out of the general debt formula

$$\Phi(D_d) = d(p + q).$$

30.1 Center mismatch directions

From $C = D_0$:

- a mismatch 1/0 moves to $D_1 = L$;
- a mismatch 0/1 moves to $D_{-1} = R$.

Thus the general SEARCH rail reproduces the original side-state directions.

The original left state is the positive debt state.

The original right state is the negative debt state.

This gives the old L/R asymmetry a uniform arithmetic interpretation.

30.2 Left cancellation rail

At $L = D_1$, matching 0/0 is simply MATCH:

$$(L, 0T) \xrightarrow{0} (L, T).$$

This is exactly the original cancellation transition that generated the entire Fuss–Catalan branch.

Thus the transition that later produced correction geometry is also a literal specialization of the general MATCH rule inside positive debt.

30.3 Right locality-repair rail

At $R = D_{-1}$, the general rules reproduce the right-side locality repair behavior.

The state is not an unrelated special-purpose gadget.

It is the negative SEARCH-debt sheet required when a one-cell blocker has been displaced from the canonical demand stack.

This is the general explanation of why the original machine needed a right side at all.

30.4 Historical WRITE arithmetic

The primitive kernel is

$$K = (1, 2).$$

The minimal kernel-carry WRITE equation now reproduces the characteristic replacement Parikh vectors of the historical table.

Depending on:

- source debt;
- target debt;
- input;
- removed top symbol;

the first nonnegative kernel shift produces the historical families represented by words such as

10,

11,

00,

000,

11111,

111111,

together with the corresponding bottom-WRITE forms.

The precise ordering is the selected canonical literal-demand order.

Thus the long WRITE strings are recovered from the general arithmetic compiler rather than copied into the generalization by hand.

30.5 Specialization and recovery theorem

Theorem 30.1 — 2:1 specialization and recovery.

Specializing the centered real-time $p:q$ construction to

$$(p, q) = (2, 1)$$

recovers:

1. the three-state

$$C/L/R$$

skeleton;

2. the state potentials

$$0, +3, -3;$$

3. the center mismatch directions;
4. the left cancellation rail;
5. the right locality-repair rail;
6. the arithmetic of the historical WRITE words.

The only reachable transition families not reproduced literally are two right-state matching families.

Those differences are purely normalization-schedule differences.

This theorem closes the first large loop of the monograph:

original machine $\rightarrow$ general theory $\rightarrow$ original architecture reconstructed.

30.6 Lazy versus eager discharge

Consider the reachable configuration

$$(R, 0).$$

The centered specialization performs the lazy matching pop

$$(R, 0) \xrightarrow{0}_{\text{lazy}} (R, \varepsilon).$$

The historical machine instead immediately materializes the represented demand and moves to the center:

$$(R, 0) \xrightarrow{0}_{\text{eager}} (C, 111).$$

Both configurations represent the same minimum demand.

Likewise, for a unary right stack

$$(R, 1^k),$$

the centered specialization uses

$$(R, 1^k) \xrightarrow{1}_{\text{lazy}} (R, 1^{k-1}),$$

whereas the historical machine uses

$$(R, 1^k) \xrightarrow{1}_{\text{eager}} (C, 1^{k+2}).$$

Again, the semantic demand agrees.

Only the time at which virtual debt is rematerialized differs.

30.7 Canonical materialization map

Define

$$\mathcal{N}(c) = (C, \eta(\Theta(c))).$$

This map forgets the physical distribution of demand between state and stack and returns the canonical center representation of the same shortest demand.

The historical eager successor is exactly the canonical materialization of the centered lazy successor on the two differing right-state families.

30.8 Eager–lazy semantic equivalence

Theorem 30.2 — Eager–lazy semantic equivalence.

Let

$$c_{\text{hist}}(w)$$

and

$$c_{\text{cent}}(w)$$

be the core configurations reached after the same binary prefix w by:

- the historical 2:1 machine;
- the centered 2:1 specialization.

Then

$$\Theta(c_{\text{hist}}(w)) = \Theta(c_{\text{cent}}(w))$$

for every prefix w .

On every transition family where the two tables differ,

$$\delta_{\text{hist}}(c, a) = \mathcal{N}(\delta_{\text{cent}}(c, a)).$$

Thus:

$$\text{historical 2:1} = \text{eager discharge},$$

while

$$\text{centered 2:1} = \text{lazy discharge}.$$

The semantics is the same.

The normalization chronology is different.

30.9 Historical machine as an eager refinement

Corollary 30.3 — Historical eager refinement.

Start with the centered 2:1 specialization.

Retain every specialized transition except on the two reachable right-state matching families.

On exactly those families, replace the lazy successor by its canonical materialization $\mathcal{N}$.

The resulting core is the historical 2:1 machine.

Therefore the return from the general theory is exact at the level of:

- control architecture;
- state potentials;
- WRITE arithmetic;
- shortest-demand semantics;

and becomes literal transition-table identity after choosing the historical eager normalization schedule.

Direct specialization alone is not claimed to be literal table identity.

The remaining difference is precisely the eager/lazy schedule choice.

30.10 Computational audit

The reference implementation of the centered family was independently rerun during the monograph audit for every ordered coprime pair

$$1 \leq p, q \leq 20,$$

including

(1, 1).

The run covered:

255 ratios,

1,909,191 distinct reached configurations,

and

3,738,864 binary edges.

It checked:

1. residual invariant;
2. residual-collision consistency;
3. two-cell qualitative visibility;
4. zero-residual correctness.

Result:

PASS.

These computations are consistency checks.

The symbolic theorems remain the proof.

30.11 The real discovery at the end of the $p:q$ generalization

The general machine has done more than extend 2:1.

It has exposed a structural nonuniqueness.

The semantic evolution is rigid:

one residual $\rightarrow$ one minimum demand.

But the physical normalization chronology is not rigid:

same semantics $\nrightarrow$ unique physical schedule.

This is the hinge of the remainder of the monograph.

The next paper does not ask for another DPDA construction.

It asks:

Why do WRITE, SEARCH, and CANCEL appear as the right primitive roles at all?

And more sharply:

Which parts are forced by semantic arithmetic, and which arise only because a one-top LIFO representation cannot follow the semantic update directly?

That is the beginning of the **Structure Theory**.

From this point onward the main sequence is uninterrupted:

Structure Theory $\rightarrow$ Uniqueness/Lift Theory $\rightarrow$ General Theory.

PART IX

Write–Search–Cancel as a Structural Theory

Research provenance. The WSC decomposition, debt channels, physical materialization level, and structural factorization theory are developed in [9]. Congruence and minimization theory for visibly pushdown systems and classical deterministic-pushdown equivalence are related but distinct background [14, 35, 61, 45].

Chapter 31

Semantic Arithmetic Before Representation

Chapter 30 ended with a fact that the original 2:1 machine could not reveal by inspection alone:

the same semantic demand can be normalized
on different physical schedules.

The centered generalization is lazy on two historical right-state families.

The original machine is eager.

Both are correct.

The next task is therefore no longer to construct a machine.

It is to separate what is forced by arithmetic from what is forced by LIFO locality and what is merely a normalization choice.

This is the beginning of the Structure Theory.

31.1 Dominating kernel level

Let a prefix contain

$$a = N_0(w), \quad b = N_1(w).$$

The primitive balanced kernel is

$$K = (q, p).$$

Define

$$\kappa(w) = \max \left(\left\lceil \frac{a}{q} \right\rceil, \left\lceil \frac{b}{p} \right\rceil \right).$$

Theorem 31.1 — Dominating kernel level.

The unique minimum completion demand is

$$\mu(w) = (q\kappa(w) - a, p\kappa(w) - b).$$

Equivalently,

$$(a, b) + \mu(w) = \kappa(w)K.$$

Proof.

If (x, y) completes w , then

$$q(b + y) = p(a + x).$$

Coprimality gives an integer k such that

$$a + x = qk, \quad b + y = pk.$$

Thus

$$x = qk - a, \quad y = pk - b.$$

Nonnegativity requires

$$k \geq \left\lceil \frac{a}{q} \right\rceil$$

and

$$k \geq \left\lceil \frac{b}{p} \right\rceil.$$

The least feasible k is therefore $\kappa(w)$.

The completion length is

$$(p + q)k - (a + b),$$

which is strictly increasing in k .

Hence the first dominating kernel level gives the unique minimum demand.

□

Geometrically, κ is the first point

$$kK = (kq, kp)$$

on the kernel ray that dominates the current count point (a, b) coordinate-wise.

The minimum demand is the slack vector from the current count point to that first dominating kernel point.

31.2 One input symbol: CANCEL plus optional carry

Let

$$e_0 = (1, 0), \quad e_1 = (0, 1).$$

Write

$$\mu_t$$

and

$$\kappa_t$$

for the minimum demand and dominating kernel level after t input symbols.

Theorem 31.2 — Unique semantic CANCEL–carry factorization.

If the next input symbol is

$$a \in \{0, 1\},$$

then

$$\mu_{t+1} = \mu_t - e_a + \chi_t K,$$

where

$$\chi_t = \kappa_{t+1} - \kappa_t \in \{0, 1\}.$$

Moreover:

- on input 0,

$$\chi_t = 1 \iff \mu_t^{(0)} = 0;$$

- on input 1,

$$\chi_t = 1 \iff \mu_t^{(1)} = 0.$$

Proof.

Reading one symbol increments exactly one prefix-count coordinate.

Each ceiling in the formula for κ therefore changes by at most one.

Hence

$$\chi_t \in \{0, 1\}.$$

Before the input,

$$(N_0, N_1) + \mu_t = \kappa_t K.$$

After the input,

$$(N_0, N_1) + e_a + \mu_{t+1} = \kappa_{t+1} K.$$

Subtracting gives the factorization.

For input 0, the zero slack is

$$q\kappa_t - N_0.$$

The kernel level rises precisely when that coordinate is already saturated,

i.e. when the zero-demand coordinate is 0.

The input-1 case is symmetric. $\square$

Thus every semantic input step has exactly two roles:

CANCEL one unit of the incoming symbol

and, if the relevant demand coordinate would become negative,

add one primitive kernel packet.

This is realization-independent.

It is the semantic reason WRITE exists at all.

31.3 Semantic WRITE count

Every semantic carry increments κ by one.

Corollary 31.3 — Semantic WRITE count.

For every prefix,

$$W_{\text{sem}}(w) = \kappa(w).$$

If

$$w \in L_{p,q}$$

and

$$N_0(w) = qn, \quad N_1(w) = pn,$$

then

$$W_{\text{sem}}(w) = n.$$

The number of semantic kernel arrivals is fixed by the input.

What can vary is when those arrivals are physically materialized.

31.4 Demand-length clock

Let

$$r = p + q.$$

Corollary 31.4 — Demand-length clock.

At prefix length t ,

$$|\mu_t| = r\kappa_t - t.$$

Hence

$$|\mu_t| \equiv -t \pmod{r}.$$

The minimum-demand length carries a rigid modular clock.
This clock will later have a physical analogue.

31.5 Semantic chronology is unique

Corollary 31.5 — Semantic chronology uniqueness.

For a fixed input word, the entire sequences

$$(\rho_t), \quad (\kappa_t), \quad (\mu_t), \quad (\chi_t)$$

are uniquely determined by the input.

They do not depend on which correct pushdown representation is used.

This is the semantic backbone of the Structure Theory.

It says:

$$\boxed{\text{the semantic chronology is unique}}$$

before we say anything about:

- stack order;
- side states;
- SEARCH debt;
- eager/lazy materialization.

The next chapter introduces the first mature abstraction separating this semantic process from its physical realization.

Chapter 32

What One-Top LIFO Locality Forbids

We now formalize a physical realization of the unique demand dynamics.

The abstraction appears here, not at the beginning of the monograph, because the reader has already seen why it is needed.

32.1 Exact literal-demand lift

Let

$$\mathcal{M}_{p,q} = \{\mu(\rho) : \rho \in \mathbb{Z}\}$$

be the semantic minimum-demand system.

Definition 32.1 — Exact real-time literal-demand lift.

An exact literal-demand lift consists of:

- a reachable physical configuration set

$$\mathcal{C};$$

- deterministic binary transitions

$$\delta_a;$$

- a semantic projection

$$\Theta : \mathcal{C} \rightarrow \mathcal{M}_{p,q};$$

- a canonical section

$$\sigma : \mathcal{M}_{p,q} \rightarrow \mathcal{C};$$

such that

$$\Theta \circ \sigma = \text{id}$$

and

$$\Theta(\delta_a(c)) = U_a(\Theta(c)),$$

where U_a is the unique semantic update from Chapter 31.

The section chooses one preferred physical representative of each semantic demand.

The other physical representatives form the semantic fiber.

That vocabulary will become central later.

32.2 Fixed block orientation

For the present Structure Theory, choose the literal canonical encoding

$$\eta(x, y) = 1^y 0^x Z_0,$$

with stack top on the left.

A one-top real-time transition has the form

$$XT \longrightarrow \gamma T.$$

The suffix T is immutable during that transition.

This one fact gives an exact direct-realizability criterion.

32.3 Top-rewrite suffix criterion

Theorem 32.2 — Top-rewrite suffix criterion.

Suppose a one-top transition begins from

$$XT.$$

A desired target word S^* can be reached in that one transition if and only if

$$T \text{ is a suffix of } S^*.$$

Proof.

Every one-top successor has the form

$$\gamma T.$$

Therefore direct realization requires T to be a suffix of the target.

Conversely, if

$$S^* = \gamma T,$$

replace X by γ . $\square$

This theorem is one of the simplest and most important results in the entire structural line.

It converts one-top LIFO locality into an exact string condition.

32.4 Complete obstruction classification

Apply the suffix criterion to the canonical word

$$1^y 0^x Z_0.$$

Theorem 32.3 — Canonical direct-realizability classification.

For the fixed orientation $1^y 0^x Z_0$:

1. input 1 is directly canonical-realizable for every canonical demand;
2. input 0 fails to be directly canonical-realizable exactly in two families:

Buried-CANCEL obstruction

$$\boxed{x > 0, \quad y > 0.}$$

Carry-order obstruction

$$\boxed{x = 0, \quad y \geq 2, \quad q \geq 2.}$$

For the reverse block orientation, exchange 0 and 1.

Proof — input 1. If $y > 0$, pop the top 1.

If $y = 0$ and $x > 0$, the semantic target is

$$1^{p-1} 0^{x+q} Z_0.$$

The untouched suffix after removing the top zero is

$$0^{x-1} Z_0,$$

which is a suffix of the target.

Thus a single top replacement suffices.

The bottom case is direct as well.

Proof — input 0. If $y = 0$ and $x > 0$, pop the top 0.

Now let $x > 0, y > 0$.

The semantic target is

$$1^y 0^{x-1} Z_0.$$

But after removing the top 1, the untouched suffix is

$$1^{y-1} 0^x Z_0.$$

That suffix is not a suffix of the target.

Direct canonicalization is impossible.

Finally let $x = 0$.

The carry target is

$$1^{y+p} 0^{q-1} Z_0.$$

For $y = 0$, write at the bottom.

For $y = 1$, the untouched suffix is empty and direct replacement is possible.

For $y \geq 2$, the untouched suffix is

$$1^{y-1} Z_0.$$

It is a suffix of the target exactly when

$$q - 1 = 0,$$

that is,

$$q = 1.$$

So for $q \geq 2$ the second obstruction family is exact. $\square$

32.5 Two fundamentally different failures

The classification reveals that “cannot stay canonical” has two different causes.

Buried cancellation

The symbol that semantic CANCEL wants to remove exists in the physical word, but a blocker lies above it.

This is a visibility/locality failure.

Carry-order obstruction

A semantic carry can create the correct **number** of zeros and ones in one transition, yet the immutable suffix forces those letters into the wrong order.

This is not a quantity failure.

It is an ordering failure.

Thus:

count-exact $\nrightarrow$ canonical.

The next chapter separates the resulting representation debt into three channels.

Chapter 33

Quantity, Order, and Control Debt

A configuration may leave the canonical section in more than one way.

Parikh counts see only one of them.

The Structure Theory therefore distinguishes:

1. quantity debt;
2. order debt;
3. control/sheet debt.

These channels will remain separate through the Uniqueness and General Theory Parts.

33.1 Quantity defect

Let $S(c)$ be the literal work-stack word of a rich configuration c .

Define its physical Parikh vector

$$P(c) = \text{Par}(S(c)).$$

Definition 33.1 — Quantity defect.

$$Q(c) = \Theta(c) - P(c).$$

A canonical configuration has

$$Q = 0.$$

But $Q = 0$ does not imply canonicity because Parikh counts forget literal order.

33.2 Universal representation conservation

Proposition 33.2 — Universal representation conservation.

For every exact transition on input a ,

$$\boxed{\Delta P + \Delta Q = -e_a + \chi K.}$$

Proof.

Since

$$Q = \Theta - P,$$

$$\Delta Q = \Delta \Theta - \Delta P.$$

The semantic update is

$$\Delta \Theta = -e_a + \chi K.$$

Rearrange. $\square$

This equation holds before we impose the later WSC lattice restrictions.

It is the universal accounting identity.

33.3 Order defect

For literal word

$$S = S_1 \cdots S_m,$$

define

$$\boxed{\Omega(S) = \#\{(i, j) : i < j, S_i = 0, S_j = 1\}.}$$

Lemma 33.3 — Block-order test.

$$\boxed{\Omega(S) = 0 \iff S \in 1^*0^*}.$$

Thus Ω measures departure from the chosen canonical block order.

33.4 Three-channel canonicity

A third channel is finite control.

A configuration can have:

- correct counts;
- correct literal order;

but still lie on a noncanonical control sheet.

Proposition 33.4 — Three-channel canonicity.

A rich configuration is canonical exactly when:

1. quantity defect is zero;
2. order defect is zero;
3. the control state is the canonical section state.

Thus:

$$\text{canonicity} = \text{quantity} + \text{order} + \text{sheet/control}.$$

This is the first mature answer to the question:

Where can normalization debt live?

Proposition 33.4A — Configuration-level separation of the three defect channels.

At the level of exact semantic representations used in this theory, none of the three canonicity channels is implied by the other two.

Each channel has a witness for which the other two vanish.

- **Quantity-only defect.** In the 2:1 schedule interaction studied later, the kernel-excess configuration

$$(C, 1110)$$

is compared with canonical

$$(C, 1).$$

The control sheet and block order are canonical, but the quantity defect is one nonzero kernel packet.

- **Order-only defect.** The carry-order repair of Theorem 33.6 preserves the exact canonical Parikh vector while forcing positive inversion count

$$\Omega > 0.$$

- **Control-only defect.** Take any exact literal lift and adjoin a semantically ignored two-sheet control decoration. A balanced input can return the base machine to a canonical stack while leaving the auxiliary sheet noncanonical. Quantity and order are then canonical while control defect remains nonzero.

Therefore

$$\boxed{Q = 0, \quad \Omega = 0, \quad \text{canonical control}}$$

are genuinely independent physical requirements.

This point becomes important in the Uniqueness Theory: a count-shadow invariant alone cannot reconstruct the full physical semantic fiber.

33.5 Kernel economy

Define a literal-demand representation to be **kernel-economical** if

$$\boxed{|P(c)| \leq |\Theta(c)|}$$

for every reachable configuration.

This prohibits adding complete zero-residual kernel packets merely as redundant workspace.

Under economy, the buried-CANCEL obstruction forces an exact quantity debt.

33.6 Forced unit SEARCH

Theorem 33.5 — Forced unit SEARCH for buried cancellation.

Assume:

- one-top real-time locality;
- kernel economy;
- canonical mixed demand

$$\Theta(c) = (x, y)$$

with

$$x > 0, \quad y > 0;$$

- literal encoding

$$1^y 0^x Z_0.$$

On input 0, the one-step successor has

$$Q' = (-1, 1) = -(1, -1).$$

Proof.

The semantic target is

$$(x - 1, y),$$

whose length is

$$x + y - 1.$$

A one-top transition can replace only the blocking top 1 by some word γ .
The successor physical length is

$$x + y - 1 + |\gamma|.$$

Kernel economy requires this to be at most the semantic target length.
Hence

$$|\gamma| = 0.$$

The top 1 must be popped.

The physical successor counts are

$$(x, y - 1),$$

while the semantic target is

$$(x - 1, y).$$

Therefore

$$Q' = (-1, 1).$$

□

Thus SEARCH is not merely one convenient implementation trick.

Under the stated locality and economy assumptions, the buried cancellation produces exactly one forced exchange defect.

33.7 Carry obstruction with zero quantity debt

The second obstruction behaves differently.

Assume:

$$x = 0, \quad y \geq 2, \quad q \geq 2.$$

On input 0, the semantic target is

$$1^{y+p}0^{q-1}Z_0.$$

A one-top transition can replace the top 1 by

$$\gamma = 1^{p+1}0^{q-1}.$$

The resulting word is

$$1^{p+1}0^{q-1}1^{y-1}Z_0.$$

Theorem 33.6 — Count-exact carry-order repair.

The successor has exactly the target Parikh counts:

$$\boxed{Q' = 0.}$$

But it has nonzero order defect.

Thus the representation can be semantically count-exact and still noncanonical.

33.8 Minimum order debt

Theorem 33.7 — Minimum carry-order debt.

Under the same hypotheses, among all one-top successors with exact target counts,

$$\boxed{\Omega_{\min} = (q-1)(y-1).}$$

Proof.

The untouched suffix is

$$1^{y-1}Z_0.$$

Exact target counts force the replacement word to contain:

$$q-1$$

zeros and

$$p + 1$$

ones.

Every one of those $q - 1$ zeros lies before all $y - 1$ untouched suffix ones.
Therefore at least

$$(q - 1)(y - 1)$$

inversions are forced.

The replacement

$$1^{p+1}0^{q-1}$$

introduces no additional inversions inside the replacement and attains the bound. $\square$

The obstruction/debt principle is now clear:

LIFO obstruction forces departure from the canonical section,

but the debt may be stored in:

- quantity;
- order;
- finite control;
- or several channels at once.

The next chapter isolates a count-shadow class in which the quantity arithmetic becomes rigid.

Chapter 34

The WSC Count Shadow

The preceding debt theory was general within the literal-demand setting.

We now impose an explicit representation policy to obtain a sharper arithmetic structure.

This policy is not claimed to describe every correct DPDA.

The scope is part of the theorem.

Set

$$\boxed{\mathbf{h} = (1, -1), \quad K = (q, p), \quad r = p + q.}$$

34.1 Phase coherence

Definition 34.1 — Phase coherence.

A reachable configuration is phase-coherent if

$$\boxed{|P(c)| \equiv |\Theta(c)| \pmod{r}.}$$

A realization is phase-coherent if every reachable configuration satisfies this congruence.

Since

$$Q = \Theta - P,$$

phase coherence is equivalent to

$$Q_0 + Q_1 \equiv 0 \pmod{r}.$$

This congruence produces the WSC lattice.

34.2 WSC defect lattice

Define

$$\mathcal{D}_{p,q} = \mathbb{Z}\mathbf{h} + \mathbb{Z}K.$$

Theorem 34.2 — WSC defect lattice.

$$\mathcal{D}_{p,q} = \{(u, v) \in \mathbb{Z}^2 : u + v \equiv 0 \pmod{r}\}.$$

Consequently every phase-coherent quantity defect has a unique decomposition

$$Q = s\mathbf{h} + wK, \quad s, w \in \mathbb{Z}.$$

Proof.

Both generators have coordinate sum divisible by r :

$$(\mathbf{h})_0 + (\mathbf{h})_1 = 0,$$

$$K_0 + K_1 = r.$$

Conversely, if

$$u + v \equiv 0 \pmod{r},$$

define

$$w = \frac{u + v}{r},$$

and

$$s = \frac{pu - qv}{r}.$$

The first is integral by assumption.

Since

$$p \equiv -q \pmod{r},$$

the numerator of s is also divisible by r .

Direct substitution gives

$$(u, v) = s\mathbf{h} + wK.$$

Uniqueness follows because

$$\det(\mathbf{h}, K) = p + q = r \neq 0.$$

□

34.3 Defect coordinates

Define

$$\Phi(Q) = pQ_0 - qQ_1,$$

$$\Lambda(Q) = Q_0 + Q_1.$$

Corollary 34.3 — Defect coordinates.

For

$$Q = s\mathbf{h} + wK,$$

$$\boxed{\Phi(Q) = rs,}$$

$$\boxed{\Lambda(Q) = rw.}$$

Hence

$$\boxed{s = \frac{\Phi(Q)}{r}, \quad w = \frac{\Lambda(Q)}{r}.}$$

The coordinate s measures the SEARCH/exchange sheet in the quantity shadow.

The coordinate w measures deferred kernel quantity.

34.4 Defect quotient

Corollary 34.4 — Defect quotient.

$$\boxed{\mathbb{Z}^2 / \mathcal{D}_{p,q} \cong \mathbb{Z}_r.}$$

The homomorphism is

$$(u, v) \mapsto u + v \pmod{r}.$$

This quotient is a useful arithmetic shadow.

It is not the physical semantic fiber.

34.5 WSC lattice is not the full LIFO fiber

The map

$$S \mapsto \text{Par}(S)$$

forgets stack order.

The quantity defect also forgets finite-control sheets.

Therefore two physical configurations can have the same coordinates

$$(s, w)$$

while having:

- different literal stack words;
- different control states;
- different future locality properties.

Thus:

WSC defect lattice $\neq$ full physical representation fiber.

This warning is essential.

Chapter 38 onward will develop the missing lift/fiber theory.

34.6 Kernel economy in lattice coordinates

Since

$$|\Theta| - |P| = Q_0 + Q_1 = rw,$$

we get:

Corollary 34.5 — Kernel economy.

Within the phase-coherent class,

$$\text{kernel economy} \iff w \geq 0.$$

Thus the deferred-kernel coordinate acquires a direct nonnegative backlog interpretation under economy.

34.7 Abstract count-shadow fiber region

Suppose the semantic demand is

$$\Theta(c) = \mu = (x, y)$$

and

$$Q = s\mathbf{h} + wK.$$

Then

$$P = \mu - Q = (x, y) - s(1, -1) - w(q, p).$$

So

$$P_0 = x - s - qw,$$

$$P_1 = y + s - pw.$$

Proposition 34.6 — Count-shadow fiber region.

Physical nonnegativity is equivalent to

$$\boxed{pw - y \leq s \leq x - qw,}$$

with

$$\boxed{0 \leq w \leq \left\lfloor \frac{x + y}{r} \right\rfloor.}$$

For fixed w , the number of integer s -positions is

$$\boxed{x + y - rw + 1.}$$

The count-shadow fiber therefore has a finite triangular/layered shape over each semantic demand.

Again, this is an abelian shadow of the full physical fiber, not the complete fiber itself.

The next chapter gives these coordinates a temporal interpretation.

Chapter 35

Physical Materialization, Backlog, and WSC Factorization

The semantic carry schedule

$$\chi_t$$

is unique.

The physical machine may materialize those kernel packets at different times.

We now make that chronology explicit.

Let

$$Q_t = s_t h + w_t K$$

along a phase-coherent run.

35.1 Materialization level

Definition 35.1 — Physical materialization level.

Define

$$\lambda_t = \kappa_t - w_t.$$

The semantic system has generated κ_t kernel packets.

The quantity shadow still defers w_t packets.

Therefore λ_t is the number of kernel packets already materialized physically at the count-shadow level.

35.2 Physical length formula

Theorem 35.2 — Physical length formula.

At prefix length t ,

$$\boxed{|P_t| = r\lambda_t - t.}$$

Proof.

$$|P_t| = |\Theta_t| - (Q_{t,0} + Q_{t,1}).$$

Chapter 31 gives

$$|\Theta_t| = r\kappa_t - t.$$

Chapter 34 gives

$$Q_{t,0} + Q_{t,1} = rw_t.$$

Subtract and use

$$\lambda_t = \kappa_t - w_t.$$

□

The physical stack length therefore has its own modular clock:

$$\boxed{|P_t| \equiv -t \pmod{r}.}$$

35.3 Materialization amount

Define

$$\boxed{M_t = \lambda_{t+1} - \lambda_t.}$$

Theorem 35.3 — Materialization amount from stack length.

For every phase-coherent one-top transition,

$$\boxed{M_t = \frac{\Delta|P_t| + 1}{r}.}$$

In particular,

$$\boxed{M_t \in \mathbb{N}.}$$

Proof.

Subtract the physical length formula at t and $t + 1$:

$$\Delta|P_t| = rM_t - 1.$$

Thus the formula follows.

Because a one-top rewrite can delete at most the visible literal cell,

$$\Delta|P_t| \geq -1.$$

Hence

$$M_t \geq 0.$$

□

The number of kernel packets physically materialized by a transition can therefore be reconstructed from the physical stack-length change alone.

35.4 Backlog law

Theorem 35.4 — Backlog law.

$$w_{t+1} - w_t = \chi_t - M_t.$$

Equivalently, if

$$A_t = \kappa_t$$

is cumulative semantic arrival and

$$B_t = \lambda_t$$

is cumulative physical materialization, then

$$w_t = A_t - B_t.$$

This is a queue-like law:

$$\text{backlog change} = \text{semantic arrivals} - \text{physical service}.$$

The analogy is structural, not stochastic.

35.5 Canonical carry dichotomy

Assume phase coherence and kernel economy, and begin a transition from canonical quantity state

$$w_t = 0.$$

Theorem 35.5 — Canonical carry dichotomy.

If

$$\chi_t = 0,$$

then

$$M_t = 0.$$

If

$$\chi_t = 1,$$

then

$$M_t \in \{0, 1\}.$$

Thus a semantic carry is either:

$$\text{eagerly materialized}$$

or

$$\text{deferred as exactly one kernel packet of backlog.}$$

This is the structural generalization of the eager/lazy distinction discovered concretely in Chapter 30.

35.6 Unique physical WSC role factorization

Let

$$\alpha_t = s_{t+1} - s_t.$$

Theorem 35.6 — Unique physical WSC count-shadow factorization.

Every phase-coherent one-top transition on input a satisfies

$$\Delta P = -e_a - \alpha_t \mathbf{h} + M_t K.$$

The coefficients

$$\alpha_t \in \mathbb{Z}$$

and

$$M_t \in \mathbb{N}$$

are unique.

Proof.

The defect change is

$$\Delta Q = \alpha_t \mathbf{h} + (w_{t+1} - w_t)K.$$

Use the backlog law:

$$\Delta Q = \alpha_t \mathbf{h} + (\chi_t - M_t)K.$$

Insert this into the universal conservation identity

$$\Delta P + \Delta Q = -e_a + \chi_t K.$$

The $\chi_t K$ terms cancel:

$$\Delta P = -e_a - \alpha_t \mathbf{h} + M_t K.$$

Uniqueness follows from linear independence of $\mathbf{h}$ and K . $\square$

Meaning

This is a **role decomposition**, not a syntactic transition taxonomy.

The terms mean:

$$-e_a$$

= semantic CANCEL forced by the input;

$$-\alpha \mathbf{h}$$

= SEARCH/exchange effect in the abelian quantity shadow;

$$MK$$

= physically materialized kernel packets.

One concrete PDA transition may combine several of these roles.

35.7 Reconstructing the charges

Corollary 35.7 — Reconstruction.

$$M = \frac{\Delta|P| + 1}{r},$$

and

$$\alpha = -\frac{p(\Delta P_0 + e_{a,0}) - q(\Delta P_1 + e_{a,1})}{r}.$$

Thus the WSC charges can be recovered from the physical count effect.

Corollary 35.7A — General physical phase clock.

Every phase-coherent run satisfies

$$|P_t| \equiv -t \pmod{r}.$$

This follows immediately from

$$|P_t| = r\lambda_t - t.$$

The semantic demand clock of Chapter 31 and the physical stack clock therefore have the same modulus $r = p + q$, even though the physical materialization schedule may be eager, lazy, or mixed.

35.8 Global materialization budget

Let an accepted word have

$$N_0 = qn, \quad N_1 = pn.$$

Theorem 35.8 — Global materialization budget.

For a phase-coherent kernel-economical run starting canonically,

$$\sum_t M_t = n.$$

This equals the total semantic carry count:

$$\sum_t \chi_t = n.$$

Proof.

The accepted input length is

$$rn.$$

At acceptance, semantic demand is zero.

Kernel economy forces the physical literal count to zero as well.

The physical length formula gives

$$0 = r\lambda_{rn} - rn.$$

Hence

$$\lambda_{rn} = n.$$

Since $\lambda_0 = 0$,

$$\sum_t M_t = n.$$

The semantic equality follows from Chapter 31. $\square$

The schedule may move materialization earlier or later.

The total materialized kernel budget cannot change.

35.9 Deferral capacity

Corollary 35.9 — Deferral capacity bound.

At prefix length t ,

$$\left\lceil \frac{t}{r} \right\rceil \leq \lambda_t \leq \kappa_t,$$

and

$$0 \leq w_t \leq \kappa_t - \left\lceil \frac{t}{r} \right\rceil.$$

Thus the physical representation has an exact finite capacity for delaying kernel materialization at a given semantic prefix.

The next chapter moves from arithmetic chronology to semantic fibers and asks which parts of the physical presentation are unique.

Chapter 36

Unique Semantic Quotient, Nonunique Physical Presentation

The Structure Theory has now separated:

- the unique semantic demand trajectory;
- the physical normalization trajectory.

We can state that distinction as a quotient.

36.1 Semantic fibers

For semantic demand

$$m \in \mathcal{M}_{p,q},$$

define

$$\boxed{F_m = \Theta^{-1}(m).}$$

Two rich configurations are semantically equivalent when

$$\boxed{c \sim c' \iff \Theta(c) = \Theta(c').}$$

A fiber can contain several physical configurations representing the same minimum demand.

36.2 Fiber congruence

Theorem 36.1 — Fiber congruence and semantic quotient.

In every exact literal-demand lift, the relation $\sim$ is preserved by binary input transitions.

Hence

$$\mathcal{C}/\sim \cong \mathcal{M}_{p,q}.$$

Proof.

If

$$\Theta(c) = \Theta(c'),$$

then exactness gives

$$\Theta(\delta_a(c)) = U_a(\Theta(c)) = U_a(\Theta(c')) = \Theta(\delta_a(c')).$$

Thus successors remain in the same semantic fiber.

The quotient state is therefore exactly the common semantic demand, and the quotient transitions are the canonical demand updates. $\square$

This is an early version of the later machine-independent semantic quotient theorem.

The stronger language-level statement will come in the Uniqueness Theory.

36.3 Canonical retraction

Define

$$N = \sigma \circ \Theta.$$

This sends every physical configuration to the canonical section representative of its semantic state.

Proposition 36.2 — Canonical absorption.

$$N^2 = N.$$

Moreover, for every input word w ,

$$N \circ \delta_w = N \circ \delta_w \circ N.$$

Proof.

Since

$$\Theta \circ \sigma = \text{id},$$

$$N^2 = \sigma\Theta\sigma\Theta = \sigma\Theta = N.$$

For one symbol,

$$N\delta_a = \sigma\Theta\delta_a = \sigma U_a\Theta.$$

Also,

$$N\delta_a N = \sigma\Theta\delta_a\sigma\Theta = \sigma U_a\Theta\sigma\Theta = \sigma U_a\Theta.$$

Induct on $|w|$. $\square$

Interpretation:

Once we normalize to the canonical section, earlier physical normalization history is erased.

This is a structural form of the eager/lazy semantic equivalence.

36.4 What is actually unique

The following are semantic and unique:

- residual;
- minimum demand;
- dominating kernel level;
- semantic carry indicator;
- canonical quotient path.

The following need not be unique:

- physical stack word;
- finite-control sheet;
- timing of materialization;
- normalization schedule;
- full pushdown presentation.

Thus the correct statement is:

unique semantic core; nonunique physical normalization lift.

36.5 Cheap nonuniqueness already exists

Proposition 36.3 — Unrestricted physical presentation nonuniqueness.

The canonical semantic quotient does not determine a unique deterministic pushdown presentation.

Proof.

Start with any exact deterministic lift with finite control Q .

For integer

$$m \geq 2,$$

adjoin an independent cyclic coordinate

$$j \in \mathbb{Z}_m.$$

On every consumed input symbol, update

$$j \mapsto j + 1 \pmod{m}.$$

Let:

- stack dynamics;
- semantic projection;
- acceptance

ignore the extra coordinate.

The result is still an exact deterministic lift.

Its control set is

$$Q \times \mathbb{Z}_m.$$

For infinitely many suitable pairwise distinct m , these presentations have different finite-control cardinalities and are therefore non-isomorphic, while their semantic quotient is unchanged. $\square$

This proposition proves physical presentation nonuniqueness.

But it is deliberately cheap.

It uses redundant control decoration.

It does **not** answer the important question:

Does nonuniqueness survive after redundant presentation freedom is removed?

That sharper question is exactly why the next paper is needed.

36.6 Why “confluence” is not yet justified

Canonical absorption says that normalizing after a fixed input word forgets earlier physical history.

It does not prove a Church–Rosser theorem for arbitrary rewrite schedules. We have not defined a general rewrite relation with arbitrary branch choices. Therefore the safe conclusion is:

canonical endpoint coherence under the chosen normalizer
--

not

global confluence.

This scope will remain important in the Uniqueness and General Theory Parts.

Chapter 37

The Original 2:1 Machine as WSC

The Structure Theory now returns to the original machine for a second reinterpretation.

Chapter 30 explained the historical transition table as an eager schedule of the centered 2:1 construction.

The present chapter goes deeper.

It interprets:

- the historical side states as SEARCH sheets;
- the historical stack-length changes as materialization bursts;
- an old transition-count identity as a global kernel budget law.

Set

$$\boxed{p = 2, \quad q = 1, \quad r = 3, \quad K = (1, 2).}$$

37.1 Historical states as SEARCH sheets

The original state potentials were

$$\Phi(C) = 0,$$

$$\Phi(L) = 3,$$

$$\Phi(R) = -3.$$

Since the WSC sheet coordinate s contributes residual quantity

$$rs = 3s,$$

we obtain:

Corollary 37.1 — Historical states as SEARCH sheets.

$$\boxed{s(C) = 0, \quad s(L) = 1, \quad s(R) = -1.}$$

Thus the original $C/L/R$ architecture is exactly a three-sheet realization of the WSC exchange coordinate.

The side states are not merely “left” and “right.”

They are the ± 1 SEARCH sheets.

37.2 Historical stack-length changes

For $r = 3$, the physical materialization formula is

$$\Delta|P| = 3M - 1.$$

Therefore:

$$M = 0 \implies \Delta|P| = -1,$$

$$M = 1 \implies \Delta|P| = +2,$$

$$M = 2 \implies \Delta|P| = +5.$$

Corollary 37.2 — Historical stack actions as materialization bursts.

The characteristic historical stack-length changes

$$\boxed{-1, \quad +2, \quad +5}$$

are exactly the WSC materialization levels

$$\boxed{M = 0, \quad M = 1, \quad M = 2.}$$

Therefore a $+5$ transition is not an anomalously long push.

It is a two-kernel-packet materialization burst.

This explains why some original transition replacements were much longer than others.

37.3 The global identity $b + 2c = n$

Suppose an accepted 2:1 run has:

- b transitions with stack-length change $+2$;
- c transitions with stack-length change $+5$.

Those transitions carry materialization amounts:

$$1 \quad \text{and} \quad 2.$$

The global materialization budget says that an accepted word with

$$N_0 = n, \quad N_1 = 2n$$

must satisfy

$$\sum_t M_t = n.$$

Hence:

Corollary 37.3 — Structural derivation of the historical identity.

$$\boxed{b + 2c = n.}$$

The identity is no longer a transition-table counting curiosity.

It is the 2:1 specialization of the universal materialization budget.

37.4 What Structure Theory has accomplished

The generalization and Structure Theory now give three successive interpretations of the same original machine.

Chapter 1

$$\boxed{\text{state+stack residual invariant.}}$$

Chapter 30

$$\boxed{C/L/R = \text{centered } p:q \text{ debt rail at } (2, 1),}$$

with historical behavior as eager discharge.

Chapter 37

$C/L/R = \text{WSC SEARCH sheets,}$

and the physical stack actions are kernel-materialization bursts.

This repeated return to the original machine is intentional.

Each return explains something that was previously only observed.

37.5 The Structure Theory endpoint

The strongest conclusion is:

CANCEL and semantic WRITE are forced by arithmetic;

SEARCH and normalization debt arise when that semantic update
meets one-top LIFO locality.

But one major question remains unresolved.

The Structure Theory has shown that physical presentations are nonunique.

Its explicit nonuniqueness construction can be dismissed as redundant decoration.

The next Part asks the sharper question:

After fixing the semantic quotient, the canonical section, and an admissible LIFO morphism category, can two **reduced** physical lifts still be genuinely different?

The answer is yes.

Moreover, in the motivating 2:1 family there are not merely two schedules.

There are exactly six admissible reduced schedules in the selected three-site family, with one directed causal conflict.

That is the Uniqueness/Lift Theory.

PART X

Semantic Uniqueness and Physical Lift Freedom

Research provenance. The language-level semantic quotient, exact sectioned lifts, reduced-lift nonuniqueness, and causal schedule analysis are developed in [10]. The terminology is kept separate from classical VPL congruence/minimization and DPDA equivalence results [14, 35, 61].

Chapter 38

The Semantic Core Belongs to the Language

The Structure Theory established a canonical semantic quotient inside exact literal-demand lifts.

The Uniqueness Theory strengthens that result.

The semantic core will now be derived from the **language itself**, not from a chosen WSC construction.

38.1 Right residuals

For a prefix u , define the right-continuation language

$$u^{-1}L_{p,q} = \{z : uz \in L_{p,q}\}.$$

Chapter 4 already proved that this continuation language depends only on the arithmetic residual

$$\rho(u) = qN_1(u) - pN_0(u).$$

We now promote that result to the canonical semantic theorem.

Theorem 38.1 — Right-residual classification.

For all prefixes u, v ,

$$u^{-1}L_{p,q} = v^{-1}L_{p,q} \iff \rho(u) = \rho(v).$$

Every integer residual occurs.

Thus the reachable right-residual classes are canonically indexed by

$$\mathbb{Z}.$$

The semantic transition system is therefore

$$\mathcal{R}_{p,q} : \quad r \xrightarrow{0} r - p, \quad r \xrightarrow{1} r + q.$$

This is the language-determined behavioral core.

Corollary 38.1A — Canonical behavioral core.

The reachable right-residual classes of $L_{p,q}$ are canonically identified with

$$\boxed{\mathbb{Z}},$$

and their induced labeled transition system is exactly

$$\mathcal{R}_{p,q} : \quad r \xrightarrow{0} r - p, \quad r \xrightarrow{1} r + q.$$

Thus the integer residual line is not a coordinate convention inherited from a particular DPDA. It is the canonical continuation-behavior core of the language.

38.2 Universal semantic quotient

Consider any correct deterministic real-time recognizer of $L_{p,q}$.

Look only at reachable configurations at input boundaries.

Map every such physical configuration c , reached by prefix u , to

$$\boxed{\Theta_{\text{res}}(c) = \rho(u)}.$$

Determinism and correctness make this map well defined on behavioral equivalence classes.

Theorem 38.2 — Universal semantic quotient.

Every correct deterministic real-time recognizer of $L_{p,q}$ admits a unique surjective transition-preserving quotient onto

$$\boxed{\mathcal{R}_{p,q}}.$$

Thus the residual system is not merely the semantic quotient of our chosen constructions.

It is the canonical behavioral quotient forced by the language.

This strengthens Chapter 36.

38.3 Minimum-demand semantics is a coordinate change

The minimum-demand system

$$\mathcal{M}_{p,q}$$

has one state

$$\mu(r)$$

for every residual r .

The map

$$r \longmapsto \mu(r)$$

is bijective.

It also intertwines the input transitions.

Theorem 38.3 — WSC semantic core equals the residual core.

$$\mathcal{R}_{p,q} \cong \mathcal{M}_{p,q}.$$

Therefore:

$$\text{language semantics} = \text{right-residual system},$$

while

$$\text{completion semantics} = \text{a useful isomorphic coordinate system}.$$

This distinction is conceptually important.

The minimum-demand representation was discovered first because it grew naturally out of the original DPDA.

But the language-level semantic object is the residual quotient.

38.4 Three levels of description

We can now separate:

Language level

$$\mathcal{R}_{p,q}.$$

This is canonical.

Semantic coordinate level

$$\mathcal{M}_{p,q}.$$

This is isomorphic to the residual system.

Physical lift level

A DPDA realization mapping onto that semantic system.

This level is not unique.

The rest of the Uniqueness Theory asks how to compare such lifts without quotienting away the very pushdown structure we want to study.

Chapter 39

Exact Sectioned Lifts and Category-Relative Minimality

If arbitrary behavioral quotienting is allowed, every correct machine collapses to the same residual core.

That is too coarse for physical comparison.

We need a category that remembers:

- the chosen semantic projection;
- the canonical section;
- the LIFO coordinates.

39.1 Reachable versus ambient configurations

Let

$$C^{\text{amb}}$$

be all syntactically possible physical configurations.

Let

$$C^{\text{reach}} \subseteq C^{\text{amb}}$$

be the configurations actually reachable from the initial condition.

All minimality, fiber, and schedule claims in this Part concern reachable configurations unless explicitly stated otherwise.

This prevents unreachable syntactic artifacts from distorting the theory.

39.2 Exact sectioned lift

Definition 39.1 — Exact sectioned lift.

An exact sectioned lift is

$$\mathcal{L} = (C^{\text{reach}}, \delta, \Theta, \sigma)$$

such that:

1. Θ is a surjective semantic projection;
2. σ is a reachable section;
- 3.

$$\Theta \circ \sigma = \text{id};$$

4. physical transitions commute exactly with semantic transitions.

The section is part of the data.

Two lifts with the same quotient but different preferred physical representatives are therefore not automatically identified.

39.3 Section-preserving morphisms

A morphism

$$F : \mathcal{L}_1 \rightarrow \mathcal{L}_2$$

is section-preserving when

$$\Theta_2 F = \Theta_1,$$

$$F \delta_a^{(1)} = \delta_a^{(2)} F,$$

and

$$F \sigma_1 = \sigma_2.$$

For concrete pushdown comparison, a **coordinate LIFO morphism** must also respect the finite-control and stack coordinates letterwise rather than replacing the pushdown by an arbitrary behavioral quotient.

This is the correct category for asking whether two physical lifts are genuinely the same representation.

39.4 Minimality is category-relative

Proposition 39.2 — Category-relative minimality.

There is no meaningful representation-independent notion of a “minimal physical lift” without specifying:

- the admissible physical policy;
- the allowed morphisms;
- whether the canonical section is preserved;
- which stack/control coordinates may be identified.

The same semantic system can have different minima in different categories. This is not a weakness of the theory.

It is the same lesson already encountered in the resource frontiers:

physical optimality is meaningful only relative to a representation category.

The next chapter asks what one-top locality forces inside one fixed literal policy.

Chapter 40

When LIFO Locality Forces Fiber Branching

The Structure Theory proved the top-rewrite suffix criterion.

We now use it at the level of an entire observation fiber.

40.1 Uniform local rule

A deterministic one-top transition cannot inspect the hidden stack suffix.

Therefore all reachable configurations sharing the same locally visible cell must be handled by the same rule.

Theorem 40.1 — Uniform one-top rewrite criterion.

Fix one deterministic local observation class.

A single one-top rule can send every source in that class to its canonical semantic target if and only if the untouched suffix of every source is compatible with that source's canonical target under the suffix criterion.

This theorem is the fiberwise form of Chapter 32.

The obstruction is no longer one bad stack.

It is incompatibility across a family of sources that the local rule cannot distinguish.

40.2 Canonical-section closure dichotomy

Use the fixed literal section

$$\sigma(x, y) = (C, 1^y 0^x Z_0).$$

Theorem 40.2 — Canonical-section closure dichotomy.

The canonical literal section is transition-closed under one-top real-time updates if and only if

$$(p, q) = (1, 1).$$

Meaning. For the balanced one-to-one language, the literal section can remain canonical after every input step.

For every nontrivial coprime ratio, some exact semantic successor cannot be realized while staying on the canonical section.

Thus the use of noncanonical physical representatives is not merely a feature of our centered construction.

Inside this literal one-top policy, it is unavoidable.

40.3 Essential physical fiber branching

Corollary 40.3 — Essential branching.

For every nontrivial ratio

$$(p, q) \neq (1, 1),$$

every exact one-top realization of the fixed literal policy must reach some noncanonical physical configuration in the same semantic fiber as a canonical one.

Therefore:

nontrivial semantic fibers are forced physically

inside the policy.

This is a much stronger statement than:

“our machine happens to use side states.”

It says the canonical section itself cannot be closed.

40.4 Quotient-protected branching

Could one simply identify the noncanonical representative with the canonical one?

Not if the quotient must preserve legal one-top transitions in the same policy.

Corollary 40.4 — Quotient-protected branching.

The forced branch cannot be quotiented away while preserving:

- the fixed literal section;
- the one-top LIFO policy;
- exact transition behavior.

Thus some physical branching is structurally protected inside the representation category.

But the claim remains policy-relative.

Chapter 44 will show that a broader residual re-coordination can remove this branching entirely.

Chapter 41

Two Reduced 2:1 Lifts

We now revisit the eager and lazy 2:1 machines.

Chapter 30 proved semantic equivalence.

The Structure Theory explained their chronology.

The Uniqueness Theory asks whether one is merely a redundant presentation of the other.

The answer is no.

41.1 Historical eager lift

The historical machine eagerly materializes two reachable right-state matching families.

It returns immediately to the canonical center representation.

41.2 Centered lazy lift

The centered specialization leaves those same semantic debts on the right sheet and postpones materialization.

Both lifts use:

- the same semantic residual core;
- the same three control labels C, L, R ;
- the same two work symbols;
- the same canonical section.

Their difference is temporal.

41.3 Section incomparability

Theorem 41.1 — Eager–lazy section incomparability.

There is no section-preserving transition morphism

$$\mathcal{L}_{\text{eager}} \rightarrow \mathcal{L}_{\text{lazy}}$$

and no such morphism

$$\mathcal{L}_{\text{lazy}} \rightarrow \mathcal{L}_{\text{eager}}.$$

The two lifts are therefore genuinely incomparable in the chosen category.

Their semantic quotient is identical.

Their physical transition structures are not.

41.4 Coordinate reducedness

Reachable witnesses force all three control states to remain distinct.

Likewise, the two work symbols cannot be merged without collapsing physically distinguishable reachable behavior.

Proposition 41.2 — Coordinate reducedness.

Both eager and lazy lifts are reduced inside the stated:

three-state / two-work-symbol coordinate category.

Thus their nonisomorphism is not caused by obvious redundant states or symbols.

41.5 Morphism-relative nonuniqueness

Corollary 41.3 — Reduced lift nonuniqueness.

The same semantic system admits multiple:

- coordinate-reduced;
- section-preserving;
- non-isomorphic

physical lifts.

Therefore:

same semantics + same coordinate budget + reducedness
 $\nRightarrow$ unique physical lift.

This strengthens Chapter 36's cheap cyclic-decoration nonuniqueness into a substantive structural theorem.

The next chapter shows that there are not merely two such schedules in the selected 2:1 family.

There are exactly six.

Chapter 42

The Six 2:1 Normalization Schedules

The historical and centered machines differ at specific reachable sites.

To understand the space of schedule choices, isolate three eager-discharge sites.

42.1 Sites

Site *A*

Left-state 0/0 matching family.

Site *B*

Right-state, top 0, input 0 family.

Site *C*

Right-state, top 1, input 1 family.

A schedule is a bit vector

$$(a, b, c) \in \{0, 1\}^3,$$

where 1 means that the corresponding eager rule is enabled.

There are eight possible schedules.

Exactly six are admissible.

42.2 Exact schedule classification

Theorem 42.1 — Exact three-site schedule classification.

$$\text{Adm}_{2,1} = \{(a, b, c) \in \{0, 1\}^3 : ab = 0\}.$$

Thus the six admissible schedules are

$$000, 001, 010, 011, 100, 101.$$

The forbidden schedules are

$$110, 111.$$

The unique minimal conflict is

$$\{A, B\}.$$

Site C is compatible with either side of that conflict.

42.3 Schedule geometry

Order schedules by inclusion of eager sites.

Corollary 42.2 — Schedule poset.

The admissible schedule poset is not a lattice.

Its two maximal elements are

$$\{A, C\}$$

and

$$\{B, C\}.$$

Their union

$$\{A, B, C\}$$

is inadmissible.

So no join exists inside the admissible poset.

The missing 110 and 111 vertices should be shown separately as forbidden.

42.4 Six reduced lifts

Proposition 42.3 — Six coordinate-reduced schedule lifts.

All six admissible schedules define pairwise non-isomorphic coordinate-reduced lifts in the selected 2:1 category.

Thus schedule nonuniqueness is not merely eager versus lazy.

The physical realization space already contains a finite family of six reduced points.




Figure 42.1: Hasse diagram of the six admissible 2:1 schedules. The schedule bits are ordered by inclusion of eager sites A, B, C . The two omitted cube vertices 110 and 111 are shown separately as forbidden because A and B cannot be simultaneously eager.

42.5 Terminal normalization coherence

Let

$\mathfrak{N}$

be the stated terminal canonical normalizer.

Theorem 42.4 — Six-schedule terminal coherence.

For every admissible schedule and every common processed prefix, applying $\mathfrak{N}$ yields the same canonical semantic endpoint.

Thus all six systems have common normalized endpoint behavior.

Terminology boundary

This is **not** a Church–Rosser theorem.

The safe term is:

normalizer-relative terminal schedule coherence.

The six lifts are physically non-isomorphic.

They nevertheless agree after the specified terminal normalization.

Corollary 42.4A — Reduced non-isomorphism with common normalized behavior.

There exist coordinate-reduced, pairwise non-isomorphic physical lifts with identical terminal-normalization behavior on every processed prefix.

Thus even after obvious coordinate redundancy is removed,

common normalized semantics $\not\Rightarrow$ unique reduced physical schedule.

This is exactly the kind of semantic/physical separation that the later General Theory will abstract.

Chapter 43

Why A Causes B

The schedule classification tells us that

$$A + B$$

is forbidden.

It does not yet explain why the conflict is directional.

The answer is one of the clearest causal mechanisms in the corpus.

43.1 Local carry visibility

A fixed eager rule is chosen from a bounded local observation.

If two sources share that observation but require different semantic carry behavior, the same eager rule cannot be backlog-neutral on both.

Theorem 43.1 — Local carry visibility criterion.

A fixed eager local rule can remain backlog-neutral throughout an observation fiber only if the semantic carry indicator is constant on that fiber.

Thus carry aliasing is a local source of schedule failure.

43.2 Portability

For an eager rule r , define its portability set

$$\text{Port}(r)$$

to be the physical source configurations on which applying that rule is safe relative to the selected canonical policy.

If F_o^α is the observation fiber exposed by schedule α , then:

Proposition 43.2 — Eager-rule portability criterion.

The eager rule is portable over that observation class exactly when

$$\boxed{F_o^\alpha \subseteq \text{Port}(r)}.$$

Schedule interactions therefore arise when enabling one rule changes which configurations another rule encounters.

43.3 A expands the B -fiber

Without A , site B sees only its baseline portable contexts.

Enabling A changes earlier normalization chronology.

That change exposes to B the infinite family

$$\boxed{(R, 0^m), \quad m \geq 2.}$$

Lemma 43.3 — Schedule-induced fiber expansion.

Enabling A enlarges the first-hit observation fiber of B by the family $(R, 0^m)$, $m \geq 2$.

This is the causal mechanism.

A does not make B 's local rule intrinsically different.

It changes which hidden suffixes reach the same local B -observation.

43.4 Direct canonicalization becomes impossible

For source

$$(R, 0^m)$$

on input 0, the canonical semantic target is

$$\boxed{10^{m-2}.}$$

A one-top rewrite must preserve the untouched suffix

$$\boxed{0^{m-1}.}$$

But 0^{m-1} is not a suffix of 10^{m-2} .

Theorem 43.4 — Schedule-induced direct-canonicity obstruction.

No one-top rule at site B can directly canonicalize every newly exposed source $(R, 0^m)$, $m \geq 2$.

This is the suffix-theorem diagnosis.

43.5 Forced kernel excess

Suppose the B -rule nevertheless insists on returning directly to the center while preserving residual.

Then it must create a nonminimal amount of literal demand.

Theorem 43.5 — Forced kernel-excess return.

Every such center-return repair introduces at least one extra kernel packet:

$$(t + 1)K$$

relative to the minimum demand.

The historical replacement

$$111$$

realizes the minimum one-packet excess.

Thus the same conflict has three exact readings.

Corollary 43.6 — Three-way diagnosis.

Within this $A \rightarrow B$ family,

suffix failure

$\Longleftrightarrow$

carry aliasing

$\Longleftrightarrow$

forced kernel excess.

This triangle is important because the three notions are not globally identical in arbitrary representations.

They coincide here because of the specific one-top literal policy and source family.

The next chapter turns this concrete explanation into a directed causal relation.

Chapter 44

Directed Causal Scheduling

A minimal conflict such as

$$\{A, B\}$$

is set-theoretic.

It does not identify which rule caused the other to become unsafe.

The first-hit fiber expansion gives a directed notion.

44.1 First-hit exposure

For target site b and previously enabled schedule set H , let

$$\boxed{\mathcal{E}_b(H)}$$

be the first-hit exposure fiber: the configurations at which a computation first reaches site b under the schedule H .

A directed cause

$$H \rightsquigarrow b$$

occurs when H is minimal among schedule sets whose first-hit exposure contains a b -source outside the portability set of the b -rule.

Thus causality is about **newly exposed unsafe contexts**.

44.2 Exact 2:1 primary dependency

Theorem 44.1 — Exact primary dependency classification.

For the selected 2:1 three-site family,

$$\boxed{\mathcal{D}_{2,1}^{\text{prim}} = \{\{A\} \rightsquigarrow B\}.$$

So the directed causal graph is

$$\boxed{A \longrightarrow B \quad \sqcup \quad C.}$$

There is:

- no $B \rightarrow A$;
- no edge involving C .

Thus the forbidden pair $\{A, B\}$ has a unique primary direction.

Corollary 44.2 — Directed cause recovers the conflict.

Forgetting orientation gives the unique minimal schedule conflict:

$$\boxed{A \rightarrow B \quad \mapsto \quad \{A, B\}.}$$

44.3 Concrete causal certificate

One first-hit witness is

$$\boxed{1111100 \rightsquigarrow (R, 00).}$$

Applying the B -rule yields

$$\boxed{(C, 1110)}$$

instead of canonical

$$\boxed{(C, 1).}$$

The quantity excess is one kernel packet in the relevant sign convention:

$$\boxed{-K.}$$

Appending one more 1 reaches residual zero at the noncanonical configuration

$$\boxed{(C, 110).}$$

This gives a finite causal certificate:

1. enabling A creates a new first hit at B ;
2. B 's eager rule is not portable there;
3. the resulting debt is behaviorally exposed by a short continuation.

Proposition 44.2A — 2:1 causal certificate.

There exists a finite word, beginning with the witness prefix

$$\boxed{1111100,}$$

whose first use of the B -site after enabling A reaches

$$(R, 00)$$

and whose eager B -successor is the noncanonical kernel-excess configuration

$$(C, 1110).$$

A finite continuation then distinguishes this successor behaviorally from the canonical semantic successor.

Hence the directed dependency

$$\boxed{A \rightsquigarrow B}$$

admits a finite first-hit certificate.

This concrete witness is the prototype for the Fatal Frontier certificates of the General Theory.

44.4 Coordinate invariance

Suppose two lift presentations are related by a site-preserving coordinate LIFO isomorphism.

Then:

- first-hit exposure fibers correspond;
- portability sets correspond;
- directed causal hyperedges correspond.

Theorem 44.3 — Causal hypergraph invariance.

The directed causal schedule structure is invariant under site-preserving coordinate isomorphism.

Thus the $A \rightarrow B$ graph is not an artifact of renaming states or stack symbols inside the same coordinate category.

44.5 But semantics does not determine causality

A raw residual-conjugacy machine can realize the same language semantic quotient without any literal A, B, C normalization sites.

Theorem 44.4 — Semantic non-determination of causal structure.

The same semantic quotient can admit physical lifts with different causal schedule structures.

Therefore:

$$\text{semantic quotient} \not\Rightarrow \text{causal schedule signature.}$$

Corollary 44.5 — Coordinate/semantic separation principle.

A property can be invariant under coordinate isomorphism and still fail to be language-intrinsic.

Thus:

$$\text{coordinate-intrinsic} \not\Rightarrow \text{semantic-intrinsic.}$$

This distinction is central to every later obstruction theorem.

Chapter 45

Which Obstructions Are Really Forced?

The Uniqueness Theory has proved several kinds of necessity.

But necessity depends on the class of allowed representations.

This final chapter of the Part makes that dependence explicit.

45.1 Nested representation classes

Consider nested classes

$$\mathfrak{C}_0 \subseteq \mathfrak{C}_1 \subseteq \mathfrak{C}_2.$$

A typical interpretation is:

- $\mathfrak{C}_0$: a narrow coordinate category;
- $\mathfrak{C}_1$: a fixed literal-demand policy with more presentation freedom;
- $\mathfrak{C}_2$: a broader realization class allowing re-coordination.

For an obstruction $\mathcal{O}$, ask whether $\mathcal{O}$ is unavoidable in each class.

45.2 Obstruction monotonicity

Theorem 45.1 — Obstruction monotonicity.

If the realization class is enlarged, an essential obstruction can disappear but cannot newly become essential.

Thus:

$$\text{Ess}(\mathfrak{C}_2) \subseteq \text{Ess}(\mathfrak{C}_1) \subseteq \text{Ess}(\mathfrak{C}_0).$$

For fully resolved three-level binary profiles, only

$$\boxed{111, \quad 110, \quad 100, \quad 000}$$

are possible.

45.3 Branching profile

For nontrivial semantic-fiber branching in the stated nested classes:

Theorem 45.2 — Policy-essential but broad-class removable branching.

$$\boxed{\Pi(\mathcal{O}_{\text{branch}}) = 110.}$$

Thus branching is:

- essential in the narrow coordinate class;
- essential in the fixed literal policy;
- removable in a broader class containing raw residual re-coordination.

This theorem connects the early raw DPDA constructions to the later lift theory.

The compact raw representation did not contradict forced branching.

It left the literal-demand policy in which branching was forced.

45.4 Current status of $A \rightarrow B$

The specific causal obstruction has only a partially resolved profile.

Proposition 45.3 — Current $A \rightarrow B$ status.

$$\boxed{\Pi(A \rightarrow B) = (1, ?, 0).}$$

We know:

- it is essential in the narrow coordinate category;
- it disappears in a broad enough class;
- whether it is essential throughout the intermediate fixed literal policy remains open.

The question mark must remain a question mark.

The monograph must not silently turn this into 110.

45.5 Representation–obstruction separation

Theorem 45.4 — Representation-obstruction separation.

Every theorem claiming that a physical obstruction is “forced” must specify the representation class in which the necessity is proved.

This is one of the main methodological conclusions of the entire project.

It applies to:

- state lower bounds;
- stack-symbol lower bounds;
- fiber branching;
- normalization debt;
- schedule conflicts;
- causal edges.

45.6 Open problems retained

The following remain open at this stage:

1. fixed-policy essentiality of the specific $A \rightarrow B$ obstruction;
2. general downward closure of arbitrary schedule families;
3. general $p:q$ causal hypergraph classification;
4. classical Church–Rosser/confluence theory for these normalization systems;
5. unrestricted DPDA lift minimality.

These questions are not defects in the monograph.

They mark the boundary between proved structure and the remaining research program.

45.7 Reproducibility note

The P11 manuscript reports a regression script

```
uniqueness_lift_validation_v0_2.py
```

with

458,776

assertion-level checks, including exhaustive words through length 14 for all eight schedules and the first common forbidden-schedule failure.

That script is not present in the provided corpus archive.

Therefore the current monograph records the computational claim as source-reported, not independently rerun from the supplied package.

Before final publication, the project should either:

- retrieve the repository artifact and rerun it; or
- rebuild an equivalent validation suite.

The symbolic proofs do not depend on the missing script.

45.8 Transition to General Theory

The Uniqueness Theory has now produced a family of concepts that are no longer specific to weighted binary balance:

- semantic transition system;
- exact physical lift;
- canonical section;
- semantic fiber;
- bounded local observation;
- forced branching;
- portability;
- causal schedule structure;
- representation-relative obstruction.

The next Part asks:

Which of these notions survive for a general semantic system and a bounded-local physical normalization policy?

The answer will organize them into:

locality $\rightarrow$ debt $\rightarrow$ behavioral safety $\rightarrow$ causality $\rightarrow$ synthesis $\rightarrow$ verification.

That is the General Theory.

PART XI

General Theory of Exact Pushdown Normalization

Research provenance. The abstract exact-lift framework, bounded-local obstruction theory, behavioral safety, causal scheduling, synthesis, and verification hierarchy are developed in [11]. Weighted-balance families are stress tests rather than the definition of the theory. Classical neighboring machinery includes visibly pushdown semantics, pushdown/finite-state equivalence frameworks, pushdown saturation, and deterministic pushdown language equivalence [15, 49, 19, 61].

Chapter 46

Semantic Fibers and Canonical Retraction

The preceding Parts reached the same conclusion repeatedly from increasingly general directions:

semantic evolution can be unique while physical normalization is not.

The General Theory now removes the weighted-balance specifics.
Its objects are:

- an abstract semantic transition system;
- an exact physical lift;
- a chosen section;
- bounded-local physical update rules;
- explicitly stated representation categories.

The word **General** is scoped accordingly.
This Part does not claim:

- a theory of every pushdown automaton;
- existence of a global section for every correct DPDA;
- unrestricted DPDA minimization;
- unrestricted infinite-lift synthesis;
- universal schedule downward closure;
- universal confluence.

46.1 Exact semantic lift

Let M be a semantic state set with deterministic transitions

$$U_a : M \rightarrow M.$$

Let C be a reachable physical configuration set with transitions

$$\delta_a : C \rightarrow C.$$

An exact lift has a surjective semantic projection

$$\boxed{\Theta : C \rightarrow M}$$

satisfying

$$\boxed{\Theta \delta_a = U_a \Theta.}$$

Choose a section

$$\boxed{\sigma : M \rightarrow C}$$

with

$$\Theta \sigma = \text{id}.$$

Define the canonical retraction

$$\boxed{N = \sigma \Theta.}$$

46.2 Canonical fibers

Theorem 46.1 — Canonical fiber and retraction.

The map N satisfies

$$\boxed{N^2 = N,}$$

$$\boxed{\Theta N = \Theta,}$$

and

$$\boxed{\text{Fix}(N) = \text{im } \sigma.}$$

Moreover,

$$\boxed{\Theta(c) = \Theta(c') \iff N(c) = N(c').}$$

Thus the fibers of the canonical retraction are exactly the semantic fibers.
Proof.

Idempotence follows from

$$\Theta\sigma = \text{id}.$$

If $N(c) = c$, then $c = \sigma(\Theta(c))$, so c lies in the section image.

Conversely every section point is fixed.

Finally, equal semantic values clearly give equal canonical representatives, and equal canonical representatives project to equal semantic states. $\square$

This is the final abstract version of the canonical-fiber idea first seen concretely in WSC.

46.3 Canonical absorption

Theorem 46.2 — Canonical absorption.

For every input word w ,

$$N\delta_w = N\delta_w N = \sigma U_w \Theta.$$

Thus the canonical endpoint after reading w depends only on the initial semantic state.

All prior physical normalization history is forgotten by the final retraction.

This theorem is the general form of:

- eager/lazy semantic equivalence;
- WSC canonical absorption;
- terminal normalization coherence.

It remains weaker than an unrestricted confluence theorem.

46.4 Zero debt and raw conjugacy

A section is **transition-closed** when

$$\delta_a \sigma = \sigma U_a$$

for every input symbol.

Theorem 46.3 — Zero-debt/raw-conjugacy equivalence.

Under the stated reachability hypotheses, the following are equivalent:

1. the chosen section is transition-closed;
2. every reachable physical configuration lies on the section:

$$C^{\text{reach}} = \sigma(M).$$

Thus zero representation debt means that the reachable physical transition graph is literally the chosen semantic section.

The reachability hypotheses matter.

Mere bijectivity of Θ on reachable configurations is not enough if the chosen section uses different representatives.

46.5 Three levels now separated permanently

From this point onward we distinguish:

semantic state

canonical section representative

arbitrary reachable physical representative in the same fiber.

The next chapter asks whether a bounded-local physical policy can keep every semantic transition on the chosen section.

Chapter 47

Bounded-Local Realizability and the Three-Stage Obstruction

The top-rewrite suffix criterion was discovered in the literal WSC setting.

We now generalize the principle.

A bounded-local physical transition observes only finite local information and must update the hidden remainder in a restricted, tail-injective manner.

The fundamental question is:

Can every semantic predecessor of one target be assigned a locally distinguishable canonical source whose physical update lands directly on the canonical target?

There are three distinct ways this can fail.

47.1 Uniform one-top portability

For ordinary one-top rewriting, sources with the same local observation must share one rewrite rule.

Lemma 47.1 — Uniform one-top portability.

A finite family of one-top sources can all be sent directly to their chosen canonical targets by one deterministic local rule exactly when their untouched suffixes are jointly compatible with that rule.

For a single source, this specializes to the suffix criterion of Chapter 32.

We do not re-prove the single-source theorem here.

47.2 Finite observation capacity

Let Ω_a be the finite set of local observations available on input a .

Suppose a semantic target m^* has many one-symbol predecessors:

$$U_a^{-1}(m^*).$$

If the chosen section is transition-closed, distinct predecessor sources that require different hidden-tail behavior must fit into the available observation classes.

Theorem 47.2 — Finite-observation collapse capacity.

If

$$\boxed{|U_a^{-1}(m^*)| > |\Omega_a|},$$

then no tail-injective bounded-local policy can realize a transition-closed section over all those predecessors.

The obstruction occurs before any detailed stack alignment is considered.

It is a pure semantic-to-local capacity overflow.

Corollary 47.3 — Infinite-collapse obstruction.

If one semantic target has infinitely many relevant predecessors while the observation set is finite, total transition closure is impossible.

47.3 Bounded- d -top pushdown capacity

For a pushdown policy that can inspect at most d stack symbols, the number of local observations is bounded by the finite combinations of:

- control state;
- top window of length at most d ;
- input symbol.

This yields an explicit observation-capacity lower bound in terms of

$$|Q|, \quad |\Gamma|, \quad d.$$

Corollary 47.4 — Bounded- d -top capacity.

A semantic collapse fiber larger than the available state/window observation capacity cannot be represented by a transition-closed bounded- d -top section.

This is the category-level generalization of the earlier finite-control debt lower bounds.

47.4 The three stages

Theorem 47.5 — Three-stage bounded-local obstruction decomposition.

Failure of direct canonical bounded-local realizability occurs at one of three ordered stages.

Stage I — Semantic capacity overflow

There are more same-target semantic predecessor roles than the observation policy can distinguish.

Stage II — Chosen-section observation collision

Capacity is sufficient in principle, but the chosen section assigns two same-target sources to the same local observation even though their canonical updates require different physical behavior.

Stage III — Tail/action alignment failure

The sources are locally separated, but the actual bounded-local rewrite cannot transform the hidden tail into the desired canonical target.

For one-top pushdown rewriting, Stage III becomes exactly the suffix/tail obstruction.

This theorem organizes what earlier papers saw separately:

- finite-control width pressure;
- fiber expansion;
- local aliasing;
- suffix failure.

47.5 One-top specialization

Corollary 47.6 — One-top specialization.

In the one-top literal setting, the Stage III test is exactly:

untouched suffix is a suffix of the canonical target.

Thus the simple theorem of Chapter 32 is the one-top boundary case of the general locality decomposition.

47.6 Locality–debt–recoordination trilemma

Theorem 47.7 — Trilemma.

When direct canonical realization fails, an exact implementation cannot keep all three of the following fixed:

1. the local observation regime;
2. the chosen canonical section;
3. direct canonical successor behavior.

At least one response is necessary:

refine observation

or

allow representation debt

or

re-coordinate the section.

This theorem summarizes a large fraction of the monograph:

- Chapters 6–8 used re-coordination;
- Chapters 26–37 used debt;
- broader observation would use more physical resources.

The next chapter asks an equally important question:

If debt is forced, is it necessarily dangerous?

The answer is no.

Chapter 48

Representation Debt and Behavioral Safety

A physical configuration can be noncanonical and still behave perfectly.

The General Theory therefore separates:

representation debt

from

behavioral incorrectness.

This distinction prevents a common mistake:

“The machine left the canonical section, therefore it is already wrong.”

That implication is false.

48.1 Terminal semantics

Let

$$f : M \rightarrow O$$

be the semantic terminal observable and

$$\alpha : C \rightarrow O$$

the physical terminal observable.

A physical configuration c is behaviorally correct for all continuations if

$\alpha(\delta_w(c)) = f(U_w(\Theta(c)))$

for every continuation w .

48.2 Immediate disagreement

Define the immediate disagreement set

$$\mathcal{D} = \{c : \alpha(c) \neq f(\Theta(c))\}.$$

A configuration can be outside $\mathcal{D}$ now but still have a future continuation reaching $\mathcal{D}$.

Thus immediate agreement is not enough.

Proposition 48.2A — Terminal factorization.

The identity

$$\alpha = f\Theta$$

holds on all physical configurations if and only if:

1. α is constant on every semantic fiber of Θ ; and
- 2.

$$\alpha\sigma = f.$$

Proof.

If $\alpha = f\Theta$, equal semantic projections have equal terminal values, and

$$\alpha\sigma = f\Theta\sigma = f.$$

Conversely, let c be any physical configuration. Fiber constancy gives

$$\alpha(c) = \alpha(\sigma(\Theta(c))).$$

Using $\alpha\sigma = f$,

$$\alpha(c) = f(\Theta(c)).$$

□

This proposition separates two requirements that are easy to conflate:

terminal correctness of the chosen section

and

terminal constancy across every physical representative of a semantic fiber.

48.3 Behavioral hull

Define

$$\mathcal{B} = \{c : \alpha(\delta_w(c)) = f(U_w(\Theta(c))) \text{ for all } w\}.$$

Theorem 48.1 — Behavioral hull characterization.

$$\mathcal{B} = \nu X. \Phi(X) = C \setminus \text{Pre}^*(\mathcal{D}).$$

That is:

- $\mathcal{B}$ is the greatest forward-safe fixed point;
- a configuration is behaviorally safe exactly when no continuation reaches immediate disagreement.

This is the exact physical correctness object.

Corollary 48.2 — Root correctness.

A fixed policy is correct from root c_0 exactly when

$$c_0 \in \mathcal{B}.$$

48.4 Forward invariance

Theorem 48.3 — Forward invariance of behavioral safety.

If

$$c \in \mathcal{B},$$

then every reachable successor

$$\delta_w(c)$$

also lies in $\mathcal{B}$.

Thus safe debt remains safe under the fixed policy.

48.5 Debt birth dichotomy

Theorem 48.4 — Debt birth dichotomy.

When a transition first leaves the canonical section, the resulting debt is born in exactly one of two regimes:

1.

$$\boxed{c \in \mathcal{B}}$$

— behaviorally harmless debt;

2.

$$\boxed{c \notin \mathcal{B}}$$

— debt that is already behaviorally unsafe.

A later continuation does not transform safe debt into unsafe debt.

If a continuation eventually reveals an error, the configuration was outside $\mathcal{B}$ from the moment the unsafe debt was created.

This corrects the tempting narrative:

“debt becomes fatal later.”

The correct statement is:

$\boxed{\text{a later suffix reveals pre-existing unsafety.}}$

48.6 Two opposite examples

Weighted-balance schedule conflict

Some noncanonical debts produced by forbidden eager schedules lie outside the behavioral hull and are exposed by short continuations.

Multi-type Dyck collapse

At a rejecting collapse boundary, bounded-locality may force noncanonical physical representatives, yet every such state can lie in an absorbing rejecting semantic fiber.

The debt is forced but harmless.

Thus:

$\boxed{\text{forced debt} \not\approx \text{fatal debt.}}$

The next chapter develops semantic certificates for proving harmlessness without constructing the exact physical behavioral hull.

Chapter 49

Semantic Containment Cores

The behavioral hull

$$\mathcal{B}$$

is exact.

But it lives in the full physical configuration space.

A coarser proof certificate can sometimes be obtained entirely at the semantic level.

49.1 Semantic trap containment

Let

$$S \subseteq M$$

be forward invariant:

$$U_a(S) \subseteq S.$$

Suppose every physical representative in every fiber over S is immediately terminally correct.

Theorem 49.1 — Semantic trap containment.

Under those hypotheses,

$$\boxed{\Theta^{-1}(S) \subseteq \mathcal{B}.}$$

Proof idea. A run beginning over S remains over S semantically.

Every reached physical representative lies in a fiber whose entire ambient fiber is terminally correct.

Therefore no continuation can reach disagreement. $\square$

This is conservative because it requires **every** representative in the full fiber to be safe.

49.2 Absorbing-fiber harmlessness

Corollary 49.2 — Absorbing-fiber harmlessness.

If a semantic state is absorbing and all physical representatives in its fiber give the correct terminal observation, then the entire fiber lies in the behavioral hull.

This is the clean explanation of multi-type Dyck collapse debt:

- the debt can be forced physically;
- the semantic state is already an absorbing reject;
- the entire fiber is harmless.

49.3 Maximal full-fiber containment core

Define the semantic good set of states whose full fibers are terminally correct.

Iteratively remove any semantic state whose successor leaves the current set.

Theorem 49.3 — Maximal fiberwise semantic containment core.

There exists a greatest forward-invariant semantic set

$$\boxed{\mathcal{T}}$$

such that

$$\boxed{\Theta^{-1}(\mathcal{T}) \subseteq \mathcal{B}.}$$

This is the maximal **ambient full-fiber** semantic containment certificate. It is deliberately conservative.

An unreachable bad representative can remove an otherwise useful semantic state from $\mathcal{T}$.

49.4 Reachable-fiber refinement

Let

$$F_m^{\text{reach}}$$

be only the representatives actually reachable under the policy.

Use those reachable fibers in the containment construction.

Theorem 49.4 — Reachable semantic containment.

There is a greatest reachable-fiber containment core

$$\boxed{\mathcal{T}^{\text{Reach}}}$$

whose reachable physical representatives are behaviorally safe.

It satisfies the conceptual hierarchy

$$\boxed{\mathcal{T} \subseteq \mathcal{T}^{\text{Reach}}.}$$

The reachable core removes ambient syntactic contamination.

49.5 Root correctness via containment**Theorem 49.5 — Root correctness via reachable containment.**

If the root semantic state belongs to

$$\mathcal{T}^{\text{Reach}},$$

then the physical root lies in

$$\mathcal{B}.$$

Thus the three correctness scales are:

$$\boxed{\mathcal{T} \longrightarrow \mathcal{T}^{\text{Reach}} \longrightarrow \mathcal{B}.}$$

They answer increasingly exact questions:

- ambient semantic certificate;
- reachable semantic certificate;
- exact physical behavioral safety.

The next chapter returns to normalization schedules and uses the behavioral hull to distinguish merely structural unsafe targets from genuinely fatal ones.

Chapter 50

Safe Reset and the Fatal Frontier

The 2:1 causal theory introduced first-hit exposure and portability.

The General Theory now separates two notions:

1. a structural reset certificate;
2. exact behavioral safety.

This separation is essential because leaving a convenient reset domain need not be fatal.

50.1 Wordwise safe-reset domain

Let $\mathcal{R} \subseteq C$ be a set of physical configurations.

Informally, $\mathcal{R}$ is a safe-reset domain if every run that stays within the certified reset behavior can be normalized correctly word by word.

Proposition 50.1 — Safe-reset certificate.

If $\mathcal{R}$ is a wordwise safe-reset domain for a schedule family $\mathfrak{H}$, then every schedule whose eager activations remain within the certified reset behavior is correct.

The domain is a proof certificate.

It need not equal the exact behavioral hull.

50.2 Structural unsafe targets

For schedule H , let

$$U_{\mathcal{R}}(H)$$

be the set of actual eager activation targets that lie outside $\mathcal{R}$.
 These are structurally uncertified targets.
 Some may still lie in the behavioral hull.
 Define that difference as behavioral slack.
 The exact badness criterion therefore needs both layers.

50.3 Fatal Frontier

Theorem 50.2 — Fatal Frontier.

Assume $\mathcal{R}$ is a wordwise safe-reset domain for the full schedule family.
 Then

$$H \notin \text{Adm} \iff U_{\mathcal{R}}(H) \setminus \mathcal{B}_H \neq \emptyset.$$

Thus a schedule is bad exactly when it actually produces an eager target that is:

1. outside the reset certificate;
2. outside the exact behavioral hull.

This is the abstract form of what Chapter 44 saw concretely.

A structural deviation becomes a **fatal frontier event** only when no behavioral slack saves it.

50.4 Activation tightness

Corollary 50.3 — Activation tightness.

If every actual eager target outside $\mathcal{R}$ is behaviorally unsafe, then

$$H \notin \text{Adm} \iff U_{\mathcal{R}}(H) \neq \emptyset.$$

In that case the structural unsafe frontier is exact.

Without activation tightness, structural detection is only sufficient, not necessary.

50.5 Site-free diagnostics

Sometimes one can certify fatality by a continuation that uses no additional optional eager site.

Theorem 50.4 — Site-free diagnostic frontier.

Under the stated safe-reset hypotheses, a finite site-free continuation can certify that an activation target is already outside the behavioral hull.

This gives a causal witness local to the activation event.

The 2:1 certificate

$$1111100 \rightsquigarrow (R, 00)$$

followed by the short exposing suffix is the motivating example.

50.6 Fatal causality

Proposition 50.5 — Fatal causality implies conflict.

If a minimal enabled site set H creates a fatal first-hit target at site b , then

$$H \cup \{b\}$$

is an inadmissible schedule set.

Thus directed fatal cause induces set-theoretic schedule conflict.

The next chapter develops the combinatorics of those conflicts and stress-tests the theory on 3:2 and an infinite odd- p :2 family.

Chapter 51

Conflict Complexes and Weighted-Balance Stress Tests

The Fatal Frontier tells us when one schedule is wrong.

We now ask how the whole family of schedules is organized.

51.1 Minimal conflicts

A conflict is a set of optional sites that cannot be enabled together.

A conflict is minimal if all of its proper subsets are admissible.

Lemma 51.1 — Minimal-conflict participation.

Every site belonging to a minimal conflict participates in some fatal mechanism for that conflict.

No member can be completely irrelevant.

Otherwise the conflict would not be minimal.

51.2 Robust fatal supports

Suppose a family of minimal site sets has fatal certificates that remain valid whenever additional sites outside the support are enabled.

These are robust fatal supports.

Theorem 51.2 — Robust fatal supports give an independence complex.

Under the robust-certificate hypotheses, the admissible schedule family is exactly the family of site subsets containing no minimal fatal support.

Therefore it is the independence complex of the hypergraph of minimal fatal supports.

The phrase **independence complex** is combinatorial packaging.

The project-specific content is the derivation of the fatal supports from physical normalization causality.

Corollary 51.3 — Finite diagnostic basis.

If the robust minimal fatal support family is finite, the entire finite schedule classification has a finite diagnostic basis.

51.3 Reset-ray fork

Theorem 51.4 — Reset-ray fork template.

Suppose one site D changes the reachable reset-ray geometry so that it exposes fatal first hits at two victim sites A and B , while schedules not containing D remain reset-safe and schedules containing D but neither victim remain safe.

Then the minimal fatal supports are

$$\boxed{\{A, D\}, \quad \{B, D\}.$$

The admissible schedules are the independence complex avoiding those two pairs.

This template abstracts the directed fork mechanism later verified uniformly.

51.4 Imported weighted-balance foundations

The General Theory does **not** re-prove:

- primitive-kernel minimum demand;
- centered residual exactness;
- centered canonical-episode terminal correctness.

Those results were already established in the $p:q$ construction.

Here they are used as a stress-test environment for the new causal machinery.

51.5 Exact 3:2 case

A finite reset analysis produces an exact certificate for the centered 3:2 schedule family.

Lemma 51.5A — Exact 3:2 finite reset certificate.

For the centered 3:2 schedule family there is an explicit finite reset repertoire whose episode closure certifies every schedule avoiding the two fatal pairs

$$\boxed{\{A, D\}, \quad \{B, D\}.$$

The same finite analysis produces first-hit fatal witnesses when either forbidden pair is enabled.

This lemma is the finite reset engine behind the exact schedule count; the detailed reset catalogue is retained for the technical appendix rather than repeated in the narrative.

Theorem 51.5 — Exact 3:2 schedule classification.

There are

40 admissible schedules

and

24 inadmissible schedules.

The minimal fatal fork is

$D \rightarrow A, \quad D \rightarrow B.$

Equivalently, the minimal fatal supports are

$\{A, D\}, \quad \{B, D\}.$

The technical finite reset catalogue supporting this result belongs in the appendix; the classification remains in the main narrative.

51.6 Infinite odd- p :2 family

Now let

$$q = 2$$

and

$$p \geq 3$$

be odd.

The source paper derives:

- a detailed arithmetic catalogue;
- source geometry;
- reset closure;
- uniform D -induced fiber expansion;
- reachability of a common length-three fatal frontier;
- kernel-suffix fatality.

The long arithmetic/source catalogues are appendix material.
 The structural conclusion is main-text material.

Lemma 51.6A — Odd- p :2 reset closure.

Let

$$q = 2, \quad p = 2h - 1 \geq 3.$$

There are two explicit safe reset repertoires:

- $\mathcal{R}_{\neg D}$, used when site D is absent;
- $\mathcal{R}_D$, used when D is present but A and B are absent.

They satisfy

$$D \notin H \implies H \in \text{Adm}_{p,2},$$

and

$$D \in H, \quad A, B \notin H \implies H \in \text{Adm}_{p,2}.$$

This accounts for

$$32 + 8 = 40$$

safe schedules before the fatal pairs are analyzed.

Lemma 51.6B — Uniform D -induced fiber expansion.

For every

$$x \geq h + 1,$$

the D -activated canonical ray satisfies

$$(C, \eta(x, h)) \xrightarrow{0^{2h-1}} (D_{-h}, 0^{x-h+1}).$$

Therefore enabling D expands the common A/B observation fiber from the singleton unary source

$$\{0\}$$

to

$$\{0, 00, 000, \dots\}.$$

This is the uniform directed mechanism

$$\boxed{D \rightsquigarrow A, \quad D \rightsquigarrow B.}$$

Lemma 51.6C — Uniform reachability of the length-three frontier.

For every

$$h \geq 2,$$

there is an explicit prefix u_h , using no optional eager site except D , that reaches the common fatal-frontier source

$$\boxed{(D_{-h}, 000).}$$

One convenient construction is

$$\boxed{u_h = v_h 1^{2h-1} 0^{2h},}$$

where v_h is the parity-dependent prefix that first reaches

$$(D_{h-1}, \varepsilon).$$

Thus the same length-three unary source is not merely syntactically possible; it is uniformly reachable across the entire odd- $p:2$ family.

Lemma 51.6D — Uniform kernel-suffix fatality.

For

$$p:2,$$

the primitive kernel is

$$\boxed{K = (2, p).}$$

The zero-weight physical kernel word

$$\boxed{1^p 00}$$

appears as a suffix in the eager A - and B -targets generated from the common source

$$(D_{-h}, 000).$$

A site-free continuation removes the non-kernel prefix and exposes this residual-zero kernel excess as behavioral disagreement.

Hence both A and B are uniformly fatal after the D -induced fiber expansion.

Appendix migration note

Two long supporting results are intentionally not expanded here:

- the **odd- p :2 arithmetic catalogue**;
- the **odd- p :2 source geometry catalogue**.

They are retained as appendix-level proof infrastructure. Their purpose is to justify the closed forms and source restrictions used in Lemmas 51.6A–51.6D, not to create a second narrative inside the main theorem.

Theorem 51.6 — Odd- p :2 fatal-fork family.

For every odd

$$p \geq 3,$$

$$\text{Adm}_{p,2} = \text{Ind}\{\{A, D\}, \{B, D\}\}.$$

Thus every member of the infinite family has:

$$40 \text{ admissible,} \quad 24 \text{ inadmissible}$$

schedules and the same directed fatal fork

$$D \rightarrow A, \quad D \rightarrow B.$$

This is an important stress test.

The General Theory is not merely repackaging the 2:1 example.

Its causal mechanism persists uniformly across an infinite parameter family.

The next chapter changes direction from diagnosis to construction:

Can a transition-closed section be synthesized when one exists?

Chapter 52

Rooted Section Synthesis

The theory has so far explained why canonical sections fail.

It now becomes constructive.

Given:

- a semantic system;
- a physical transition template;
- a chosen physical representative of the root;

can we synthesize a transition-closed section?

The answer is an exact path-outcome coherence criterion.

52.1 Rooted path-outcome fibers

Fix semantic root

$$m_0$$

and physical representative

$$c_0$$

with

$$\Theta(c_0) = m_0.$$

For semantic state m , define

$$\boxed{\mathcal{H}_{c_0}(m)}$$

to be the set of all physical outcomes obtained by following physical transitions along semantic words that take m_0 to m .

Different semantic paths to the same state can produce different physical endpoints.

That is exactly the synthesis obstruction.

52.2 Path-outcome coherence

Theorem 52.1 — Rooted path-outcome coherence criterion.

A rooted transition-closed section through c_0 exists if and only if

$$\boxed{|\mathcal{H}_{c_0}(m)| = 1}$$

for every reachable semantic state m .

When it exists, the section is uniquely determined by the root representative.

Proof.

If a transition-closed section exists, every semantic path to m must end at the single section point $\sigma(m)$.

Conversely, suppose every path-outcome fiber is a singleton.

Define $\sigma(m)$ to be that unique physical outcome.

If $m \xrightarrow{a} U_a(m)$, append a to any path reaching m .

Path-outcome uniqueness gives

$$\delta_a(\sigma(m)) = \sigma(U_a(m)).$$

Thus the section is transition-closed. $\square$

Corollary 52.2 — One root representative determines the section.

For a fixed physical root candidate, there is at most one rooted transition-closed section.

52.3 Finite synthesis algorithm

For finite semantic systems, the theorem yields an exact propagation procedure.

Start from a root candidate.

Propagate its physical image along semantic edges.

Whenever a semantic state is reached again:

- if the physical endpoint agrees, continue;
- if it differs, synthesis fails.

Theorem 52.3 — Finite rooted synthesis.

For a finite semantic transition graph, this propagation decides exactly whether a rooted transition-closed section exists and constructs it when it does.

The algorithm is graph-theoretic once the physical transition template is fixed.

52.4 Short obstruction certificate

Theorem 52.4 — Short parallel-path obstruction certificate.

If finite rooted synthesis fails, there exist two semantic paths from the root to the same semantic state whose physical endpoints differ.

A shortest such pair gives a finite bounded certificate of failure.

Thus synthesis failure is witnessed by path disagreement.

52.5 Counting sections

Corollary 52.5 — Counting synthesized sections.

If the physical root fiber is finite, the number of rooted transition-closed sections equals the number of root representatives that pass the path-outcome coherence test.

The theory therefore turns section existence into a concrete search over root candidates.

52.6 Infinite semantic cores require symbolic synthesis

The coherence theorem itself does not require finiteness.

But brute-force propagation over an infinite semantic graph is not an algorithm.

The next chapter shows one important infinite case where a finite symbolic rule skeleton can be synthesized exactly.

That construction also demonstrates again that re-coordination can remove debt.

Chapter 53

Symbolic Re-Coordination and Forced-Debt Frontiers

The locality trilemma gave three responses to a failed canonical section:

refine observation, allow debt, re-coordinate.

This chapter gives a constructive example of the third option, then quantifies what remains unavoidable when only a partial section can be maintained.

53.1 Modular quotient/remainder synthesis

Consider integer translation semantics generated by fixed signed weights.

Choose a modulus large enough to absorb every one-step weight change.

Represent each semantic integer by:

- a finite sign/remainder state;
- a unary quotient stack.

Theorem 53.1 — Modular quotient/remainder synthesis.

There exists an explicit one-top zero-debt lift in which:

- the finite control stores the signed remainder;
- the stack stores the unbounded quotient;
- every semantic translation is realized directly;
- the chosen section is transition-closed.

This is a general constructive example of:

re-coordination removes normalization debt.

Scope

The construction is:

- pseudo-polynomial in the numerical weight bound;
- not polynomial in the bit length of weights;
- not a state-minimality theorem;
- not a general algorithm for arbitrary infinite lifts.

53.2 Finite symbolic skeleton

Theorem 53.2 — Finite symbolic skeleton and global path-outcome coherence.

The modular representation is governed by a finite set of symbolic local rules.

Those rules satisfy global rooted path-outcome coherence on the infinite semantic core.

Thus an infinite transition-closed section can sometimes be certified by finite symbolic arithmetic.

Corollary 53.3 — Zero-weight loop coherence.

Every semantic zero-weight loop returns the synthesized physical representation to the same section point.

This is an immediate path-outcome consequence, not a separate confluence theorem.

53.3 Partial canonical sections

Total transition closure may be impossible.

Let

$$G \subseteq M$$

be a semantic region over which a chosen section is transition-closed.

Theorem 53.4 — Canonical interior.

As long as the semantic run remains inside G , the physical run remains exactly on the chosen canonical section.

Representation debt is born only on an edge leaving the canonical interior. This turns debt creation into a boundary phenomenon.

53.4 Canonical-hit capacity

Suppose many semantic source states share one target under the same input.

Only finitely many source representatives can land directly on one canonical target if the local observation space is finite.

Theorem 53.5 — Canonical-hit capacity.

For a fixed bounded-local policy, the number of source edges that can hit one chosen canonical target directly is bounded by the available local observation capacity.

This theorem no longer assumes total transition closure.

It applies specifically at the boundary of a partial section.

53.5 Quantitative forced debt

Let $\mathcal{S}$ be a finite set of relevant same-target source states and let Ω_a be the input- a observation set.

Corollary 53.6 — Quantitative forced-debt frontier.

At least

$$|\mathcal{S}| - |\Omega_a|$$

source edges are forced to leave the chosen section.

If $\mathcal{S}$ is infinite and Ω_a is finite, infinitely many debt-producing edges are forced.

This is the quantitative form of the earlier fiber-branching theorem.

53.6 Multi-type Dyck frontier

For one bracket type,

$$\text{Dyck}_1 : \text{total canonical synthesis.}$$

For

$$k \geq 2,$$

the multi-type Dyck semantic system has an infinite collapse boundary that exceeds bounded local observation capacity.

Theorem 53.7 — Dyck forced-debt frontier.

For

$$k \geq 2,$$

there is:

zero-debt live interior + infinite forced-debt collapse boundary.

Yet that boundary lies in an absorbing rejecting semantic fiber.

Therefore the forced debt is behaviorally harmless.

This example is crucial because it separates:

forced physical noncanonicity

from

behavioral failure.

The theory can therefore quantify unavoidable debt without equating debt with error.

Chapter 54

Effective Verification Hierarchy

The General Theory ends with algorithms.

The objective is not to claim new classical automata algorithms.

The contribution is a hierarchy of reductions from normalization correctness to established decision problems.

There are three principal verification levels.

54.1 Level I — Structural safe-reset saturation

Suppose a regular set Bad_{str} contains every structurally unsafe eager target outside a certified safe-reset domain.

Theorem 54.1 — Regular structural safe-reset certification.

If pushdown reachability saturation proves that no configuration in Bad_{str} is reachable under the schedule, then the safe-reset certificate proves policy correctness.

This test is:

sufficient in general.

Under activation tightness, it becomes exact.

This is the cheapest level when a good structural reset certificate is available.

54.2 Level II — Exact immediate-disagreement saturation

Suppose the exact immediate disagreement set

$$\mathcal{D}_H$$

is regular and effectively represented.

Theorem 54.2 — Regular immediate-disagreement saturation.

Then

$$\mathcal{B}_H = C \setminus \text{Pre}_H^*(\mathcal{D}_H)$$

is regular and effectively constructible by pushdown predecessor saturation.
Thus the exact behavioral hull can be computed.

Corollary 54.3 — Exact root classification.

The fixed-policy root is correct exactly when it lies outside the saturated predecessor set of $\mathcal{D}_H$.

Corollary 54.4 — Exact finite schedule classification.

For a finite optional-site family, if the regular disagreement construction is effective for every schedule, the full schedule family is exactly decidable by enumerating schedules and applying the saturation test.

This turns the Fatal Frontier theory into a complete finite classifier in the regular subclass.

54.3 Reachable-containment shortcut

Sometimes one only needs to prove correctness of one root, not construct the entire exact hull.

Corollary 54.5 — Reachable-containment root shortcut.

A root can be certified through the reachable semantic containment core without explicitly constructing the whole physical behavioral hull.

This can be substantially cheaper when semantic containment is small.

54.4 Finite semantic containment computation

If the semantic system M is finite and the initial semantically good set G_0 is known:

Theorem 54.6 — Finite containment-core computation.

The maximal forward-invariant semantic containment core is computed by iterative removal of states with outgoing edges leaving the current set.

Once G_0 is known, the edge-processing complexity is

$$O(|M||\Sigma|).$$

Corollary 54.7 — Finite containment-failure certificate.

If a semantic root is removed, there is a semantic witness path of length at most

$$\boxed{|M| - 1}$$

leading from the root to a containment failure.

Thus semantic containment failure has a short certificate.

54.5 Level III — DPDA-equivalence fallback

Regular immediate disagreement is powerful but not universal.

When no useful regular representation of $\mathcal{D}_H$ is available, compare continuation languages directly.

For physical configuration c , construct:

- the physical continuation language from c ;
- the semantic continuation language from $\Theta(c)$.

If both have effective deterministic pushdown presentations, classical deterministic pushdown equivalence can decide equality [61, 45].

Theorem 54.8 — DPDA-equivalence behavioral fallback.

Behavioral safety of a fixed physical root reduces exactly to equivalence between its physical and semantic deterministic pushdown continuation languages.

Corollary 54.9 — Exact fixed-policy root verification.

Whenever those two continuation languages are effectively represented as DPDAs, fixed-policy root correctness is decidable.

The imported equivalence algorithm is classical.

The contribution here is the reduction from normalization safety to that equivalence problem.

54.6 Verification hierarchy

The three levels are therefore:

structural safe-reset certification

$\Downarrow$

exact regular disagreement saturation



general DPDA-equivalence fallback.

They trade:

- structural specificity;
- computational convenience;
- generality.

54.7 General Theory synthesis

The final theoretical chain is now complete:

semantic fibers



bounded-local realizability



representation debt



behavioral hull / containment



safe reset / fatal causality



conflict complexes



section synthesis / forced debt

⇓

effective verification.

But the monograph should not end on an abstract verification theorem.
The final Part will return one last time to the two 2023 DPDA figures.
The point of the General Theory is not to replace the original machine.
It is to explain it.

PART XII

Closing the Loop

Chapter 55

The 2023 DPDA Revisited

The monograph began with two diagrams.

They represented the same $C/L/R$ recognition core but different empty-word wrappers:

$$L_{2,1} \setminus \{\varepsilon\}$$

and

$$L_{2,1}.$$

At the beginning, the figures were concrete machines to be understood.
We now return to them after developing:

- shortest-completion geometry;
- exact resource frontiers;
- representation-independent hierarchy methods;
- higher-dimensional completion geometry;
- cancellation/defect geometry;
- chronological arithmetic refinement;
- the general real-time $p:q$ machine;
- WSC Structure Theory;
- Uniqueness/Lift Theory;
- General Theory.

The purpose of this chapter is not to redraw the machine as an abstract object.

It is to show what every part of the original design means after the full theory has been developed.

55.1 The two origin figures remain distinct




Figure 55.1: The empty-input-excluded original DPDA revisited after the general theory. This is the same formalized machine as Figure 1.1, deliberately repeated so that the final interpretation can be read against the original object.

This version recognizes

$$L_{2,1} \setminus \{\varepsilon\}.$$




Figure 55.2: The empty-input-included original DPDA revisited. This is the same recognition/completion core as Figure 1.2; the wrapper accepts ε .

This version recognizes

$$L_{2,1}.$$

The common $C/L/R$ operational core is the same.

The difference is the acceptance/initialization wrapper.

This is the first lesson of the completed theory:

$$\boxed{\text{acceptance convention} \neq \text{semantic core.}}$$

55.2 What the stack was storing

The original stack looked like an operational memory structure.

Chapter 1 proved that its reachable form is highly restricted.

The state-potential invariant then gave

$$\boxed{\rho(w) = D(S) + \Phi(q).}$$

For 2:1,

$$D(S) = 2N_0(S) - N_1(S)$$

under the original machine's stack-weight convention, while the state potentials are

$$\Phi(C) = 0, \quad \Phi(L) = 3, \quad \Phi(R) = -3.$$

Thus the machine stores the arithmetic imbalance **distributed between control and stack**.

That observation became the prototype for every later semantic/physical decomposition.

55.3 What the completion rules were doing

The right-endmarker completion rules do not create new semantic information.

They materialize information already present in the final state/stack configuration.

After normalization, the physical word decodes a shortest completion demand.

The machine therefore contains an implicit map

$$\boxed{\begin{array}{l} \text{physical configuration} \rightarrow \text{residual} \rightarrow \text{minimum demand} \\ \rightarrow \text{canonical shortest continuation.} \end{array}}$$

This was the discovery that opened the entire completion-geometry branch.

55.4 Why depth two appeared

The original normalized stores lie in only a few structural families.

Therefore the qualitative question

Which demand symbols occur anywhere in the remaining stack?

is determined by the first two cells.

The general centered $p:q$ machine later recovered this two-cell qualitative visibility through a completely different canonical ordering.

Thus the original depth-two property was not merely a graphical accident.

It was an early manifestation of a broader principle:

unbounded quantitative demand can have bounded qualitative visibility.

55.5 The cancellation transition

The transition

$$(L, 0S) \xrightarrow{0} (L, S)$$

acquired three meanings.

Machine meaning

Delete one matching stack symbol.

Geodesic meaning

Consume one required zero along a shortest completion.

Combinatorial meaning

In the nonnegative 2-Dyck sector, mark a down-step landing at an exact odd-height family.

That third interpretation generated:

correction signatures $\rightarrow$ defect $\rightarrow$ Catalan boundary
 $\rightarrow$ binary skeleton $\rightarrow$ enumeration.

Thus one local transition produced an entire independent research branch.

55.6 Why the improved raw DPDA looked so different

The three-symbol raw residual-conjugate DPDA of Chapters 7–8 is not a cleaned-up drawing of the original machine.

It solves a different physical representation problem.

It asks:

If the residual itself is the semantic state, what is the most efficient stable raw encoding within the stated model?

That representation moves finite phase information into control and leaves unary magnitude on the stack.

It achieves the exact raw residual-faithful frontier.

The original machine instead stores a richer literal completion-demand process with normalization chronology.

Therefore the two machines are not rivals for one undefined notion of optimality.

They are optimal or explanatory in different representation categories.

55.7 Why the general $p:q$ machine returned to the original architecture

The centered real-time $p:q$ construction deliberately preserved the operational idea of the original machine:

- literal demand on the stack;
- one-cell LIFO blocker displacement;
- signed debt in finite control;
- later WRITE materialization.

When

$$(p, q) = (2, 1),$$

the general debt rail becomes

$$D_{-1}, D_0, D_1$$

with

$D_{-1} = R,$	$D_0 = C,$	$D_1 = L.$
---------------	------------	------------

The general potentials become

$$-3, 0, +3.$$

The historical WRITE arithmetic is recovered from the kernel-carry compiler.

Thus the original architecture was not merely generalized by analogy. It was reconstructed from the general arithmetic.

55.8 Why the historical table was not literally the centered specialization

Two reachable right-state matching families remained different.

The centered machine keeps debt lazily in R .

The historical machine immediately rematerializes the same semantic demand in C .

Therefore:

$$\boxed{\text{historical} = \text{eager normalization},}$$

$$\boxed{\text{centered specialization} = \text{lazy normalization}.}$$

This small transition-table discrepancy turned out to be one of the most important facts in the research program.

It proved concretely that:

$$\boxed{\text{same semantics} \not\equiv \text{unique physical chronology}.}$$

That observation led directly to Structure Theory and Uniqueness Theory.

55.9 C/L/R after Structure Theory

For 2:1,

$$r = p + q = 3.$$

The WSC exchange-sheet coordinate gives

$$\boxed{s(C) = 0, \quad s(L) = 1, \quad s(R) = -1.}$$

Thus the original states become SEARCH sheets.

The characteristic stack-length changes

$$-1, +2, +5$$

become materialization levels

$$M = 0, 1, 2.$$

And the historical identity

$$b + 2c = n$$

becomes a specialization of the global materialization budget.

The original table has therefore been reinterpreted from local syntax into semantic/physical accounting.

55.10 C/L/R after Uniqueness Theory

The original eager machine is one reduced lift.

The centered lazy machine is another.

They are:

- semantically equivalent;
- section-preserving incomparable;
- coordinate-reduced;
- non-isomorphic.

Within the selected three-site family, there are six admissible reduced normalization schedules.

The historical machine is therefore not “the” unique physical realization of the semantics.

It is a distinguished point in a nontrivial lift space.

55.11 The original machine after General Theory

The General Theory gives a final vocabulary.

The original configuration space is a physical lift.

The residual is semantic state.

The canonical completion word is a section choice.

The side-state configurations occupy noncanonical semantic fibers.

The one-top LIFO rule creates locality constraints.


Leaving the canonical section creates representation debt.

Some schedule choices preserve the behavioral hull.

Some expose fatal frontiers.

Alternative coordinate systems can remove obstructions that are essential in the literal policy.

Thus the full research program can be read backward into the original figure.



Semantic invariant: $\rho = D(S) + \Phi(q)$ • physical lift $\Theta : C_{\text{phys}} \rightarrow M$ • locality $\rightarrow$ debt $\rightarrow$ safe/fatal scheduling

Figure 55.3: The original $C/L/R$ architecture annotated by the complete theory. The figure is interpretive rather than a replacement transition table: it overlays the residual/state-potential invariant, WSC SEARCH-sheet coordinates, literal-demand magnitude, kernel materialization, end-of-input completion, schedule freedom, and the semantic-lift viewpoint on the original machine architecture.

Suggested annotations:

- C : canonical sheet / zero SEARCH charge;
- L : positive SEARCH sheet;
- R : negative SEARCH sheet;
- unary stack tails: physical demand magnitude;
- $0, 0/\varepsilon$: CANCEL / geodesic descent / correction event;
- long replacement words: kernel-materialization bursts;
- right-endmarker completion: final canonical materialization;
- eager right-state families: schedule choice;
- wrapper: acceptance convention, not semantic core.

The final visual object of the monograph should be this reinterpreted original machine.

Chapter 56

What Was Forced, What Was Chosen, What Became General

The entire monograph can now be organized by one question:

Which features were mathematically forced, and at what level?

The answer requires several categories.

56.1 Language-semantic necessities

These are determined by $L_{p,q}$ itself.

Residual quotient

$$u^{-1}L_{p,q} = v^{-1}L_{p,q} \iff \rho(u) = \rho(v).$$

The right-residual system is canonical.

Minimum completion demand

Each residual determines one unique minimum nonnegative completion Parikh vector.

Primitive kernel

$$K = (q, p).$$

Every equivalent completion vector differs by a kernel packet.

Semantic carry chronology

For a fixed input word,

$$(\rho_t), \quad (\kappa_t), \quad (\mu_t), \quad (\chi_t)$$

are unique.

These facts do not depend on a pushdown representation.

56.2 Necessities of a chosen locality category

These are not language-intrinsic, but become forced once the physical category is fixed.

Examples include:

One-top suffix obstruction

A hidden suffix cannot be arbitrarily rearranged by one top rewrite.

Literal-policy fiber branching

For nontrivial $p:q$, the fixed canonical literal section cannot remain transition-closed.

Deferred-SEARCH control width

$$2 \left\lfloor \frac{p+q-1}{2} \right\rfloor + 1$$

states are exact in the canonical deferred-SEARCH model.

Bounded-local collapse capacity

Too many semantic predecessors cannot be distinguished by too few local observations.

These are **representation-category necessities**.

They must always be stated with their category.

56.3 Coordinate artifacts

Some difficulties disappear when the physical coordinates change.

The clearest example is the four-symbol normalized lift.

Its

$$p + q - 1$$

raw boundary defects were exact in those coordinates.

But a boundary-compatible coordinate shift removed them and produced a three-symbol raw DPDA.

Therefore:

locality defect $\nRightarrow$ semantic necessity.

A defect may belong to the coordinate system.

56.4 Normalization-schedule choices

The semantic chronology can be unique while physical materialization chronology varies.

Examples:

- historical eager 2:1;
- centered lazy 2:1;
- the six admissible reduced 2:1 schedules.

The total kernel budget remains rigid even when its timing changes.

Thus:

semantic WRITE arrivals are forced;

physical WRITE timing may be chosen.

56.5 Special 2:1 features

Several concrete facts are specializations rather than universal principles.

For

$$p = 2, \quad q = 1 :$$

$$K = (1, 2),$$

$$m = 3,$$

$$h = 1,$$

so the centered debt rail has exactly three states.

These become:

$$R, C, L.$$

The state potentials are

$$-3, 0, +3.$$

The historical materialization changes are

$$-1, +2, +5.$$

The correction transition in the nonnegative path sector lands at

$$3, 5, 7, \dots$$

These numerical values are 2:1-specific.

The principles they instantiate are broader.

56.6 What generalized to all coprime $p:q$

The following extend uniformly:

- residual semantics;
- shortest completion demand;
- primitive kernel;
- geodesic rectangles;
- bounded canonical action visibility;
- raw residual-conjugate DPDA;
- exact raw residual-faithful frontier;
- centered real-time literal-demand DPDA;
- SEARCH debt;
- minimal kernel-carry WRITE;
- canonical-segment invariant;
- WSC semantic carry factorization;
- materialization/backlog accounting.

The original 2:1 machine was therefore a genuine first member of a general structural family.

56.7 What generalized beyond rational binary counting

Several branches moved beyond $L_{p,q}$.

Matrix completion geometry

The unique minimum pair became an optimal Parikh fiber.

The rectangle became a cubical quotient by the equal-length kernel.

Topological hierarchy

Kernel rank controlled local/global shortest-schedule topology.

General exact-lift theory


Semantic fibers, canonical sections, bounded-local obstruction, behavioral hulls, containment, causality, synthesis, and verification no longer depend specifically on weighted binary balance.

Thus the project eventually generalized not just the machine, but the distinction that the machine exposed:

what must be remembered semantically <i>versus</i> how a LIFO machine can physically remember it.
--

56.8 The research story in one chain

The complete causal story is:



This is the story the final publication version must preserve even if the page count is later reduced.

Chapter 57

Open Problems and the Continuing Research Program

The complete research version should end by separating proved results from genuine remaining questions.

The following list is intentionally conservative.

No open problem is declared solved merely because a related scoped theorem exists.

57.1 Unrestricted recognition complexity of $L_{p,q}$

The monograph proves exact frontiers for several representation categories, including raw residual conjugacy and deferred literal-demand SEARCH.

It does not determine the unrestricted minimum DPDA resources for recognizing every $L_{p,q}$.

Open directions include:

- unrestricted state complexity;
- unrestricted stack-alphabet complexity;
- complete state/stack Pareto frontiers;
- interactions with acceptance conventions;
- exact tradeoffs beyond residual-faithful categories.

57.2 Broader resource hierarchies

The representation-independent hierarchy theorem proves the fixed-state adjacent stack-alphabet hierarchy for the stated general DPDA model with $m \geq 2$.

Open extensions include:

- bounded-push variants;

- complete Pareto frontiers for the binary witnesses;
- related hierarchy questions under other initial-store conventions;
- transfer of completion-slope methods to natural language families rather than engineered witnesses.

57.3 General causal schedule theory for $p:q$

The 2:1 three-site family is classified exactly.

The 3:2 and odd- p :2 stress tests give larger exact families.

A full causal hypergraph classification for arbitrary coprime $p:q$ remains open.

Questions include:

- which sites are structurally canonical across all ratios?
- how do minimal fatal supports depend on p, q ?
- when are schedule families downward closed?
- when do robust fatal supports fully determine admissibility?

57.4 Fixed-policy status of $A \rightarrow B$

The specific 2:1 obstruction has persistence profile

$$(1, ?, 0).$$

Its essentiality in the intermediate fixed literal policy remains open.

This question mark must not be erased in later editing.

57.5 Unrestricted lift minimality

The Uniqueness Theory proves:

- reduced non-isomorphic lifts;
- category-relative minimality;
- forced branching inside fixed policies.

It does not give a universal minimal-lift theory for all deterministic pushdown representations.

A complete theory would need a carefully chosen morphism category that preserves enough physical structure to remain meaningful without being representation-arbitrary.

57.6 Confluence and schedule rewriting

The monograph proves:

- canonical absorption;
- normalizer-relative terminal schedule coherence;
- behavioral equivalence of several schedules.

It does not prove a general Church–Rosser theorem.

A classical confluence theory for normalization schedules remains open.

Such a theory would require:

- a precise rewrite relation on physical normalization choices;
- critical-pair analysis or an alternative invariant;
- separation of terminal semantic coherence from full rewrite confluence.

57.7 Higher-dimensional completion topology

The current theory gives:

- rank-zero boxes;
- rank-one global S^1 ;
- rank-two local holes;
- exact three-letter cubical formulas;
- eventual local-geometry periods.

Open questions include:

- higher homotopy types at kernel rank ≥ 2 ;
- classification of possible squarefree-divisor action complexes arising from minimum-completion fibers;
- exact global quotients beyond rank one;
- relations between affine-semigroup invariants and geodesic topology.

57.8 Higher-order correction geometry

The correction-signature theory is complete for its current defect functional.

Natural extensions include:

- other distinguished DPDA transitions;
- other rational ratios;
- multitype cancellation histories;
- higher-dimensional path analogues;
- refined statistics beyond correction landing height.

The binary cancellation skeleton suggests that local pushdown events may encode unexpectedly rigid combinatorial objects.

57.9 Fourth-order chronological arithmetic

CCR-3 is intentionally third-order.

At order four, a new connected chronological motif may appear.

The next problem is therefore to derive the genuine fourth-order interaction before writing any CCR-4 recurrence.

Other questions include:

- which local chronological deformations have Stirling shadows?
- what multiple-root lifting patterns can occur?
- when does singular extinction lead to higher ramification?
- which parts of the CCR-3 family-specific chain have exact analogues at higher order?

A full external novelty audit of those future families would be required independently.

57.10 General section synthesis

The rooted path-outcome criterion is exact.

Finite synthesis is effective.

One infinite integer-translation family admits finite symbolic synthesis.

But unrestricted synthesis over arbitrary infinite semantic systems is not solved.

Open directions include:

- decidable symbolic synthesis classes;
- minimal-section synthesis;
- synthesis under bounded stack alphabets or state budgets;
- quantitative optimization of forced-debt frontiers.

57.11 Verification complexity

The General Theory gives a hierarchy of reductions:

safe-reset $\rightarrow$ regular disagreement saturation $\rightarrow$ DPDA equivalence.

Further questions include:

- complexity of these reductions for specific normalization classes;
- efficient construction of disagreement automata;
- symbolic containment cores in infinite semantic systems;
- proof-certificate generation suitable for machine-checked verification.

57.12 Reproducibility tasks before publication

The complete research archive should contain every claimed computational artifact.

Current audit notes include:

- P9 centered-family audit script present and independently rerun successfully;
- P10 validation suite reconstructed from the authoritative manuscript/checkpoint audit specification and rerun successfully at the manuscript-reported total of 477,033 assertion-level checks;
- P11 validation suite reconstructed from the retained v0.1 core plus the later schedule/causality audit specification and rerun successfully at the manuscript-reported total of 458,776 assertion-level checks.

Before final publication:

1. preserve the reconstructed P10 and P11 scripts together with their provenance note;
2. preserve exact versioned outputs;
3. place computer-assisted certificates beside the theorems they support;
4. distinguish recovered historical artifacts from reconstructed regression artifacts;
5. distinguish symbolic proof from executable consistency evidence.

57.13 Final perspective

The research program began with a concrete deterministic pushdown automaton.

Its most productive feature was not simply that it recognized

$$L_{2,1}.$$

It distributed one arithmetic quantity across:

- finite control;
- a LIFO store;
- completion operations.

Trying to understand that distribution revealed:

semantic information $\neq$ physical representation.

Everything that followed can be traced to that separation.

The final lesson is therefore not one specific DPDA construction or one specific theorem.

It is a method:

extract the semantic object from the machine,

identify which physical constraints are truly forced,

change coordinates when the obstruction is representational,

retain debt when locality forces it,

and distinguish semantic correctness from normalization history.

The original 2023 DPDA remains the first concrete object in that story—and, after the full theory, the final object through which the story can be read.

Appendix A

Complete Transition Audit of the Original 2023 DPDA

A.1 Purpose and scope

This appendix closes the finite verification step that is deliberately compressed in Chapter 1.

The object audited here is the **original 2023 completion-aware 2:1 DPDA**, in the formalized presentation used throughout this monograph. The language convention is

$$L_{2,1} = \{ w \in \{0, 1\}^* : N_1(w) = 2N_0(w) \},$$

with signed prefix residual

$$\boxed{\rho(w) = N_1(w) - 2N_0(w).}$$

For a logical stack word $S \in \{0, 1\}^*$, define

$$\boxed{D(S) = 2N_0(S) - N_1(S)}$$

and assign the core-state potentials

$$\boxed{\Phi(C) = 0, \quad \Phi(L) = 3, \quad \Phi(R) = -3.}$$

The Chapter 1 invariant is

$$\boxed{\rho(w) = D(S) + \Phi(q)} \tag{A.1}$$

for every processed prefix w , where (q, S) is the corresponding core-projected configuration.

Since reading a binary symbol changes the residual by

$$\Delta\rho = \begin{cases} -2, & a = 0, \\ +1, & a = 1, \end{cases} \quad (\text{A.2})$$

it is enough to verify, for every reachable recognition transition,

$$\boxed{\Delta D + \Delta\Phi = \Delta\rho.} \quad (\text{A.3})$$

The audit also checks that every transition preserves the state-dependent reachable-store families

$$\boxed{\mathcal{S}_C = 0^* \cup 1^* \cup \{10\}, \quad \mathcal{S}_L = 0^*, \quad \mathcal{S}_R = 1^* \cup \{0\}.} \quad (\text{A.4})$$

The bottom marker Z_0 is omitted from logical-stack words below. The leftmost logical symbol is the top.

A.2 Wrapper projection

The two formal variants differ only in how initialization and empty-input acceptance are wrapped around the semantic core.

The projection used throughout the monograph is

$$\pi(C, SZ_0) = (C, S), \quad \pi(L, SZ_0) = (L, S), \quad \pi(R, SZ_0) = (R, S),$$

together with

$$\boxed{\pi(q_s, \varepsilon) = (C, \varepsilon), \quad \pi(q_f, Z_0) = (C, \varepsilon).} \quad (\text{A.5})$$

Thus the wrappers carry no additional residual semantics.

Empty-input-excluded wrapper

The first binary symbol is consumed directly at the source wrapper.

input	pro- jected source	pro- jected target	ΔD	$\Delta\Phi$	$\Delta(D + \Phi)$	$\Delta\rho$
1	(C, ε)	$(C, 10)$	+1	0	+1	+1
0	(C, ε)	$(C, 11)$	-2	0	-2	-2

Indeed,

$$D(10) = 2 - 1 = 1, \quad D(11) = -2.$$

Empty-input-included wrapper

The initialization move

$$\varepsilon, \varepsilon / Z_0$$

takes the physical source state to q_f . Under (A.5), both endpoints represent (C, ε) , so this move has semantic change 0.

When the first binary symbol is later consumed from q_f ,

$$1, Z_0 / 10Z_0 \quad \text{or} \quad 0, Z_0 / 11Z_0,$$

the same two projected checks above apply.

Finally, the core-to-final move

$$\varepsilon, Z_0 / Z_0$$

is available only at the balanced empty logical stack and projects

$$(C, \varepsilon) \longrightarrow (C, \varepsilon).$$

It therefore also preserves (A.1).

A.3 Center-state recognition transitions

A.3.1 Matching 1-cancellation

$$(C, 1T) \xrightarrow{1} (C, T).$$

One 1 is popped, hence

$$\Delta D = +1, \quad \Delta \Phi = 0.$$

Therefore

$$\Delta(D + \Phi) = +1 = \Delta \rho.$$

The target remains in $0^* \cup 1^* \cup \{10\}$: from a homogeneous 1-stack it remains homogeneous, while the exceptional source 10 becomes 0.

A.3.2 Matching 0-cancellation

$$(C, 0T) \xrightarrow{0} (C, T).$$

One 0 is popped, so

$$\Delta D = -2, \quad \Delta \Phi = 0,$$

and therefore

$$\Delta(D + \Phi) = -2 = \Delta\rho.$$

Reachability forces the source to lie in 0^+ , so the target lies in 0^* .

A.3.3 Center-to-left mismatch

$$(C, 0T) \xrightarrow{1} (L, T).$$

Popping the top 0 gives

$$\Delta D = -2.$$

The state potential changes by

$$\Delta \Phi = \Phi(L) - \Phi(C) = 3.$$

Thus

$$\boxed{\Delta(D + \Phi) = -2 + 3 = +1 = \Delta\rho.}$$

Because the source is a reachable center configuration with top 0, it belongs to 0^+ ; hence $T \in 0^*$, exactly the required left-state family.

A.3.4 Center-to-right mismatch

$$(C, 1T) \xrightarrow{0} (R, T).$$

Popping one 1 yields

$$\Delta D = +1,$$

while

$$\Delta \Phi = \Phi(R) - \Phi(C) = -3.$$

Therefore

$$\boxed{\Delta(D + \Phi) = 1 - 3 = -2 = \Delta\rho.}$$

There are two reachable source shapes:

$$1^{m+1} \mapsto 1^m,$$

which enters the 1^* part of $\mathcal{S}_R$, and

$$10 \mapsto 0,$$

which produces the exceptional right-state store 0.

These are exactly the two right-state families in (A.4).

A.4 Left-state recognition transitions

Every reachable left-state logical store lies in 0^* .

A.4.1 Left-state 0-cancellation

$$(L, 0T) \xrightarrow{0} (L, T).$$

The stack change is one popped 0:

$$\Delta D = -2, \quad \Delta \Phi = 0.$$

Hence

$$\Delta(D + \Phi) = -2 = \Delta\rho.$$

The target remains in 0^* .

A.4.2 Empty-left state on input 0

The bottom-marker rule is

$$0, Z_0/10Z_0,$$

so logically

$$(L, \varepsilon) \xrightarrow{0} (C, 10).$$

The source value is

$$D(\varepsilon) + \Phi(L) = 3.$$

The target value is

$$D(10) + \Phi(C) = 1.$$

Thus

$$\Delta(D + \Phi) = 1 - 3 = -2 = \Delta\rho.$$

The target 10 is the exceptional center word permitted by $\mathcal{S}_C$.

A.4.3 Nonempty-left state on input 1

For $T \in 0^*$,

$$(L, 0T) \xrightarrow{1} (C, 000T).$$

Replacing one 0 by three 0's adds two zeros, so

$$\Delta D = +4.$$

The state potential drops by 3:

$$\Delta\Phi = 0 - 3 = -3.$$

Hence

$$\Delta(D + \Phi) = 4 - 3 = +1 = \Delta\rho.$$

The target is again homogeneous 0^* .

A.4.4 Empty-left state on input 1

The bottom-marker rule

$$1, Z_0/00Z_0$$

gives

$$(L, \varepsilon) \xrightarrow{1} (C, 00).$$

Now

$$\Delta D = +4, \quad \Delta\Phi = -3,$$

so again

$$\Delta(D + \Phi) = +1 = \Delta\rho.$$

The target lies in 0^* .

A.5 Right-state recognition transitions

Every reachable right-state logical store lies in

$$1^* \cup \{0\}.$$

A.5.1 Nonempty homogeneous-right state on input 0

For $T \in 1^*$,

$$(R, 1T) \xrightarrow{0} (C, 111111T).$$

Replacing one 1 by six 1's adds five ones, therefore

$$\Delta D = -5.$$

The state potential rises from -3 to 0 :

$$\Delta \Phi = +3.$$

Thus

$$\boxed{\Delta(D + \Phi) = -5 + 3 = -2 = \Delta \rho.}$$

The target is in 1^* .

A.5.2 Empty-right state on input 0

The bottom-marker rule

$$0, Z_0/11111Z_0$$

gives

$$(R, \varepsilon) \xrightarrow{0} (C, 11111).$$

The stack contribution changes by

$$\Delta D = -5,$$

and the potential changes by $+3$. Hence

$$\Delta(D + \Phi) = -2 = \Delta \rho.$$

A.5.3 Exceptional right store 0 on input 0

The rule

$$0, 0/111$$

gives

$$(R, 0) \xrightarrow{0} (C, 111).$$

The stack weight changes from

$$D(0) = 2$$

to

$$D(111) = -3,$$

so

$$\Delta D = -5.$$

Again $\Delta \Phi = +3$, and

$$\Delta(D + \Phi) = -2 = \Delta \rho.$$

A.5.4 Exceptional right store 0 on input 1

The rule

$$1, 0/\varepsilon$$

gives

$$(R, 0) \xrightarrow{1} (C, \varepsilon).$$

The stack contribution changes by

$$\Delta D = -2,$$

while the state-potential change is +3. Therefore

$$\boxed{\Delta(D + \Phi) = -2 + 3 = +1 = \Delta \rho.}$$

A.5.5 Empty-right state on input 1

The bottom-marker rule

$$1, Z_0/11Z_0$$

gives

$$(R, \varepsilon) \xrightarrow{1} (C, 11).$$

Here

$$\Delta D = -2, \quad \Delta \Phi = +3,$$

and therefore

$$\Delta(D + \Phi) = +1 = \Delta\rho.$$

A.5.6 Nonempty homogeneous-right state on input 1

For $T \in 1^*$,

$$(R, 1T) \xrightarrow{1} (C, 111T).$$

Replacing one 1 by three 1's adds two ones:

$$\Delta D = -2.$$

Together with $\Delta \Phi = +3$,

$$\Delta(D + \Phi) = +1 = \Delta\rho.$$

The target remains homogeneous 1^* .

A.6 Complete recognition-phase audit table

The preceding cases may be compressed into one table.

source	input / top	logical action	target	ΔD	$\Delta \Phi$	total	required
C	1/1	pop 1	C	+1	0	+1	+1
C	0/0	pop 0	C	-2	0	-2	-2
C	1/0	pop 0	L	-2	+3	+1	+1
C	0/1	pop 1	R	+1	-3	-2	-2
L	0/0	pop 0	L	-2	0	-2	-2

source	input / top	logical action	target	ΔD	$\Delta \Phi$	total	required
L	$0/Z_0$	write 10	C	+1	-3	-2	-2
L	$1/0$	$0 \mapsto$ 000	C	+4	-3	+1	+1
L	$1/Z_0$	write 00	C	+4	-3	+1	+1
R	$0/1$	$1 \mapsto$ 111111	C	-5	+3	-2	-2
R	$0/Z_0$	write 11111	C	-5	+3	-2	-2
R	$0/0$	$0 \mapsto$ 111	C	-5	+3	-2	-2
R	$1/0$	pop 0	C	-2	+3	+1	+1
R	$1/Z_0$	write 11	C	-2	+3	+1	+1
R	$1/1$	$1 \mapsto$ 111	C	-2	+3	+1	+1

Thus every reachable recognition transition satisfies (A.3).

This proves the finite transition step omitted from the short proof in Chapter 1.

A.7 Reachable-store closure

The same table proves the state-dependent store normal form.

Center

The center can contain only

$$0^*, \quad 1^*, \quad 10.$$

Its two matching loops shorten an existing block. A mismatch can leave the center only by popping the visible blocking symbol. Every side-to-center transition returns to a homogeneous block except the left-empty/input-0 case, which creates exactly 10.

Left

The only transition entering L is

$$(C, 0T) \xrightarrow{1} (L, T).$$

Reachability forces the source center stack to belong to 0^+ , so

$$T \in 0^*.$$

Therefore no reachable left configuration has top 1.

This proves, in particular, that the listed completion rule

$$\neg, 1/0$$

has an unreachable source condition.

Right

The only transition entering R is

$$(C, 1T) \xrightarrow{0} (R, T).$$

If the source is 1^{m+1} , then $T = 1^m$.

If the source is the exceptional center word 10 , then $T = 0$.

Hence

$$\mathcal{S}_R = 1^* \cup \{0\}.$$

No other right-state store can be produced.

A.8 Audit of the end-of-input completion rules

The endmarker $\neg$ is not a binary input symbol, so completion normalization must preserve the residual already accumulated by the prefix.

The side-state potential is transferred into literal stack content.

A.8.1 Left-state completion

Every reachable left store is 0^m .

Empty left store.

$$(L, \varepsilon) \xrightarrow{\neg} 100.$$

Before completion,

$$D(\varepsilon) + \Phi(L) = 0 + 3 = 3.$$

After completion,

$$D(100) = 4 - 1 = 3.$$

Nonempty left store. For $m \geq 1$,

$$(L, 0^m) \xrightarrow{\perp} 10^{m+2}.$$

Indeed the physical rule replaces the top 0 by 1000, leaving 0^{m-1} untouched.
Before completion,

$$D(0^m) + \Phi(L) = 2m + 3.$$

After completion,

$$D(10^{m+2}) = 2(m+2) - 1 = 2m + 3.$$

Thus left completion preserves the represented residual exactly.
The listed left rule with top 1 is unreachable and plays no role in any run.

A.8.2 Right-state completion

There are three reachable right forms.

Empty right store.

$$(R, \varepsilon) \xrightarrow{\perp} 111.$$

Before completion,

$$D(\varepsilon) + \Phi(R) = -3,$$

and after completion,

$$D(111) = -3.$$

Nonempty homogeneous right store. For $m \geq 1$,

$$(R, 1^m) \xrightarrow{\perp} 1^{m+3}.$$

Before completion,

$$D(1^m) + \Phi(R) = -m - 3.$$

After completion,

$$D(1^{m+3}) = -m - 3.$$

Exceptional right store.

$$(R, 0) \xrightarrow{-1} 1.$$

Before completion,

$$D(0) + \Phi(R) = 2 - 3 = -1.$$

After completion,

$$D(1) = -1.$$

Therefore every reachable endmarker normalization rule preserves the prefix residual.

A.9 From residual preservation to the canonical shortest completion

After the side-state correction, every normalized logical stack belongs to

$$0^* \cup 1^* \cup 10^+.$$

The three residual cases are therefore:

$$\widehat{S}(w) = \begin{cases} 0^m, & \rho(w) = 2m, \quad m \geq 0, \\ 10^{m+1}, & \rho(w) = 2m + 1, \quad m \geq 0, \\ 1^m, & \rho(w) = -m, \quad m \geq 1. \end{cases} \quad (\text{A.6})$$

For any continuation containing x zeros and y ones,

$$wz \in L_{2,1} \iff 2x - y = \rho(w). \quad (\text{A.7})$$

The words in (A.6) realize the unique minimum nonnegative Parikh solution of (A.7):

- $2m$: $(x, y) = (m, 0)$;
- $2m + 1$: $(x, y) = (m + 1, 1)$;
- $-m$: $(x, y) = (0, m)$.

Thus the transition audit not only proves recognition-phase arithmetic consistency; together with the endmarker audit it recovers the shortest-completion semantics of the original design.

A.10 Two representative runs

A.10.1 Accepted word 110

Using the core projection and omitting Z_0 ,

$$(C, \varepsilon) \xrightarrow{1} (C, 10) \xrightarrow{1} (C, 0) \xrightarrow{0} (C, \varepsilon).$$

The residual sequence is

$$0, 1, 2, 0,$$

while $D(S) + \Phi(q)$ gives the same sequence:

$$0, 1, 2, 0.$$

The final empty center stack therefore represents zero residual and the word is accepted.

A.10.2 Prefix 111

The run is

$$(C, \varepsilon) \xrightarrow{1} (C, 10) \xrightarrow{1} (C, 0) \xrightarrow{1} (L, \varepsilon).$$

The residual is

$$\rho(111) = 3.$$

The final logical stack is empty, but

$$D(\varepsilon) + \Phi(L) = 3.$$

The endmarker rule transfers this finite-control contribution into the literal completion word:

$$(L, \varepsilon) \xrightarrow{1} 100.$$

Since

$$2N_0(100) - N_1(100) = 4 - 1 = 3,$$

the normalized stack represents exactly the outstanding residual. Appending 100 gives

$$111100 \in L_{2,1}.$$

This is the canonical shortest completion selected by the machine.

A.11 Determinism of the completion extension

The completion callouts should not be interpreted as ordinary ε -moves.

For a fixed state/top pair:

- binary recognition transitions are selected by next input 0 or 1;
- completion is selected only by the distinct right endmarker $\neg$.

Direct all completion rules to a dedicated completion sink with no outgoing transitions.

Then no configuration has two enabled successors.

Hence the completion-extended formalization is deterministic in the standard DPDA sense.

A.12 Exhaustive computational cross-check

The symbolic audit above is complete: it covers every reachable transition family.

The accompanying P1 validation suite provides an independent bounded regression layer. In the final research corpus it checks, among other items,

- 524,287 original-machine binary inputs;
- 8,912,898 prefix transition checks;
- the residual invariant;
- the reachable-store normal form;
- completion normalization;
- independent shortest-completion length;
- depth-two support visibility.

All checks pass.

The executable audit is supporting evidence only; the proof of (A.1) is the finite symbolic transition audit in Sections A.2–A.8.

A.13 Conclusion

The original 2023 design is not merely a correct recognizer.

Its state and stack form an exact split representation of the prefix residual:

$$\boxed{\rho = D + \Phi.}$$

The center state stores the residual entirely in the stack.

The side states store part of the same residual in finite control:

$$\boxed{\Phi(L) = +3, \quad \Phi(R) = -3.}$$

Every binary transition preserves that semantic quantity with the correct residual increment, and every reachable endmarker rule transfers the finite-control offset back into literal stack form.

Thus the three operational ideas visible in the original machine—

$$\boxed{\text{cancellation,} \quad \text{side-state mismatch,} \quad \text{completion,}}$$

—are already an exact semantic normalization mechanism.

This is the local machine-level fact from which the later completion geometry, residual quotients, WRITE-SEARCH-CANCEL theory, lift theory, and General Theory are developed.

Appendix B

General Construction Transition Schemas

B.1 Purpose and model separation

The main text deliberately develops the general constructions conceptually before presenting every transition rule. This appendix collects the exact machine syntax in one place so that the book remains reproducible without forcing long transition tables into the narrative chapters.

Fix coprime integers $p, q \geq 1$ and write

$$\rho(w) = qN_1(w) - pN_0(w), \quad T_0(r) = r - p, \quad T_1(r) = r + q.$$

The primitive zero-weight kernel is

$$K = (q, p).$$

Three constructions must be kept separate.

1. The *paired-sign four-symbol lift* realizes the residual graph exactly after a deterministic normalization macro-step. It is a normalized residual-configuration conjugacy, not a raw one-step realization.
2. The *boundary-compatible three-symbol lift* changes coordinates so that every residual edge is one ordinary binary-input transition. It is ε -free and gives raw residual-configuration conjugacy with $\max(p, q)$ states.
3. The *centered literal-demand construction* stores the unique shortest completion demand itself, uses a finite-control debt rail to repair the LIFO ordering obstruction, consumes exactly one binary symbol per transition, and is real-time during binary input.

The exact resource-frontier theorems later compare these constructions only inside their stated representation models. In particular, no table below should be read as an unrestricted minimum-state theorem for every DPDA recognizing $L_{p,q}$.

B.2 Four-symbol normalized paired-sign lift

This is the first compact general physical lift used for the residual-geometry and configuration-rectangle results.

Put

$$m_{\text{st}} = \max(p, q), \quad Q_4 = \{R_0, \dots, R_{m_{\text{st}}-1}\}, \quad \Gamma_4 = \{0, 1, X, Z_0\}.$$

For $0 \leq s < p$ and $0 \leq t < q$, define the finite residual bases

$$a_s = \begin{cases} 0, & s = 0, \\ s - p, & 1 \leq s < p, \end{cases} \quad b_t = \begin{cases} 0, & t = 0, \\ q - t, & 1 \leq t < q. \end{cases}$$

The stable encoding $\eta_4 : \mathbb{Z} \rightarrow Q_4 \times \Gamma_4^*$ is

$$\eta_4(0) = (R_0, Z_0).$$

If $r > 0$, let $s = r \bmod p$ and write uniquely

$$r = pk + a_s, \quad k \geq 1;$$

then

$$\eta_4(r) = (R_s, 0^k Z_0).$$

If $r < 0$, let $t = (-r) \bmod q$ and write uniquely

$$r = -qk + b_t, \quad k \geq 1;$$

then

$$\eta_4(r) = (R_t, 1^k X).$$

The two sign roles may share a control state; their work symbols and sentinels keep the stable encodings disjoint.

Positive stable rules

For $0 \leq s < p$, put

$$s' = (s + q) \bmod p, \quad c_s = \frac{a_s + q - a_{s'}}{p}.$$

Then

$$\delta(R_s, 0, 0) = (R_s, \varepsilon)$$

and

$$\delta(R_s, 1, 0) = (R_{s'}, 0^{c_s+1}).$$

When the last positive work symbol is removed, a unique forced boundary normalization is used:

$$\delta(R_s, \varepsilon, Z_0) = \eta_4(a_s) \quad (1 \leq s < p).$$

For $s = 0$, popping the final positive 0 already exposes the zero encoding (R_0, Z_0) , so no boundary normalization is required.

Negative stable rules

For $0 \leq t < q$, put

$$t' = (t + p) \bmod q, \quad d_t = \frac{p - b_t + b_{t'}}{q}.$$

Then

$$\delta(R_t, 1, 1) = (R_t, \varepsilon)$$

and

$$\delta(R_t, 0, 1) = (R_{t'}, 1^{d_t+1}).$$

When the negative sentinel is exposed,

$$\delta(R_t, \varepsilon, X) = \eta_4(b_t).$$

Zero rules and acceptance

At residual zero,

$$\delta(R_0, 1, Z_0) = \eta_4(q), \quad \delta(R_0, 0, Z_0) = \eta_4(-p).$$

A mandatory right endmarker is accepted only at the stable zero configuration $\eta_4(0)$.

An expression of the form $\delta(\cdot) = \eta_4(r')$ is semantic shorthand: the exposed boundary symbol is replaced by the stack word of the stable encoding while control enters the corresponding state. Long fixed writes may be compiled into deterministic forced ε -chains if a bounded-write PDA syntax is required. Such a compilation preserves the stable target but is no longer a raw one-binary-transition realization.

Define a normalized input macro-step $c \xRightarrow{a} c'$ to mean one consumed binary symbol followed by the unique forced normalization chain. The transition schema satisfies

$$\boxed{\eta_4(r) \xRightarrow{0} \eta_4(r-p), \quad \eta_4(r) \xRightarrow{1} \eta_4(r+q).}$$

This is the exact normalized residual-configuration conjugacy used in the configuration-rectangle chapters.

B.3 One-work-symbol raw conjugacy

The one-work-symbol construction is included because it is the matching upper bound for the exact $p+q$ -state result in the permanent-bottom raw-conjugacy model.

Let

$$\Gamma_2 = \{A, Z_0\},$$

with disjoint state families

$$P_0, \dots, P_{p-1}, \quad N_0, \dots, N_{q-1}.$$

For $r \geq 0$, write

$$r = pk + s, \quad k \geq 0, \quad 0 \leq s < p,$$

and define

$$\xi(r) = (P_s, A^k Z_0).$$

For $r < 0$, write uniquely

$$r = u - qh, \quad 0 \leq u < q, \quad h \geq 1,$$

and define

$$\xi(r) = (N_u, A^{h-1} Z_0).$$

For $0 \leq s < p$, set

$$s^+ = (s + q) \bmod p, \quad c_s = \frac{s + q - s^+}{p} \geq 0,$$

and write

$$s - p = u_s - qh_s, \quad 0 \leq u_s < q, \quad h_s \geq 1.$$

The positive and boundary rules are

$$\boxed{\delta(P_s, 0, A) = (P_s, \varepsilon)}, \quad \boxed{\delta(P_s, 1, A) = (P_{s^+}, A^{c_s+1})},$$

$$\boxed{\delta(P_s, 1, Z_0) = (P_{s^+}, A^{c_s} Z_0)}, \quad \boxed{\delta(P_s, 0, Z_0) = (N_{u_s}, A^{h_s-1} Z_0)}.$$

For $0 \leq u < q$, put

$$u^- = (u - p) \bmod q, \quad d_u = \frac{u^- - u + p}{q} \geq 0.$$

The negative rules are

$$\begin{aligned} \boxed{\delta(N_u, 1, A) = (N_u, \varepsilon)}, \quad & \boxed{\delta(N_u, 0, A) = (N_{u^-}, A^{d_u+1})}, \\ \boxed{\delta(N_u, 1, Z_0) = \xi(u)}, \quad & \boxed{\delta(N_u, 0, Z_0) = (N_{u^-}, A^{d_u} Z_0)}. \end{aligned}$$

The mandatory right endmarker is accepted only at

$$\xi(0) = (P_0, Z_0).$$

There are no ε -transitions. Direct substitution gives

$$\boxed{\delta(\xi(r), 0) = \xi(r - p), \quad \delta(\xi(r), 1) = \xi(r + q)}.$$

The construction uses exactly $p + q$ states and two total stack symbols, one of which is the permanent bottom marker.

B.4 Boundary-compatible three-symbol raw DPDA

The four-symbol normalized lift shows that the residual graph has a compact physical realization, but it has a finite set of boundary edges that require normalization. The following coordinate change removes those defects completely.

Assume first that $p \geq q$. Let

$$Q_3 = \{R_0, \dots, R_{p-1}\}, \quad \Gamma_3 = \{0, 1, Z_0\}.$$

For $r \geq 0$, write uniquely

$$r = pk + s, \quad k \geq 0, \quad 0 \leq s < p,$$

and set

$$\boxed{\zeta(r) = (R_s, 0^k Z_0)}.$$

For $r < 0$, write uniquely

$$r = u - qh, \quad h \geq 1, \quad 0 \leq u < q,$$

and set

$$\boxed{\zeta(r) = (R_u, 1^h Z_0)}.$$

The state sets overlap deliberately. The sign is visible from the top of stack: negative encodings have top 1, whereas nonnegative encodings have top 0 or Z_0 .

For $0 \leq s < p$, define

$$s^+ = (s + q) \bmod p, \quad c_s = \frac{s + q - s^+}{p} \in \{0, 1\},$$

$$u_s = (s - p) \bmod q, \quad h_s = \frac{u_s + p - s}{q} \geq 1.$$

For $0 \leq u < q$, define

$$u^- = (u - p) \bmod q, \quad d_u = \frac{u^- - u + p}{q} \geq 1.$$

Complete transition table for $p \geq q$

region	input/top	target	replacement
nonnegative interior	0/0	R_s	ε
nonnegative interior	1/0	R_{s+}	0^{c_s+1}
nonnegative boundary	1/ Z_0	R_{s+}	$0^{c_s} Z_0$
nonnegative boundary	0/ Z_0	R_{u_s}	$1^{h_s} Z_0$
negative	1/1	R_u	ε
negative	0/1	R_{u-}	1^{d_u+1}

Equivalently,

$$\begin{aligned} \delta(R_s, 0, 0) &= (R_s, \varepsilon), \\ \delta(R_s, 1, 0) &= (R_{s+}, 0^{c_s+1}), \\ \delta(R_s, 1, Z_0) &= (R_{s+}, 0^{c_s} Z_0), \\ \delta(R_s, 0, Z_0) &= (R_{u_s}, 1^{h_s} Z_0), \\ \delta(R_u, 1, 1) &= (R_u, \varepsilon), \\ \delta(R_u, 0, 1) &= (R_{u-}, 1^{d_u+1}). \end{aligned}$$

The right endmarker is accepted only at

$$\zeta(0) = (R_0, Z_0).$$

No ε -transition is used. Every binary transition consumes exactly one input symbol, and

$$\delta(\zeta(r), 0) = \zeta(r - p), \quad \delta(\zeta(r), 1) = \zeta(r + q).$$

Hence the raw binary-input configuration graph is label-isomorphic to the residual graph.

If $q > p$, use the symmetric construction obtained by simultaneously exchanging

$$0 \leftrightarrow 1, \quad p \leftrightarrow q, \quad r \leftrightarrow -r.$$

The resulting state count is always

$$\boxed{|Q_3| = \max(p, q)},$$

with three total stack symbols including the permanent bottom marker.

B.5 Centered real-time literal-demand construction

The raw residual encodings above store a signed residual coordinate. The centered construction solves a different problem: it keeps a physical representation whose semantic projection is the unique shortest completion demand and repairs one-cell LIFO ordering conflicts by moving a blocker into finite control.

Let

$$m = p + q, \quad h_{\text{ctr}} = \left\lfloor \frac{m-1}{2} \right\rfloor.$$

For a residual r , let

$$\mu(r) = (x_r, y_r)$$

be the unique minimum nonnegative solution of

$$px - qy = r.$$

For a demand vector (x, y) use the canonical two-visible ordering

$$\eta(x, y) = \begin{cases} (10)^y 0^{x-y}, & x \geq y, \\ (10)^x 1^{y-x}, & y > x. \end{cases}$$

For debt $d \in \mathbb{Z}$, define the virtual demand

$$V_d = \mu(dm), \quad \Phi(D_d) = dm.$$

The centered state set is

$$\boxed{Q_{\text{ctr}} = \{D_{-h_{\text{ctr}}}, \dots, D_{-1}, D_0, D_1, \dots, D_{h_{\text{ctr}}}\}},$$

with $D_0 = C$, and

$$\Gamma_{\text{ctr}} = \{0, 1, Z_0\}.$$

For a logical stack S , the semantic projection is

$$\Theta(D_d, S) = \text{red}_{p,q}(P(S) + V_d).$$

The residual invariant is

$$\boxed{ho = W(P(S)) + dm.}$$

MATCH

For $z \in \{0, 1\}$ and every $|d| \leq h_{\text{ctr}}$,

$$(D_d, zT) \xrightarrow{z} (D_d, T).$$

Interior SEARCH

For $d < h_{\text{ctr}}$,

$$(D_d, 0T) \xrightarrow{1} (D_{d+1}, T),$$

and for $d > -h_{\text{ctr}}$,

$$(D_d, 1T) \xrightarrow{0} (D_{d-1}, T).$$

The displaced blocker is not restored immediately. Its signed effect is represented by the debt coordinate d .

Outer-boundary WRITE

Put

$$U_+ = \eta(\mu((h_{\text{ctr}} + 1)m)), \quad U_- = \eta(\mu(-(h_{\text{ctr}} + 1)m)).$$

Then

$$(D_{h_{\text{ctr}}}, 0T) \xrightarrow{1} (C, U_+T),$$

and

$$(D_{-h_{\text{ctr}}}, 1T) \xrightarrow{0} (C, U_-T).$$

Bottom WRITE

For every $|d| \leq h_{\text{ctr}}$,

$$(D_d, Z_0) \xrightarrow{1} (C, \eta(\mu(dm + q))Z_0),$$

$$(D_d, Z_0) \xrightarrow{0} (C, \eta(\mu(dm - p))Z_0).$$

Acceptance

A mandatory right endmarker is permitted only at the balanced empty configuration:

$$(C, Z_0) \xrightarrow{1} (C, \varepsilon).$$

There is no right-endmarker rule from any other reachable configuration. Every binary transition consumes exactly one input symbol; the centered construction has no ε -move during binary processing.

B.6 Minimal kernel-carry compiler identity

The WRITE strings in the centered construction are not ad hoc. Suppose a transition consumes input symbol a , removes visible top symbol z , changes state from D_s to D_t , and writes a replacement word u . Let

$$e_0 = (1, 0), \quad e_1 = (0, 1).$$

Ignoring the untouched suffix, invariant preservation forces

$$P(u) = e_z - e_a + V_s - V_t + \kappa K.$$

At the bottom marker, omit the term e_z . If

$$B = e_z - e_a + V_s - V_t,$$

then the smallest kernel carry making the replacement nonnegative is

$$\kappa_* = \max\left(\left\lceil -\frac{B_0}{q} \right\rceil, \left\lceil -\frac{B_1}{p} \right\rceil\right).$$

The unique shortest compatible replacement vector is

$$R_* = B + \kappa_* K.$$

Every other nonnegative invariant-preserving replacement is $R_* + \ell K$ for some $\ell \geq 1$ and is longer by exactly $\ell(p + q)$ symbols.

This identity is the common arithmetic compiler behind the large literal WRITE words of the original 2:1 machine and the general centered family.

B.7 The 2:1 specialization and the eager–lazy distinction

For $(p, q) = (2, 1)$,

$$m = 3, \quad h_{\text{ctr}} = 1,$$

so

$$D_{-1} = R, \quad D_0 = C, \quad D_1 = L,$$

with potentials

$$-3, \quad 0, \quad +3.$$

The centered schema reproduces the historical center/left/right control skeleton, the center mismatch pops, the left cancellation family, and the arithmetic of the characteristic WRITE words.

Two reachable right-state matching families differ in schedule, not semantics. The centered machine keeps the resulting debt in finite control (lazy materialization), whereas the historical machine immediately materializes the same canonical demand on the stack (eager materialization). If $\mathcal{N}$ denotes canonical materialization, then on exactly those families

$$\delta_{\text{hist}}(c, a) = \mathcal{N}(\delta_{\text{cent}}(c, a)).$$

Accordingly, direct specialization is not claimed to be literal transition-table identity. The remaining difference is precisely the eager-lazy normalization-schedule choice.

B.8 Determinism and implementation conventions

The constructions in this appendix use three different transition disciplines and should not be conflated.

- In the four-symbol paired-sign lift, forced ε -normalization is enabled only on transient boundary pairs where binary input is disabled. The normalization macro-step is therefore deterministic.
- In the one-work-symbol and boundary-compatible three-symbol raw lifts, there are no ε -transitions at all; the right endmarker is the only nonbinary input symbol.
- In the centered literal-demand construction, every binary step is real-time and the right endmarker is enabled only at the balanced empty configuration.

All replacement words are fixed finite constants once (p, q) is fixed. If a separate PDA convention imposes a bound on replacement length, long fixed writes can be compiled using auxiliary deterministic microstates. Such compilation changes the physical transition granularity and must not be silently substituted into a theorem whose statement concerns raw one-input transitions or exact state counts in a fixed model.

B.9 Summary table

construction		states	stack alpha- bet	exact role
paired-sign normalized lift	normal-	$\max(p, q)$	$\{0, 1, X, Z_0\}$	exact residual conjugacy after a unique normalization macro- step
one-work-symbol raw lift		$p + q$	$\{A, Z_0\}$	exact raw residual conjugacy; matching upper bound for the one-work-symbol state theorem
boundary- compatible lift	raw	$\max(p, q)$	$\{0, 1, Z_0\}$	ε -free raw residual conjugacy by ordinary binary transitions
centered demand lift	literal-	$2\lfloor(p+q-1)/2\rfloor + 1$	$\{0, 1, Z_0\}$	real-time shortest-demand nor- malization with finite-control SEARCH debt; exact width only in the canonical deferred- SEARCH model

The purpose of the table is organizational. The constructions optimize different interfaces, and their state counts are not interchangeable unrestricted complexity claims.

Appendix C

Scoped Lower Bounds and Finite Certificates

C.1 Why this appendix is necessary

Several chapters use the word *optimal*, but they do so in different representation categories. The purpose of this appendix is to make the scope of every principal lower bound explicit and to place the one genuinely computer-assisted finite elimination in a reproducible form.

The central distinction is:

representation-relative lower bound $\neq$ unrestricted recognition lower bound.

A machine may recognize the same language while abandoning the residual coordinates, stack-tail form, literal-demand interface, or normalization schedule used by another construction. A lower bound obtained by counting coordinates in one architecture therefore cannot automatically be promoted to all DPDAs.

C.2 Lower-bound scope table

result	opponent class	exact bound	excluded overclaim
one-work-symbol raw residual conjugacy	permanent bottom, one work symbol, ε -free raw residual-configuration conjugacy	$ Q \geq p + q$	not an unrestricted recognition bound for $L_{p,q}$
homogeneous-tail raw frontier	permanent bottom, sign-separated homogeneous tails, at least two work symbols	$ Q \geq \max(p, q)$	not a theorem about arbitrary stack encodings

result	opponent class	exact bound	excluded overclaim
arbitrary three-symbol ratios 2:1, 3:1, 3:2	permanent bottom, three total stack symbols, arbitrary injective raw residual encoding	$ Q \geq \max(p, q)$ for these ratios	general coprime (p, q) case remains open
canonical deferred-SEARCH width	literal shortest-demand representation with deferred SEARCH interface	$2\lfloor(p+q-1)/2\rfloor + 1$ states	not an unrestricted state minimum for every DPDA recognizing $L_{p,q}$
phase-local capacity	specified local observation depth and remainder architecture	$ Q g^\lambda \geq BH(B)$	not a universal Pareto frontier
binary stack-alphabet hierarchy witness	arbitrary general DPDA in the stated single-initial-symbol model, deterministic ε -moves allowed	$ Q \Gamma \geq mk + 1$ for $B_{m,k}$	applies to the engineered witness language, not automatically to $L_{p,q}$

C.3 Raw residual-conjugacy frontier

For the permanent-bottom raw residual-conjugacy models developed in the resource chapters, the exact structured frontier is:

$$S_{\text{raw}}(g) = \begin{cases} \infty, & g = 0, \\ p + q, & g = 1, \\ \max(p, q), & g \geq 2, \end{cases}$$

where g denotes the number of work symbols above the permanent bottom marker in the homogeneous-tail model.

The upper bounds are realized by the constructions collected in Appendix B. The lower bounds rely on the raw-conjugacy interface: every residual must have one injective physical configuration, and each binary residual edge must be realized by one ordinary input transition.

The $g = 1$ lower bound is exact because a one-work-symbol store can encode unbounded magnitude only through height, while the p positive-end and q negative-end arithmetic phases require disjoint recurrent state resources. The resulting state requirement is

$$|Q| \geq p + q.$$

With two distinct work symbols, sign can be placed in the visible top symbol. The boundary-compatible construction then attains

$$|Q| = \max(p, q), \quad |\Gamma| = 3$$

including the permanent bottom marker.

The homogeneous-tail hypothesis is removed completely only for selected ratios in the present corpus. Exact arbitrary-encoding three-symbol state complexity is proved for

$$(2, 1), \quad (3, 1), \quad (3, 2),$$

and, by the input/sign symmetry, for

$$(1, 2), \quad (1, 3), \quad (2, 3).$$

The corresponding arbitrary-encoding theorem for every coprime pair remains open.

C.4 The 3:2 finite reduction

The first genuinely two-sided arbitrary-encoding case is $(p, q) = (3, 2)$. The theorem concerns the full permanent-bottom three-symbol model

$$\Gamma_3 = \{A, B, Z_0\}$$

with no ε -moves and an injective raw residual-configuration conjugacy. No homogeneous-tail or sign-separated assumption is imposed.

Assume, for contradiction, that two control states suffice. Choose an encoded configuration of minimum work-stack height and write the two parity chains relative to its residual r_* as

$$A_k = \eta(r_* + 2k), \quad B_k = \eta(r_* + 1 + 2k).$$

Let G denote the input-1 action and F the input-0 action. Since input 1 adds 2 to the residual and input 0 subtracts 3,

$$G(A_k) = A_{k+1}, \quad G(B_k) = B_{k+1},$$

$$F(A_{k+2}) = B_k, \quad F(B_{k+1}) = A_k.$$

Therefore, on the encoded residual graph,

$$\boxed{F^2 G^3 = G^3 F^2 = \text{id}.} \tag{C.1}$$

Equation (C.1) is the key finiteness mechanism. One consuming PDA transition deletes at most one work symbol. If a work-top rule writes m work symbols in place of one, its first step increases work height by $m - 1$. The remaining four moves of either five-step inverse identity can delete at most four symbols before the original configuration must be recovered. Hence

$$\boxed{m \leq 5}$$

for every work-top replacement encountered in the search. From the bare bottom, the first step creates rather than replaces its work word, so every bottom write has length at most four.

Thus an a priori infinite transition search reduces to a finite deterministic rule-synthesis problem.

C.5 Normalization of the minimum configuration

There are exactly two minimum-height cases.

Positive minimum height. The suffix strictly below the visible top of a minimum encoded configuration can never be exposed by another encoded configuration; it is inert and may be stripped. Rename states and work symbols so that

$$A_0 = (0, A\#),$$

where $\#$ denotes the fixed inert suffix.

Zero minimum height. Normalize instead to

$$A_0 = (0, \#),$$

where $\#$ is the bare bottom marker.

These two cases exhaust the possible minimum work heights.

C.6 Legal branching and pruning rules

The finite certificate assigns a transition rule only when a state/top pair is encountered for the first time. At that first occurrence it branches over every legal target state and every replacement word satisfying the proven length bounds. A repeated state/top pair must reuse the rule already fixed by determinism.

Branches are rejected only by consequences already forced by the raw conjugacy and the inverse identities:

1. a prescribed F -relation or G -relation has the wrong target;
2. suffix preservation makes the required replacement impossible;
3. two distinct residuals obtain the same encoded configuration, violating injectivity;
4. a configuration is too tall to return to a prescribed target in the remaining number of one-pop transitions.

Representative deletion-capacity bounds are

$$F^2(A_3) = A_0, \quad F(B_1) = A_0,$$

which imply

$$|A_3| \leq |A_0| + 2, \quad |B_1| \leq |A_0| + 1.$$

More generally,

$$F(B_j) = A_{j-1}, \quad F(A_{j+2}) = B_j,$$

and the corresponding two-deletion relations

$$F^2(B_j) = B_{j-3}, \quad F^2(A_{j+2}) = A_{j-1}$$

provide the remaining legal height-pruning bounds.

No heuristic pruning is permitted: every discarded branch violates one of these proved necessary conditions.

C.7 Exact survivor vectors

For positive minimum height, the exact numbers of surviving partial rule systems at successive synthesis stages are

$$\boxed{1, 123, 11643, 1691, 2120, 1591, 348, 43, 13, 0.} \quad (\text{C.2})$$

The search examines exactly

$$\boxed{1, 130, 430}$$

branch attempts in this case.

For minimum height zero, the survivor vector is

$$\boxed{1, 61, 7556, 1134, 2146, 2966, 6846, 7032, 2376, 650, 190, 26, 0.} \quad (\text{C.3})$$

This case examines exactly

$$\boxed{1, 696, 120}$$

branch attempts.

Both exhaustive searches terminate with zero survivors. Since the symbolic reduction proves that every possible two-state raw encoding must occur in one of the two finite searches, two states are impossible. The explicit three-state boundary-compatible construction supplies the matching upper bound. Hence

$$\boxed{|Q|_{\min} = 3}$$

for the arbitrary permanent-bottom three-symbol raw-conjugacy model at $(3, 2)$, and symmetrically at $(2, 3)$.

C.8 Certificate reproducibility

The authoritative P2 validation suite was rerun from the final research corpus. The complete run terminates with

80,717,789 checks – PASS.

Among the independently checked components are

component	count
brute minimum-demand cases	11,011
broad arithmetic cases	1,548,547
compact normalized macro cases	667,110
normalized rectangle vertices	1,859,269
normalized rectangle diamonds	1,286,746
raw four-symbol immediate cases	1,111,110
one-work-symbol raw cases	1,111,110
three-symbol raw cases	1,111,110
raw rectangle classifications	61,455
three-symbol raw rectangle vertices	731,091
three-symbol raw rectangle diamonds	474,684

The script also asserts the exact branch counts and survivor vectors in (C.2)–(C.3).

The later resource-frontier validation suite was likewise rerun and terminates with

79,606,679 checks – PASS.

These executable audits support the symbolic results and certify the bounded search. Except for the explicitly computer-assisted (3, 2) elimination, they are regression evidence rather than substitutes for proof.

C.9 Architecture-relative deferred-SEARCH width

For the centered literal-demand machine, put

$$h_{\text{ctr}} = \left\lfloor \frac{p + q - 1}{2} \right\rfloor.$$

In the canonical deferred-SEARCH model, after a witness demand block has been consumed, the remaining work stack is the same bare bottom marker for every debt value

$$-h_{\text{ctr}}, \dots, h_{\text{ctr}}.$$

Distinct residuals are continuation-distinguishable. Determinism therefore forces distinct control states for all these debt values. The exact width is

$$|Q|_{\min}^{\text{deferred}} = 2h_{\text{ctr}} + 1 = 2 \left\lfloor \frac{p + q - 1}{2} \right\rfloor + 1.$$

The centered construction attains the bound.

The deferred condition is essential: the final consumed witness cell is a pure pop, and materialization is postponed to a later transition. Eager encodings, alternate stack coordinates, or machines that do not store literal minimum demand are outside this lower-bound category.

C.10 Representation-independent hierarchy product bound

A different chapter deliberately leaves the counting-language representation class. For every $m \geq 2$ and $k \geq 1$, the binary witness language $B_{m,k}$ is recognized by a realtime DPDA with exactly m states and $k + 1$ pushdown symbols. Every competing DPDA in the stated *general* deterministic model—including opponents with arbitrary deterministic ε -moves and unrelated internal stack encodings—satisfies

$$|Q||\Gamma| \geq mk + 1. \tag{C.4}$$

Consequently the fixed-state stack-alphabet classes obey the strict hierarchy


$$\mathcal{D}_{m,k} \subsetneq \mathcal{D}_{m,k+1} \quad (m \geq 2, k \geq 1).$$

The proof does not count residual phases in a prescribed code. It isolates $mk + 1$ externally distinguishable affine completion slopes. Repetition, height-escape, and step-pump arguments show that one recurrent state/top-symbol mode cannot realize two distinct slopes. Since a DPDA has at most $|Q||\Gamma|$ nonempty modes, (C.4) follows.

This is representation-independent for the witness language $B_{m,k}$, but it must not be retroactively read as an unrestricted product lower bound for the rational counting languages $L_{p,q}$.

C.11 Final scope rule

The lower-bound hierarchy in this monograph is intentionally layered:



The arrows describe the historical research route, not logical implication between the theorem classes. Each lower bound remains attached to the opponent model stated in its theorem.

Appendix D

Reproducibility and Validation Ledger

D.1 Purpose

The symbolic proofs are the mathematical basis of the monograph. Executable audits serve three narrower purposes:

1. detect transcription errors in long transition tables and closed forms;
2. test boundary cases across large finite parameter ranges;
3. certify a finite bounded elimination when a theorem explicitly has a computer-assisted final step.

This appendix records exactly which validation artifacts were recovered from the final paper corpus, which were independently rerun, and which later artifacts had to be reconstructed from archived audit specifications.

D.2 Provenance classes

Every executable artifact is assigned one of the following provenance classes.

Recovered and rerun.

The exact script is present in the final paper corpus and was independently executed during the monograph audit.

Reconstructed and rerun.

The historical script named by a manuscript is absent from the final corpus. A replacement regression script was reconstructed from retained earlier code, the authoritative manuscript, and archived checkpoint specifications. The replacement reproduces the reported assertion total but is not claimed to be byte-for-byte identical to the missing historical file.

Source-reported only.

A manuscript or checkpoint reports a computation, but no independent rerun is being claimed in the monograph package.

The distinction is substantive. A reconstructed artifact must never be described as a recovered original.

D.3 Recovered and independently rerun suites**P1 — completion/frontier/original-machine suite**

Artifact in the final corpus:

`paper_validation_suite.py`

The independent rerun passes and includes:

check family	count
original-machine binary inputs	524,287
prefix transition checks	8,912,898
independent quotient-distance checks	25,281,628
compact-DPDA runs	1,282,610
parameter formula audits	3,931

The suite checks the original-machine residual invariant, reachable-store structure, completion normalization, shortest-completion formulas, compact realization, and related finite-frontier claims.

P2 — structure/resource suite

Artifact in the final corpus:

`second_paper_validation_suite.py`

The independent rerun terminates with

80,717,789 checks – PASS.

Principal reported subcounts are:

check family	count
brute minimum-demand cases	11,011
broad arithmetic cases	1,548,547
compact macro cases	667,110

configuration-rectangle vertices	1,859,269
configuration-rectangle diamonds	1,286,746
phase samples	44,047
raw four-symbol cases	1,111,110
raw four-symbol defects encountered	16,247
one-work-symbol cases	1,111,110
raw three-symbol cases	1,111,110
raw rectangle cases	61,455
raw three-symbol rectangle vertices	731,091
raw three-symbol rectangle diamonds	474,684
ladder paths	11,778

The same suite contains the bounded $(3, 2)$ rule-synthesis certificate recorded in Appendix C.

P5 — exact resource-frontier suite

Artifact in the final source package:

`DPDA4_validation_reference.py`

The independent rerun terminates with

79,606,679 checks – PASS.

The suite independently rechecks the minimum-demand arithmetic, normalized and raw configuration rectangles, raw-conjugacy boundaries, and the state-alphabet resource formulas used in the exact frontier paper.

P9 — centered real-time family audit

Artifact in the final corpus:

`centered_dpda_audit.py`

The independent rerun covers every ordered coprime pair

$$1 \leq p, q \leq 20$$

and reports

255 ordered coprime ratios

with

1,909,191 reached configurations

and the distributed result file records

3,738,864 binary edges.

The audit checks transition totality on reached binary configurations, the residual invariant, residual-collision consistency, two-cell visibility, and zero-residual correctness. The rerun passes.

D.4 Reconstructed P10 validation suite

The WSC manuscript names

`wsc_validation_v1_1.py`

and reports

477,033 assertion-level checks.

The exact historical file was not present in the final paper corpus or retained Library artifacts.

A replacement regression suite was therefore reconstructed on 19 August 2026 from the authoritative manuscript’s explicit audit scope and archived research checkpoints. The replacement checks:

- all coprime $1 \leq p, q \leq 8$ and all binary prefixes of length at most nine for the closed-form minimum-demand formula, semantic carry criterion, and demand clock;
- the bounded canonical one-top direct-realizability classification;
- WSC defect-lattice coordinate formulas;
- count-shadow fiber inequalities;
- physical-factorization reconstruction formulas;
- bounded carry-order debt instances;
- supplementary kernel-fiber identities.

The reconstructed suite reproduces the manuscript-reported total exactly:

477,033 checks – PASS.

It is stored in the monograph source package under the historical filename, but its file header explicitly records that it is a reconstruction rather than a recovered original.

D.5 Reconstructed P11 validation suite

The semantic-uniqueness/lift manuscript names

`uniqueness_lift_validation_v0_2.py`

and reports

458,776 assertion-level checks.

The exact v0.2 file was absent from the final paper corpus. However, the earlier artifact

`uniqueness_lift_validation_v0_1.py`

was recovered from the retained Library.

The replacement v0.2 suite preserves the exhaustive v0.1 core and adds the later schedule/causality checks stated in the authoritative v0.2/v0.6 manuscript. In particular it checks:

- the six admissible and two inadmissible 2:1 schedules;
- pairwise distinct reduced schedule fingerprints;
- the two-maxima/non-lattice schedule poset structure;
- the directed $A \rightsquigarrow B$ causal certificate;
- the one-kernel excess at the eager B successor;
- the first forbidden recognition witness.

The reconstructed suite reproduces the reported total exactly:

458,776 checks – PASS.

Again, this is a functional reconstruction, not a byte-for-byte recovery claim.

D.6 P12 and executable support

The General Theory is formulated and proved symbolically. Its verification chapter deliberately imports classical decision mechanisms only after reducing correctness to an exact structural target. No missing external script is required to establish the main P12 proofs.

Finite catalogues used as stress tests, including the 3:2 and odd- p :2 schedule families, are recorded separately in the General-Theory catalogue appendix rather than being treated as a substitute for the abstract theorems.

D.7 Preserved monograph code package

The LaTeX monograph project preserves the following executable files under `code/`:

- `P1_paper_validation_suite.py`;
- `P2_second_paper_validation_suite.py`;
- `P5_DPDA4_validation_reference.py`;
- `P9_centered_dpda_audit.py`;
- `wsc_validation_v1_1.py` (reconstructed);
- `uniqueness_lift_validation_v0_2.py` (reconstructed);
- a provenance README for the reconstructed artifacts;
- exact text outputs from the independently rerun P1/P2/P5/P9 suites.

The final archival release should additionally contain a machine-readable manifest of cryptographic hashes for these files and the compiled monograph source tree.

D.8 Interpretation rule

Throughout the book, the phrase “validated computationally” must be read according to the following hierarchy:

symbolic theorem first, executable regression evidence second.

The sole exception is an explicitly labeled computer-assisted theorem whose symbolic argument reduces the remaining obligation to a finite exhaustive search. In that case, the finite certificate is part of the proof and its pruning assumptions, survivor vectors, and terminal zero-survivor result must be preserved together.

Appendix E

Lagrange–Bürmann Derivation of the Exact Correction Triangle

E.1 The marked first-return grammar

This appendix supplies the coefficient derivation deferred from Chapter 20. Fix an integer $d \geq 1$. Let $F(z, t)$ denote the correction-marked excursion generating function at baseline zero, and let $A(z, t)$ and $B(z, t)$ be the two periodic tail classes introduced by the first-return decomposition. They satisfy

$$\boxed{F = 1 + zA^dF}, \tag{E.1}$$

$$\boxed{A = 1 + zA^dB}, \tag{E.2}$$

$$\boxed{B = 1 + ztA^{d-1}B^2}. \tag{E.3}$$

The variable z marks down-steps and t marks the distinguished correction events.

Define

$$\boxed{Q_d(n, k) = [z^n t^k]F(z, t)}.$$

The goal is to derive the closed form for $Q_d(n, k)$ directly from (E.1)–(E.3).

E.2 The auxiliary kernel variable

Set

$$Y = zA^{d-1}B.$$

Equation (E.2) gives

$$A = 1 + AY,$$

so

$$\boxed{A = \frac{1}{1 - Y}.} \quad (\text{E.4})$$

Similarly, divide (E.3) by B and use $Y = zA^{d-1}B$:

$$1 = B^{-1} + tY,$$

whence

$$\boxed{B = \frac{1}{1 - tY}.} \quad (\text{E.5})$$

Substituting (E.4)–(E.5) into the definition of Y yields

$$Y = \frac{z}{(1 - Y)^{d-1}(1 - tY)}.$$

Equivalently,

$$\boxed{Y(1 - Y)^{d-1}(1 - tY) = z.} \quad (\text{E.6})$$

Thus Y is the unique formal power series with zero constant term satisfying

$$\boxed{Y = z \Phi(Y), \quad \Phi(y) = (1 - y)^{-(d-1)}(1 - ty)^{-1}.} \quad (\text{E.7})$$

This is exactly the form required for Lagrange inversion.

E.3 Coefficient formula for the auxiliary kernel

Define

$$K_d(z, t) = \frac{Y}{z}.$$

Since

$$[z^n t^k] K_d = [z^{n+1} t^k] Y,$$

the basic Lagrange formula applied to (E.7) gives

$$[z^{n+1} t^k] Y = \frac{1}{n+1} [y^n t^k] \Phi(y)^{n+1}.$$

Now

$$\Phi(y)^{n+1} = (1 - y)^{-(d-1)(n+1)} (1 - ty)^{-(n+1)}.$$

Extracting the t^k coefficient gives

$$[t^k] (1 - ty)^{-(n+1)} = \binom{n+k}{k} y^k.$$

Therefore

$$[z^n t^k] K_d = \frac{1}{n+1} \binom{n+k}{k} [y^{n-k}] (1-y)^{-(d-1)(n+1)}.$$

Using

$$[y^r] (1-y)^{-a} = \binom{a+r-1}{r},$$

we obtain

$$\boxed{[z^n t^k] K_d = \frac{1}{n+1} \binom{n+k}{k} \binom{dn+d-k-2}{n-k}}. \quad (\text{E.8})$$

For $d = 2$ this reduces to

$$\boxed{[z^n t^k] K_2 = \frac{1}{n+1} \binom{n+k}{k} \binom{2n-k}{n-k}}. \quad (\text{E.9})$$

This is the auxiliary kernel triangle identified in the main text. It is not the actual correction distribution.

E.4 Recovering the correction generating function

From (E.1),

$$F = \frac{1}{1 - zA^d}.$$

Using (E.4),

$$\boxed{F = \sum_{m \geq 0} z^m A^{dm} = \sum_{m \geq 0} z^m (1-Y)^{-dm}}. \quad (\text{E.10})$$

This expansion is the convenient starting point for the exact correction triangle.

For fixed $n \geq 1$,

$$Q_d(n, k) = \sum_{m=0}^n [z^{n-m} t^k] (1-Y)^{-dm}. \quad (\text{E.11})$$

The term $m = n$ contains no positive power of Y , so it contributes

$$\boxed{\mathbf{1}_{\{k=0\}}}. \quad (\text{E.12})$$

The term $m = 0$ contributes nothing when $n \geq 1$. It remains to treat $1 \leq m \leq n-1$.

E.5 Lagrange–Bürmann extraction for a fixed summand

Fix $1 \leq m \leq n - 1$ and set

$$N = n - m > 0.$$

Apply Lagrange–Bürmann to

$$f(y) = (1 - y)^{-dm}.$$

Since

$$f'(y) = dm(1 - y)^{-dm-1},$$

we obtain

$$[z^N](1 - Y)^{-dm} = \frac{dm}{N}[y^{N-1}](1 - y)^{-dm-1}\Phi(y)^N.$$

Substitute (E.7):

$$[z^N](1 - Y)^{-dm} = \frac{dm}{N}[y^{N-1}](1 - y)^{-dm-1-(d-1)N}(1 - ty)^{-N}. \quad (\text{E.13})$$

Because $N = n - m$,

$$dm + 1 + (d - 1)N = (d - 1)n + m + 1.$$

Thus

$$[z^N t^k](1 - Y)^{-dm} = \frac{dm}{N}[y^{N-1} t^k](1 - y)^{-((d-1)n+m+1)}(1 - ty)^{-N}.$$

Now

$$[t^k](1 - ty)^{-N} = \binom{N + k - 1}{k} y^k,$$

so

$$[z^N t^k](1 - Y)^{-dm} = \frac{dm}{N} \binom{N + k - 1}{k} [y^{N-k-1}](1 - y)^{-((d-1)n+m+1)}.$$

A final binomial extraction gives

$$[y^{N-k-1}](1 - y)^{-((d-1)n+m+1)} = \binom{dn - k - 1}{N - k - 1}.$$

Hence

$$\boxed{[z^{n-m} t^k](1 - Y)^{-dm} = \frac{dm}{n - m} \binom{n - m + k - 1}{k} \binom{dn - k - 1}{n - m - k - 1}}. \quad (\text{E.14})$$

The coefficient is automatically zero unless

$$n - m \geq k + 1.$$

E.6 Exact correction triangle

Insert (E.12) and (E.14) into (E.11). Make the change of variable

$$j = n - m.$$

Then

$$m = n - j,$$

and the nonzero range is

$$k + 1 \leq j \leq n - 1.$$

Therefore

$$Q_d(n, k) = \mathbf{1}_{\{k=0\}} + \sum_{j=k+1}^{n-1} \frac{d(n-j)}{j} \binom{j+k-1}{k} \binom{dn-k-1}{j-k-1}. \quad (\text{E.15})$$

This is the formula stated in Chapter 20.

The upper support bound follows immediately. If $k > n - 2$, then there is no integer j with

$$k + 1 \leq j \leq n - 1,$$

and the indicator term vanishes because $k > 0$. Thus

$$Q_d(n, k) = 0 \quad (k > n - 2). \quad (\text{E.16})$$

E.7 Equivalent elimination through $U = A - 1$

For completeness, the scalar equation quoted in the main chapter follows directly from the same grammar. Put

$$U = A - 1.$$

From (E.2),

$$U = zA^d B = AY.$$

Since $A = 1 + U$, this gives

$$Y = \frac{U}{1 + U}.$$

Equation (E.3) can be written as

$$B = 1 + tYB,$$

so

$$B = \frac{1}{1 - tY} = \frac{1 + U}{1 + (1 - t)U}.$$

Substituting into $U = zA^dB$ with $A = 1 + U$ gives the convenient form

$$\boxed{U(1 + (1 - t)U) = z(1 + U)^{d+1}.} \quad (\text{E.17})$$

Equation (E.17) and the kernel equation (E.6) are related by

$$\boxed{Y = \frac{U}{1 + U}.}$$

Thus the two inversion routes are exactly the same formal system in different coordinates.

E.8 Row sum

Setting $t = 1$ forgets the correction marking. Equation (E.17) becomes

$$U = z(1 + U)^{d+1}.$$

At $t = 1$, the correction generating function simplifies to

$$F(z, 1) = 1 + U.$$

Lagrange inversion gives, for $n \geq 1$,

$$[z^n]U = \frac{1}{n}[u^{n-1}](1 + u)^{(d+1)n} = \frac{1}{n} \binom{(d+1)n}{n-1}.$$

Using

$$\frac{1}{n} \binom{(d+1)n}{n-1} = \frac{1}{dn+1} \binom{(d+1)n}{n},$$

we recover the Fuss–Catalan row sum

$$\boxed{\sum_{k \geq 0} Q_d(n, k) = \frac{1}{dn+1} \binom{(d+1)n}{n}.} \quad (\text{E.18})$$

E.9 Why the kernel and correction arrays differ

The kernel coefficients (E.8) arise from the single implicit series Y and record the algebraic kernel introduced to solve the periodic grammar. The correction coefficients (E.15) arise only after the full zero-baseline excursion series

$$F = \sum_{m \geq 0} z^m (1 - Y)^{-dm}$$

is reconstructed.

Therefore

K_d is an algebraic waypoint, whereas Q_d is the path statistic.

The distinction is not terminological: the two coefficient arrays have different combinatorial meanings and diverge as soon as the correction marking becomes nontrivial.

Appendix F

General-Theory Stress-Test Catalogues

F.1 Purpose and scope

The abstract General Theory developed in Part XI is not defined by the weighted-balance family. Its semantic system may be non-additive and need not have a residual space isomorphic to $\mathbb{Z}$. The centered $p:q$ constructions are instead a deliberately rigid stress-test family in which safe-reset domains, source-fiber expansion, directed exposure, and robust fatal supports can all be calculated exactly.

This appendix records the full finite 3:2 certificate and the infinite odd- p :2 arithmetic/source catalogues used in the stress-test theorems. The purpose is reproducibility: the main chapters can state the structural consequences without interrupting the narrative with long case catalogues.

Throughout this appendix, for a canonical literal-demand pair (x, y) we use

$$\eta(x, y) = \begin{cases} (10)^y 0^{x-y}, & x \geq y, \\ (10)^x 1^{y-x}, & y > x. \end{cases}$$

For the weighted-balance residual, the primitive nonnegative kernel vector is $K = (q, p)$.

F.2 The exact centered 3:2 certificate

Set

$$p = 3, \quad q = 2, \quad m = p + q = 5,$$

and use the centered states

$$D_{-2}, D_{-1}, C, D_1, D_2.$$

The six natural eager sites are indexed by the current centered state, the next input symbol, and the visible top symbol. Their eager replacement words are

site	(d, a, X)	γ
A	$(-2, 0, 0)$	11111
B	$(-2, 1, 0)$	10111
C	$(-2, 1, 1)$	11111
D	$(1, 0, 0)$	10100
E	$(1, 1, 0)$	10000
F	$(2, 0, 0)$	10000

(F.1)

F.2.1 Finite reset repertoires

Define

$$\begin{aligned} \mathcal{R}_0^{(3,2)} = & \{(2, 2)\} \cup \{(0, y) : y \geq 4\} \cup \{(1, y) : y \geq 3\} \\ & \cup \{(x, 1) : x \geq 3\} \cup \{(x, 0) : x \geq 4\} \end{aligned}$$

(F.2)

and

$$\mathcal{R}_D^{(3,2)} = \mathcal{R}_0^{(3,2)} \cup \{(x, 2) : x \geq 3\}.$$

(F.3)

Before the first reset event of an episode begun from one of these canonical pairs, the relevant eager-source shapes are

local context	$\mathcal{R}_0^{(3,2)}$	$\mathcal{R}_D^{(3,2)}$
$D_{-2}, \text{ top } 0$	0	0^k
$D_{-2}, \text{ top } 1$	1^k	1^k
$D_1, \text{ top } 0$	0^k	0^k
$D_2, \text{ top } 0$	0^k	0^k

(F.4)

The bottom-WRITE targets are

d	0	1
-2	$(1, 8)$	$(0, 4)$
-1	$(0, 4)$	$(1, 3)$
0	$(1, 3)$	$(2, 2)$
1	$(2, 2)$	$(3, 1)$
2	$(3, 1)$	$(4, 0)$

(F.5)

and every one lies in $\mathcal{R}_0^{(3,2)}$.

Writing the unary source tail as $t \geq 0$, the remaining reset targets are

reset action	canonical target pair	
positive boundary WRITE	$(t + 5, 0)$	
negative boundary WRITE	$(1, t + 9)$	
A	$(0, 5)$	(F.6)
B	$(1, 4)$	
C	$(0, t + 5)$	
D	$(t + 3, 2)$	
E	$(t + 4, 1)$	
F	$(t + 4, 1)$	

These tables give two closed reset regimes:

- (a) with D disabled, every baseline reset and every enabled eager action among A, B, C, E, F returns to $\mathcal{R}_0^{(3,2)}$;
- (b) with D enabled and A, B disabled, every baseline reset and every enabled eager action among C, D, E, F returns to $\mathcal{R}_D^{(3,2)}$.

The only new reset ray in the second regime is

$$\{(x, 2) : x \geq 3\},$$

and this is exactly the ray that enlarges the D_{-2} -top-0 source from the singleton stack 0 to the unary family 0^k .

F.2.2 Direct verification of reset closure

The canonical words in $\mathcal{R}_0^{(3,2)}$ are respectively

$$1010, \quad 1^y, \quad 101^{y-1}, \quad 10^x, \quad 0^x,$$

while the added D -ray has

$$\eta(x, 2) = 1010^{x-1}.$$

Before the first reset, MATCH and interior SEARCH only remove the visible top cell, so the current stack is always a suffix of the starting canonical word. The source catalogue in (F.4) follows immediately. In particular, without the D -ray only the word 1010 can reach D_{-2} while a zero remains visible, and exactly its final zero remains.

For the non-bottom reset events, the centered boundary blocks and the six eager replacement words yield the canonical words

$$0^{t+5}, \quad 101^{t+8}, \quad 1^5, \quad 101^3, \quad 1^{t+5}, \quad 1010^{t+2}, \quad 10^{t+4}.$$

These are precisely the encodings listed in (F.6). Every row except D lies in $\mathcal{R}_0^{(3,2)}$; the D row is precisely the additional ray of $\mathcal{R}_D^{(3,2)}$.

F.2.3 Exact schedule classification and fatal witnesses

The admissible schedules are exactly

$$\boxed{\text{Adm}_{3,2} = \{H : \{A, D\} \not\subseteq H, \{B, D\} \not\subseteq H\}.} \quad (\text{F.7})$$

Therefore

$$\boxed{40 \text{ schedules are admissible and } 24 \text{ are inadmissible.}} \quad (\text{F.8})$$

The minimal directed fatal fork is

$$\boxed{D \rightarrow A, \quad D \rightarrow B.} \quad (\text{F.9})$$

The maximal admissible faces are

$$\{A, B, C, E, F\}, \quad \{C, D, E, F\},$$

so the admissible complex is non-pure and the schedule poset is not a lattice. Its independence polynomial is

$$\boxed{(1+x)^3(1+3x+x^2).} \quad (\text{F.10})$$

Sufficiency follows from the two reset-closed regimes above: the $D = 0$ regime contributes $2^5 = 32$ schedules, while $D = 1$ with $A = B = 0$ contributes $2^3 = 8$ more.

Necessity is certified by two complete balanced words,

$$\boxed{w_B = 1111000011}, \quad \boxed{w_A = 011111110000011}. \quad (\text{F.11})$$

Their critical portions are

branch	reaching prefix	victim step	diagnostic
$D \rightarrow B$	11110000 : ($D_{-2}, 00$)	1 : ($C, 101110$)	1 : ($C, 011110$)
$D \rightarrow A$	011111110000 : ($D_{-2}, 000$)	0 : ($C, 11111100$)	11 : ($C, 111100$)

(F.12)

In the first row, the earlier eager D step occurs on prefix 11110; in the second it occurs on prefix 011111110. Before the victim step the only eager optional site encountered is therefore D . After the victim step the listed diagnostic uses only baseline center transitions. Both terminal configurations have semantic residual zero but a nonempty center stack, so they are nonaccepting under the fixed centered acceptance convention. The witnesses are therefore extension-stable under arbitrary choices at C, E, F and at the other victim site.

F.3 The infinite odd- p :2 family

Fix

$$\boxed{p = 2h - 1, \quad q = 2, \quad h \geq 2,} \quad (\text{F.13})$$

so that

$$m = p + q = 2h + 1.$$

For $t \geq 1$ set

$$\epsilon_t = t \bmod 2, \quad k_t = \left\lceil \frac{t}{2} \right\rceil, \quad \beta_t = t + p k_t,$$

and define

$$\boxed{\Gamma_t = \eta(\epsilon_t, \beta_t).} \quad (\text{F.14})$$

The six eager sites are

$$\begin{aligned} A &: (D_{-h}, 0T) \xrightarrow{0} (\mathbb{C}, \Gamma_h T), \\ B &: (D_{-h}, 0T) \xrightarrow{1} (\mathbb{C}, \Gamma_{h-1} T), \\ C &: (D_{-h}, 1T) \xrightarrow{1} (\mathbb{C}, \Gamma_h T), \\ D &: (D_{h-1}, 0T) \xrightarrow{0} (\mathbb{C}, (10)^h 0T), \\ E &: (D_{h-1}, 0T) \xrightarrow{1} (\mathbb{C}, (10)^{h-1} 000T), \\ F &: (D_h, 0T) \xrightarrow{0} (\mathbb{C}, (10)^{h-1} 000T). \end{aligned} \quad (\text{F.15})$$

F.3.1 Arithmetic catalogue

For the minimum-demand map μ the exact identities are:

(i) for every $t \geq 1$,

$$\boxed{\mu(-tm) = (\epsilon_t, \beta_t);} \quad (\text{F.16})$$

(ii) for $0 \leq j \leq h$,

$$\boxed{\mu(jm + 2) = (j + 2, p - j - 1);} \quad (\text{F.17})$$

(iii) for $1 \leq t \leq h + 1$,

$$\boxed{\mu(-tm + 2) = (\epsilon_t, \beta_t - 1);} \quad (\text{F.18})$$

(iv) at the positive outer boundary,

$$\boxed{\mu((h + 1)m) = (h + 3, h - 2).} \quad (\text{F.19})$$

For completeness, if $t = 2r$ then $(0, \beta_t)$ has weight

$$-2(2r + pr) = -2r(p + 2) = -tm,$$

while if $t = 2r + 1$ then $(1, \beta_t)$ has weight

$$p - 2(2r + 1 + p(r + 1)) = -(2r + 1)(p + 2) = -tm.$$

In both cases the first coordinate is below the kernel step $q = 2$, so the pair is kernel-reduced. Equation (F.17) follows from

$$p(j + 2) - 2(p - j - 1) = j(p + 2) + 2 = jm + 2,$$

with second coordinate below p . Equation (F.18) is obtained by lowering the second coordinate in (F.16) by one, and (F.19) follows from

$$p(h + 3) - 2(h - 2) = (h + 1)(2h + 1).$$

The complete family of canonical baseline reset pairs is therefore

$$\boxed{\begin{aligned} \mathcal{R}_h^{\text{base}} &= \{\mu(jm + 2) : -h - 1 \leq j \leq h\} \\ &\cup \{(x, h - 2) : x \geq h + 3\} \\ &\cup \{(\epsilon_{h+1}, y) : y \geq \beta_{h+1}\}. \end{aligned}} \quad (\text{F.20})$$

The two schedule-regime repertoires are

$$\boxed{\begin{aligned} \mathcal{R}_{-D,h} &= \mathcal{R}_h^{\text{base}} \cup \{(\epsilon_h, y) : y \geq \beta_h\} \\ &\cup \{(\epsilon_{h-1}, \beta_{h-1})\} \cup \{(x, h - 1) : x \geq h + 2\}, \end{aligned}} \quad (\text{F.21})$$

and

$$\boxed{\begin{aligned} \mathcal{R}_{D,h} &= \mathcal{R}_h^{\text{base}} \cup \{(\epsilon_h, y) : y \geq \beta_h\} \\ &\cup \{(x, h - 1) : x \geq h + 2\} \cup \{(x, h) : x \geq h + 1\}. \end{aligned}} \quad (\text{F.22})$$

Every pair in either repertoire has minority size at most h .

F.3.2 Source-geometry catalogue

Consider an episode beginning at a canonical center pair in one of the two repertoires and stopping at its first reset event. The exact source geometry is:

(i) in $\mathcal{R}_{-D,h}$,

$$D_{-h} \text{ with top } 0 \implies \boxed{\text{complete remaining stack } 0}; \quad (\text{F.23})$$

(ii) in $\mathcal{R}_{D,h}$,

$$D_{-h} \text{ with top } 0 \implies \boxed{\text{complete remaining stack } 0^n, n \geq 1}, \quad (\text{F.24})$$

and lengths $n \geq 2$ arise only from the additional D -ray $\{(x, h) : x \geq h+1\}$;

(iii) in either repertoire,

$$\begin{aligned} D_{-h} \text{ with top } 1 &\implies 1^n, \\ D_{h-1} \text{ with top } 0 &\implies 0^n, \\ D_h \text{ with top } 0 &\implies 0^n \end{aligned} \quad (\text{F.25})$$

for some $n \geq 1$.

The key observation is that, before the first reset event, MATCH and interior SEARCH only delete the visible top cell, so the current stack is a suffix of the canonical reset word from which the episode began. Reaching D_{-h} requires at least h more negative than positive SEARCHes. In $\mathcal{R}_{-D,h}$ the only canonical pair with enough of both symbols to leave a zero visible at that boundary is (h, h) , whose word is $(10)^h$; all h one-cells must have been negative SEARCHes, leaving exactly its final zero. The only new family in $\mathcal{R}_{D,h}$ that can leave a longer zero suffix is (x, h) , whose word is $(10)^h 0^{x-h}$.

The remaining unary-source conclusions follow from the same minority-size bound. A suffix containing the opposite symbol beyond the visible boundary cell would force minority size at least $h+1$, contradicting the defining families in (F.21)–(F.22).

F.3.3 Reset-target catalogue and closure

Writing the unary source stack length as $n \geq 1$, all eager reset targets are

site	source stack	canonical target pair
A	0	(ϵ_h, β_h)
B	0	$(\epsilon_{h-1}, \beta_{h-1})$
C	1^n	$(\epsilon_h, \beta_h + n - 1)$
D	0^n	$(n + h, h)$
E	0^n	$(n + h + 1, h - 1)$
F	0^n	$(n + h + 1, h - 1)$

(F.26)

The identities are literal: for instance

$$(10)^h 0 0^{n-1} = \eta(n + h, h)$$

and

$$(10)^{h-1} 000 0^{n-1} = \eta(n + h + 1, h - 1).$$

Consequently:

- (a) if D is disabled, baseline resets and enabled actions among A, B, C, E, F preserve $\mathcal{R}_{\neg D, h}$;
- (b) if D is enabled and A, B are disabled, baseline resets and enabled actions among C, D, E, F preserve $\mathcal{R}_{D, h}$.

F.3.4 Uniform D -induced source-fiber expansion

For every $x \geq h + 1$,

$$\boxed{(C, \eta(x, h)) \xrightarrow{0^{2h-1}} (D_{-h}, 0^{x-h+1}).} \quad (\text{F.27})$$

Indeed,

$$\eta(x, h) = (10)^h 0^{x-h}.$$

The first $2h - 2$ zeros remove $h - 1$ complete 10 pairs, and the final zero SEARCHes through the last 1, leaving the last alternating zero together with the zero tail. Thus enabling D expands the common A/B source fiber from

$$\boxed{\{0\}}$$

to

$$\boxed{\{0, 00, 000, \dots\}.} \quad (\text{F.28})$$

F.3.5 Uniform reachability of the length-three frontier

There is, for every $h \geq 2$, a word u_h that reaches

$$\boxed{(D_{-h}, 000)} \quad (\text{F.29})$$

from the initial centered configuration and encounters no eager optional site other than D .

Define

$$v_h = \begin{cases} 1^{(h-1)m/2}, & h \text{ odd,} \\ 0 1^{2h+(h-2)m/2}, & h \text{ even.} \end{cases} \quad (\text{F.30})$$

Then

$$\boxed{u_h = v_h 1^{2h-1} 0^{2h}.} \quad (\text{F.31})$$

The useful increment block is

$$(D_d, \varepsilon) \xrightarrow{1^m} (D_{d+2}, \varepsilon), \quad 0 \leq d \leq h - 2. \quad (\text{F.32})$$

The first symbol bottom-WRITEs

$$\mu(dm + 2) = (d + 2, p - d - 1),$$

whose canonical word has length $p + 1 = 2h$; the remaining $2h$ ones erase the word, and its $d + 2$ zeros contribute the required positive SEARCHes. This builds v_h so that the machine reaches (D_{h-1}, ε) . The next 1^{2h-1} segment leaves $(D_{h-1}, 00)$, the following zero invokes D and produces $(C, (10)^h 00)$, and the remaining 0^{2h-1} reaches $(D_{-h}, 000)$. No site among A, B, C, E, F is encountered before that endpoint.

F.3.6 Uniform kernel-suffix fatality

For every $t \geq 1$, Γ_t has suffix 1^p . Hence there is a unique word z_t such that

$$\boxed{\Gamma_t = z_t 1^p.} \quad (\text{F.33})$$

If $t = 2r$ then

$$\Gamma_t = 1^{\beta_t}, \quad \beta_t = rm \geq m > p,$$

while if $t = 2r + 1$ then

$$\Gamma_t = 10 1^{\beta_t-1}, \quad \beta_t - 1 = p + rm \geq p.$$

Thus the suffix claim is uniform in t .

At the common source $(D_{-h}, 000)$, eager A and B yield

$$A : (D_{-h}, 000) \xrightarrow{0} (C, \Gamma_h 00) = (C, z_h 1^p 00),$$

$$B : (D_{-h}, 000) \xrightarrow{1} (C, \Gamma_{h-1} 00) = (C, z_{h-1} 1^p 00).$$

Reading z_h or z_{h-1} from the center uses only MATCHes and reaches

$$\boxed{(C, 1^p 00).} \quad (\text{F.34})$$

The stack has two zeros and p ones, hence weight

$$p \cdot 2 - 2 \cdot p = 0,$$

but is nonempty. Under centered acceptance it is therefore a site-free fatal diagnostic.

F.3.7 Exact odd- p :2 classification

For every odd $p \geq 3$,

$$\boxed{\text{Adm}_{p,2} = \text{Ind}\{\{A, D\}, \{B, D\}\}.} \quad (\text{F.35})$$

Equivalently,

$$\boxed{H \in \text{Adm}_{p,2} \iff \neg(A \wedge D) \wedge \neg(B \wedge D).} \quad (\text{F.36})$$

Every member of the family therefore has exactly 40 admissible and 24 inadmissible schedules, with facets

$$\{A, B, C, E, F\}, \quad \{C, D, E, F\},$$

and the same fatal fork

$$\boxed{D \rightarrow A, \quad D \rightarrow B.} \quad (\text{F.37})$$

Sufficiency is the reset-closure induction: $D = 0$ gives $2^5 = 32$ schedules, and $D = 1$ with $A = B = 0$ gives another $2^3 = 8$. Necessity uses the common reaching prefix u_h from (F.31) together with the complete counterexample words

$$\boxed{w_A^{(h)} = u_h 0 z_h,} \quad \boxed{w_B^{(h)} = u_h 1 z_{h-1}.} \quad (\text{F.38})$$

The prefix u_h encounters no optional eager site except D , and the post-victim suffixes are site-free. The two forbidden supports are therefore robust under arbitrary additions of C, E, F and of the other victim site.

F.4 Interpretive boundary

The finite 3:2 calculation and the infinite odd- p :2 family have a single role in the monograph: they certify that the abstract mechanisms of Part XI are not artifacts of one numerical instance. They exhibit, in exact families,

$$\boxed{\begin{array}{l} \text{safe reset} \longrightarrow \text{fiber expansion} \\ \longrightarrow \text{directed exposure} \longrightarrow \text{robust fatal support.} \end{array}}$$

They do *not* define the General Theory. The General Theory applies to an exact physical lift over a deterministic semantic system through

$$\Theta \delta_a = U_a \Theta,$$

with no requirement that the semantic system be a weighted-balance residual space or even an additive group.

Bibliography

Bibliography

- [1] A. E. Bütün. *Completion Frontiers and Quotient Geometry in Deterministic Pushdown Automata for Rational Binary Counting*. Manuscript, 2026.
- [2] A. E. Bütün. *Exact Structure and Resource Bounds in Deterministic Pushdown Automata for Binary Counting Languages*. Manuscript, 2026.
- [3] A. E. Bütün. *A Binary Stack-Alphabet Hierarchy for Deterministic Pushdown Automata*. Manuscript, 2026.
- [4] A. E. Bütün. *Exact Resource Frontiers for Residual-Faithful Deterministic Pushdown Automata: State-Stack Tradeoffs for Rational Binary Counting Quotient Graphs*. Manuscript, 2026.
- [5] A. E. Bütün. *Higher-Dimensional Completion Geometry for Linear Counting Languages*. Manuscript, 2026.
- [6] A. E. Bütün. *From Pushdown Cancellation to Defect Geometry on Fuss-Catalan Paths*. Manuscript, 2026.
- [7] A. E. Bütün. *Chronological Carry Refinement and p -Adic Root Geometry*. Manuscript, 2026.
- [8] A. E. Bütün. *Real-Time $p:q$ Deterministic Pushdown Automata: A Uniform Centered Construction*. Manuscript, 2026.
- [9] A. E. Bütün. *Write-Search-Cancel: A Structural Theory of Real-Time Literal-Demand Normalization*. Manuscript, 2026.
- [10] A. E. Bütün. *Semantic Uniqueness and LIFO Lift Theory for Real-Time Pushdown Normalization*. Manuscript, 2026.
- [11] A. E. Bütün. *A General Theory of Real-Time Pushdown Normalization*. Manuscript, 2026.
- [12] A. Abrams and R. Ghrist. State complexes for metamorphic robots. *The International Journal of Robotics Research*, 23(7–8):811–826, 2004.

-
- [13] A. Adelberg and T. Lengyel, New results on the p -adic valuation of Stirling numbers, arXiv:2111.08766 [math.NT], 2021.
 - [14] R. Alur, V. Kumar, P. Madhusudan, and M. Viswanathan, Congruences for visibly pushdown languages, in *ICALP 2005*, Lecture Notes in Computer Science, vol. 3580, Springer, 2005, pp. 1102–1114. doi:10.1007/11523468_89.
 - [15] R. Alur and P. Madhusudan, Visibly pushdown languages, in *Proceedings of the 36th Annual ACM Symposium on Theory of Computing (STOC 2004)*, ACM, 2004, pp. 202–211. doi:10.1145/1007352.1007390.
 - [16] A. Asinowski, B. Hackl, and S. J. Selkirk, Down-step statistics in generalized Dyck paths, *Discrete Mathematics & Theoretical Computer Science* 24 (2022), no. 1, article 7163.
 - [17] J. Autry, P. Graves, J. Loucks, C. O’Neill, V. Ponomarenko, and S. Yih. Squarefree divisor complexes of certain numerical semigroup elements. *Involve*, 14(1):1–9, 2021.
 - [18] J.-C. Aval, Multivariate Fuss–Catalan numbers, *Discrete Mathematics* 308 (2008), 4660–4669.
 - [19] A. Bouajjani, J. Esparza, and O. Maler, Reachability analysis of pushdown automata: Application to model-checking, in *CONCUR 1997: Concurrency Theory*, Lecture Notes in Computer Science, vol. 1243, Springer, 1997, pp. 135–150. doi:10.1007/3-540-63141-0_10.
 - [20] J.-L. Baril, P. E. Harris, K. J. Harry, M. McClinton, and J. L. Ramírez, Enumerating runs, valleys, and peaks in Catalan words, arXiv:2404.05672 [math.CO], 2024.
 - [21] S. Böhm, S. Göller, and P. Jančár. Equivalence of deterministic one-counter automata is NL-complete. In *Proceedings of the 45th Annual ACM Symposium on Theory of Computing (STOC 2013)*, pp. 131–140, 2013. doi:10.1145/2488608.2488626.
 - [22] A. Butoi, B. DuSell, T. Vieira, R. Cotterell, and D. Chiang. Algorithms for weighted pushdown automata. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 9669–9680. Association for Computational Linguistics, 2022. doi:10.18653/v1/2022.emnlp-main.656.
 - [23] A. Carayol and O. Serre. Marking shortest paths on pushdown graphs does not preserve MSO decidability. *Information Processing Letters*, 116(10):638–643, 2016. doi:10.1016/j.ipl.2016.04.015.

- [24] D. Chistikov. Notes on counting with finite machines. In *34th International Conference on Foundation of Software Technology and Theoretical Computer Science (FSTTCS 2014)*, LIPIcs 29, pp. 339–350, 2014. doi:10.4230/LIPIcs.FSTTCS.2014.339.
- [25] D. Chistikov, W. Czerwiński, P. Hofman, M. Pilipczuk, and M. Wehar. Shortest paths in one-counter systems. *Logical Methods in Computer Science*, 15(1):19:1–19:28, 2019. doi:10.23638/LMCS-15(1:19)2019.
- [26] D. Chistikov and R. Majumdar. Unary pushdown automata and straight-line programs. In *Automata, Languages, and Programming: 41st International Colloquium (ICALP 2014), Part II*, LNCS 8573, pp. 146–157. Springer, 2014. doi:10.1007/978-3-662-43951-7_13.
- [27] P. Diaconis and J. Fulman, Carries, shuffling, and symmetric functions, arXiv:0902.0179 [math.PR], 2009.
- [28] P. Diaconis and B. Sturmfels. Algebraic algorithms for sampling from conditional distributions. *The Annals of Statistics*, 26(1):363–397, 1998.
- [29] M. Elder, M. Kambites, and G. Ostheimer. On groups and counter automata. *International Journal of Algebra and Computation*, 18(8):1345–1364, 2008. arXiv:math/0611188.
- [30] H. Fernau and P. Wolf, Synchronization of deterministic visibly pushdown automata, in *FSTTCS 2020*, LIPIcs, vol. 182, Article 45, 2020. doi:10.4230/LIPIcs.FSTTCS.2020.45.
- [31] H. Fernau, P. Wolf, and T. Yamakami, Synchronizing deterministic pushdown automata can be really hard, *Information and Computation* 295(B) (2023), 105089. doi:10.1016/j.ic.2023.105089.
- [32] P. Flajolet and A. Odlyzko, Singularity analysis of generating functions, *SIAM Journal on Discrete Mathematics* 3 (1990), no. 2, 216–240.
- [33] T. Gagie, G. Manzini, and J. Sirén, Wheeler graphs: A framework for BWT-based data structures, *Theoretical Computer Science* 698 (2017), 67–78. doi:10.1016/j.tcs.2017.06.016.
- [34] P. A. García-Sánchez, I. Ojeda, and A. Sánchez-R.-Navarro. Factorization invariants in half-factorial affine semigroups. *International Journal of Algebra and Computation*, 23(1):111–122, 2013.
- [35] O. Gauwin, A. Muscholl, and M. Raskin. Minimization of visibly pushdown automata is NP-complete. *Logical Methods in Computer Science*, 16(1):14, 2020. doi:10.23638/LMCS-16(1:14)2020.

-
- [36] J. Goldstine, M. Kappes, C. M. Kintala, H. Leung, A. Malcher, and D. Wotschke. Descriptive complexity of machines with limited resources. *Journal of Universal Computer Science*, 8(2):193–234, 2002. DOI: 10.3217/jucs-008-02-0193.
 - [37] J. Goldstine, J. K. Price, and D. Wotschke. On reducing the number of stack symbols in a PDA. *Mathematical Systems Theory*, 26:313–326, 1993. DOI: 10.1007/BF01189852.
 - [38] J. E. Graver. On the foundations of linear and integer linear programming I. *Mathematical Programming*, 9:207–226, 1975.
 - [39] V. Guba and M. Sapir. *Diagram Groups*. Memoirs of the American Mathematical Society, vol. 130, no. 620, 1997.
 - [40] A. Hatcher. *Algebraic Topology*. Cambridge University Press, 2002.
 - [41] C. Heuberger, S. J. Selkirk, and S. Wagner, Enumeration of generalized Dyck paths based on the height of down-steps modulo k , *Electronic Journal of Combinatorics* 30 (2023), no. 1, Paper 1.26, 18 pp.
 - [42] S. Hong and M. Qiu, On the p -adic properties of Stirling numbers of the first kind, *Acta Mathematica Hungarica* 161 (2020), 366–395.
 - [43] O. H. Ibarra and I. McQuillan, On store languages of language acceptors, *Theoretical Computer Science* 745 (2018), 114–132. doi:10.1016/j.tcs.2018.05.036.
 - [44] O. H. Ibarra and I. McQuillan. Language acceptors with a pushdown: characterizations and complexity. *International Journal of Foundations of Computer Science*, 36(3):345–370, 2025. doi:10.1142/S0129054124430044.
 - [45] P. Jančár. A short decidability proof for DPDA language equivalence via first-order grammars. CoRR abs/1010.4760, 2011.
 - [46] L. Kiss and B. Rován. Descriptive complexity of push down automata. In *Information Technologies – Applications and Theory (ITAT 2020)*, CEUR Workshop Proceedings, Vol. 2718, pp. 59–66, 2020.
 - [47] D. Kozlov. *Combinatorial Algebraic Topology*. Algorithms and Computation in Mathematics, vol. 21, Springer, 2008.
 - [48] D. Krob, J. Mairesse, and I. Michos. Computing the average parallelism in trace monoids. *Discrete Mathematics*, 273(1–3):131–162, 2003.

- [49] A. Kučera and R. Mayr, A generic framework for checking semantic equivalences between pushdown automata and finite-state automata, *Journal of Computer and System Sciences* 91 (2018), 82–103. doi:10.1016/j.jcss.2017.09.004.
- [50] M. Kutrib and A. Malcher, Reversible pushdown automata, *Journal of Computer and System Sciences* 78(6) (2012), 1814–1827. doi:10.1016/j.jcss.2011.12.004.
- [51] M. Kutrib and A. Malcher. Reversible computations of one-way counter automata. In *Proceedings of NCMA 2022*, EPTCS 367, pp. 126–142, 2022. arXiv:2208.14720.
- [52] T. Mansour and M. Shattuck, Avoidance of vincular patterns by Catalan words, arXiv:2405.12435 [math.CO], 2024.
- [53] O. Martynova. Exact descriptonal complexity of determinization of input-driven pushdown automata. In *Implementation and Application of Automata (CIAA 2024)*, pp. 249–260, 2024. arXiv:2404.10516.
- [54] T. Masopust. A note on limited pushdown alphabets in stateless deterministic pushdown automata. *International Journal of Foundations of Computer Science*, 24(3):319–328, 2013. DOI: 10.1142/S0129054113500068; arXiv:1208.5002.
- [55] D. A. Moj, Biorthogonal eigenvectors of the Holte carry matrix and cascade-free enumeration, arXiv:2604.04591 [math.CO], 2026.
- [56] M. Moschetti and C. O’Neill. Counting edges in factorization graphs of numerical semigroup elements. *Involve*, 18(5):861–871, 2025. DOI: 10.2140/involve.2025.18.861.
- [57] Walter D. Neumann and Michael Shapiro. Regular geodesic normal forms in virtually abelian groups. *Bulletin of the Australian Mathematical Society*, 55(3):517–519, 1997. doi:10.1017/S0004972700034171.
- [58] The OEIS Foundation, The On-Line Encyclopedia of Integer Sequences, A371395, triangle $T(n, k) = \binom{n+k}{k} \binom{2n-k}{n-k} / (n+1)$, accessed August 2026.
- [59] M. Pétréolle, A. D. Sokal, and B.-X. Zhu, Lattice paths and branched continued fractions: An infinite sequence of generalizations of the Stieltjes–Rogers and Thron–Rogers polynomials, with coefficientwise Hankel-total positivity, *Memoirs of the American Mathematical Society* 291 (2023), no. 1450.

- [60] G. Pighizzini. Deterministic pushdown automata and unary languages. *International Journal of Foundations of Computer Science*, 20(4):629–645, 2009. doi:10.1142/S0129054109006784.
- [61] G. Sénizergues. $L(A) = L(B)$? Decidability results from complete formal systems. *Theoretical Computer Science*, 251(1–2):1–166, 2001. doi:10.1016/S0304-3975(00)00285-1.
- [62] C. C. Squier. Word problems and a homological finiteness condition for monoids. *Journal of Pure and Applied Algebra*, 49(1–2):201–217, 1987.
- [63] B. Sturmfels. On vector partition functions. *Journal of Combinatorial Theory, Series A*, 72(2):302–309, 1995.
- [64] L. G. Valiant and M. S. Paterson, Deterministic one-counter automata, *Journal of Computer and System Sciences* 10(3) (1975), 340–350. doi:10.1016/S0022-0000(75)80005-5.